

Dataset import

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from statsmodels.tsa.arima.model import ARIMA
from statsmodels.tsa.statespace.sarimax import SARIMAX
from sklearn.metrics import mean_absolute_error, mean_squared_error
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import LSTM, Dense, Dropout
from sklearn.preprocessing import MinMaxScaler
import warnings
warnings.filterwarnings("ignore")

# Load dataset
df = pd.read_csv("C:\\\\Users\\\\marat\\\\OneDrive\\\\Desktop\\\\project\\\\global
temperature.csv")

df['Date'] = pd.to_datetime(df[['Year', 'Month', 'Day']])
df.set_index('Date', inplace=True)
df = df[['Temperature', 'Anomaly']]

# Resample to monthly data
df.tail()

      Temperature  Anomaly
Date
2022-07-27      10.229    1.639
2022-07-28      10.221    1.631
2022-07-29      10.164    1.574
2022-07-30      10.167    1.577
2022-07-31      10.219    1.629
```

Descriptive Statistics

```
import pandas as pd

df = pd.read_csv(r"C:\Users\marat\OneDrive\Desktop\project\global
temperature.csv")
# Keep only relevant columns
df1 = df[['Year', 'Month', 'Temperature']]

# Descriptive statistics for each year
yearly_stats = df1.groupby('Year')[['Temperature']].describe()
```

```

# Descriptive statistics for each month across all years
monthly_stats = df1.groupby('Month')['Temperature'].describe()

# Grouping into specific 20-year periods
bins = list(range(df1['Year'].min(), df1['Year'].max() + 1, 20))
bin_labels = [f"{bins[i]}-{bins[i+1]-1}" for i in range(len(bins)-1)]
df1['Year Group'] = pd.cut(df1['Year'], bins=bins, labels=None) #
Changed to None to avoid label issues
grouped_20yr_stats = df1.groupby('Year Group')
['Temperature'].describe()

# Splitting the dataset into two equal time periods
mid_year = df1['Year'].median()
df1['Half'] = df1['Year'].apply(lambda x: 'First Half' if x <=
mid_year else 'Second Half')
half_stats = df1.groupby('Half')['Temperature'].describe()

# Overall statistics
overall_stats = df1['Temperature'].describe()

# Compute average temperature for each month of each year
monthly_avg_per_year = df1.groupby(['Year', 'Month'])
['Temperature'].mean().unstack()

# Compute average temperature for each year
yearly_avg_temp = df1.groupby('Year')['Temperature'].mean()

# Calculate 10-year averages
bins_10yr = list(range(df1['Year'].min(), df1['Year'].max() + 1, 10))
bin_labels_10yr = [f"{bins_10yr[i]}-{bins_10yr[i+1]-1}" for i in
range(len(bins_10yr)-1)]
df1['Year Group 10'] = pd.cut(df1['Year'], bins=bins_10yr,
labels=bin_labels_10yr, right=False)
grouped_10yr_avg = df1.groupby('Year Group 10')['Temperature'].mean()

# ---- Descriptive Statistics for 10-Year Groups ----
bins_10yr = list(range(df1['Year'].min(), df1['Year'].max() + 10, 10))
labels_10yr = [f"{bins_10yr[i]}-{bins_10yr[i+1]-1}" for i in
range(len(bins_10yr)-1)]
df1['Year Group 10'] = pd.cut(df1['Year'], bins=bins_10yr,
labels=labels_10yr, right=False)
grouped_10yr_stats = df1.groupby('Year Group 10')
['Temperature'].describe()

# ---- Descriptive Statistics for 20-Year Groups ----
bins_20yr = list(range(df1['Year'].min(), df1['Year'].max() + 20, 20))
labels_20yr = [f"{bins_20yr[i]}-{bins_20yr[i+1]-1}" for i in
range(len(bins_20yr)-1)]
df1['Year Group 20'] = pd.cut(df1['Year'], bins=bins_20yr,
labels=labels_20yr, right=False)

```

```

grouped_20yr_stats = df1.groupby('Year Group 20')
['Temperature'].describe()

# ---- Descriptive Statistics for 35-Year Groups ---
bins_35yr = list(range(df1['Year'].min(), df1['Year'].max() + 35, 35))
labels_35yr = [f'{bins_35yr[i]}-{bins_35yr[i+1]-1}' for i in
range(len(bins_35yr)-1)]
df1['Year Group 35'] = pd.cut(df1['Year'], bins=bins_35yr,
labels=labels_35yr, right=False)
grouped_35yr_stats = df1.groupby('Year Group 35')
['Temperature'].describe()

# Save results to an Excel workbook
with pd.ExcelWriter("temperature_analysis_new1.xlsx") as writer:
    yearly_stats.to_excel(writer, sheet_name="Yearly Statistics")
    monthly_stats.to_excel(writer, sheet_name="Monthly Statistics")
    grouped_20yr_stats.to_excel(writer, sheet_name="20-Year Groups")
    half_stats.to_excel(writer, sheet_name="Two Halves")
    overall_stats.to_excel(writer, sheet_name="Overall Statistics")
    monthly_avg_per_year.to_excel(writer, sheet_name="Monthly Avg Per
Year")
    yearly_avg_temp.to_excel(writer, sheet_name="Yearly Avg Temp")
    grouped_10yr_avg.to_excel(writer, sheet_name="10-Year Averages")
    grouped_10yr_stats.to_excel(writer, sheet_name="10-Year Stats")
    grouped_20yr_stats.to_excel(writer, sheet_name="20-Year Stats")
    grouped_35yr_stats.to_excel(writer, sheet_name="35-Year Stats")
# Save 10-year averages

# Display results
selected_years = list(range(1880, 2021, 10))
print("Descriptive Statistics by Selected Years (Every 10 Years):\n",
yearly_stats.loc[selected_years])
print("\nDescriptive Statistics by Month:\n", monthly_stats)
print("\nDescriptive Statistics by 20-Year Groups:\n",
grouped_20yr_stats)
print("\nDescriptive Statistics for Two Equal Halves:\n", half_stats)
print("\nOverall Statistics:\n", overall_stats)
print("\nMonthly Average Temperature per Year (First 5 Years):\n",
monthly_avg_per_year.head())
print("\nYearly Average Temperature (First 5 Years):\n",
yearly_avg_temp.head())
print("\n10-Year Averages:\n", grouped_10yr_avg) # Display 10-year
averages
print("\nDescriptive Statistics by 10-Year Groups:\n",
grouped_10yr_stats)
print("\nDescriptive Statistics by 20-Year Groups:\n",
grouped_20yr_stats)
print("\nDescriptive Statistics by 35-Year Groups:\n",
grouped_35yr_stats)

```

Descriptive Statistics by Selected Years (Every 10 Years):								
	count	mean	std	min	25%	50%	75%	
max								
Year								
1880	366.0	7.995557	0.370984	6.945	7.76175	8.063	8.29675	
8.666								
1890	365.0	7.989942	0.300809	6.933	7.77400	8.003	8.20300	
8.770								
1900	365.0	8.528134	0.426486	6.916	8.34500	8.572	8.82600	
9.367								
1910	365.0	8.211362	0.395085	6.751	8.03100	8.280	8.47400	
8.974								
1920	366.0	8.305962	0.436845	6.956	8.14125	8.366	8.55550	
9.524								
1930	365.0	8.508871	0.430060	6.978	8.26500	8.479	8.74800	
9.909								
1940	366.0	8.680478	0.345768	7.759	8.48550	8.681	8.92300	
9.566								
1950	365.0	8.292293	0.390372	7.121	8.06000	8.355	8.58000	
9.142								
1960	366.0	8.484902	0.500478	7.228	8.14100	8.556	8.79875	
9.672								
1970	365.0	8.636427	0.397900	7.476	8.38200	8.619	8.88500	
9.854								
1980	366.0	8.936541	0.410590	7.656	8.68350	8.962	9.17225	
10.136								
1990	365.0	9.236841	0.530523	8.025	8.87500	9.169	9.56300	
10.847								
2000	366.0	9.169997	0.490043	8.017	8.83100	9.166	9.44050	
10.752								
2010	365.0	9.713625	0.471451	8.292	9.40800	9.693	10.00400	
10.958								
2020	366.0	10.058536	0.431781	9.002	9.78600	10.032	10.29725	
11.327								

Descriptive Statistics by Month:								
	count	mean	std	min	25%	50%	75%	
max								
Month								
1	4433.0	8.606167	0.836344	5.862	8.00400	8.578	9.17500	
11.055								
2	4039.0	8.622804	0.842865	6.117	8.00800	8.539	9.18250	
11.545								
3	4433.0	8.738859	0.826366	6.821	8.13500	8.625	9.26800	
11.491								
4	4290.0	8.902750	0.690573	7.156	8.40825	8.817	9.33800	
11.352								
5	4433.0	8.764846	0.591479	7.246	8.35600	8.668	9.10100	

10.924								
6	4290.0	8.720883	0.542304	7.567	8.33700	8.615	9.00475	
10.518								
7	4433.0	8.742596	0.476396	7.612	8.41300	8.633	8.95600	
10.499								
8	4402.0	8.675235	0.506688	7.594	8.32125	8.554	8.93800	
10.566								
9	4260.0	8.628015	0.538040	7.448	8.25075	8.509	8.92000	
10.351								
10	4402.0	8.754561	0.586063	7.054	8.34200	8.695	9.09600	
10.623								
11	4260.0	8.526160	0.674768	6.751	8.05900	8.454	8.94225	
10.945								
12	4402.0	8.572129	0.740657	6.412	8.07400	8.489	9.06275	
11.020								

Descriptive Statistics by 20-Year Groups:

	count	mean	std	min	25%	50%
--	-------	------	-----	-----	-----	-----

75% \
Year Group 20

1880-1899	7305.0	8.112854	0.433809	5.862	7.881	8.163
8.4010						
1900-1919	7304.0	8.258179	0.431921	6.412	8.015	8.286
8.5390						
1920-1939	7305.0	8.478008	0.434693	6.597	8.245	8.482
8.7220						
1940-1959	7305.0	8.594258	0.433365	6.803	8.341	8.622
8.8690						
1960-1979	7305.0	8.582274	0.458677	6.596	8.290	8.581
8.8860						
1980-1999	7305.0	9.025537	0.502520	7.257	8.694	9.036
9.3560						
2000-2019	7305.0	9.606187	0.503268	7.906	9.283	9.602
9.9210						
2020-2039	943.0	9.944686	0.410681	8.162	9.674	9.916
10.1725						

max

Year Group 20

1880-1899	9.583
1900-1919	9.704
1920-1939	10.209
1940-1959	10.152
1960-1979	10.282
1980-1999	10.847
2000-2019	11.545
2020-2039	11.352

Descriptive Statistics for Two Equal Halves:

		count	mean	std	min	25%	50%	75%
max								
Half								
First Half	26297.0	8.337875	0.469145	5.862	8.070	8.363	8.633	
10.209								
Second Half	25780.0	9.046162	0.661319	6.596	8.567	9.011	9.503	
11.545								
Overall Statistics:								
count	52077.000000							
mean	8.688503							
std	0.673085							
min	5.862000							
25%	8.250000							
50%	8.603000							
75%	9.088000							
max	11.545000							
Name: Temperature, dtype: float64								
Monthly Average Temperature per Year (First 5 Years):								
Month	1	2	3	4	5	6		
7	\							
Year								
1880	7.690903	7.337310	7.564581	7.773167	8.208258	8.272533		
8.381677								
1881	8.209839	8.072500	8.336581	8.603800	8.487452	8.023933		
8.467903								
1882	8.879806	8.641679	8.333516	8.068467	8.179645	8.012533		
8.102613								
1883	7.413419	7.409107	8.157935	8.131200	8.213290	8.495900		
8.212645								
1884	8.143226	7.980586	7.422839	7.490367	7.503839	7.783433		
7.944032								
Month	8	9	10	11	12			
Year								
1880	8.314419	7.988467	8.050000	8.178300	8.152032			
1881	8.531065	8.223300	8.074387	7.776267	8.227097			
1882	8.354968	8.234467	7.841419	7.505133	7.274097			
1883	7.983774	7.902900	8.142387	7.647100	8.181710			
1884	8.070097	7.807233	7.857484	7.721533	7.896355			
Yearly Average Temperature (First 5 Years):								
Year								
1880	7.995557							
1881	8.255378							
1882	8.116529							
1883	7.995145							

```
1884    7.801880
Name: Temperature, dtype: float64
```

10-Year Averages:

```
Year Group 10
1880-1889    8.051050
1890-1899    8.174676
1900-1909    8.260697
1910-1919    8.255661
1920-1929    8.416542
1930-1939    8.539491
1940-1949    8.647585
1950-1959    8.540916
1960-1969    8.553706
1970-1979    8.610850
1980-1989    8.896096
1990-1999    9.155014
2000-2009    9.485445
2010-2019    9.726963
```

```
Name: Temperature, dtype: float64
```

Descriptive Statistics by 10-Year Groups:

	count	mean	std	min	25%	50%
75% \						
Year Group 10						
1880-1889	3653.0	8.051050	0.423024	6.529	7.81400	8.0750
8.33200						
1890-1899	3652.0	8.174676	0.435706	5.862	7.96700	8.2420
8.45125						
1900-1909	3652.0	8.260697	0.427165	6.662	8.01400	8.2800
8.53200						
1910-1919	3652.0	8.255661	0.436670	6.412	8.01500	8.2905
8.54700						
1920-1929	3653.0	8.416542	0.420882	6.597	8.21600	8.4300
8.63900						
1930-1939	3652.0	8.539491	0.439623	6.793	8.28775	8.5540
8.79300						
1940-1949	3653.0	8.647585	0.400426	6.803	8.42800	8.6640
8.89700						
1950-1959	3652.0	8.540916	0.457861	7.024	8.25400	8.5675
8.83400						
1960-1969	3653.0	8.553706	0.455693	6.596	8.25200	8.5620
8.85100						
1970-1979	3652.0	8.610850	0.459933	6.610	8.32675	8.5980
8.91000						
1980-1989	3653.0	8.896096	0.470858	7.352	8.57400	8.9000
9.23200						
1990-1999	3652.0	9.155014	0.499890	7.257	8.83775	9.1560
9.45225						

2000-2009	3653.0	9.485445	0.483732	7.906	9.17200	9.4830
9.79800						
2010-2019	3652.0	9.726963	0.493415	8.156	9.41400	9.7050
10.02700						

2020-2029 943.0 9.944686 0.410681 8.162 9.67400 9.9160

10.17250

Year Group 10	max
1880-1889	9.550
1890-1899	9.583
1900-1909	9.704
1910-1919	9.671
1920-1929	10.209
1930-1939	10.166
1940-1949	10.015
1950-1959	10.152
1960-1969	10.282
1970-1979	10.007
1980-1989	10.506
1990-1999	10.847
2000-2009	11.142
2010-2019	11.545
2020-2029	11.352

Descriptive Statistics by 20-Year Groups:

75% \ Year Group 20	count	mean	std	min	25%	50%
1880-1899	7305.0	8.112854	0.433809	5.862	7.881	8.163
8.4010						
1900-1919	7304.0	8.258179	0.431921	6.412	8.015	8.286
8.5390						
1920-1939	7305.0	8.478008	0.434693	6.597	8.245	8.482
8.7220						
1940-1959	7305.0	8.594258	0.433365	6.803	8.341	8.622
8.8690						
1960-1979	7305.0	8.582274	0.458677	6.596	8.290	8.581
8.8860						
1980-1999	7305.0	9.025537	0.502520	7.257	8.694	9.036
9.3560						
2000-2019	7305.0	9.606187	0.503268	7.906	9.283	9.602
9.9210						
2020-2039	943.0	9.944686	0.410681	8.162	9.674	9.916
10.1725						

Year Group 20	max
1880-1899	9.583

1900-1919	9.704
1920-1939	10.209
1940-1959	10.152
1960-1979	10.282
1980-1999	10.847
2000-2019	11.545
2020-2039	11.352

Descriptive Statistics by 35-Year Groups:

	count	mean	std	min	25%	50%
75% \						
Year Group 35						
1880-1914	12783.0	8.177261	0.436764	5.862	7.935	8.215
8.4610						
1915-1949	12784.0	8.492936	0.445525	6.412	8.251	8.511
8.7600						
1950-1984	12784.0	8.612535	0.474145	6.596	8.308	8.613
8.9140						
1985-2019	12783.0	9.378631	0.566654	7.257	9.012	9.378
9.7480						
2020-2054	943.0	9.944686	0.410681	8.162	9.674	9.916
10.1725						

max

Year Group 35	9.704
1880-1914	9.704
1915-1949	10.209
1950-1984	10.282
1985-2019	11.545
2020-2054	11.352

```
# ---- Descriptive Statistics for 10-Year Groups ----
bins_10yr = list(range(df1['Year'].min(), df1['Year'].max() + 10, 10))
labels_10yr = [f"{bins_10yr[i]}-{bins_10yr[i+1]-1}" for i in range(len(bins_10yr)-1)]
df1['Year Group 10'] = pd.cut(df1['Year'], bins=bins_10yr,
                                labels=labels_10yr, right=False)
grouped_10yr_stats = df1.groupby('Year Group 10')[['Temperature']].describe()

# ---- Descriptive Statistics for 20-Year Groups ----
bins_20yr = list(range(df1['Year'].min(), df1['Year'].max() + 20, 20))
labels_20yr = [f"{bins_20yr[i]}-{bins_20yr[i+1]-1}" for i in range(len(bins_20yr)-1)]
df1['Year Group 20'] = pd.cut(df1['Year'], bins=bins_20yr,
                                labels=labels_20yr, right=False)
grouped_20yr_stats = df1.groupby('Year Group 20')[['Temperature']].describe()
```

```

# ---- Descriptive Statistics for 35-Year Groups ----
bins_35yr = list(range(df1['Year'].min(), df1['Year'].max() + 35, 35))
labels_35yr = [f"{bins_35yr[i]}-{bins_35yr[i+1]-1}" for i in
range(len(bins_35yr)-1)]
df1['Year Group 35'] = pd.cut(df1['Year'], bins=bins_35yr,
labels=labels_35yr, right=False)
grouped_35yr_stats = df1.groupby('Year Group 35')
['Temperature'].describe()

with pd.ExcelWriter("temperature_analysis_new.xlsx") as writer:
    yearly_stats.to_excel(writer, sheet_name="Yearly Statistics")
    monthly_stats.to_excel(writer, sheet_name="Monthly Statistics")
    grouped_20yr_stats.to_excel(writer, sheet_name="20-Year Groups")
    half_stats.to_excel(writer, sheet_name="Two Halves")
    overall_stats.to_excel(writer, sheet_name="Overall Statistics")
    monthly_avg_per_year.to_excel(writer, sheet_name="Monthly Avg Per
Year")
    yearly_avg_temp.to_excel(writer, sheet_name="Yearly Avg Temp")
    grouped_10yr_avg.to_excel(writer, sheet_name="10-Year Averages") #
Save 10-year averages

```

Distributional Study

```

from scipy.stats import kstest, probplot
from scipy import stats

dist_names = ['norm', 'expon', 'gamma', 'beta', 'lognorm',
'weibull_min', 'weibull_max']
temperature_data = df['Temperature'].dropna()

best_fit = {}
kstest_results = {}

def evaluate_fit(data, dist_name, params):
    dist = getattr(stats, dist_name)
    ks_stat, ks_pval = kstest(data, dist_name, args=params)
    return ks_stat, ks_pval

for dist_name in dist_names:
    dist = getattr(stats, dist_name)
    params = dist.fit(temperature_data)
    best_fit[dist_name] = params
    ks_stat, ks_pval = evaluate_fit(temperature_data, dist_name,
params)
    kstest_results[dist_name] = {'KS Statistic': ks_stat, 'P-Value': ks_pval}

# Convert results to a DataFrame

```

```

kstest_df = pd.DataFrame.from_dict(kstest_results, orient='index')

# Plot histogram with fitted distributions
sns.histplot(temperature_data, kde=False, bins=30, stat="density",
label="Data")

x = np.linspace(temperature_data.min(), temperature_data.max(), 1000)

for dist_name, params in best_fit.items():
    dist = getattr(stats, dist_name)
    pdf_fitted = dist.pdf(x, *params)
    plt.plot(x, pdf_fitted, label=dist_name)

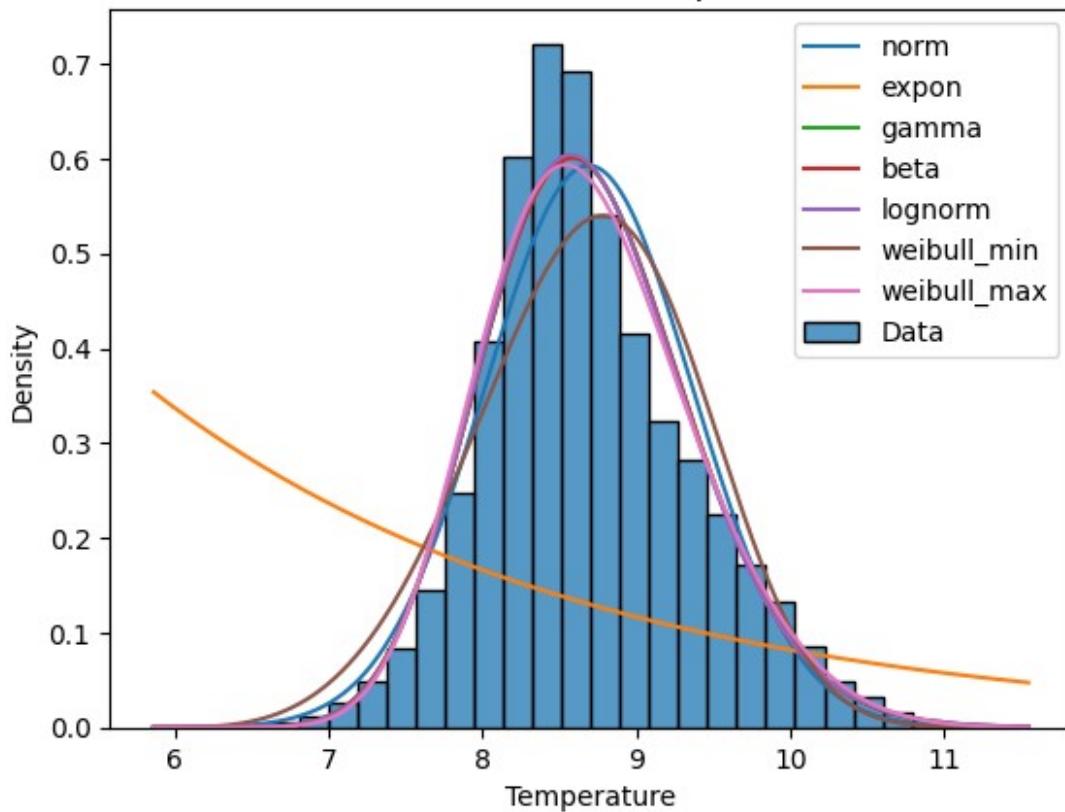
plt.legend()
plt.xlabel("Temperature")
plt.ylabel("Density")
plt.title("Fitted Distributions on Temperature Data")
plt.show()

# Print KS Test results
print("\nKolmogorov-Smirnov Test Results:")
print(kstest_df)

# Print best-fit parameters for each distribution
print("\nBest Fit Distribution Parameters:")
for dist_name, params in best_fit.items():
    print(f"{dist_name}: {params}")

```

Fitted Distributions on Temperature Data



Kolmogorov-Smirnov Test Results:

	KS Statistic	P-Value
norm	0.059938	4.460529e-163
expon	0.427401	0.000000e+00
gamma	0.039049	1.968802e-69
beta	0.039049	1.967192e-69
lognorm	0.037038	1.654976e-62
weibull_min	0.074443	1.982387e-251
weibull_max	0.036362	2.928558e-60

Best Fit Distribution Parameters:

norm: (8.688502563511724, 0.6730782024854579)
expon: (5.862, 2.826502563511724)
gamma: (39.011663603959136, 4.497700761300052, 0.10742432871731672)
beta: (39.011788306791566, 47129810.30498446, 4.497692042424953,
5062886.853759905)
lognorm: (0.11628643083088369, 2.932839197933386, 5.716837978733398)
weibull_min: (4.415801985131122, 5.859482916670485, 3.093586658959958)
weibull_max: (5.760986512463964, 12.050427154466213,
3.6289129305249883)

```

import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
from scipy import stats
from scipy.stats import shapiro

# List of years to test
target_years = [1880, 1900, 1925, 1950, 1975, 2000, 2020]

# Iterate through each year
for year in target_years:
    year_data = df[df['Year'] == year]['Temperature'].dropna()

    if len(year_data) < 3:
        print(f"Year {year}: Not enough data for testing.")
        continue

    print(f"\n===== Year {year} =====")

    # Normality test (Shapiro-Wilk is good for small to medium sample sizes)
    stat, p_value = shapiro(year_data)
    print(f"Shapiro-Wilk test statistic: {stat:.4f}, p-value: {p_value:.4f}")

    if p_value > 0.05:
        print("Data appears to be normally distributed.\n")
        transformed_data = year_data # No need to transform
    else:
        print("Data is NOT normally distributed. Applying Johnson transformation...\n")

        # Fit Johnson SU distribution
        fitted_params = stats.johnsonsu.fit(year_data)
        transformed_data =
stats.johnsonsu(*fitted_params).ppf(stats.johnsonsu(*fitted_params).cdf(year_data))

        # Re-test for normality after transformation
        stat_t, p_t = shapiro(transformed_data)
        print(f"Post-transformation Shapiro-Wilk: stat = {stat_t:.4f}, p = {p_t:.4f}")

    # Plot the final data (original or transformed)
    sns.histplot(transformed_data, bins=30, kde=True, stat="density",
color='skyblue')
    plt.title(f"Histogram (Transformed if Applied) - Year {year}")
    plt.xlabel("Temperature")
    plt.ylabel("Density")

```

```

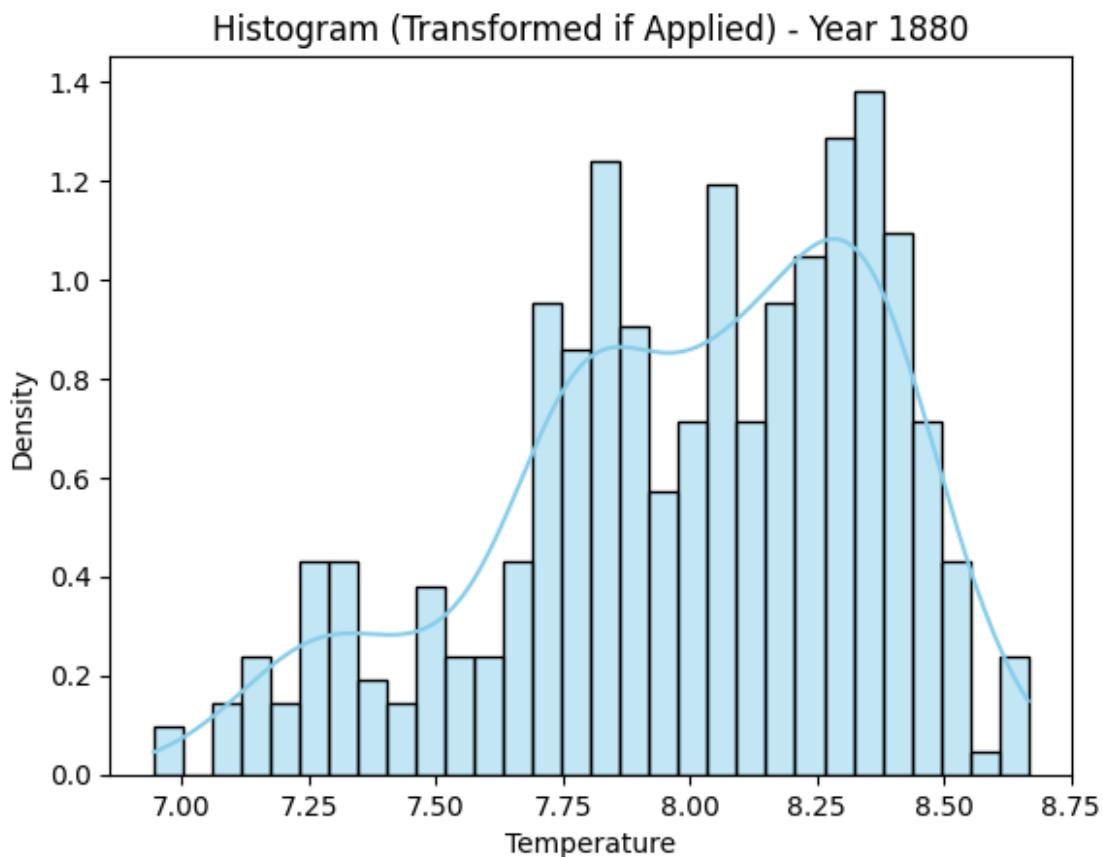
plt.show()

# Q-Q Plot to visually check normality
stats.probplot(transformed_data, dist="norm", plot=plt)
plt.title(f"Q-Q Plot (Transformed if Applied) - Year {year}")
plt.show()

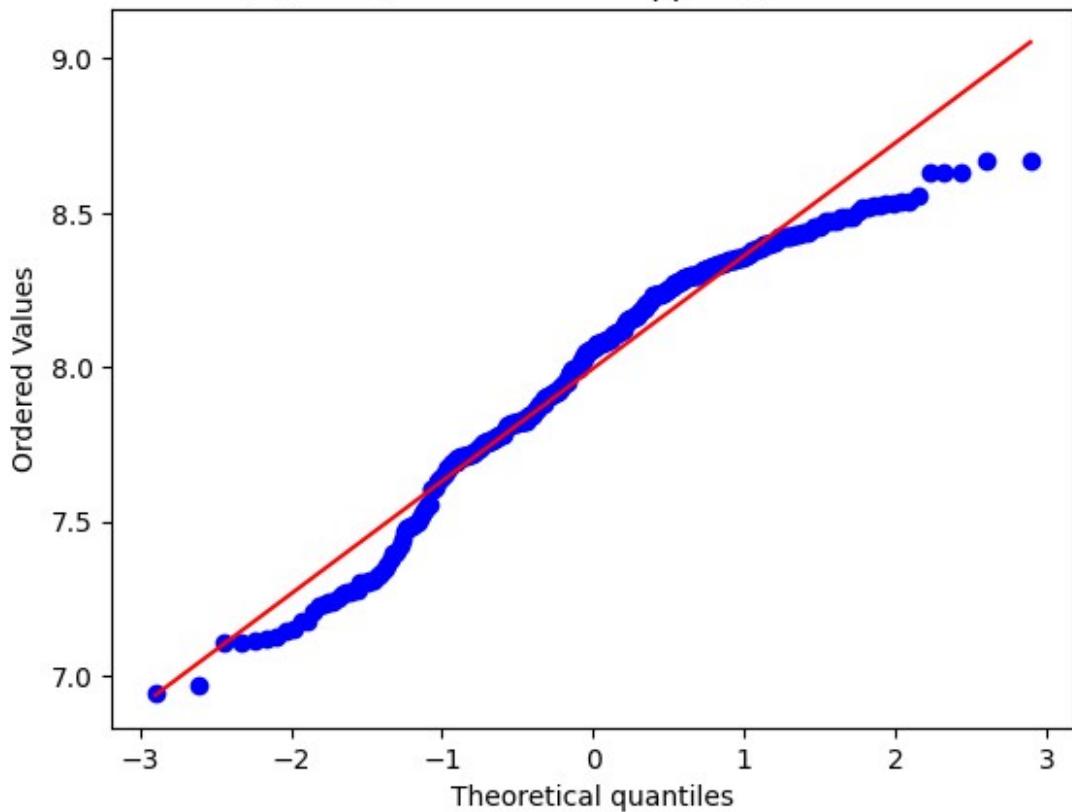
===== Year 1880 =====
Shapiro-Wilk test statistic: 0.9572, p-value: 0.0000
Data is NOT normally distributed. Applying Johnson transformation...

Post-transformation Shapiro-Wilk: stat = 0.9572, p = 0.0000

```



Q-Q Plot (Transformed if Applied) - Year 1880



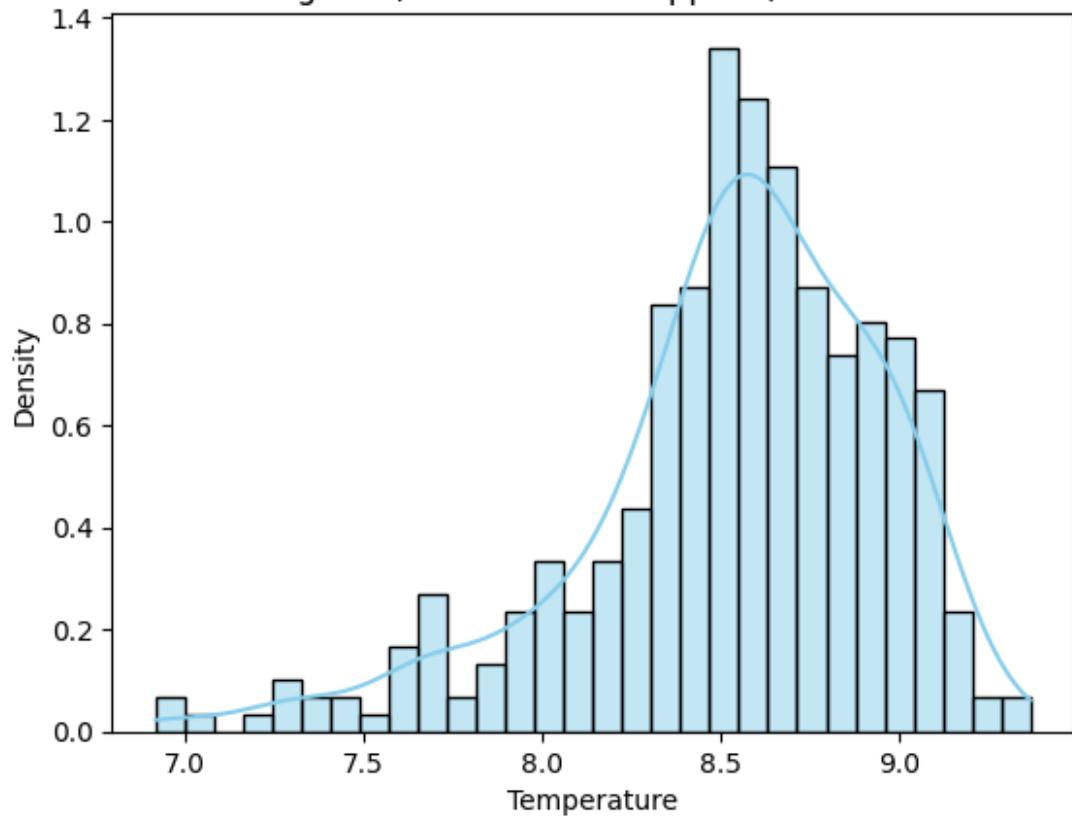
===== Year 1900 =====

Shapiro-Wilk test statistic: 0.9372, p-value: 0.0000

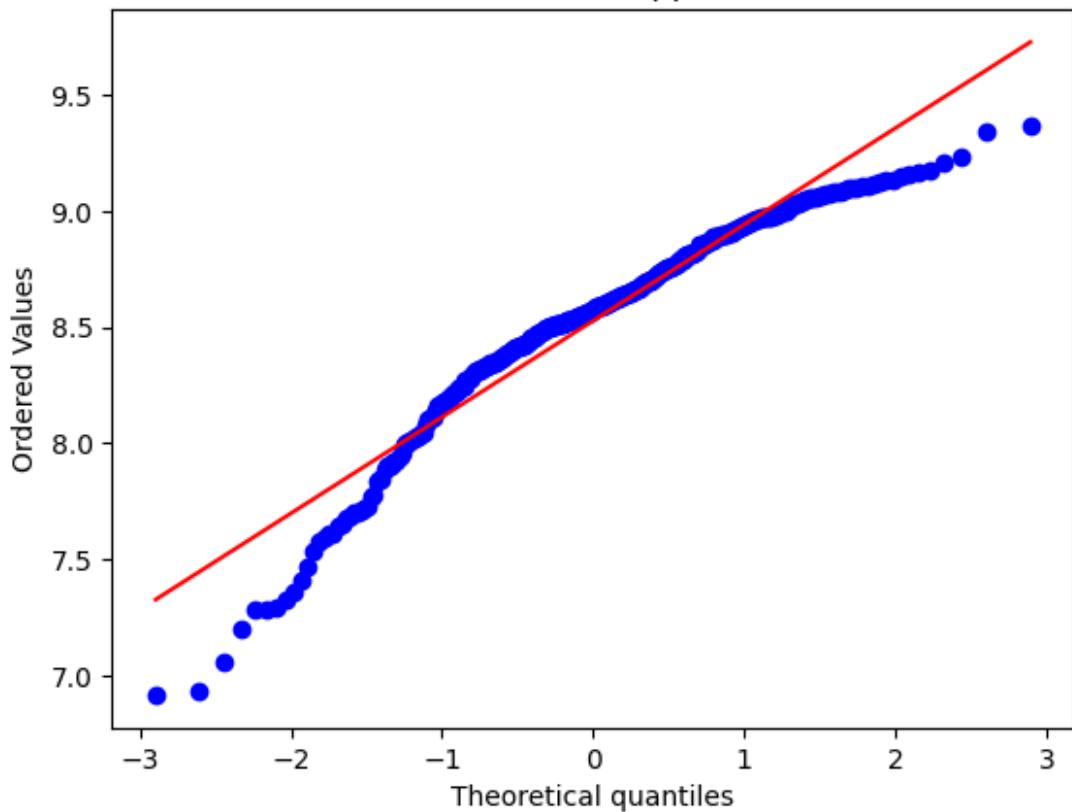
Data is NOT normally distributed. Applying Johnson transformation...

Post-transformation Shapiro-Wilk: stat = 0.9372, p = 0.0000

Histogram (Transformed if Applied) - Year 1900



Q-Q Plot (Transformed if Applied) - Year 1900



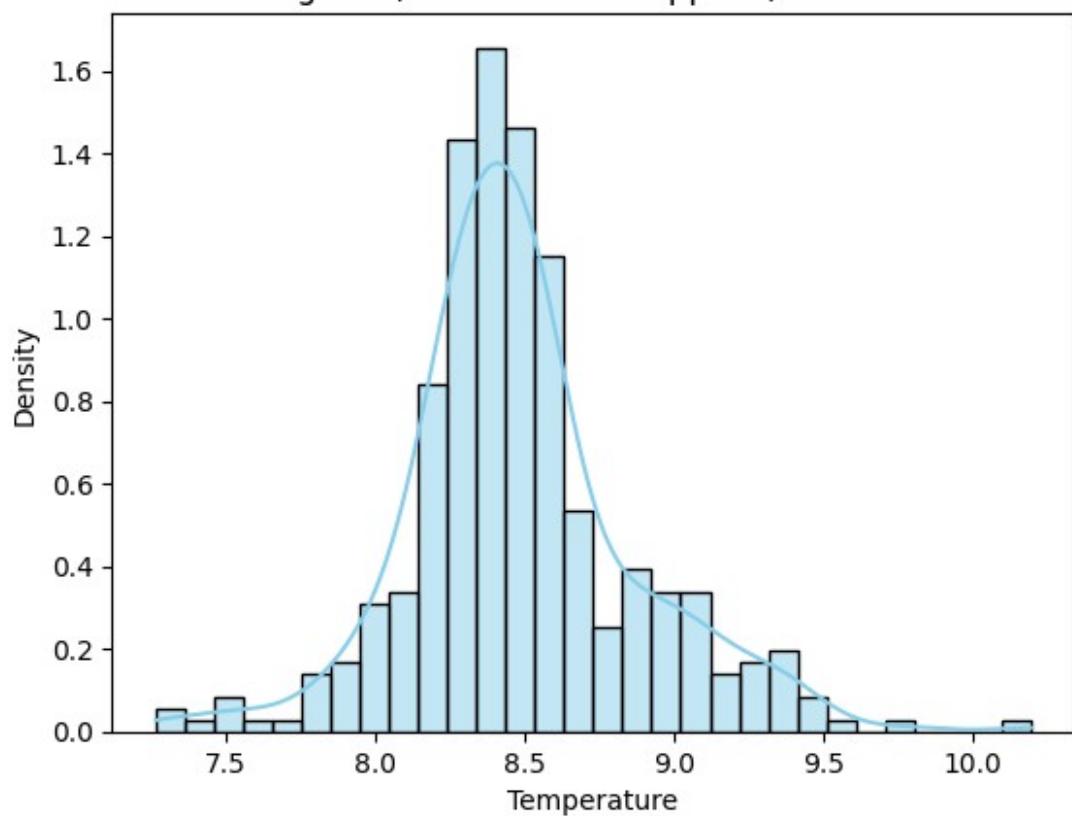
===== Year 1925 =====

Shapiro-Wilk test statistic: 0.9560, p-value: 0.0000

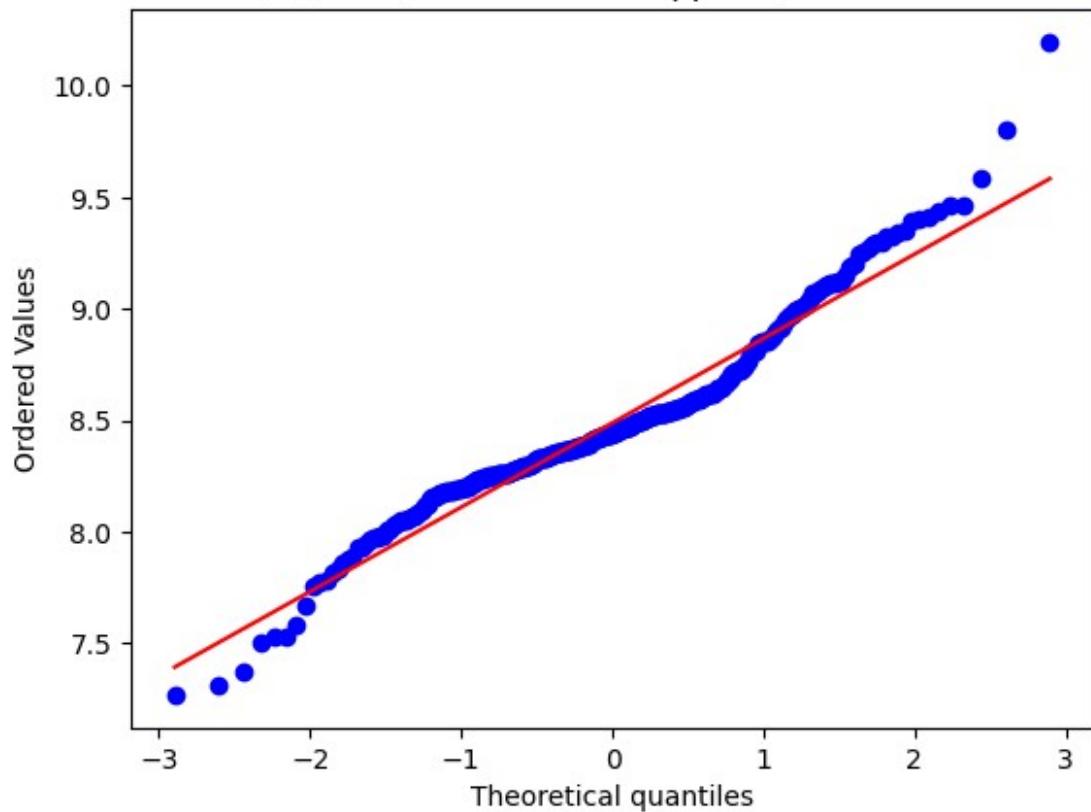
Data is NOT normally distributed. Applying Johnson transformation...

Post-transformation Shapiro-Wilk: stat = 0.9560, p = 0.0000

Histogram (Transformed if Applied) - Year 1925



Q-Q Plot (Transformed if Applied) - Year 1925



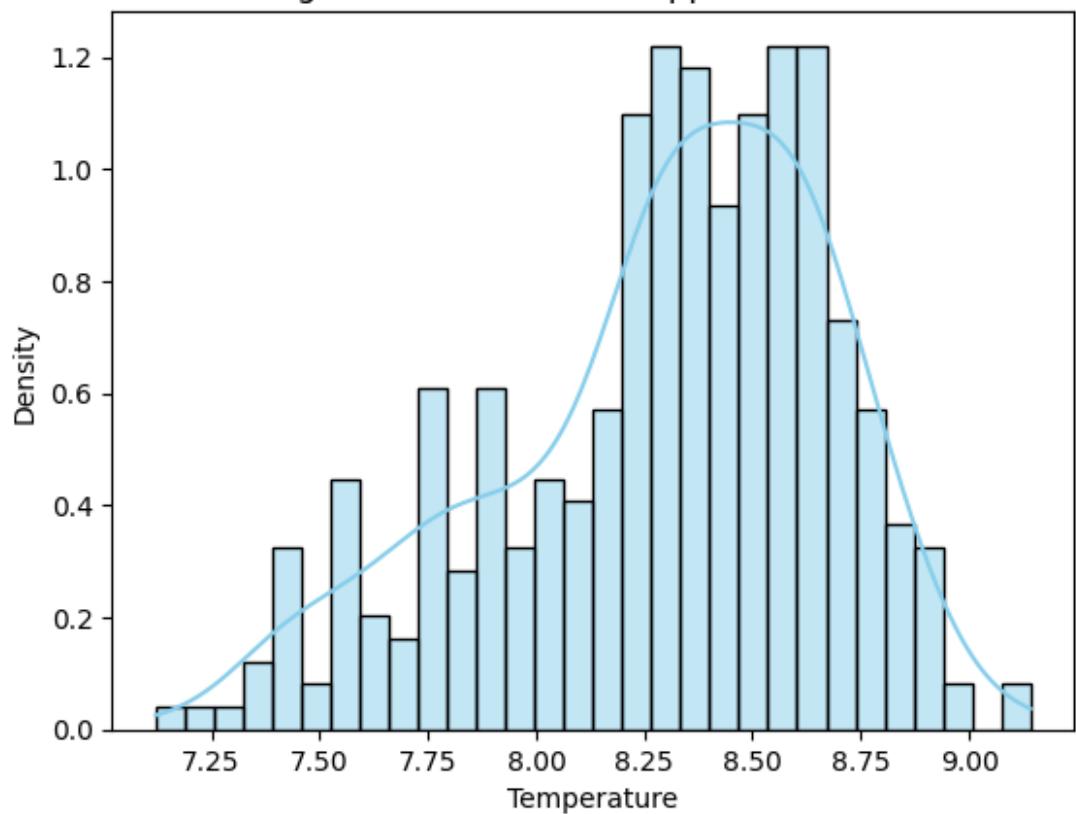
===== Year 1950 =====

Shapiro-Wilk test statistic: 0.9588, p-value: 0.0000

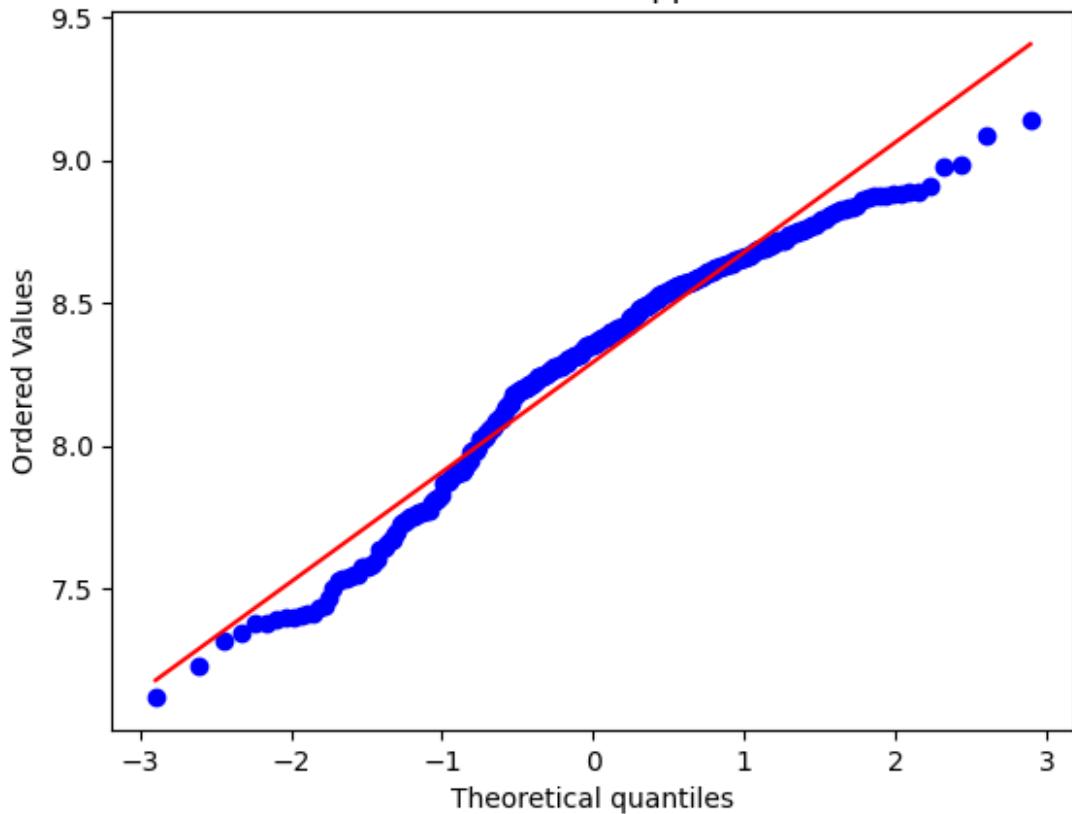
Data is NOT normally distributed. Applying Johnson transformation...

Post-transformation Shapiro-Wilk: stat = 0.9588, p = 0.0000

Histogram (Transformed if Applied) - Year 1950



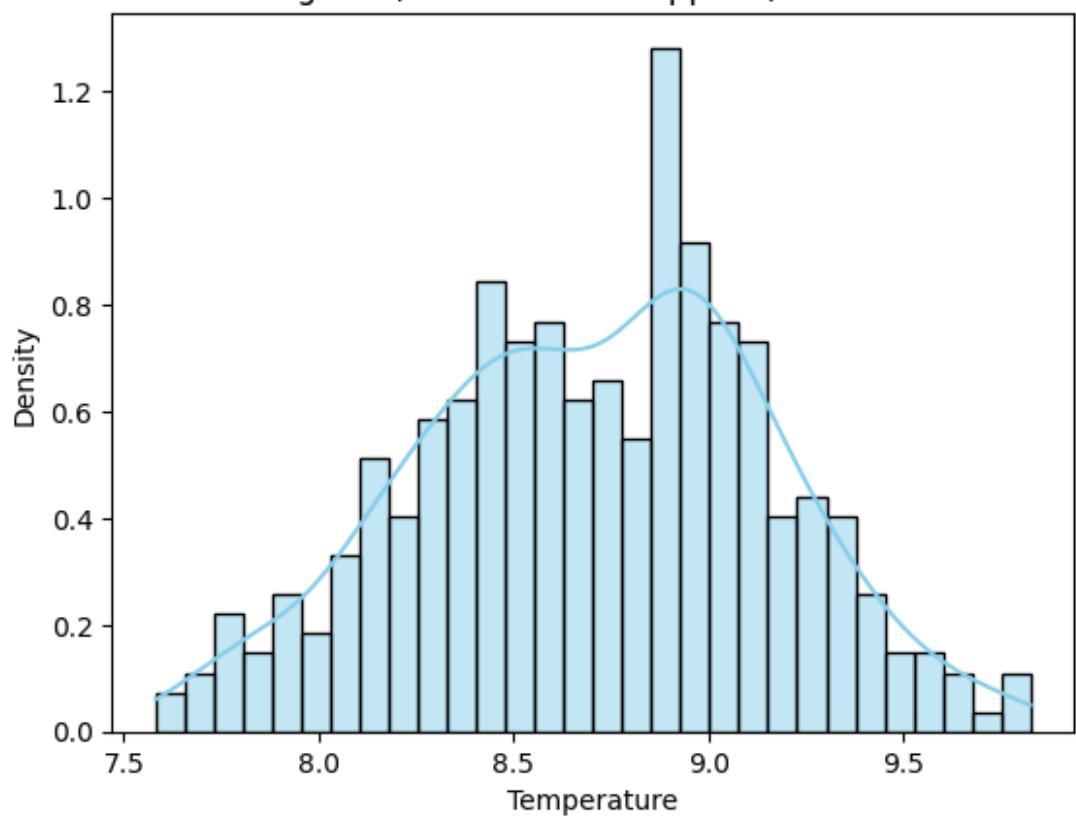
Q-Q Plot (Transformed if Applied) - Year 1950



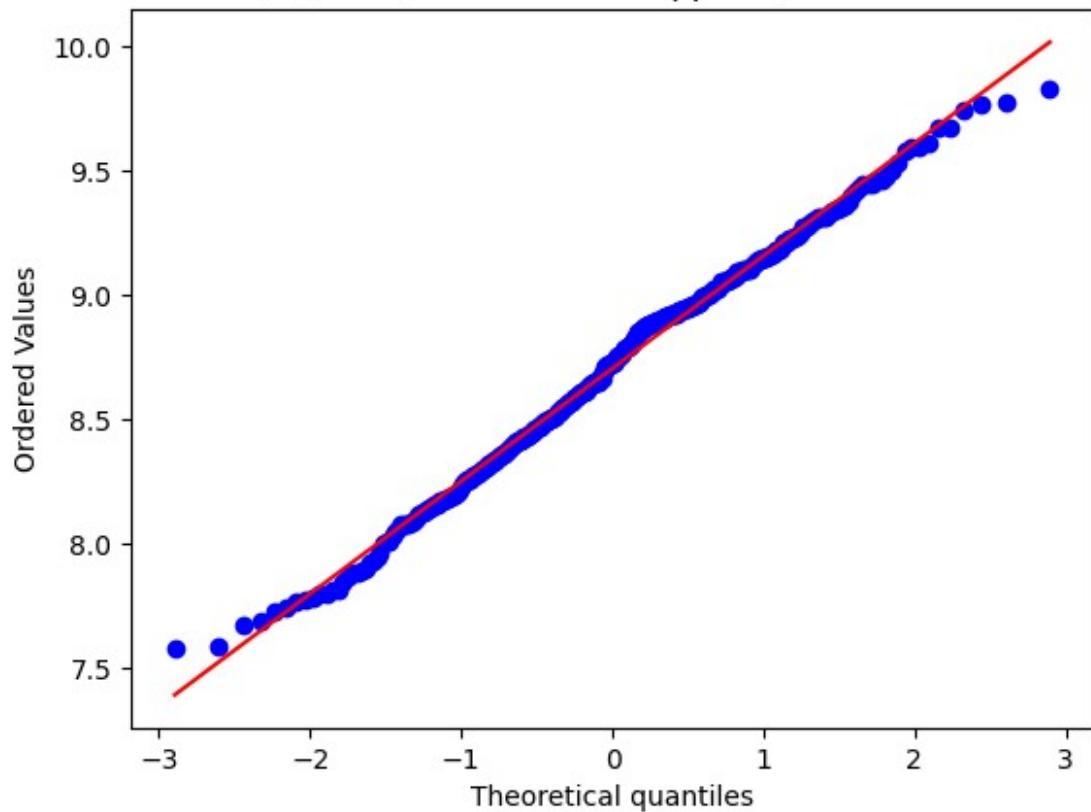
===== Year 1975 =====

Shapiro-Wilk test statistic: 0.9933, p-value: 0.1057
Data appears to be normally distributed.

Histogram (Transformed if Applied) - Year 1975



Q-Q Plot (Transformed if Applied) - Year 1975



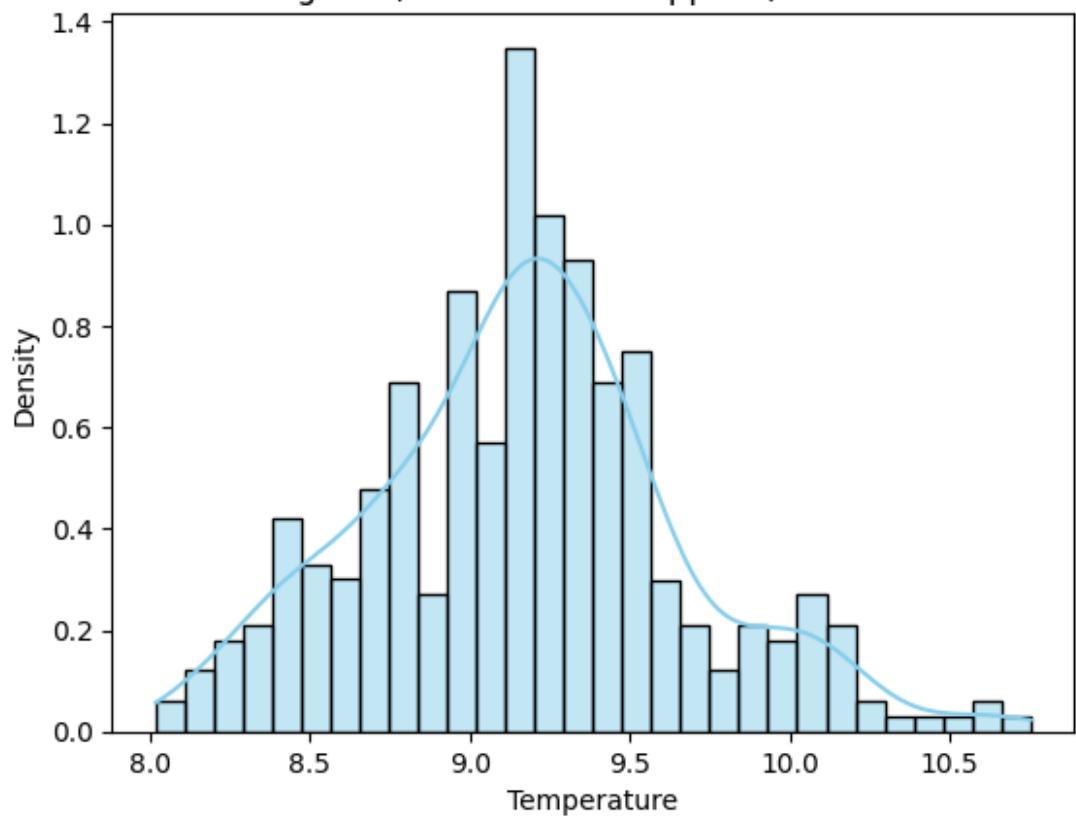
===== Year 2000 =====

Shapiro-Wilk test statistic: 0.9862, p-value: 0.0015

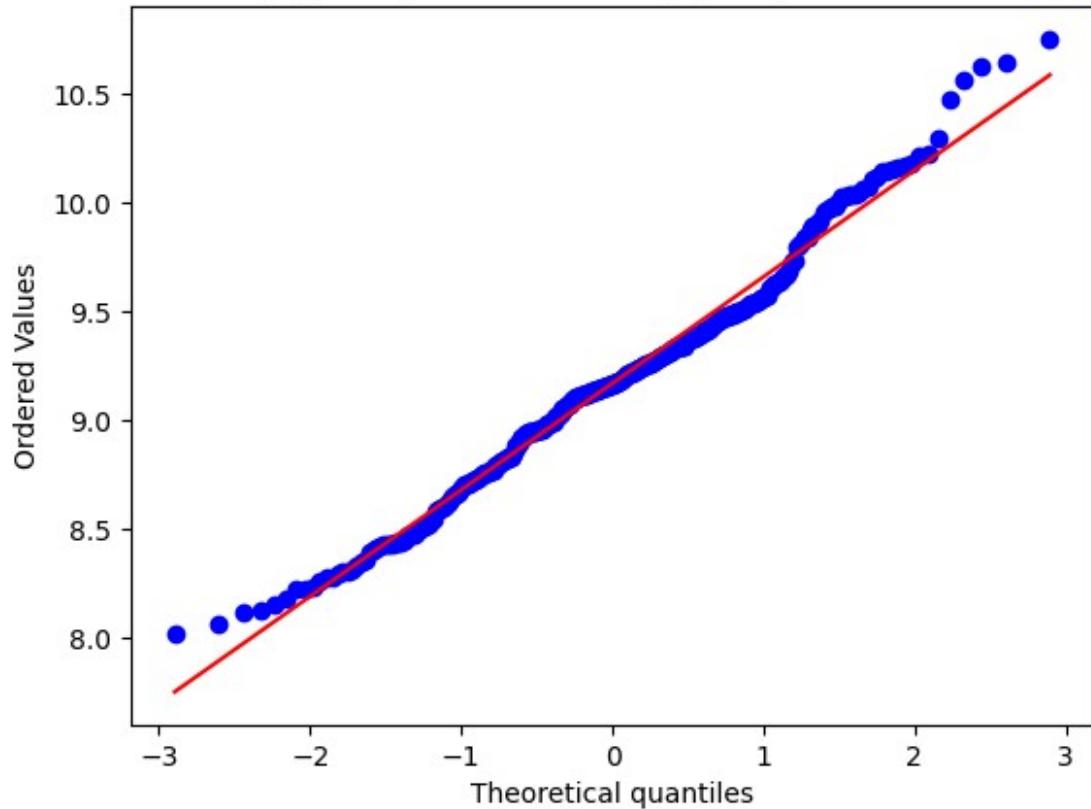
Data is NOT normally distributed. Applying Johnson transformation...

Post-transformation Shapiro-Wilk: stat = 0.9862, p = 0.0015

Histogram (Transformed if Applied) - Year 2000



Q-Q Plot (Transformed if Applied) - Year 2000



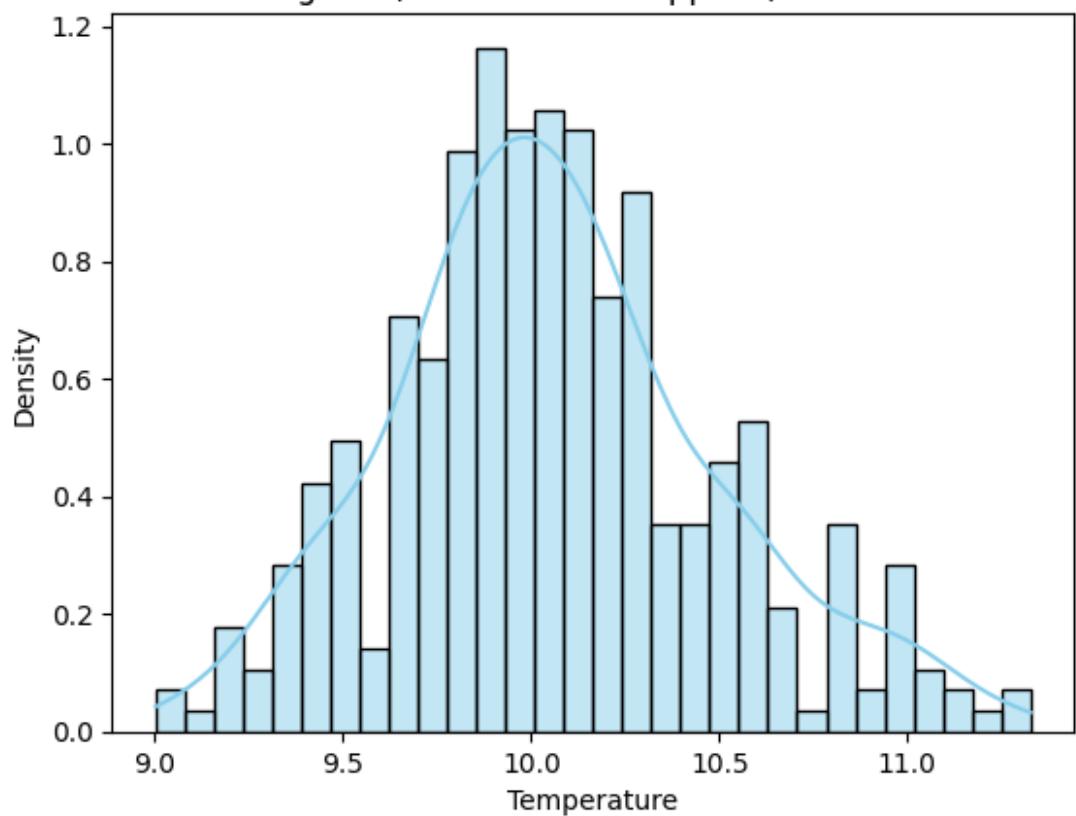
===== Year 2020 =====

Shapiro-Wilk test statistic: 0.9865, p-value: 0.0017

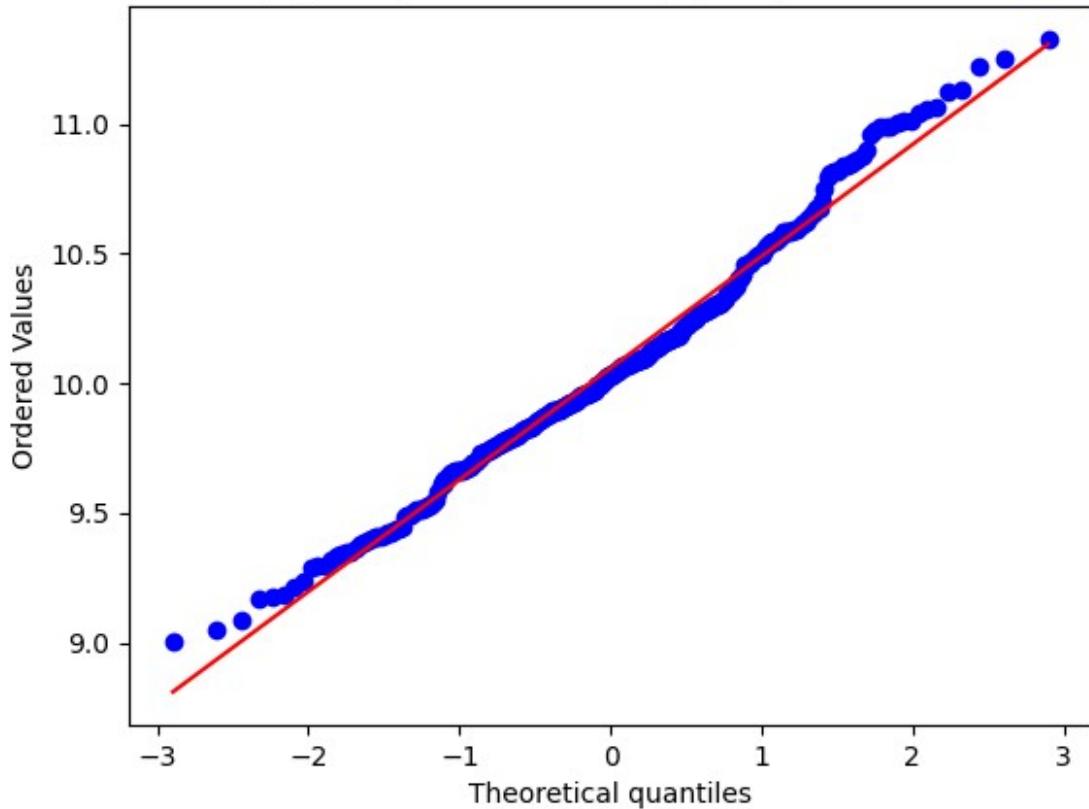
Data is NOT normally distributed. Applying Johnson transformation...

Post-transformation Shapiro-Wilk: stat = 0.9865, p = 0.0017

Histogram (Transformed if Applied) - Year 2020



Q-Q Plot (Transformed if Applied) - Year 2020



```
import pandas as pd
import numpy as np
from scipy import stats

def detect_outliers_std(df, column, threshold=2.5):
    """
    Detects outliers in a DataFrame column based on standard deviations from the mean.

    Args:
        df (pd.DataFrame): The DataFrame to analyze.
        column (str): The name of the column to check for outliers.
        threshold (float): The number of standard deviations from the mean to consider as an outlier.
                           Defaults to 2.5, as used in Hansen et al. (2010).

    Returns:
        pd.DataFrame: A DataFrame containing the outliers and their values.
    """
    mean = df[column].mean()
    std = df[column].std()
```

```

        outliers = df[(df[column] > mean + threshold * std) | (df[column]
< mean - threshold * std)]
        return outliers

def detect_inconsistent_neighbors(df, column='Temperature',
year_col='Year', threshold=2.5):
    """
    Detects data points that are inconsistent with neighboring years
    based on standard deviations.
    This is a simplified version of the neighbor comparison mentioned
    in Hansen et al. (2010).

    Args:
        df (pd.DataFrame): The DataFrame to analyze, containing 'Year'
        and the temperature column.
        column (str): The name of the column to check for
        inconsistencies (e.g., 'Temperature').
        year_col (str): The name of the year column.
        threshold (float): Threshold for considering a point as
        inconsistent.

    Returns:
        pd.DataFrame: DataFrame of inconsistent data points.
    """

    df_shifted = df.copy()
    df_shifted['Temp_lag1'] = df_shifted[column].shift(1)
    df_shifted['Temp_lead1'] = df_shifted[column].shift(-1)

    df_shifted['Mean_Neighbor'] = df_shifted[['Temp_lag1',
'Temp_lead1']].mean(axis=1)
    df_shifted['Std_Neighbor'] = df_shifted[['Temp_lag1',
'Temp_lead1']].std(axis=1)

    df_shifted['Outlier_High'] = (df_shifted[column] >
df_shifted['Mean_Neighbor'] + threshold * df_shifted['Std_Neighbor'])
    df_shifted['Outlier_Low'] = (df_shifted[column] <
df_shifted['Mean_Neighbor'] - threshold * df_shifted['Std_Neighbor'])

    inconsistent_points = df_shifted[(df_shifted['Outlier_High'] ==
True) | (df_shifted['Outlier_Low'] == True)]

    return inconsistent_points[['Year', column]] # Return only Year
and Temperature

# Example Usage:
# Overall outlier detection
overall_temp_outliers = detect_outliers_std(df1, 'Temperature')

```

```

print("Overall Temperature Outliers (Std Dev):\n",
overall_temp_outliers)

# Outlier detection by year (e.g., for each year's temperature)
yearly_outliers = df1.groupby('Year').apply(detect_outliers_std,
column='Temperature')
print("\nYearly Temperature Outliers (Std Dev):\n", yearly_outliers)

# Detect points inconsistent with neighboring years
inconsistent_neighbor_points = detect_inconsistent_neighbors(df1)
print("\nInconsistent Neighboring Year Points:\n",
inconsistent_neighbor_points)

```

Overall Temperature Outliers (Std Dev):

	Year	Month	Temperature		Year Group		Half Year
Group 10							
72	1880	3	6.968		NaN	First Half	
1880-1889							
73	1880	3	6.945		NaN	First Half	
1880-1889							
1063	1882	11	6.950	(1880.0, 1900.0]	First Half		
1880-1889							
1064	1882	11	6.921	(1880.0, 1900.0]	First Half		
1880-1889							
1070	1882	12	6.856	(1880.0, 1900.0]	First Half		
1880-1889							
...
...							
51964	2022	4	10.816		NaN	Second Half	
Nan							
51965	2022	4	10.653		NaN	Second Half	
Nan							
52025	2022	6	10.385		NaN	Second Half	
Nan							
52030	2022	6	10.394		NaN	Second Half	
Nan							
52031	2022	6	10.518		NaN	Second Half	
Nan							

[885 rows x 6 columns]

Yearly Temperature Outliers (Std Dev):

	Year	Month	Temperature		Year Group
Half \					
Year					
1880	72	1880	3	6.968	NaN First Half
	73	1880	3	6.945	NaN First Half
1881	410	1881	2	7.413	(1880.0, 1900.0] First Half
	687	1881	11	7.447	(1880.0, 1900.0] First Half
	688	1881	11	7.280	(1880.0, 1900.0] First Half

...			
2021	51557	2021	2	8.957		NaN	Second	Half
2022	51960	2022	4	11.137		NaN	Second	Half
	51961	2022	4	11.352		NaN	Second	Half
	51962	2022	4	11.291		NaN	Second	Half
	51963	2022	4	11.014		NaN	Second	Half
Year Group 10								
Year								
1880	72		1880-1889					
	73		1880-1889					
1881	410		1880-1889					
	687		1880-1889					
	688		1880-1889					
...			...					
2021	51557			NaN				
2022	51960			NaN				
	51961			NaN				
	51962			NaN				
	51963			NaN				

[942 rows x 6 columns]

Inconsistent Neighboring Year Points:

	Year	Temperature
1	1880	7.998
2	1880	7.917
3	1880	7.975
5	1880	7.847
6	1880	7.944
...
52035	2022	10.074
52045	2022	9.348
52052	2022	9.786
52058	2022	9.376
52072	2022	10.229

[6943 rows x 2 columns]

Missing Values: Check for missing or inconsistent records.

Outlier Detection: Boxplots and IQR method.

Trend Analysis: Moving averages, seasonal decomposition.

Visualization:

1. Time series plots for different periods.

2. Rolling mean and standard deviation.

3. Histogram of temperature anomalies.

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from statsmodels.tsa.seasonal import seasonal_decompose

# Display basic info
print(df.info())
print(df.head())

# Check for missing values
missing_values = df.isnull().sum()
print("Missing Values:\n", missing_values)

# Handling missing values (if any)
df.dropna(inplace=True) # Remove missing records
df.reset_index(drop=True, inplace=True)

# Outlier Detection using IQR
Q1 = df['Temperature'].quantile(0.25)
Q3 = df['Temperature'].quantile(0.75)
IQR = Q3 - Q1

lower_bound = Q1 - 1.5 * IQR
upper_bound = Q3 + 1.5 * IQR

outliers = df[(df['Temperature'] < lower_bound) | (df['Temperature'] >
upper_bound)]
print(f"Number of Outliers: {len(outliers)}")

# Boxplot for outlier visualization
plt.figure(figsize=(8, 5))
```

```

sns.boxplot(x=df['Temperature'])
plt.title("Boxplot of Temperature")
plt.show()

# Trend Analysis using Moving Average
df['Rolling_Mean'] = df['Temperature'].rolling(window=365).mean()
df['Rolling_Std'] = df['Temperature'].rolling(window=365).std()

plt.figure(figsize=(12, 6))
plt.plot(df['Temperature'], label="Original Temperature", alpha=0.5)
plt.plot(df['Rolling_Mean'], label="365-day Moving Average",
color='red')
plt.plot(df['Rolling_Std'], label="365-day Rolling Std Dev",
color='green')
plt.title("Trend Analysis with Moving Average and Std Dev")
plt.legend()
plt.show()

# Seasonal Decomposition
decomposition = seasonal_decompose(df['Temperature'], period=365,
model='additive')

plt.figure(figsize=(12, 8))
plt.subplot(411)
plt.plot(df['Temperature'], label="Original", alpha=0.5)
plt.legend()

plt.subplot(412)
plt.plot(decomposition.trend, label="Trend", color='red')
plt.legend()

plt.subplot(413)
plt.plot(decomposition.seasonal, label="Seasonality", color='green')
plt.legend()

plt.subplot(414)
plt.plot(decomposition.resid, label="Residuals", color='purple')
plt.legend()

plt.tight_layout()
plt.show()

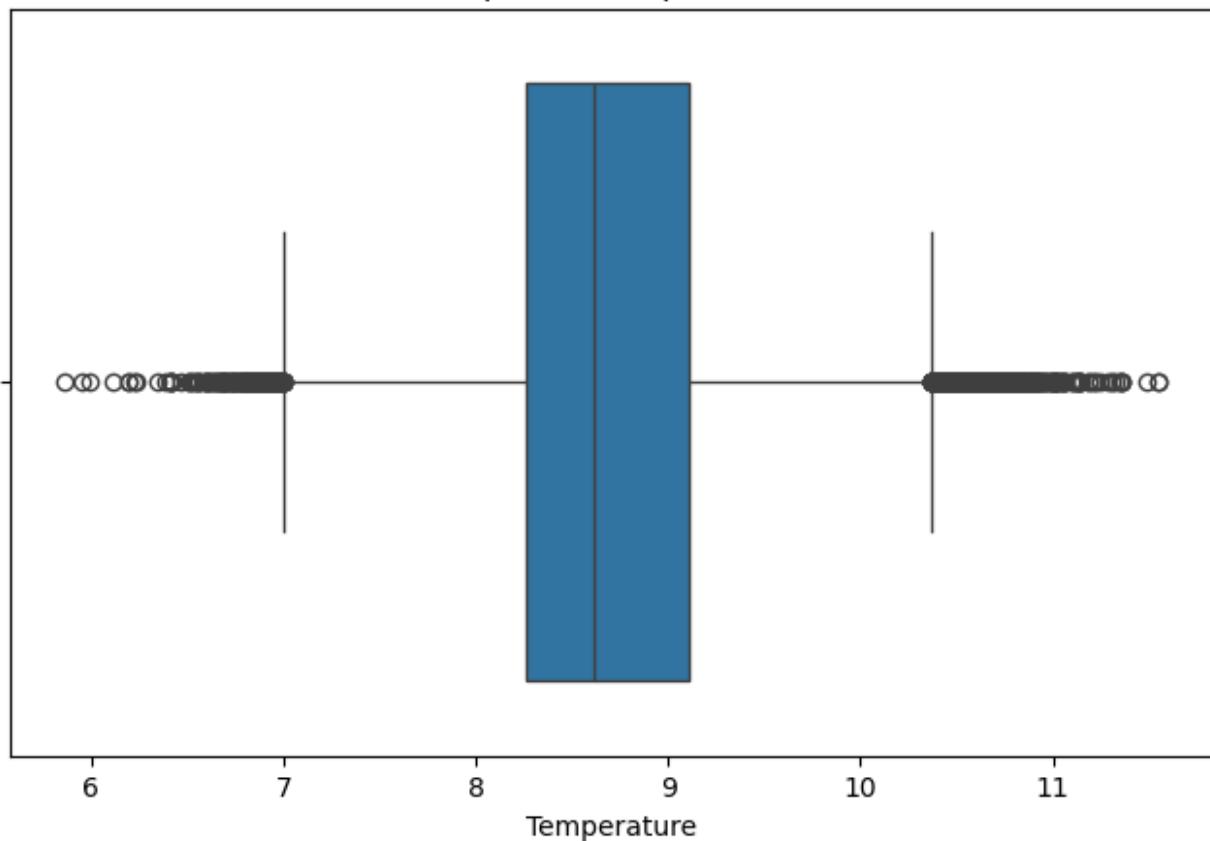
# Histogram of temperature anomalies
plt.figure(figsize=(8, 5))
sns.histplot(df['Anomaly'], bins=50, kde=True)
plt.title("Histogram of Temperature Anomalies")
plt.xlabel("Temperature Anomaly")
plt.show()

```

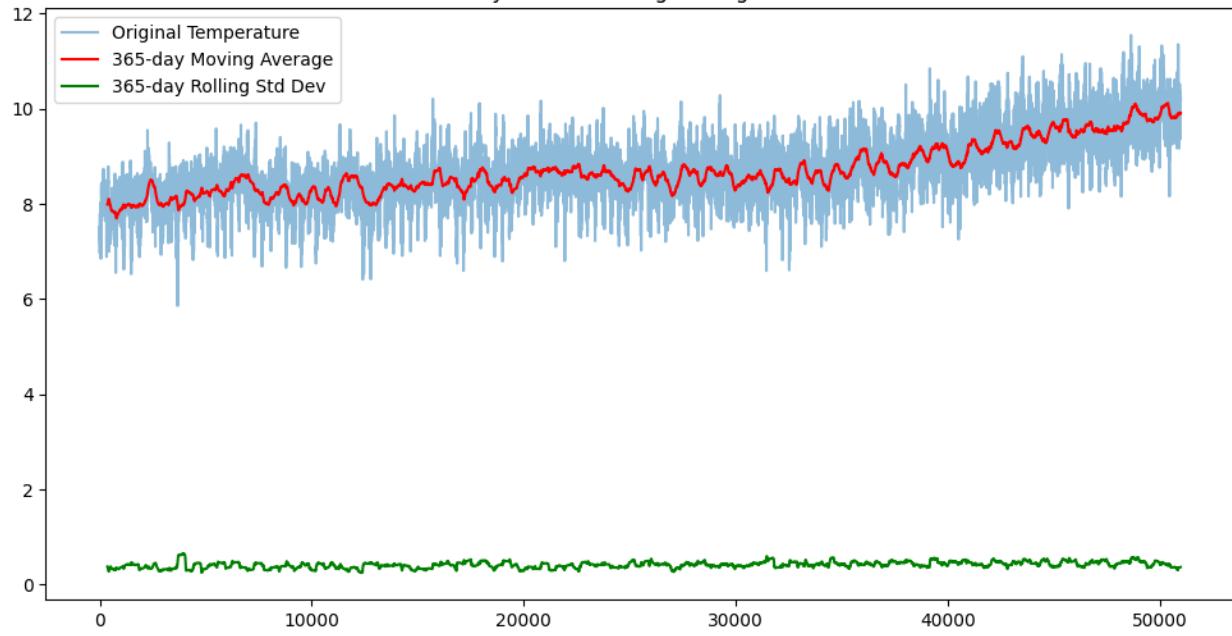
```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 51349 entries, 0 to 51348
Data columns (total 8 columns):
 #   Column           Non-Null Count  Dtype  
--- 
 0   Year              51349 non-null   int64  
 1   Month             51349 non-null   int64  
 2   Day               51349 non-null   int64  
 3   Day of Year       51349 non-null   int64  
 4   Anomaly           51349 non-null   float64 
 5   Temperature       51349 non-null   float64 
 6   Rolling_Mean      50985 non-null   float64 
 7   Rolling_Std        50985 non-null   float64 
dtypes: float64(4), int64(4)
memory usage: 3.1 MB
None
   Year  Month  Day  Day of Year  Anomaly  Temperature
Rolling_Mean \
0   1881     12    29          363    0.232     8.822      NaN
1   1881     12    30          364   -0.015     8.575      NaN
2   1881     12    31          365    0.013     8.603      NaN
3   1882      1     1            1    0.020     8.610      NaN
4   1882      1     2            2   -0.043     8.547      NaN

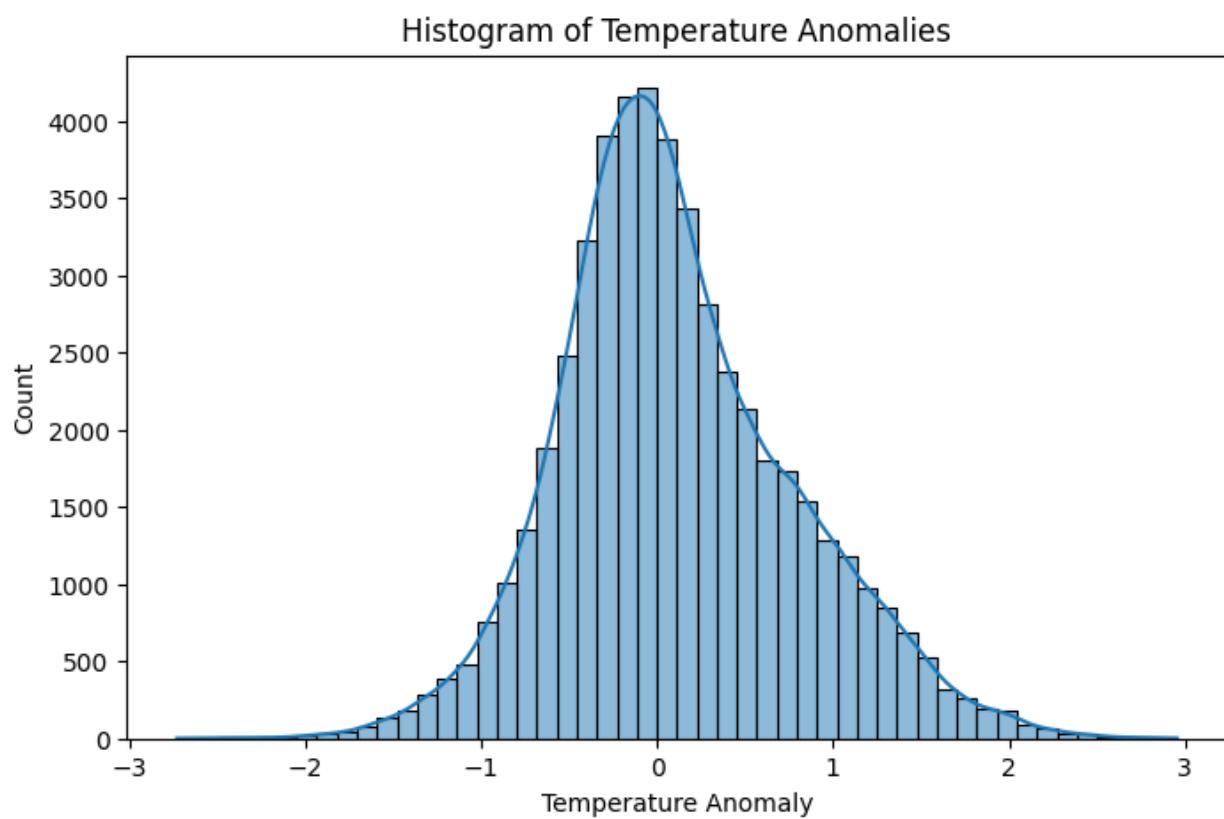
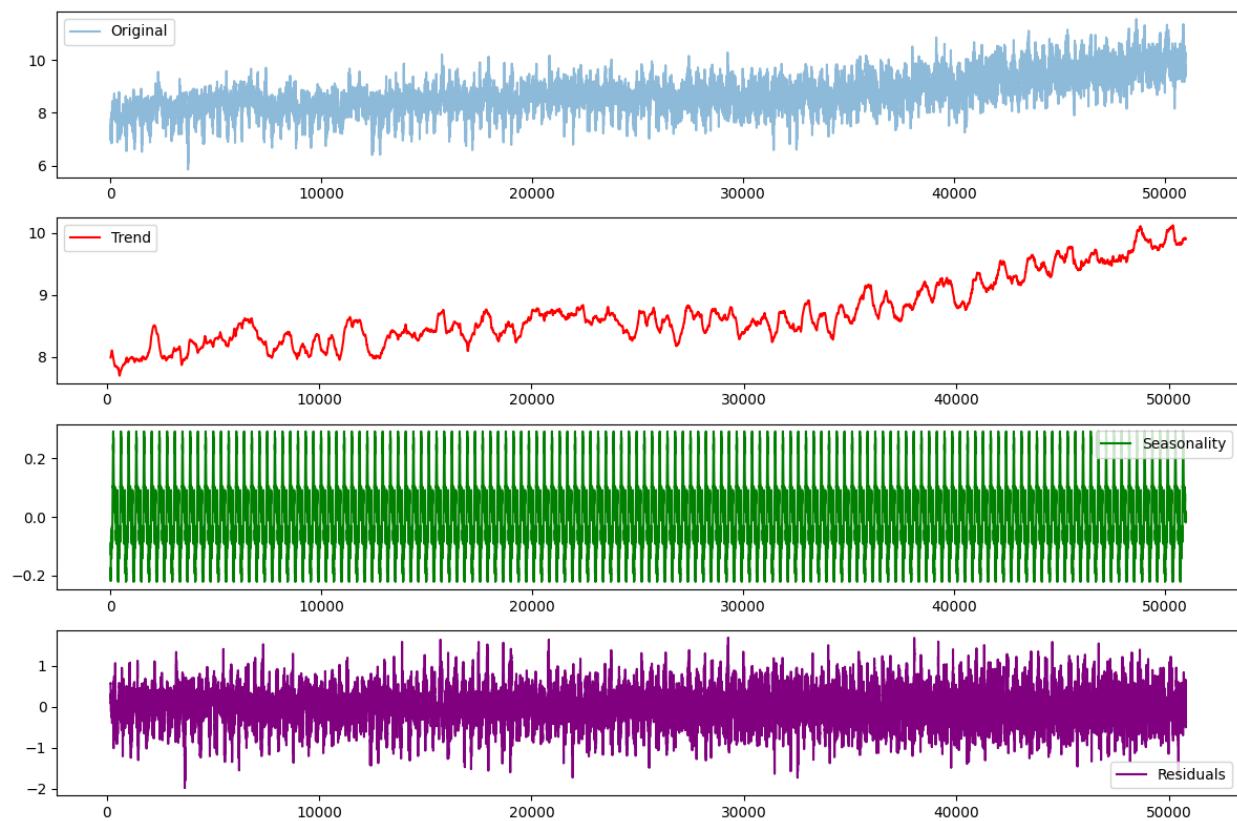
   Rolling_Std
0           NaN
1           NaN
2           NaN
3           NaN
4           NaN
Missing Values:
   Year      0
Month      0
Day        0
Day of Year 0
Anomaly    0
Temperature 0
Rolling_Mean 364
Rolling_Std  364
dtype: int64
Number of Outliers: 881
```

Boxplot of Temperature



Trend Analysis with Moving Average and Std Dev





```

import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from scipy.stats import theilslopes
from sklearn.linear_model import LinearRegression
from statsmodels.tsa.stattools import adfuller
import pymannkendall as mk

X = df[['Year']].values.reshape(-1, 1)
y = df['Temperature'].values

model = LinearRegression()
model.fit(X, y)

trend_slope = model.coef_[0]
intercept = model.intercept_
print(f"Linear Regression Trend Slope: {trend_slope:.6f}")

# Plot the trend line
plt.figure(figsize=(10, 5))
plt.scatter(df['Year'], df['Temperature'], label="Original Data",
alpha=0.5)
plt.plot(df['Year'], model.predict(X), color='red', label=f"Trend Line
(Slope: {trend_slope:.6f})")
plt.xlabel("Year")
plt.ylabel("Temperature")
plt.title("Temperature Trend (Linear Regression)")
plt.legend()
plt.show()

df['Rolling_Mean'] = df['Temperature'].rolling(window=10,
min_periods=1).mean()

plt.figure(figsize=(10, 5))
plt.plot(df['Year'], df['Temperature'], label="Original Data",
alpha=0.5)
plt.plot(df['Year'], df['Rolling_Mean'], label="10-Year Moving
Average", color='red')
plt.xlabel("Year")
plt.ylabel("Temperature")
plt.title("Temperature Trend (Moving Average)")
plt.show()

adf_test = adfuller(df['Temperature'])

```

```

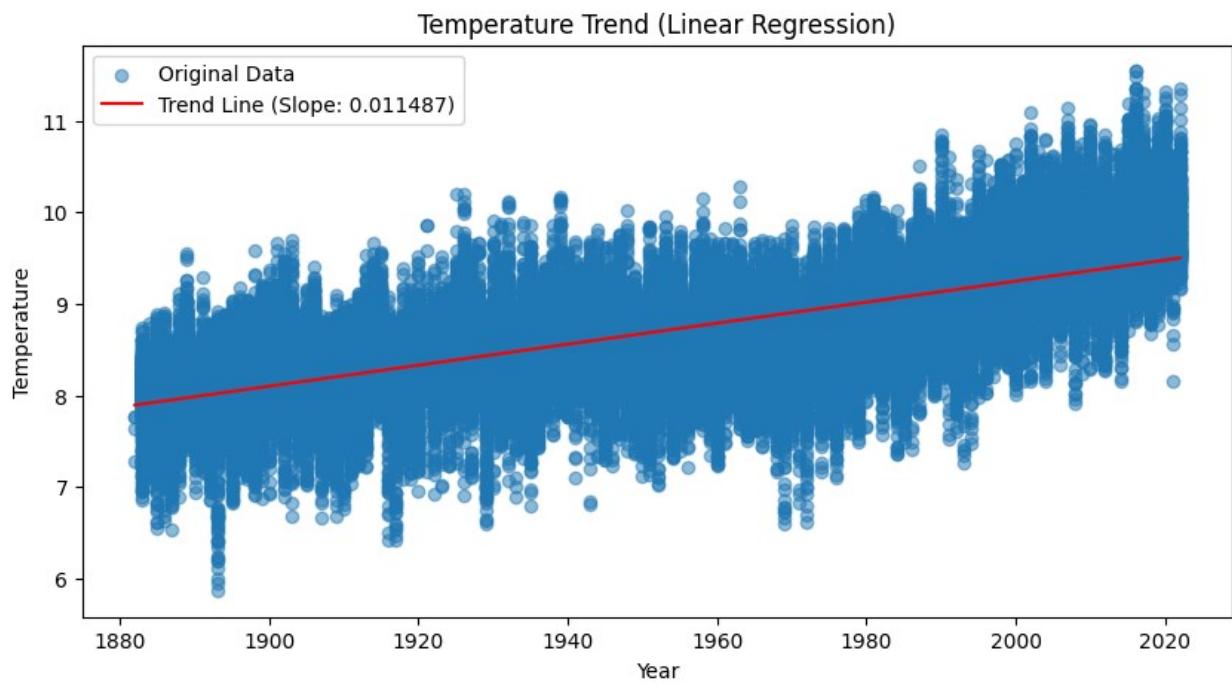
print(f"ADF Statistic: {adf_test[0]}")
print(f"p-value: {adf_test[1]}")

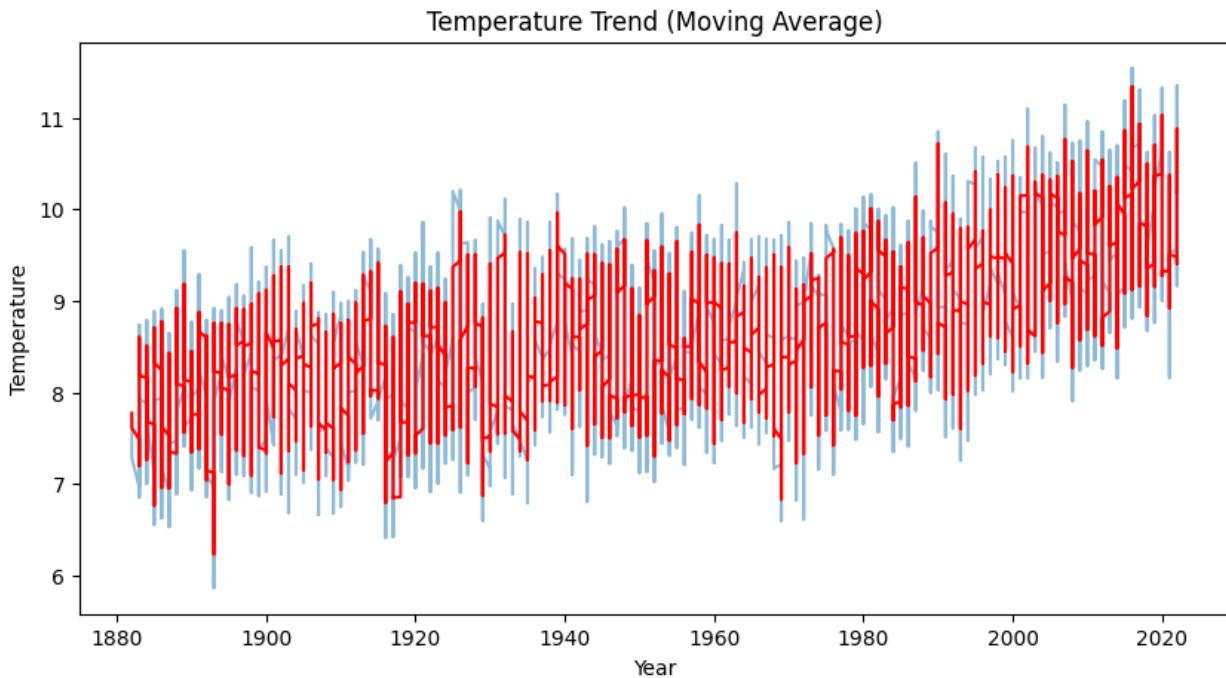
methods = {
    "Linear Regression Slope": trend_slope,
    "ADF Test p-value": adf_test[1]
}

for method, value in methods.items():
    print(f"{method}: {value:.6f}")

```

Linear Regression Trend Slope: 0.011487





```

ADF Statistic: -9.437660805688283
p-value: 4.991882507815218e-16
Linear Regression Slope: 0.011487
ADF Test p-value: 0.000000

```

```

import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from statsmodels.tsa.seasonal import seasonal_decompose

# Assuming df contains the dataset with 'Year', 'Temperature', and
# 'Anomaly' columns
group_size = 10
start_year = 1880
end_year = 2020

for start in range(start_year, end_year, group_size):
    df_group = df[(df['Year'] >= start) & (df['Year'] < start + group_size)].copy()
    print(f"Analyzing period: {start}-{start + group_size}")

    # Outlier Detection using IQR
    Q1 = df_group['Temperature'].quantile(0.25)
    Q3 = df_group['Temperature'].quantile(0.75)
    IQR = Q3 - Q1
    lower_bound = Q1 - 1.5 * IQR
    upper_bound = Q3 + 1.5 * IQR
    outliers = df_group[(df_group['Temperature'] < lower_bound) | (df_group['Temperature'] > upper_bound)]

```

```

(df_group['Temperature'] > upper_bound)
    print(f"Number of Outliers: {len(outliers)}")

# Boxplot for Outliers
plt.figure(figsize=(8, 5))
sns.boxplot(x=df_group['Temperature'])
plt.title(f"Boxplot of Temperature ({start}-{start + group_size})")
plt.show()

# Trend Analysis using Moving Average
df_group['Rolling_Mean'] =
df_group['Temperature'].rolling(window=365, min_periods=1).mean()
df_group['Rolling_Std'] =
df_group['Temperature'].rolling(window=365, min_periods=1).std()

plt.figure(figsize=(12, 6))
plt.plot(df_group['Temperature'], label="Original Temperature",
alpha=0.5)
plt.plot(df_group['Rolling_Mean'], label="365-day Moving Average",
color='red')
plt.plot(df_group['Rolling_Std'], label="365-day Rolling Std Dev",
color='green')
plt.title(f"Trend Analysis ({start}-{start + group_size})")
plt.legend()
plt.show()

# Seasonal Decomposition
decomposition = seasonal_decompose(df_group['Temperature'],
period=365, model='additive', extrapolate_trend='freq')

plt.figure(figsize=(12, 8))
plt.subplot(411)
plt.plot(df_group['Temperature'], label="Original", alpha=0.5)
plt.legend()

plt.subplot(412)
plt.plot(decomposition.trend, label="Trend", color='red')
plt.legend()

plt.subplot(413)
plt.plot(decomposition.seasonal, label="Seasonality",
color='green')
plt.legend()

plt.subplot(414)
plt.plot(decomposition.resid, label="Residuals", color='purple')
plt.legend()

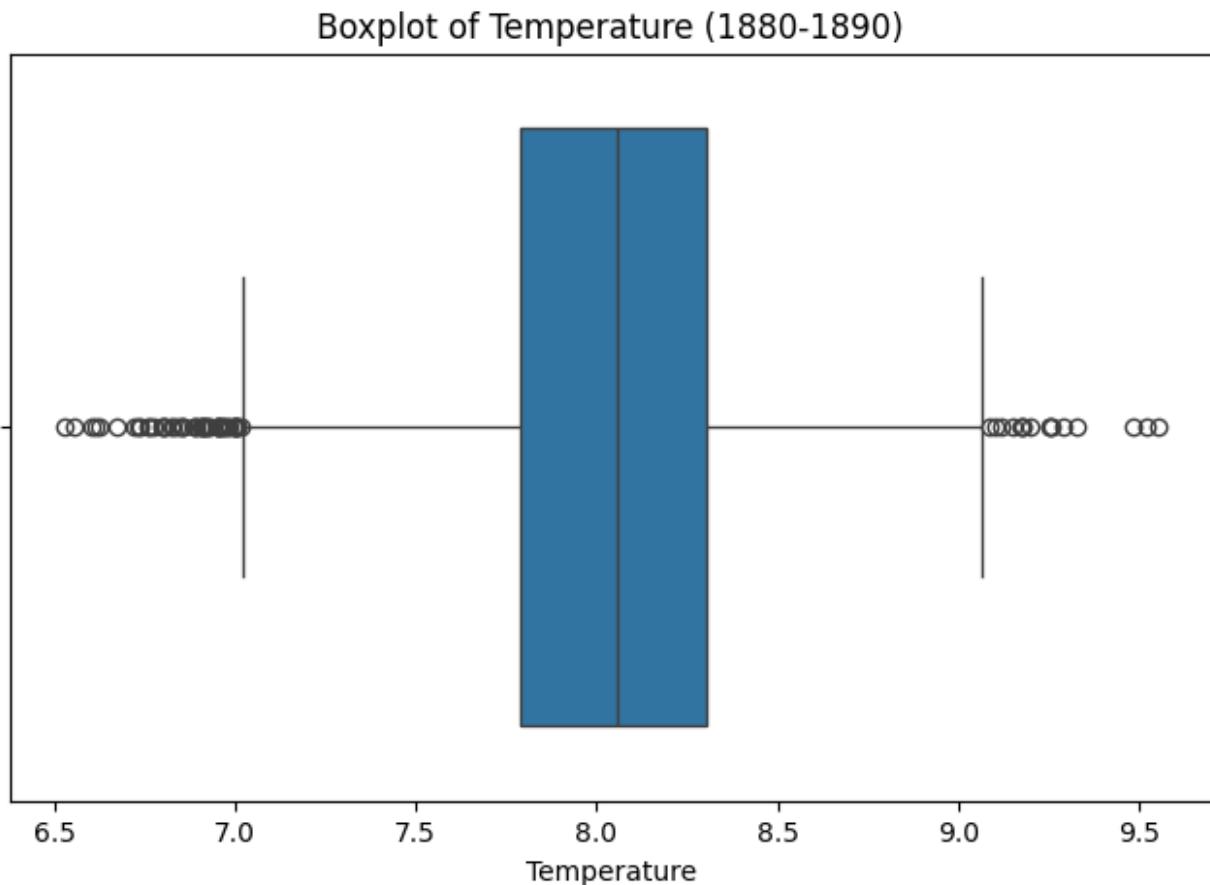
```

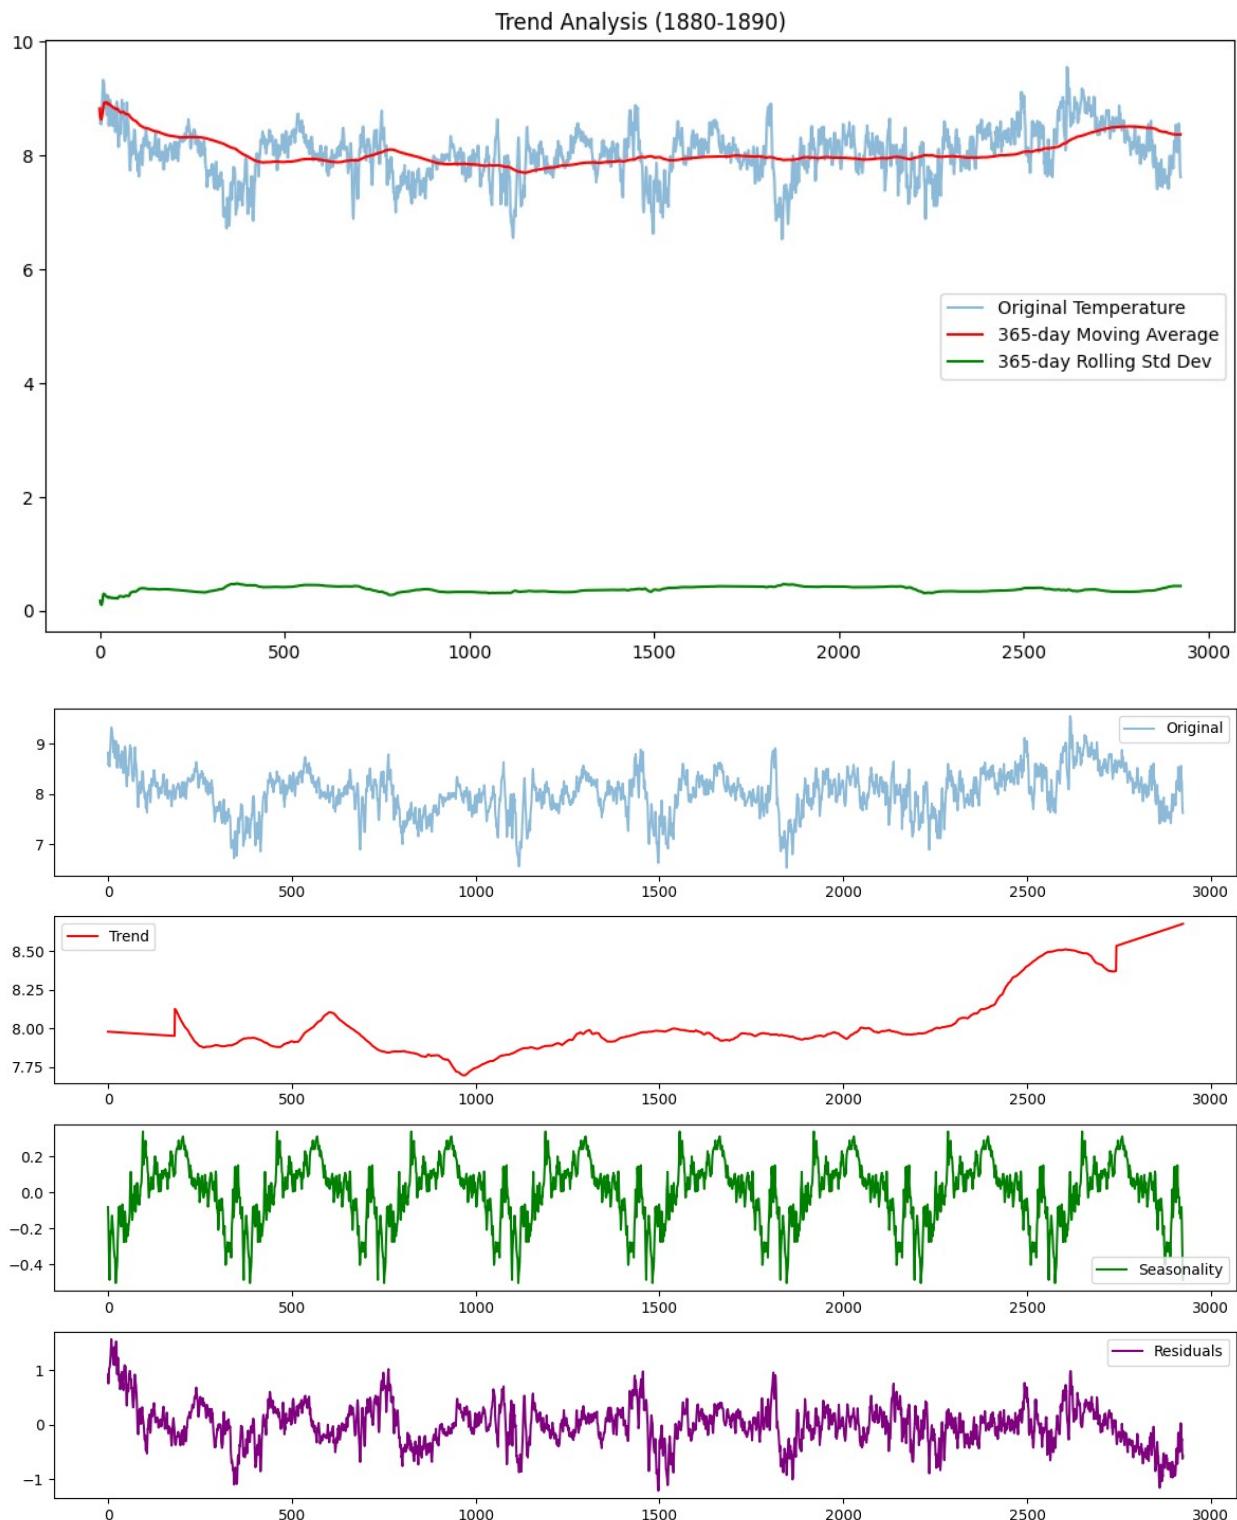
```
plt.tight_layout()
plt.show()

# Histogram of Temperature Anomalies
plt.figure(figsize=(8, 5))
sns.histplot(df_group['Anomaly'], bins=50, kde=True)
plt.title(f"Histogram of Temperature Anomalies ({start}-{start + group_size})")
plt.xlabel("Temperature Anomaly")
plt.show()
```

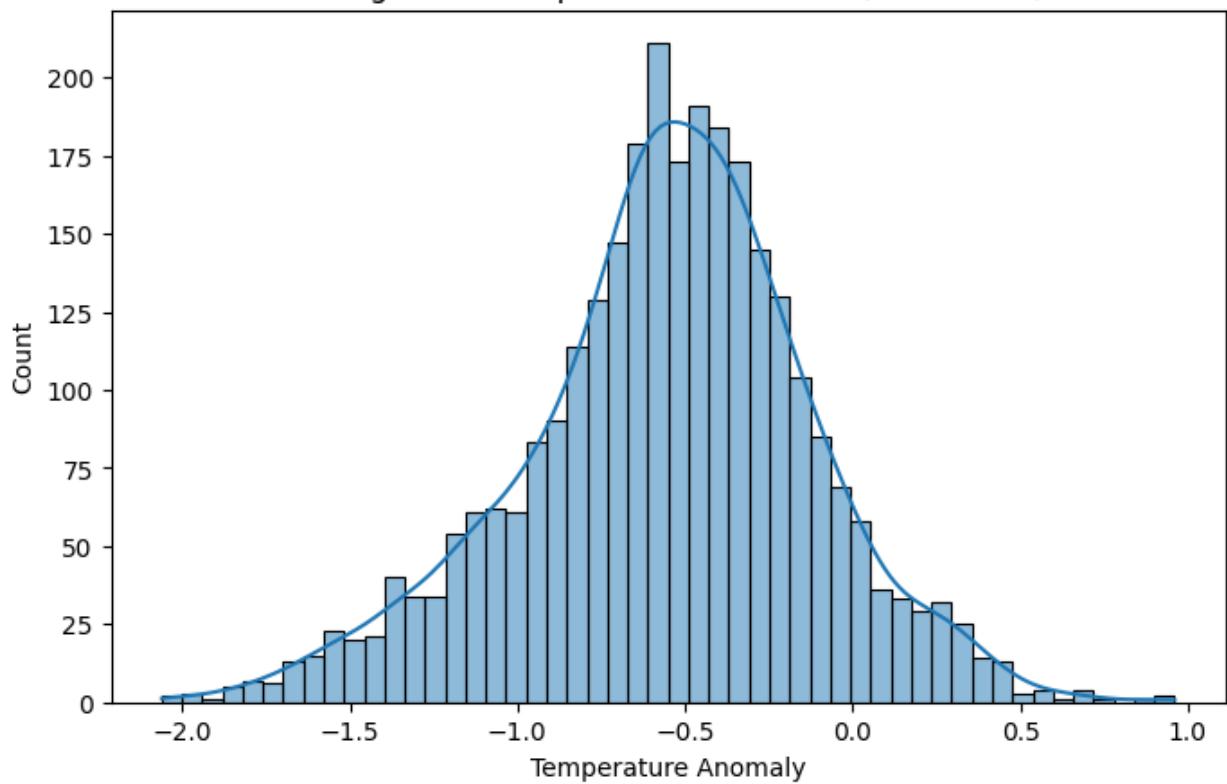
Analyzing period: 1880-1890

Number of Outliers: 69



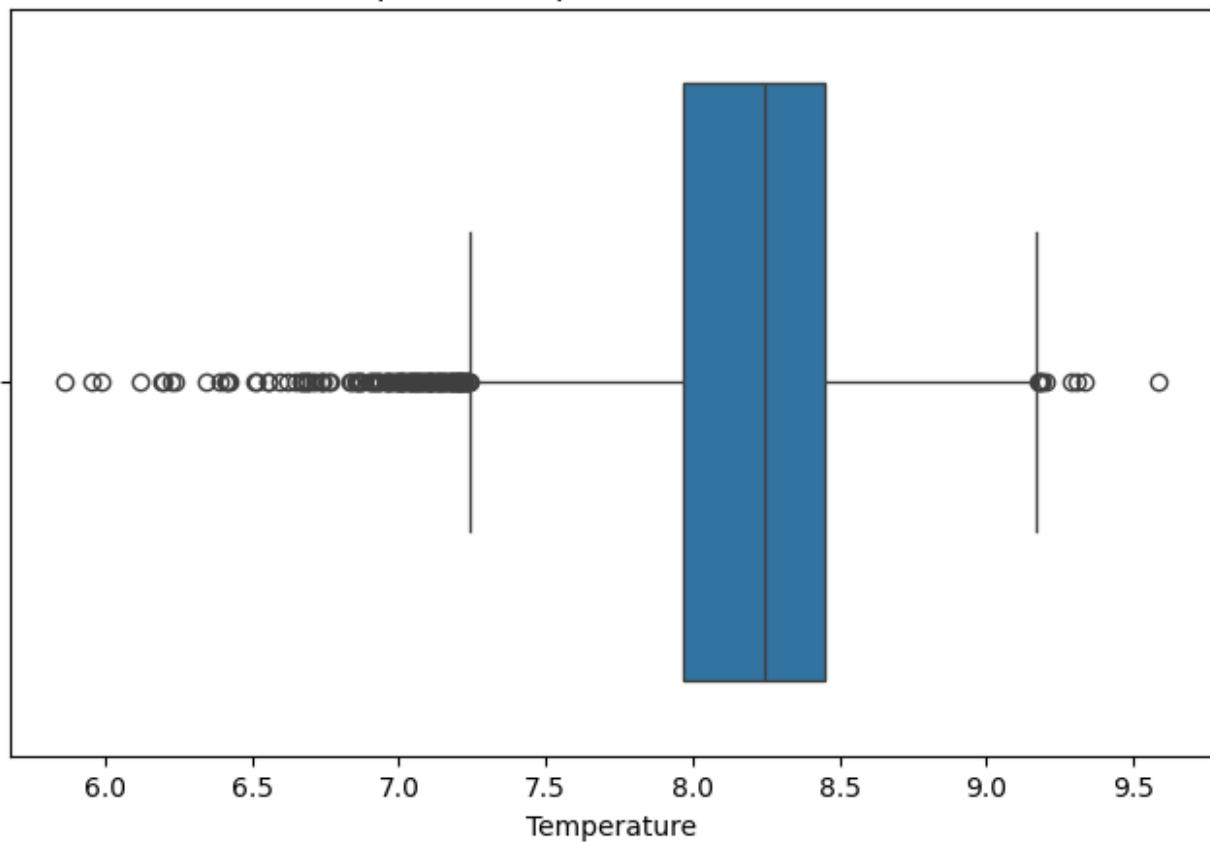


Histogram of Temperature Anomalies (1880-1890)

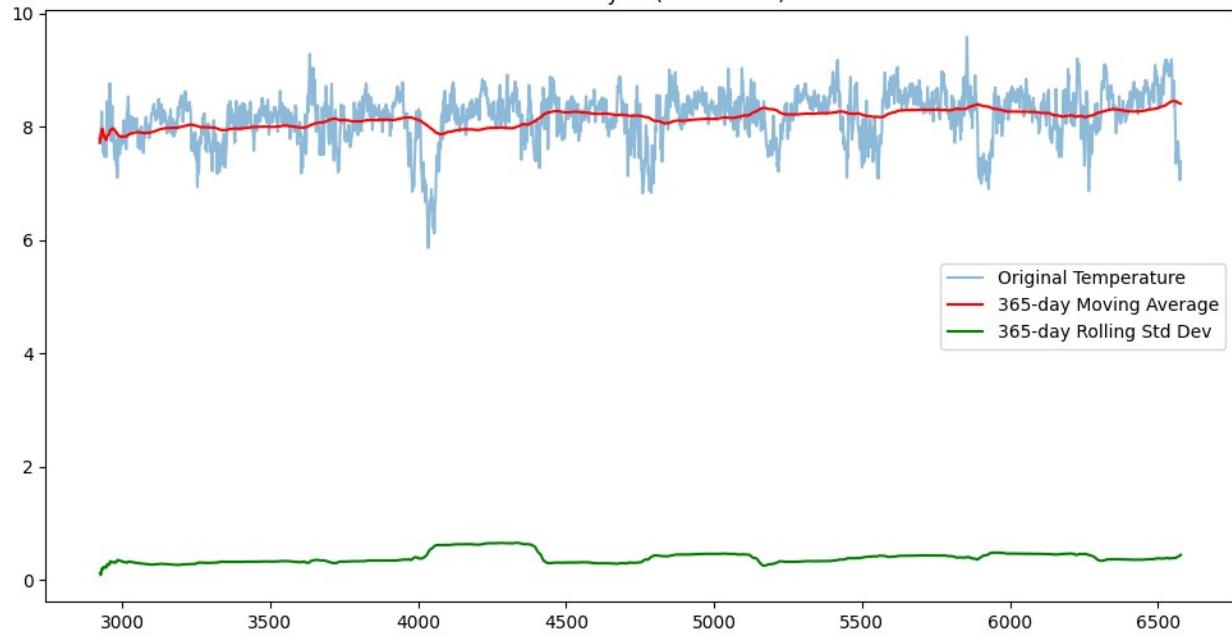


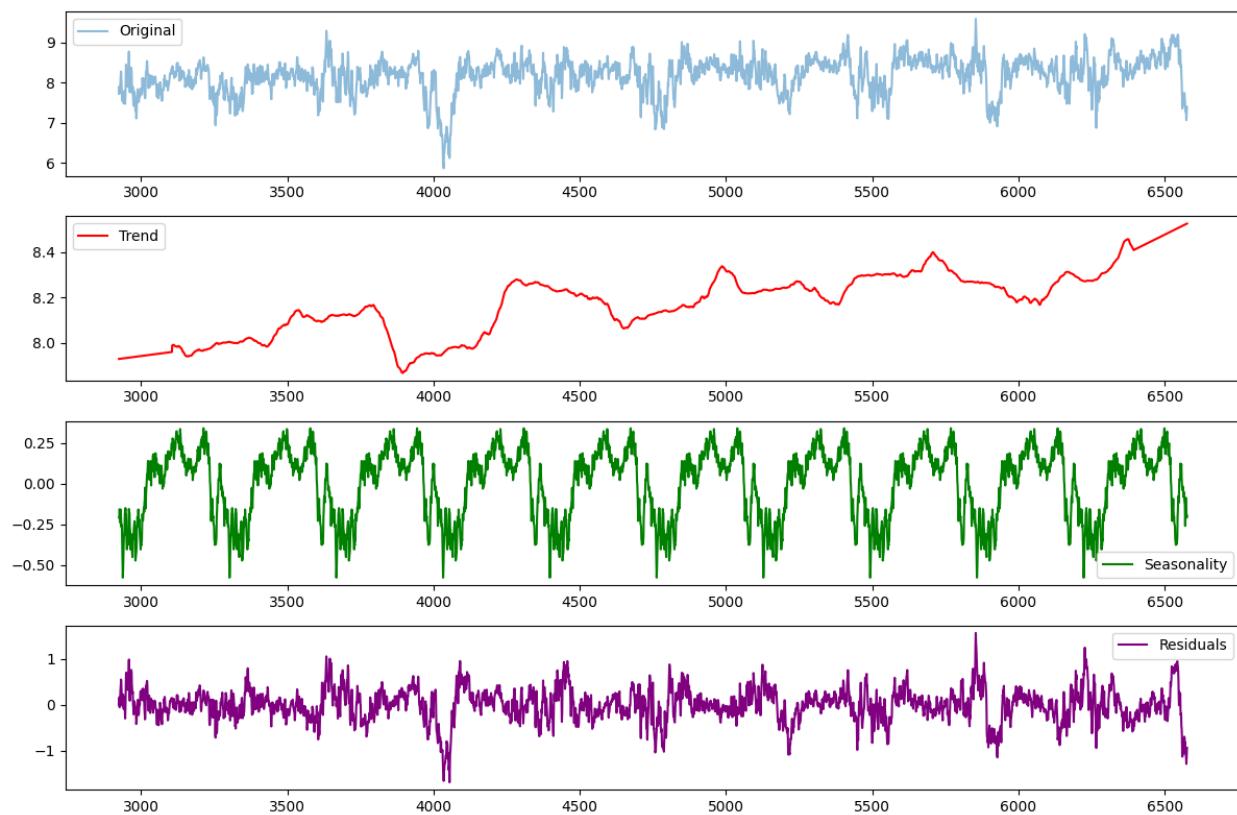
Analyzing period: 1890-1900
Number of Outliers: 137

Boxplot of Temperature (1890-1900)

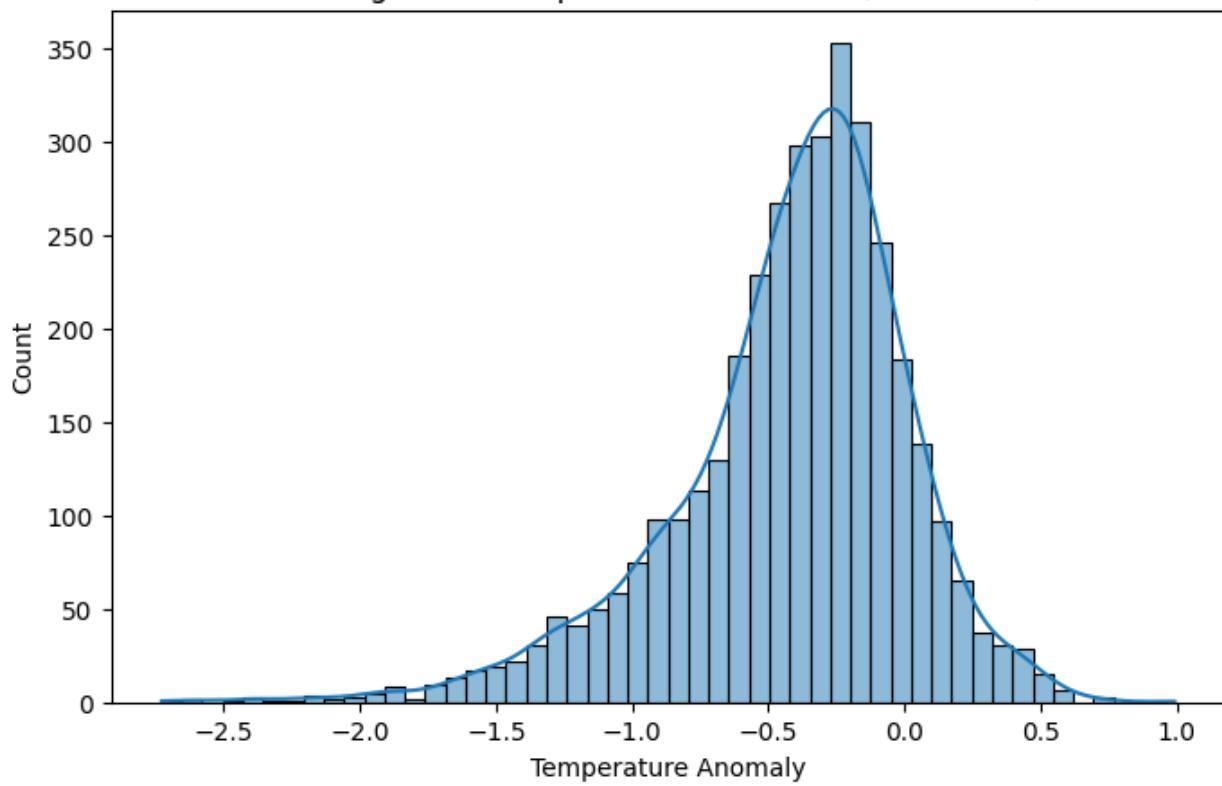


Trend Analysis (1890-1900)



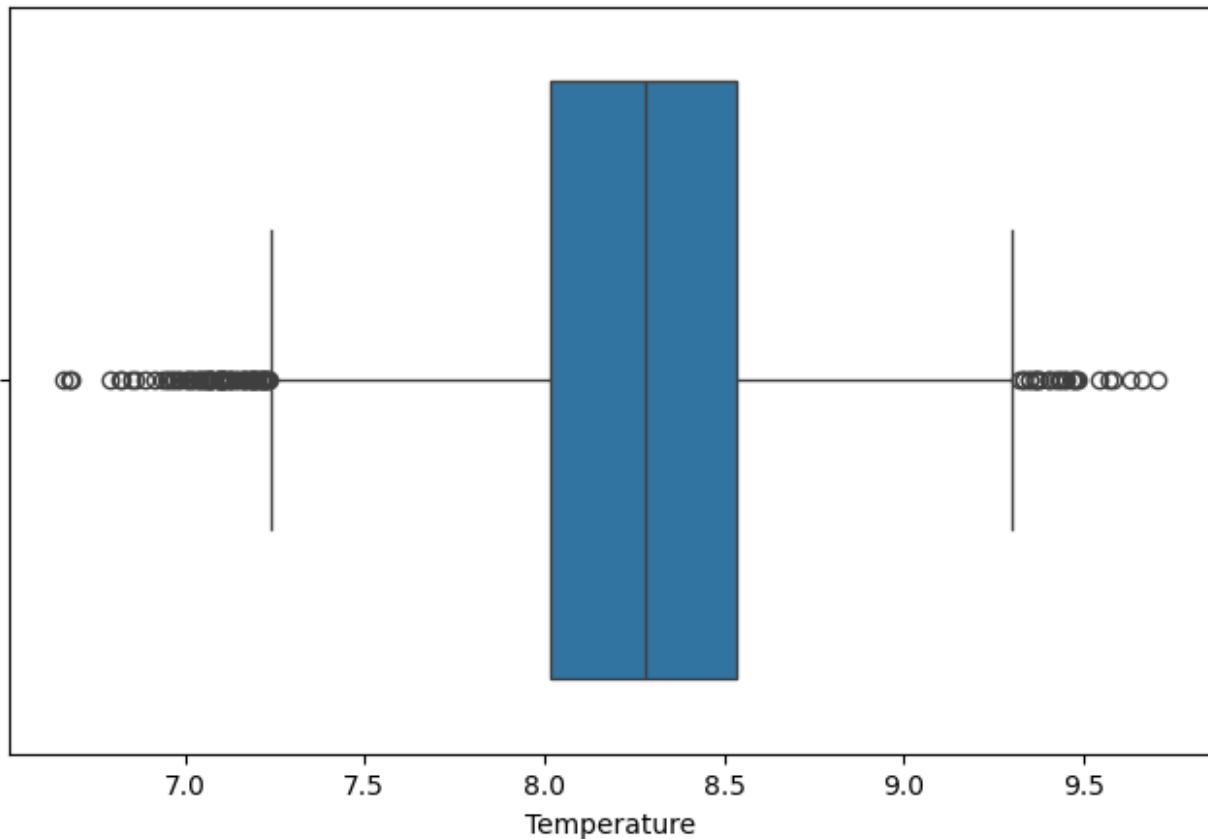


Histogram of Temperature Anomalies (1890-1900)

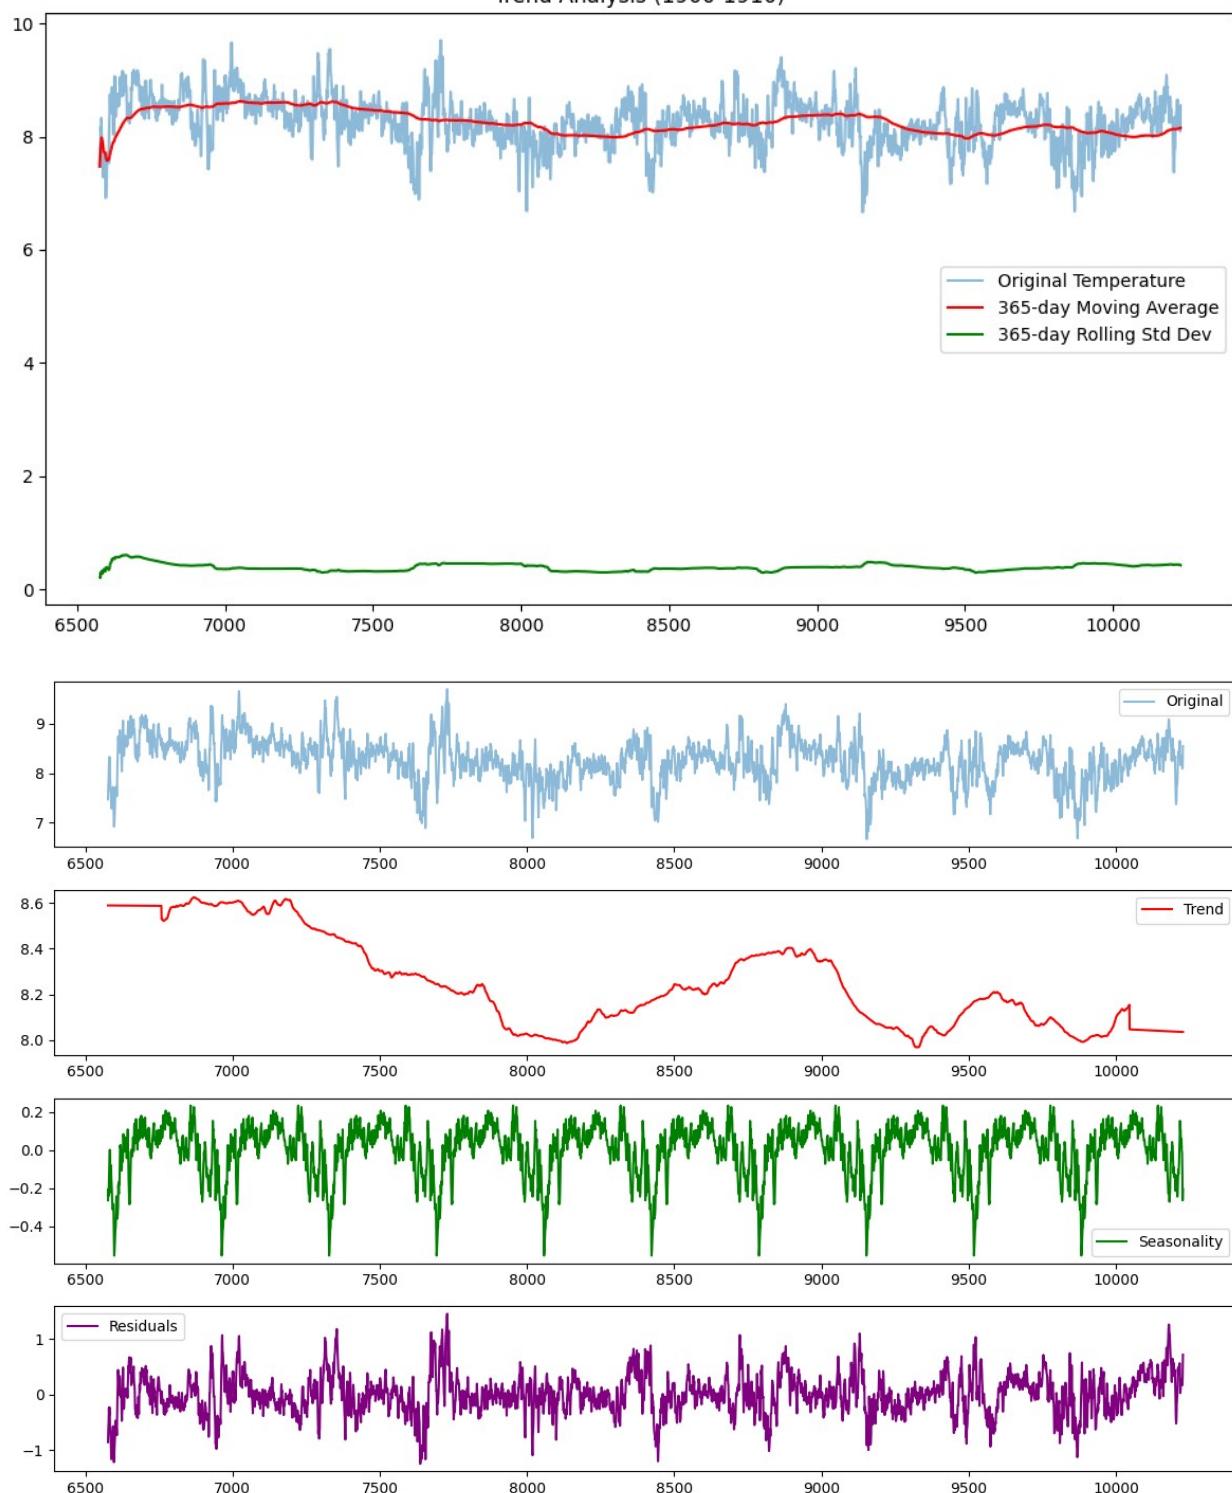


Analyzing period: 1900-1910
Number of Outliers: 105

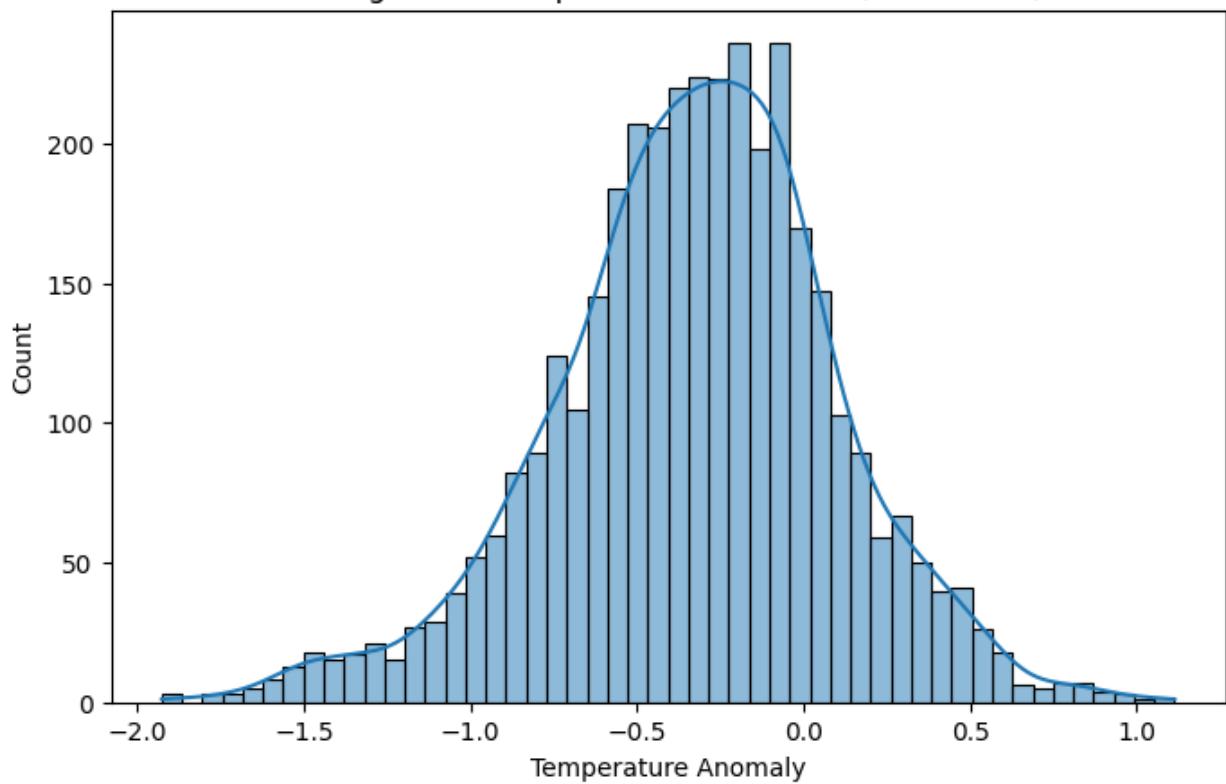
Boxplot of Temperature (1900-1910)



Trend Analysis (1900-1910)

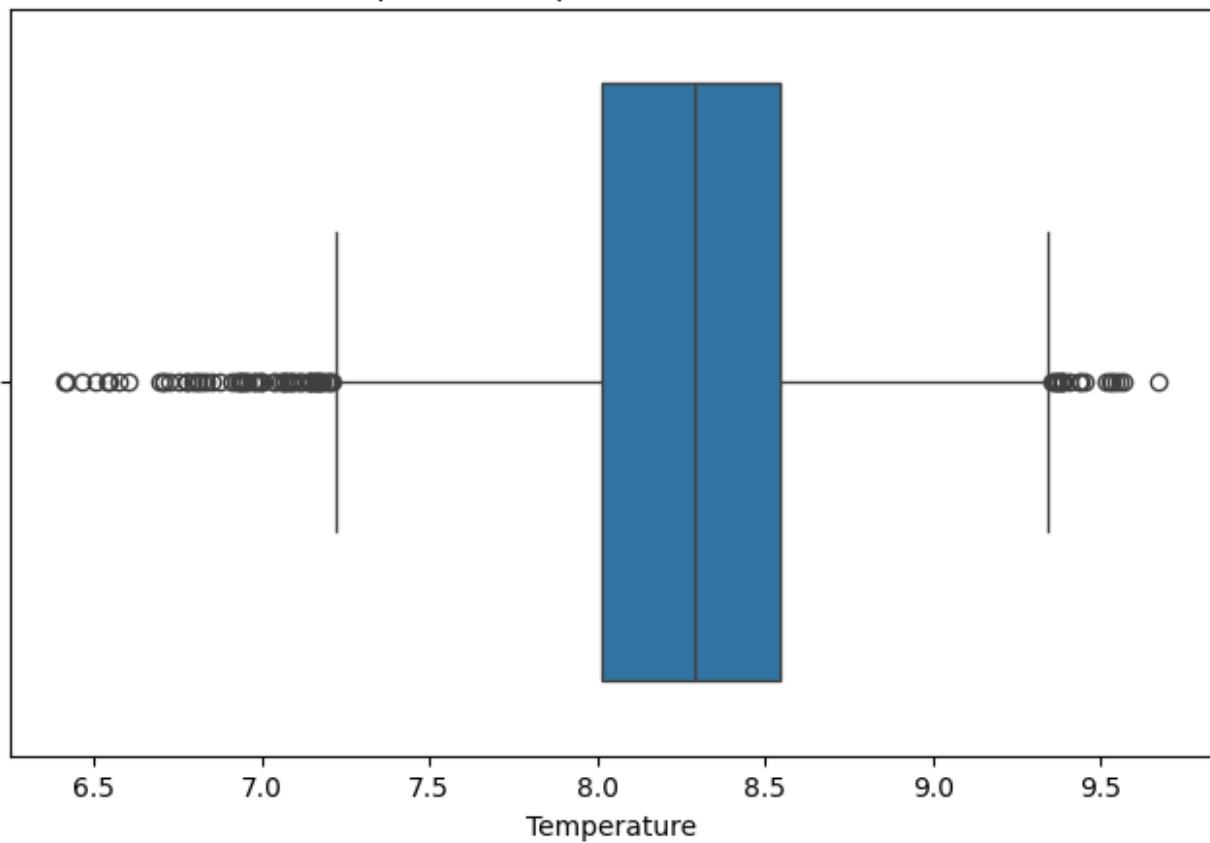


Histogram of Temperature Anomalies (1900-1910)

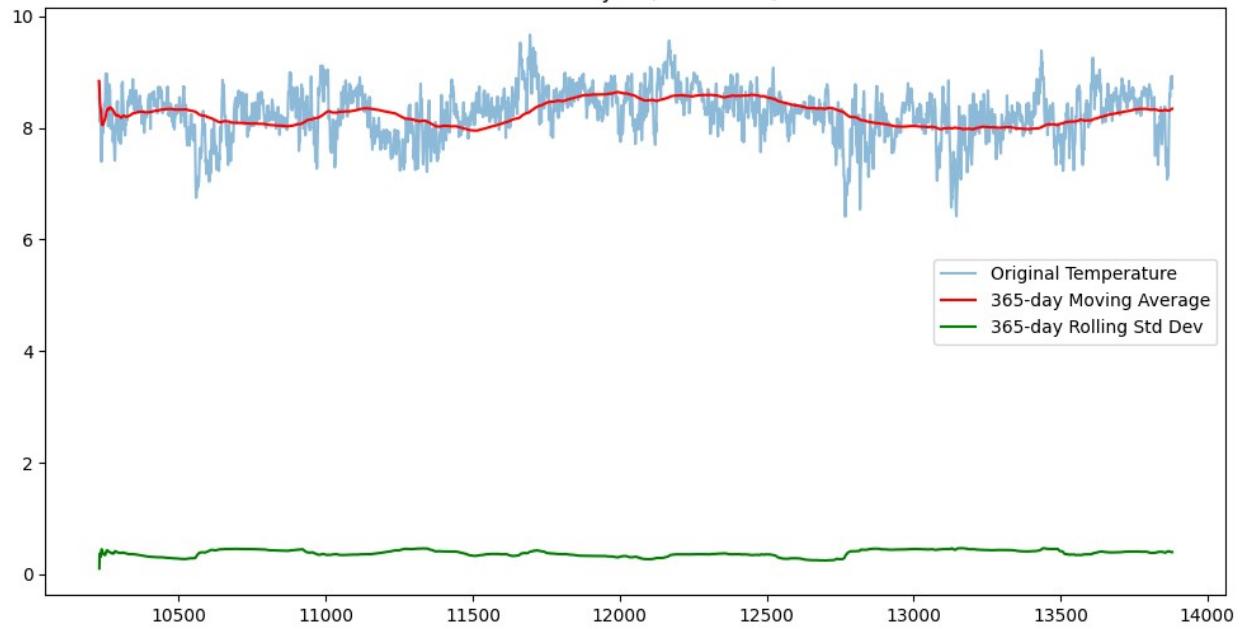


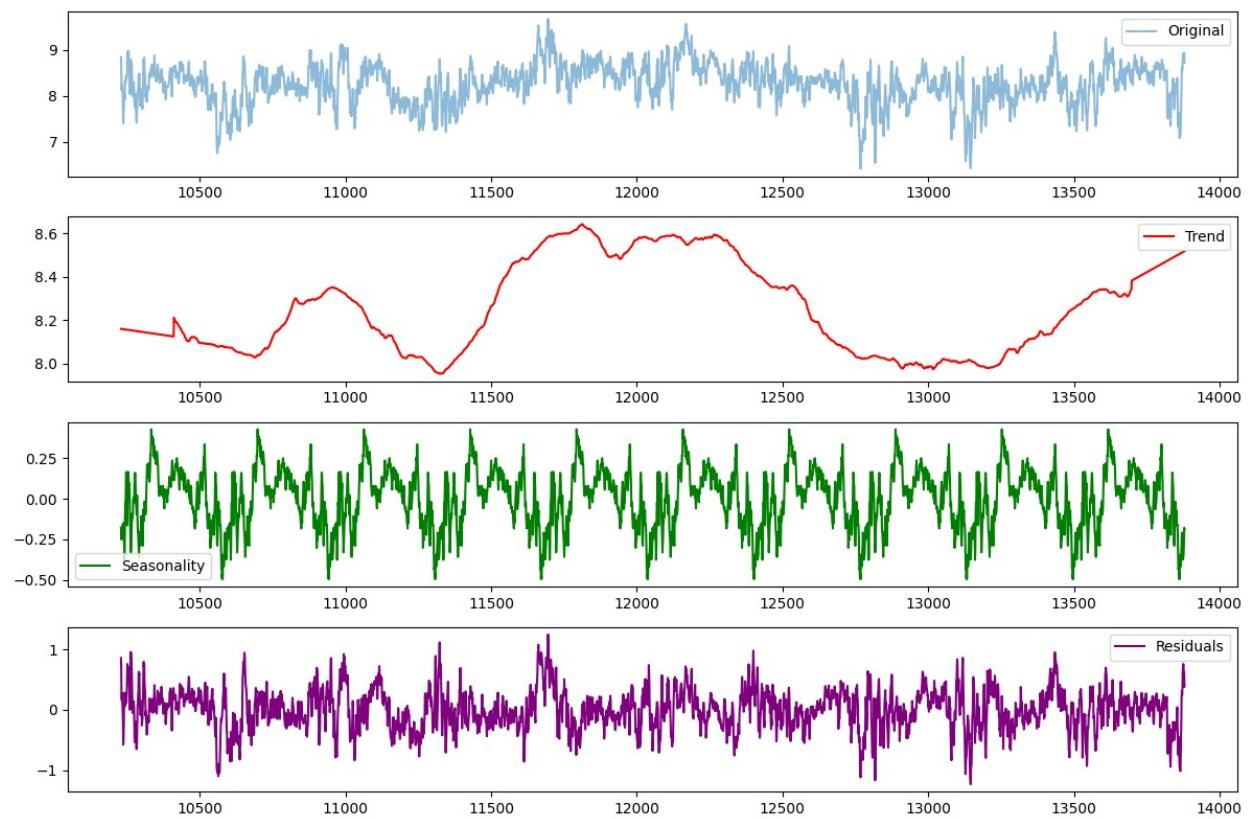
Analyzing period: 1910-1920
Number of Outliers: 88

Boxplot of Temperature (1910-1920)

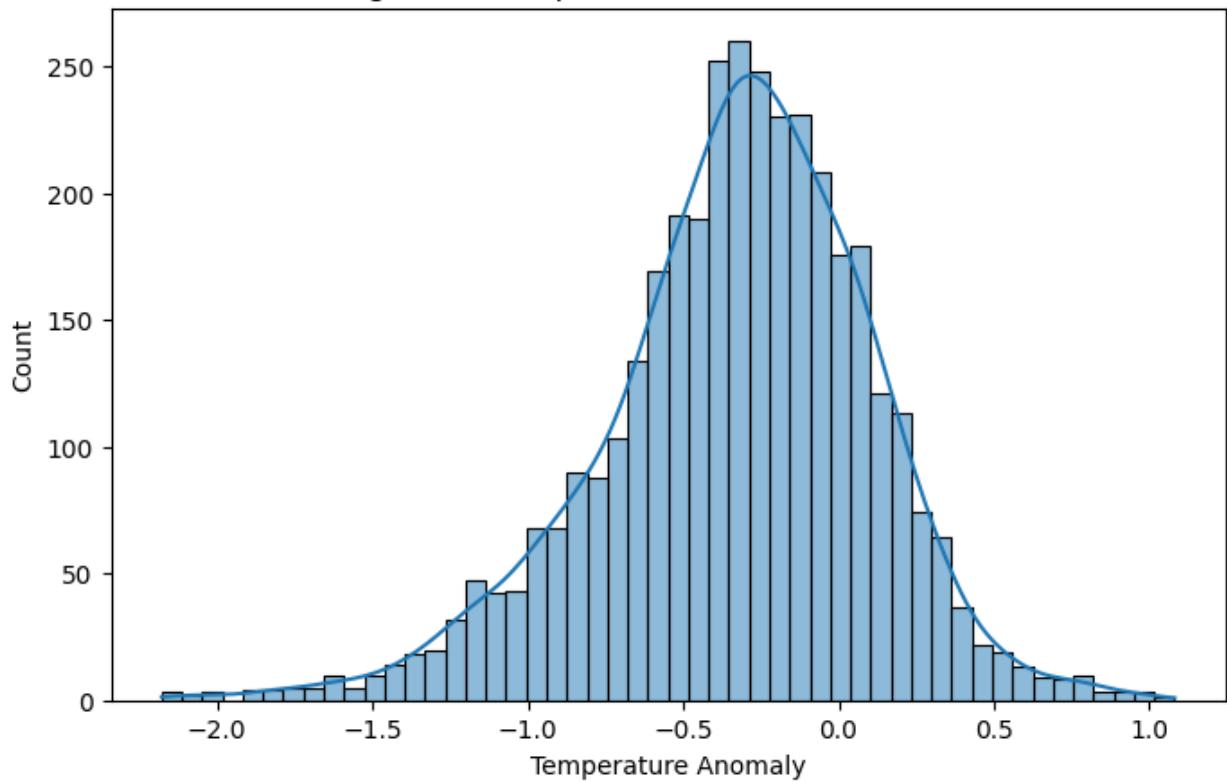


Trend Analysis (1910-1920)



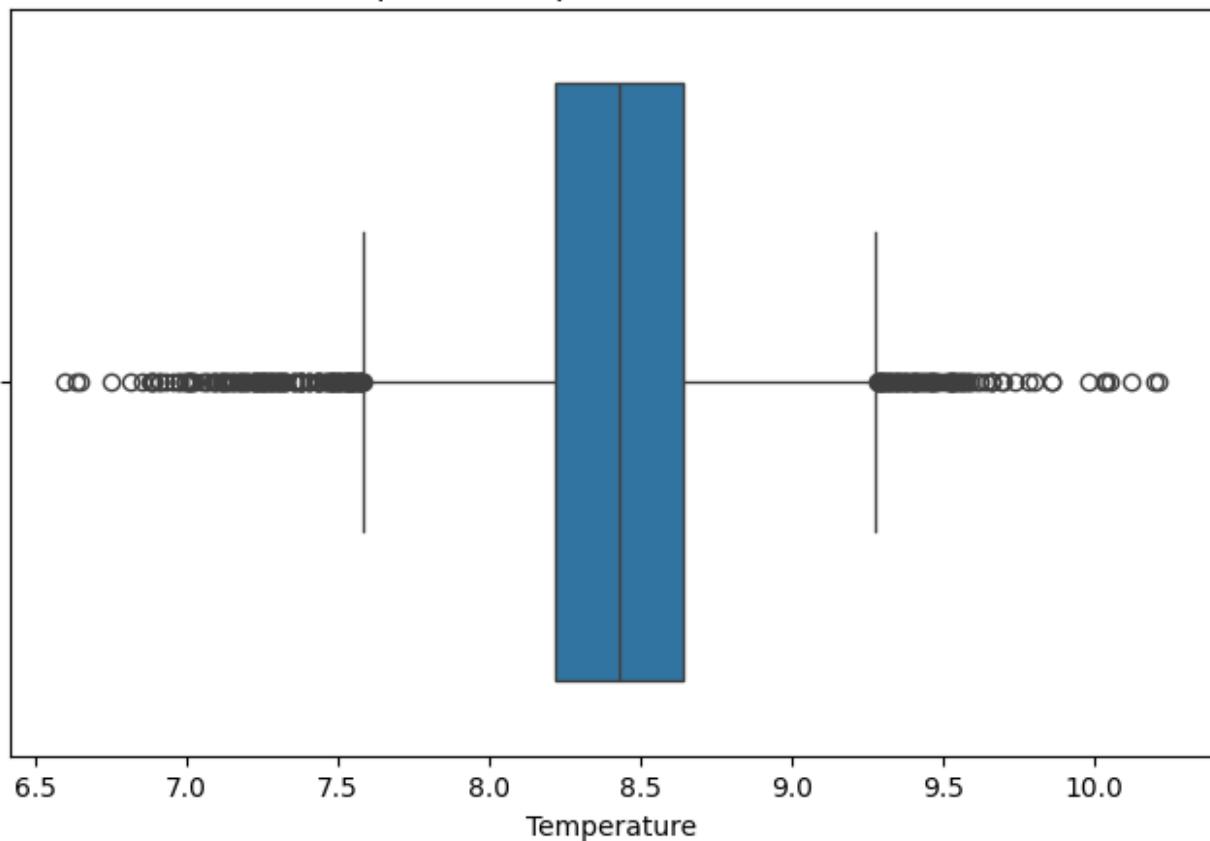


Histogram of Temperature Anomalies (1910-1920)

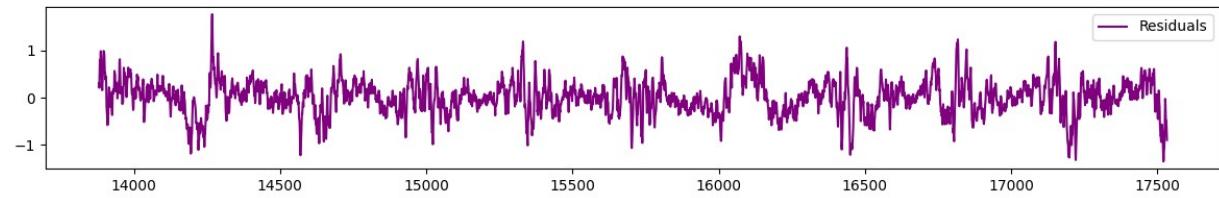
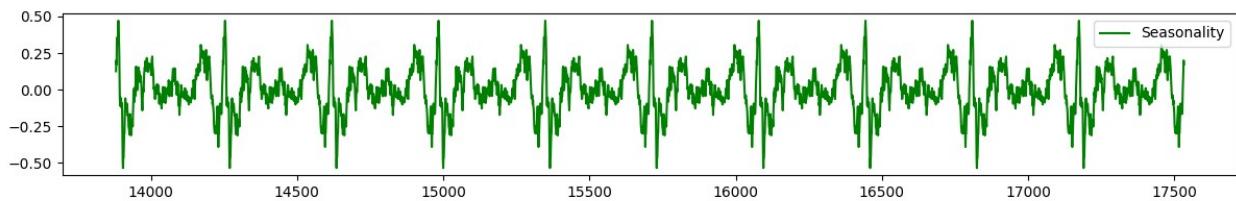
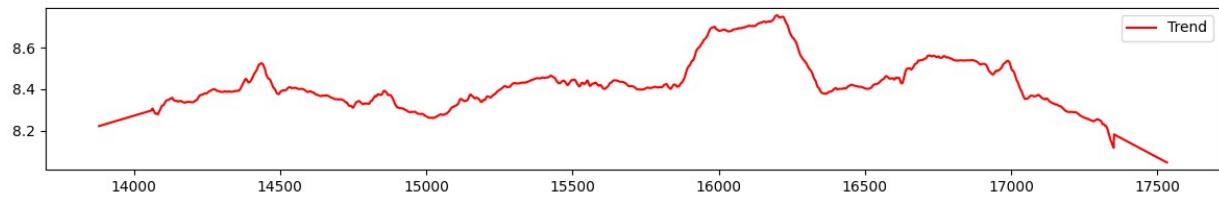
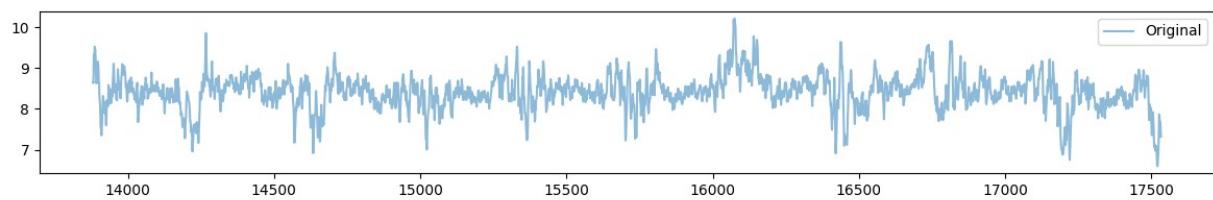
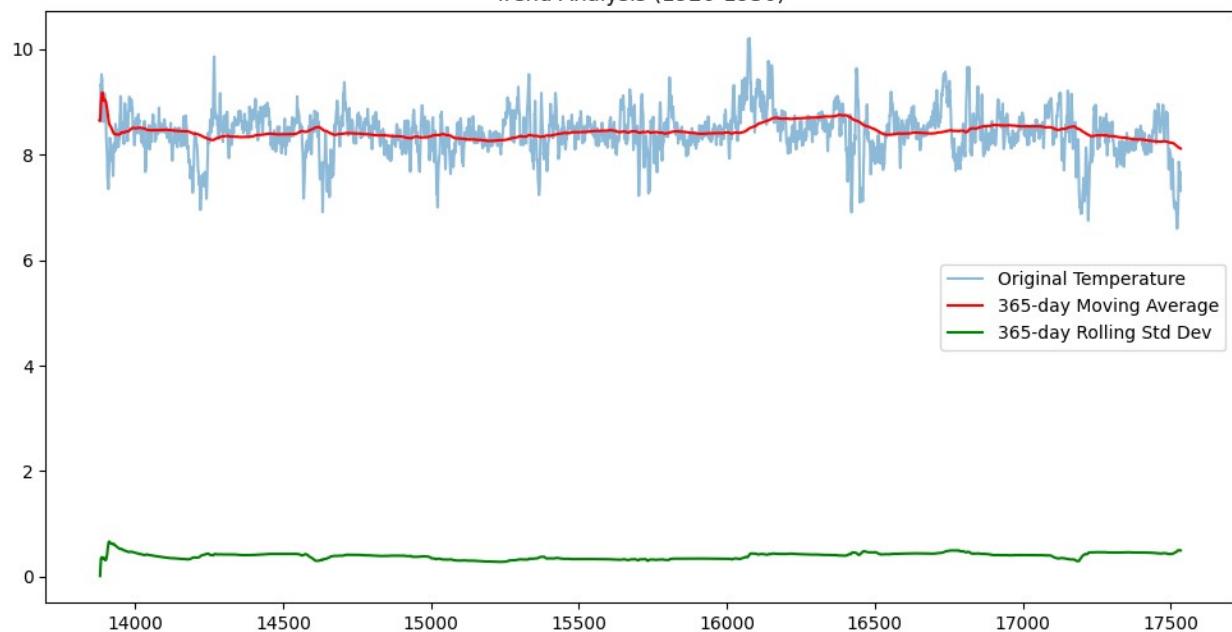


Analyzing period: 1920-1930
Number of Outliers: 236

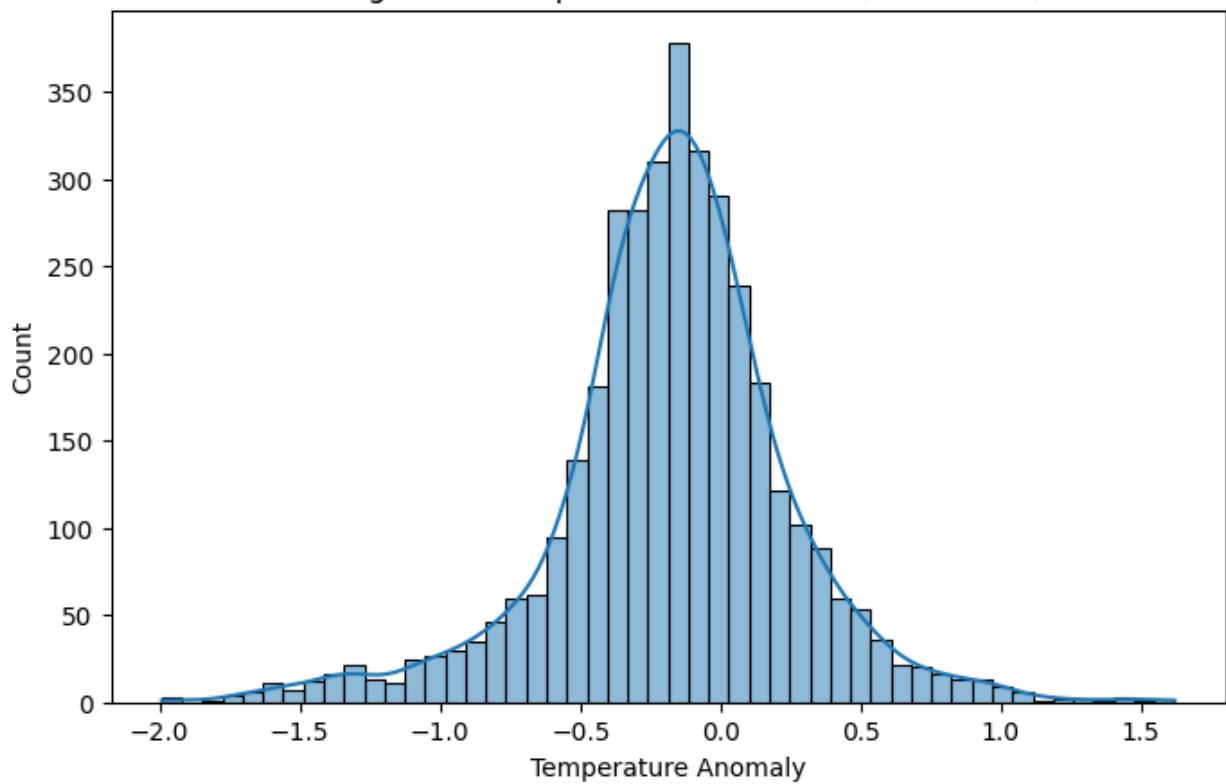
Boxplot of Temperature (1920-1930)



Trend Analysis (1920-1930)

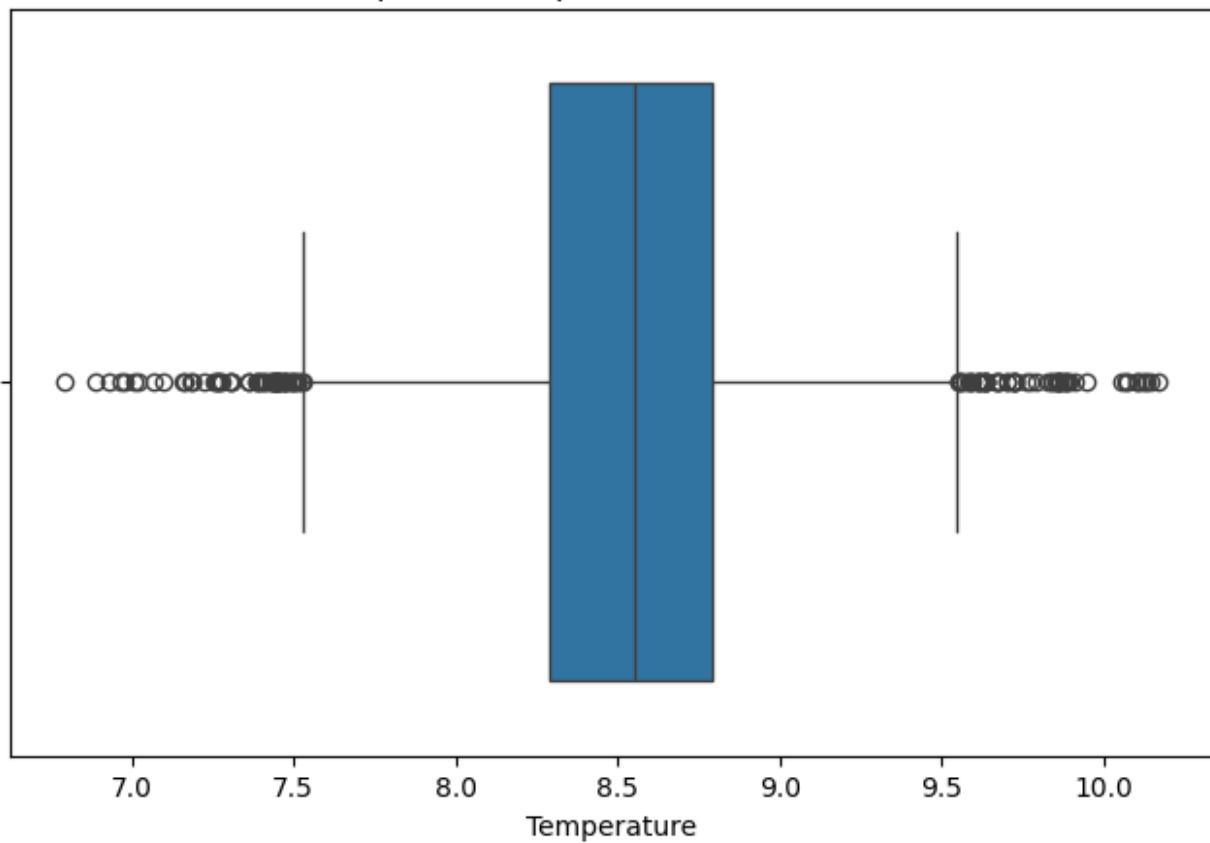


Histogram of Temperature Anomalies (1920-1930)

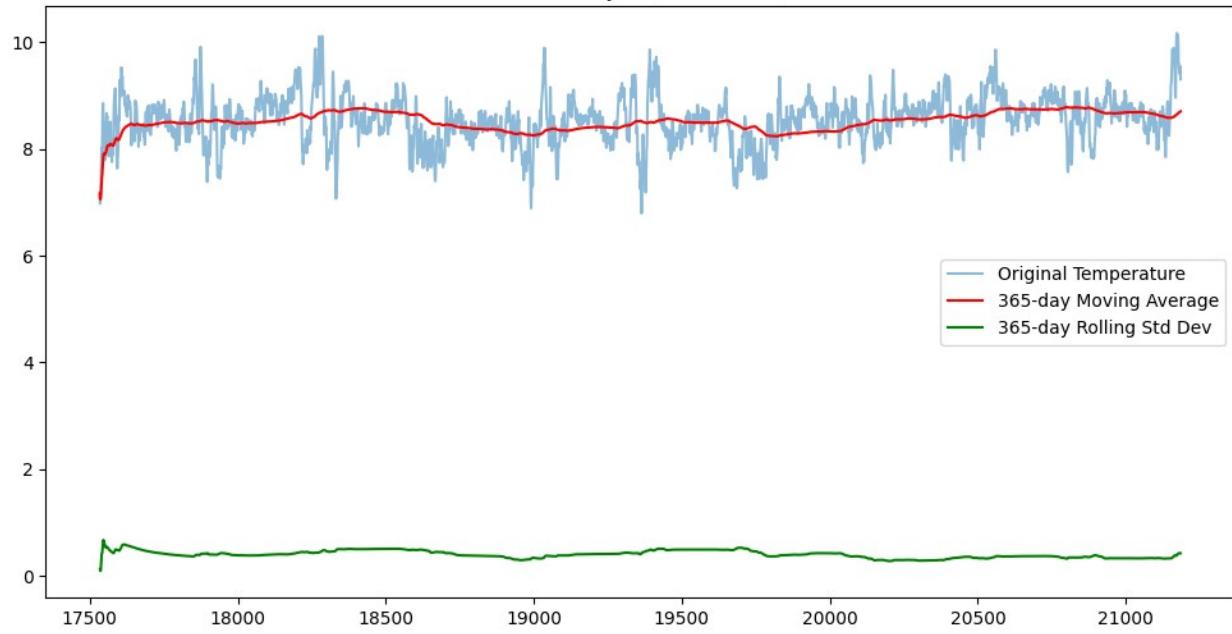


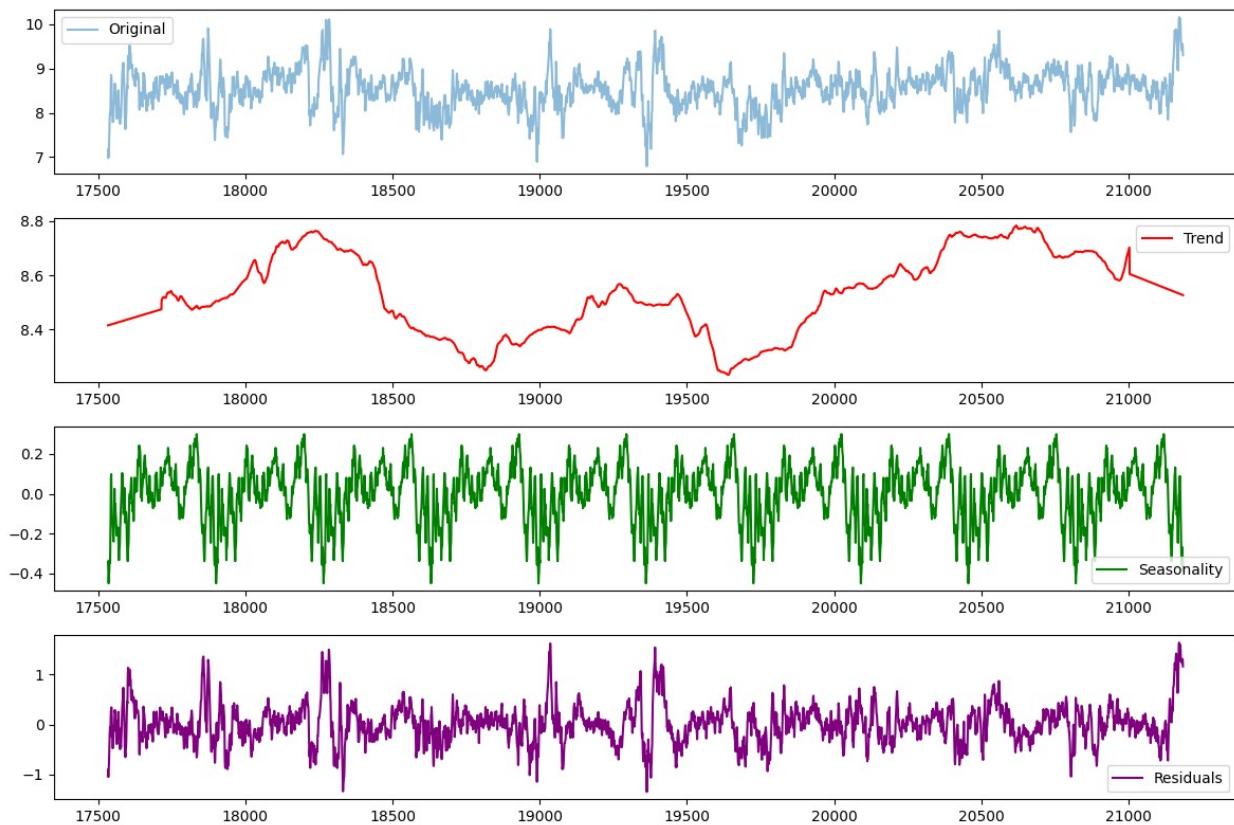
Analyzing period: 1930-1940
Number of Outliers: 121

Boxplot of Temperature (1930-1940)

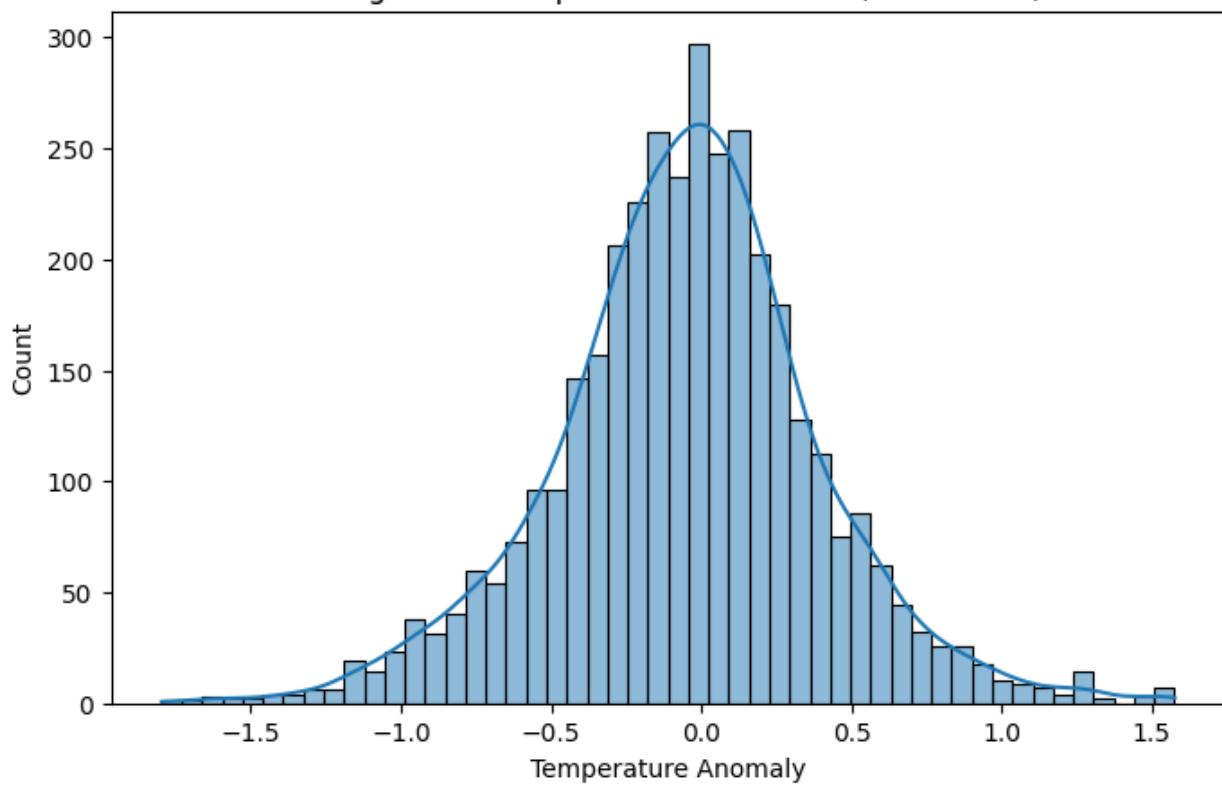


Trend Analysis (1930-1940)



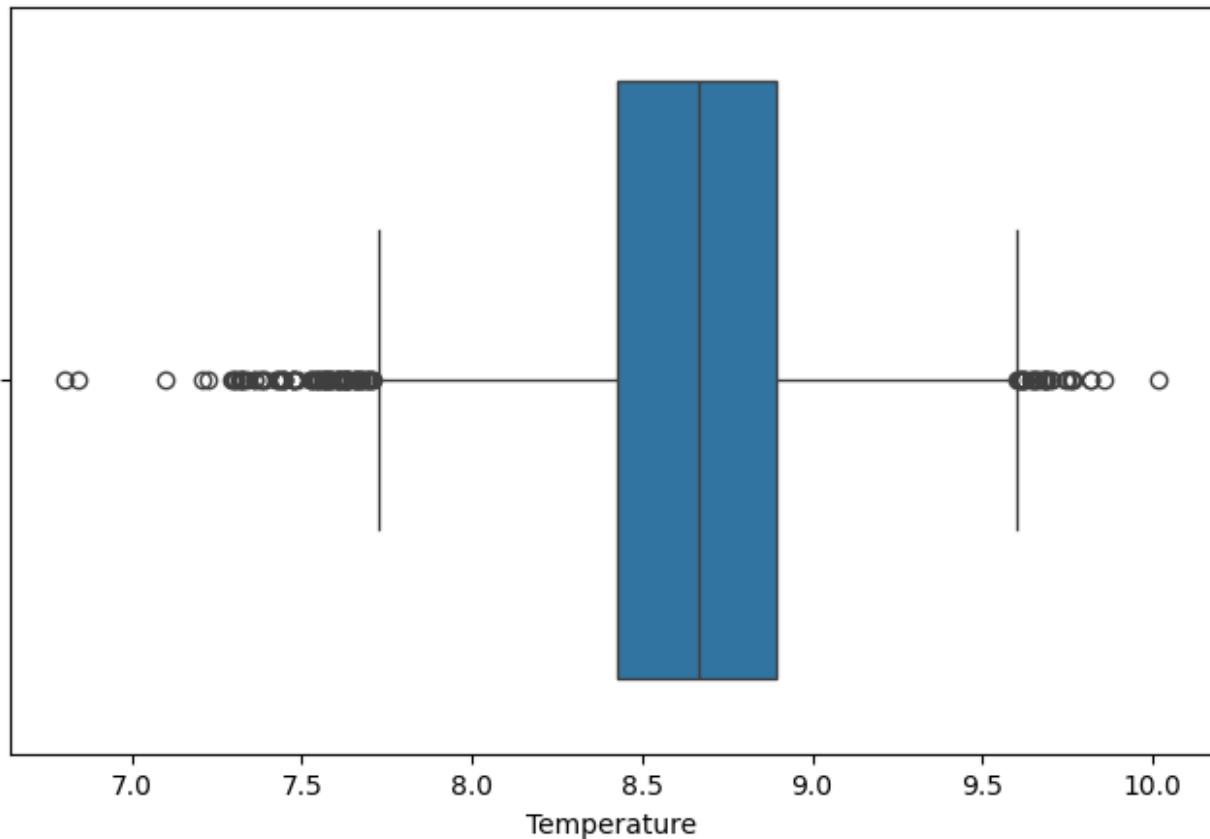


Histogram of Temperature Anomalies (1930-1940)

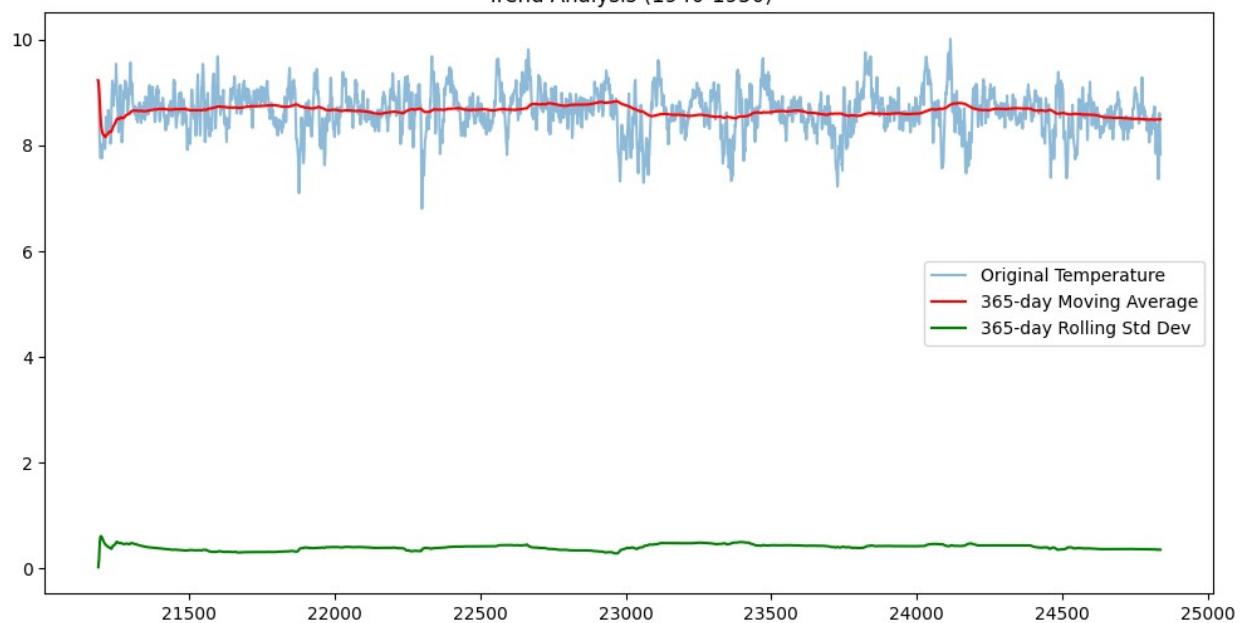


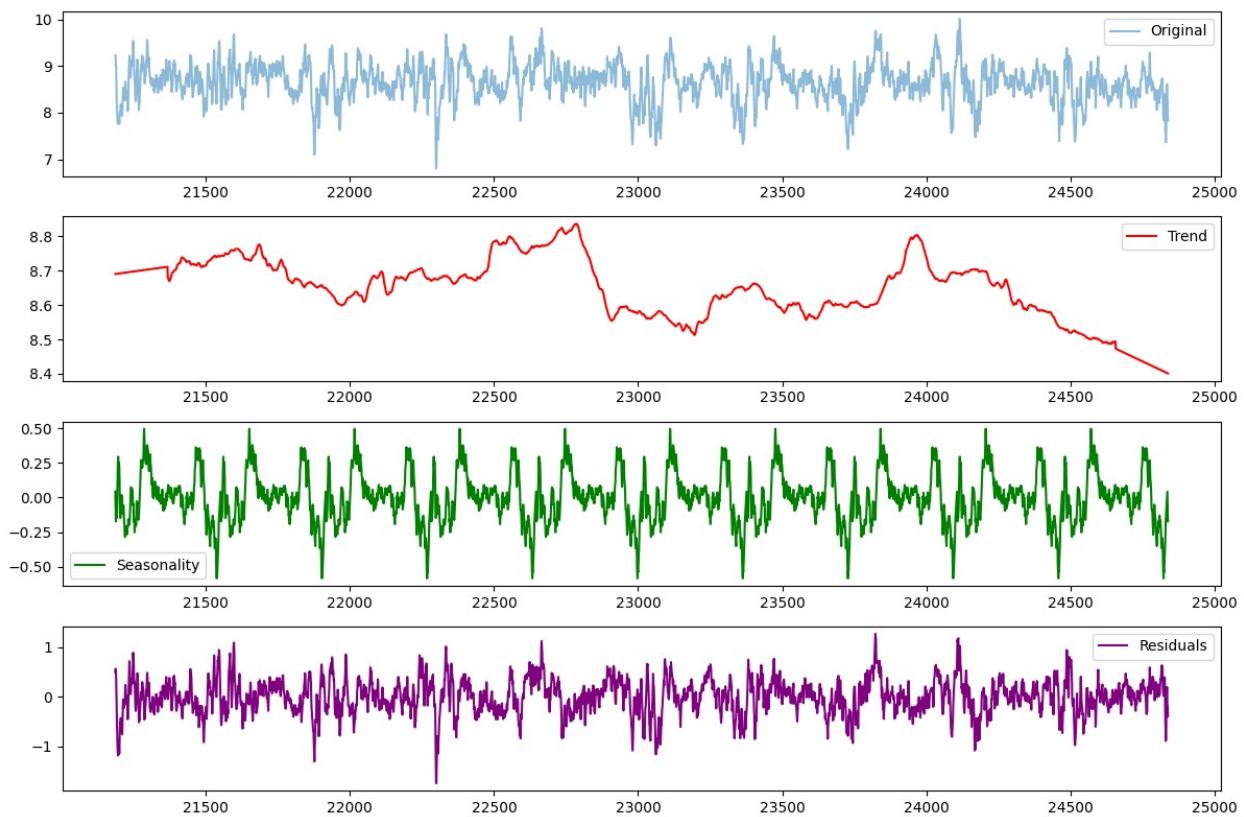
Analyzing period: 1940-1950
Number of Outliers: 120

Boxplot of Temperature (1940-1950)

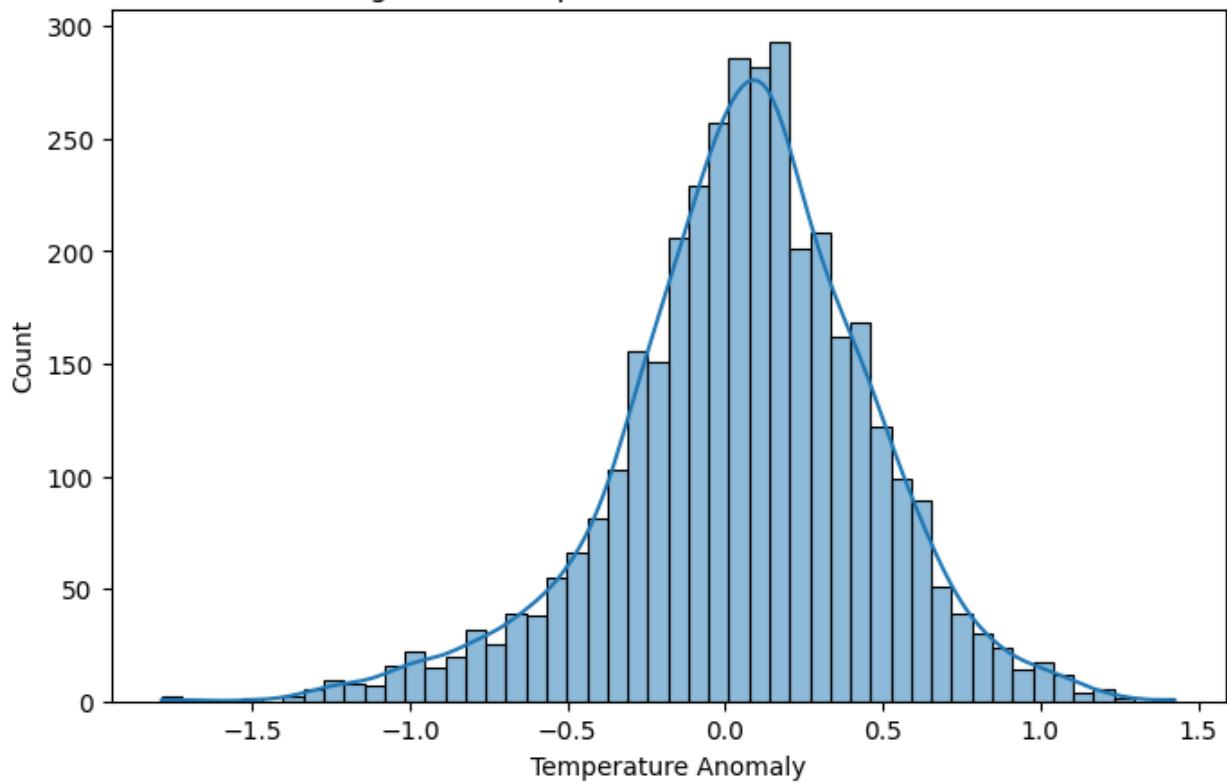


Trend Analysis (1940-1950)



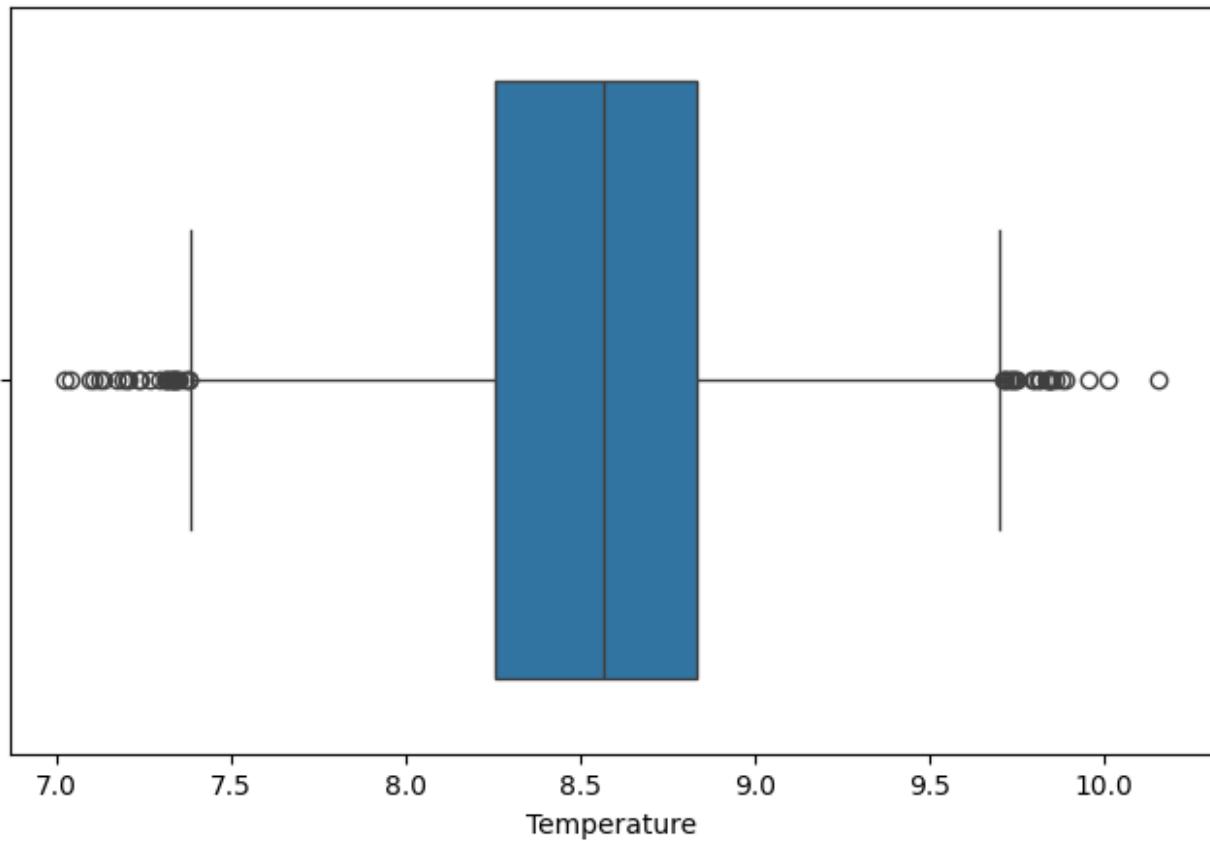


Histogram of Temperature Anomalies (1940-1950)

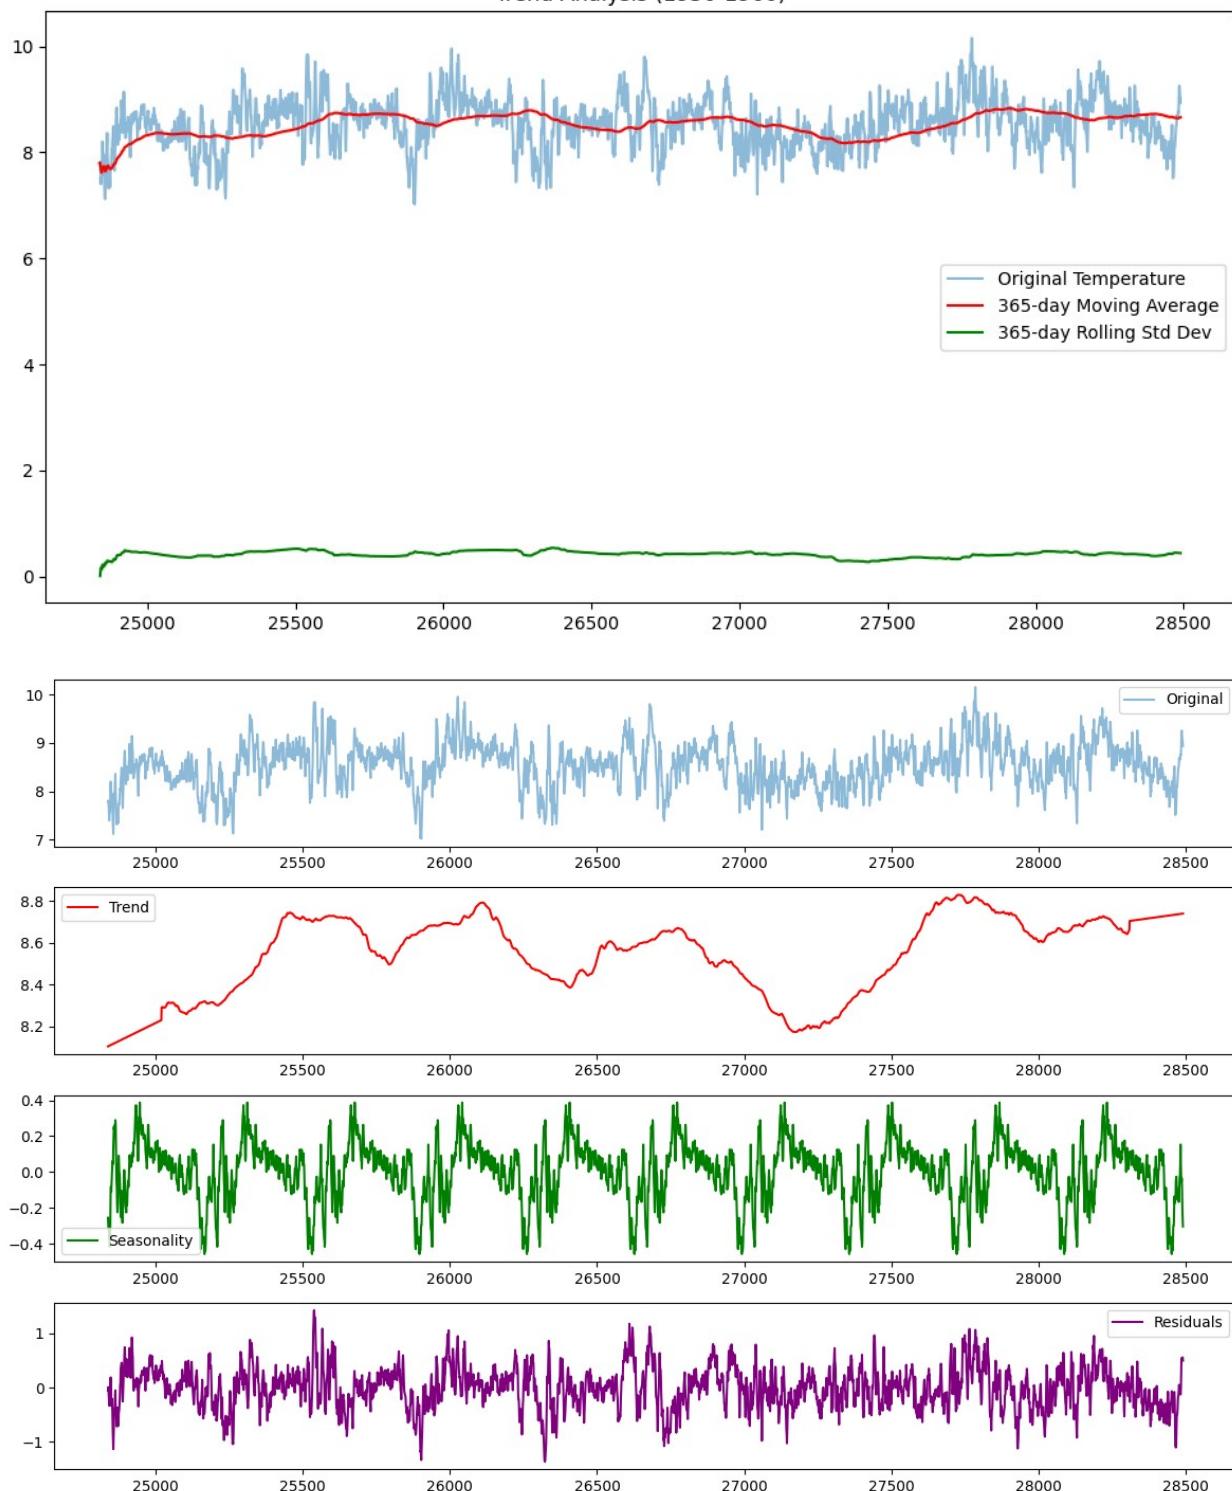


Analyzing period: 1950-1960
Number of Outliers: 52

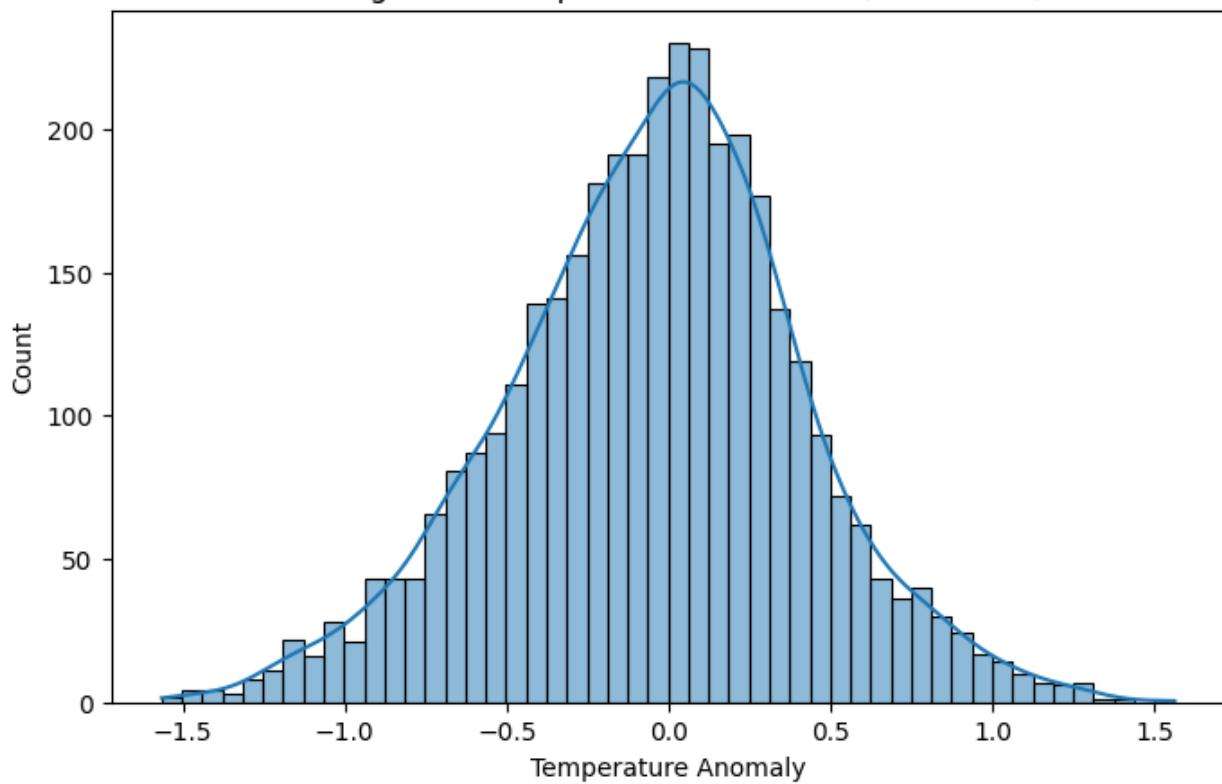
Boxplot of Temperature (1950-1960)



Trend Analysis (1950-1960)

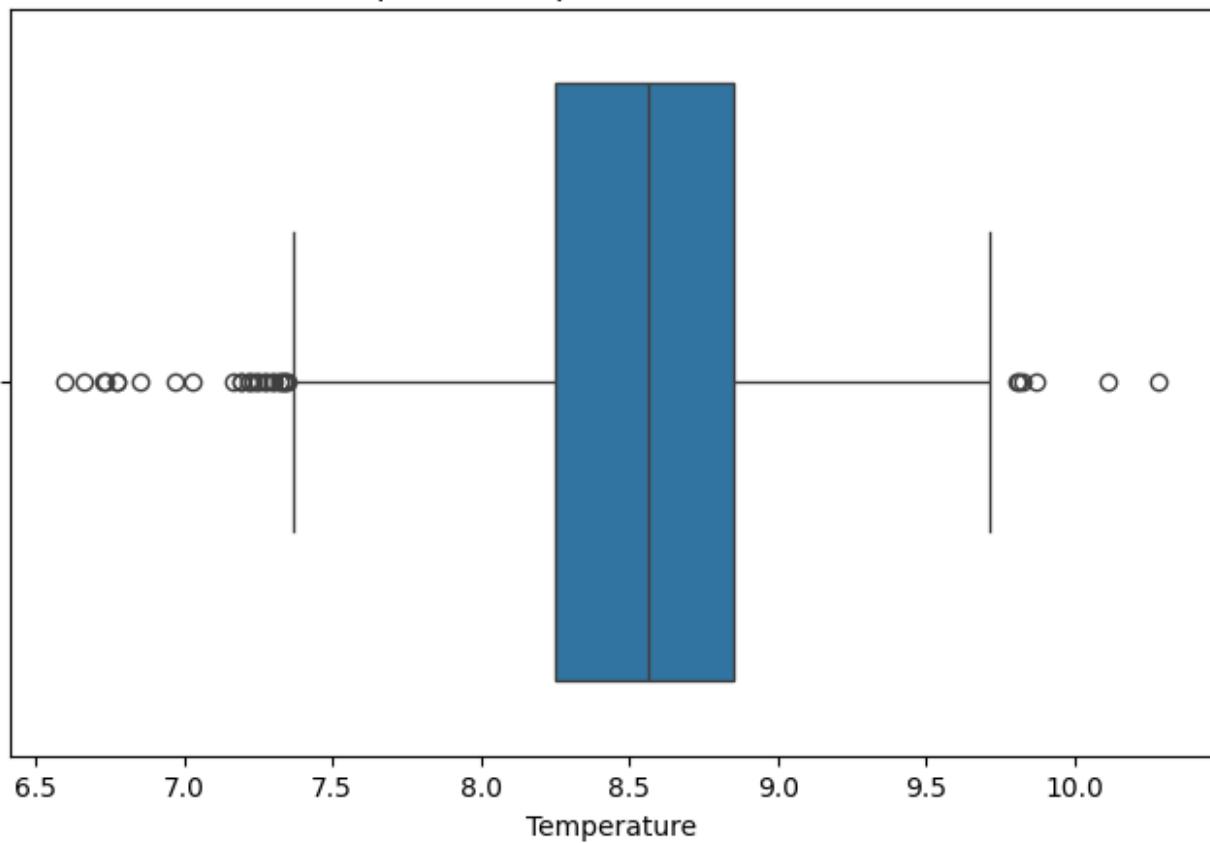


Histogram of Temperature Anomalies (1950-1960)

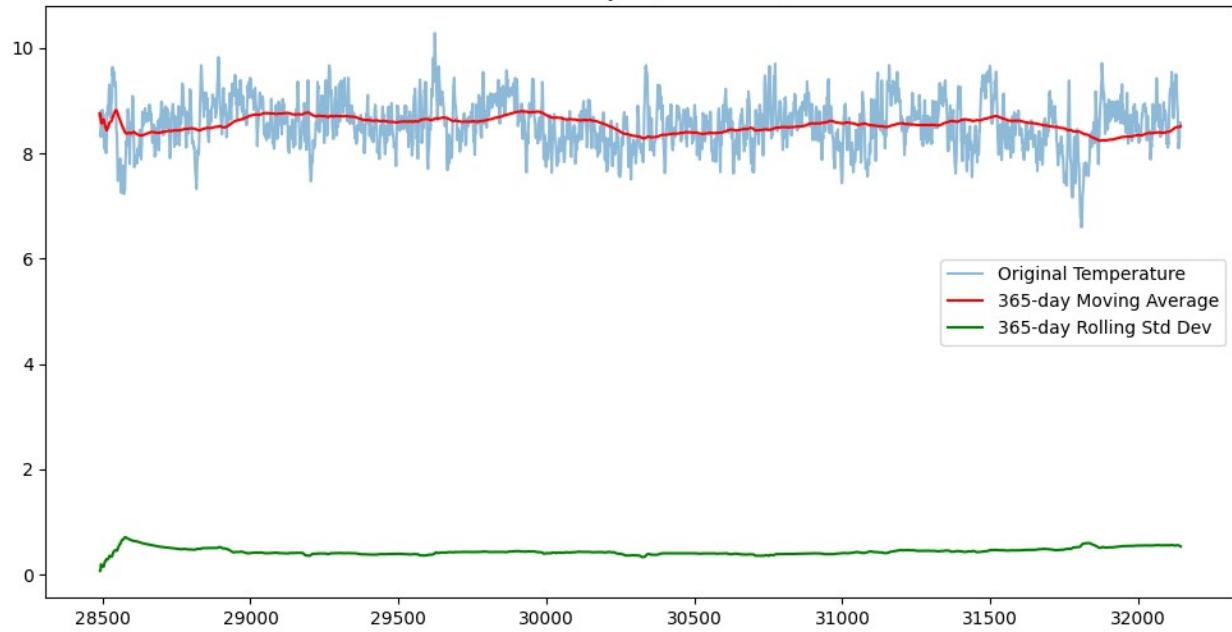


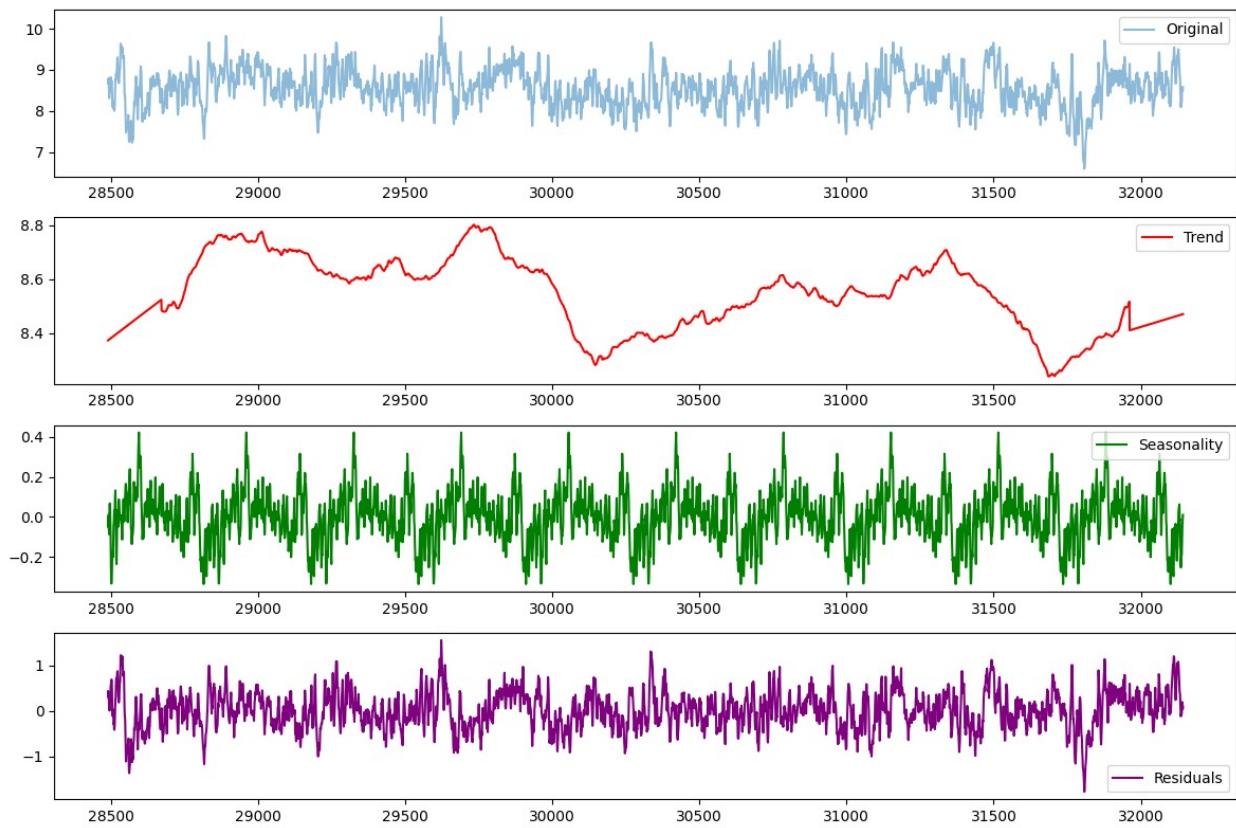
Analyzing period: 1960-1970
Number of Outliers: 33

Boxplot of Temperature (1960-1970)

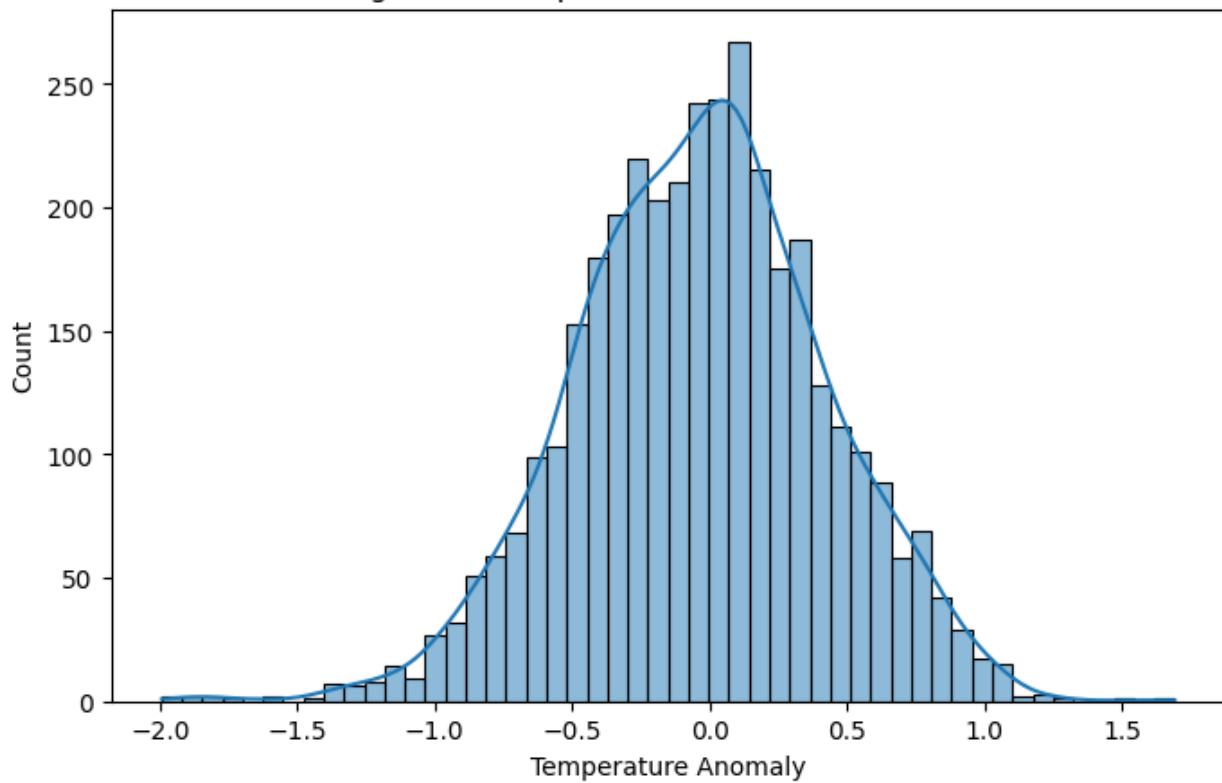


Trend Analysis (1960-1970)

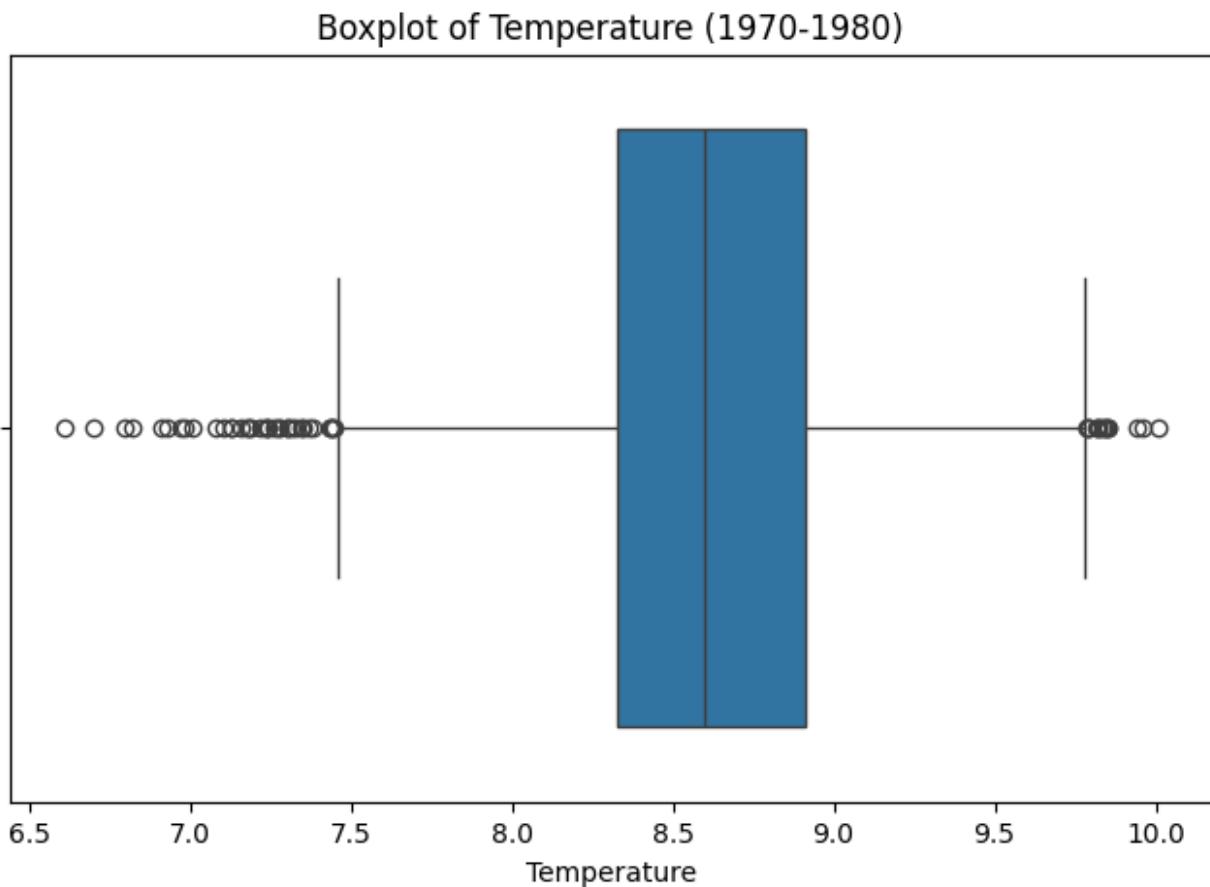




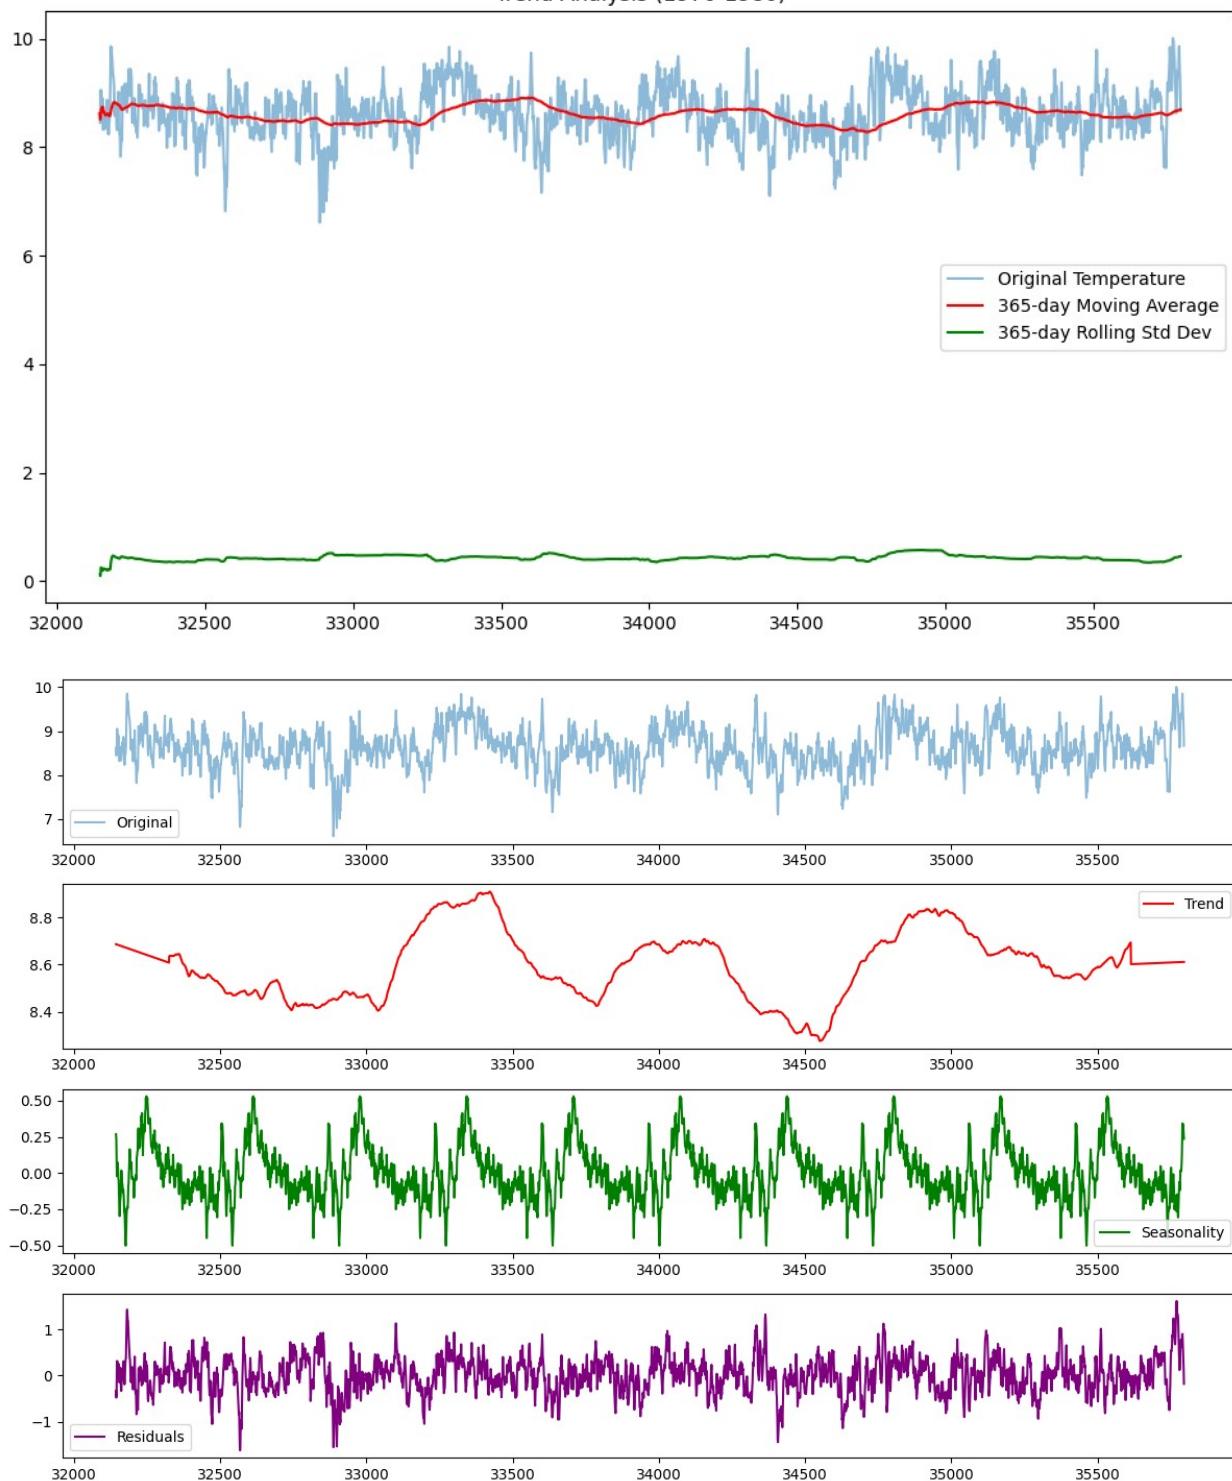
Histogram of Temperature Anomalies (1960-1970)



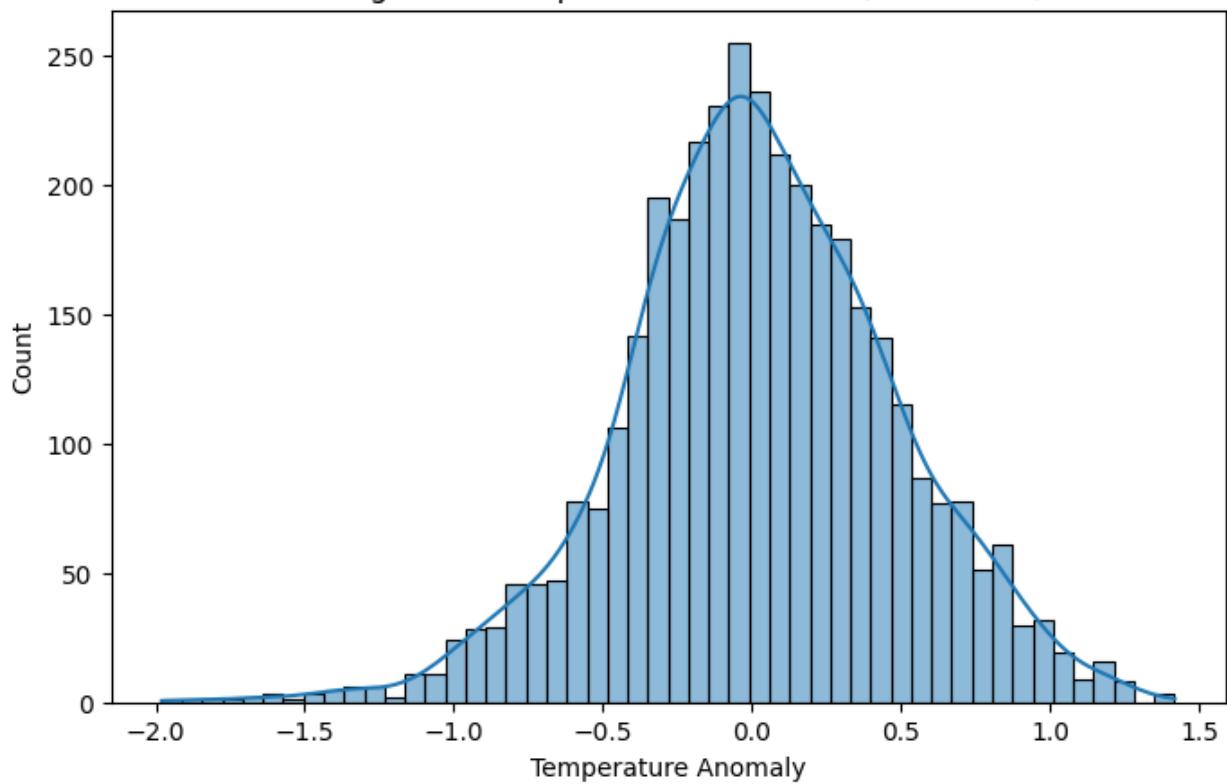
Analyzing period: 1970-1980
Number of Outliers: 50



Trend Analysis (1970-1980)

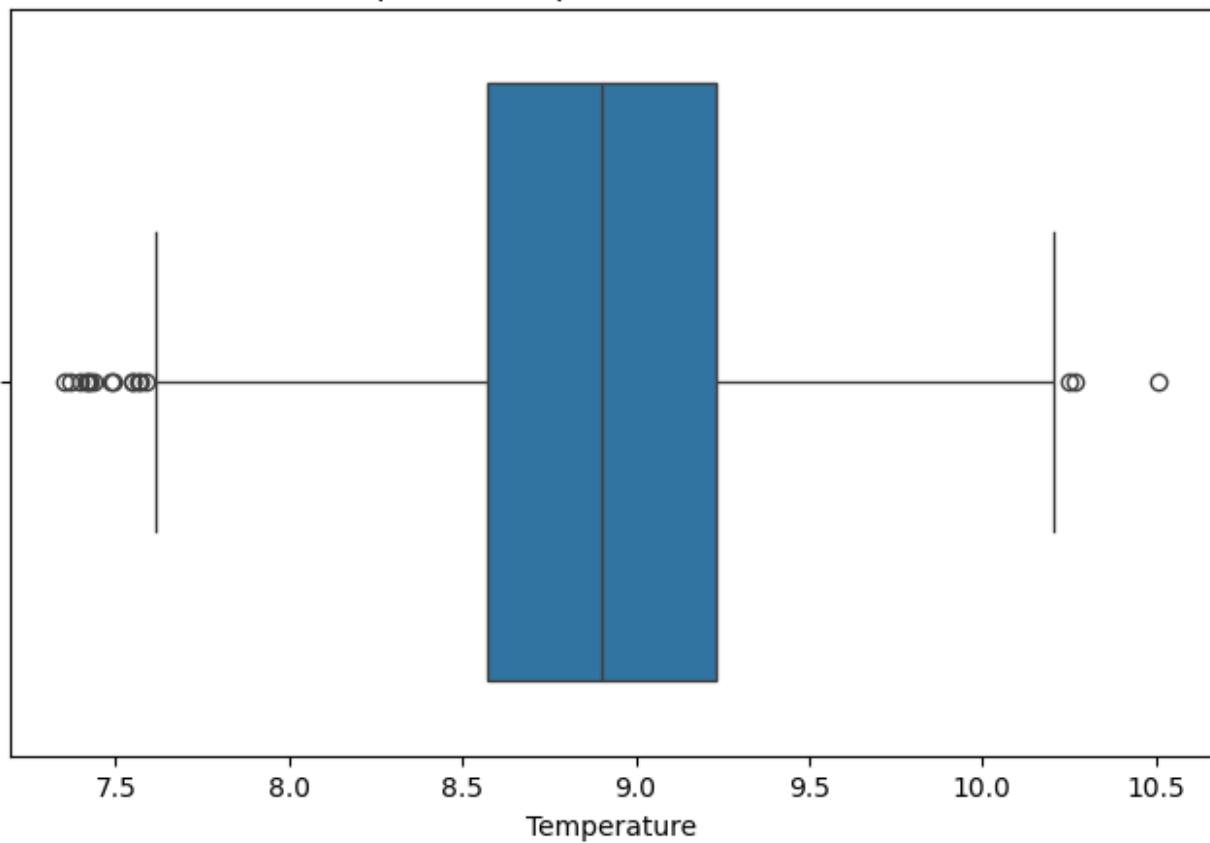


Histogram of Temperature Anomalies (1970-1980)

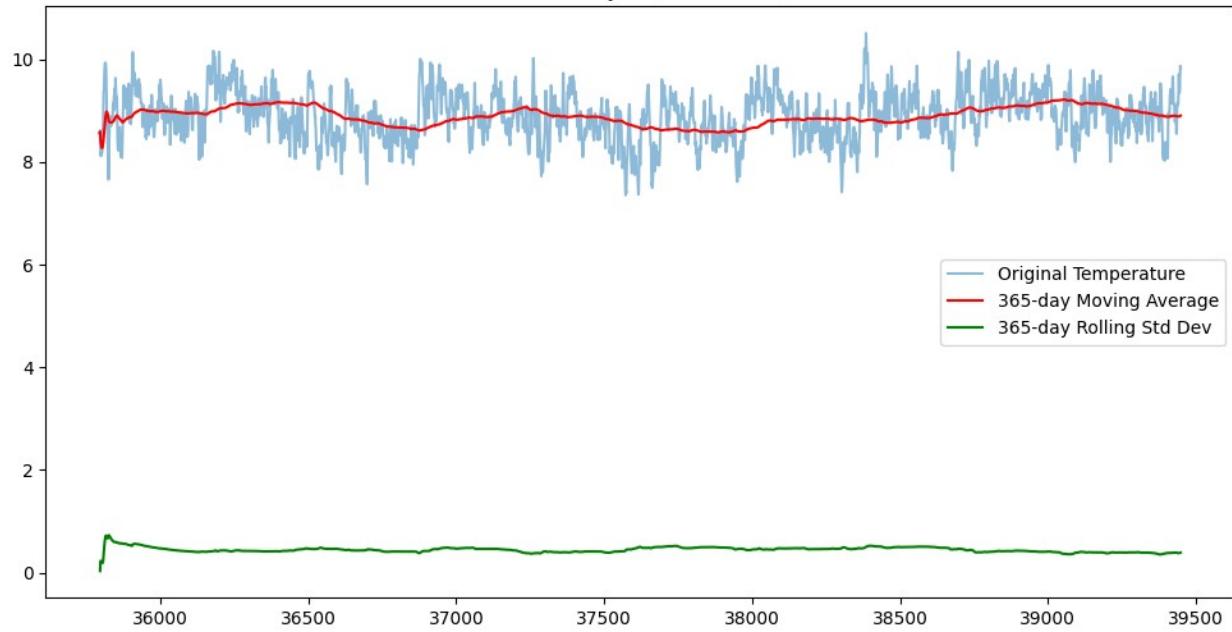


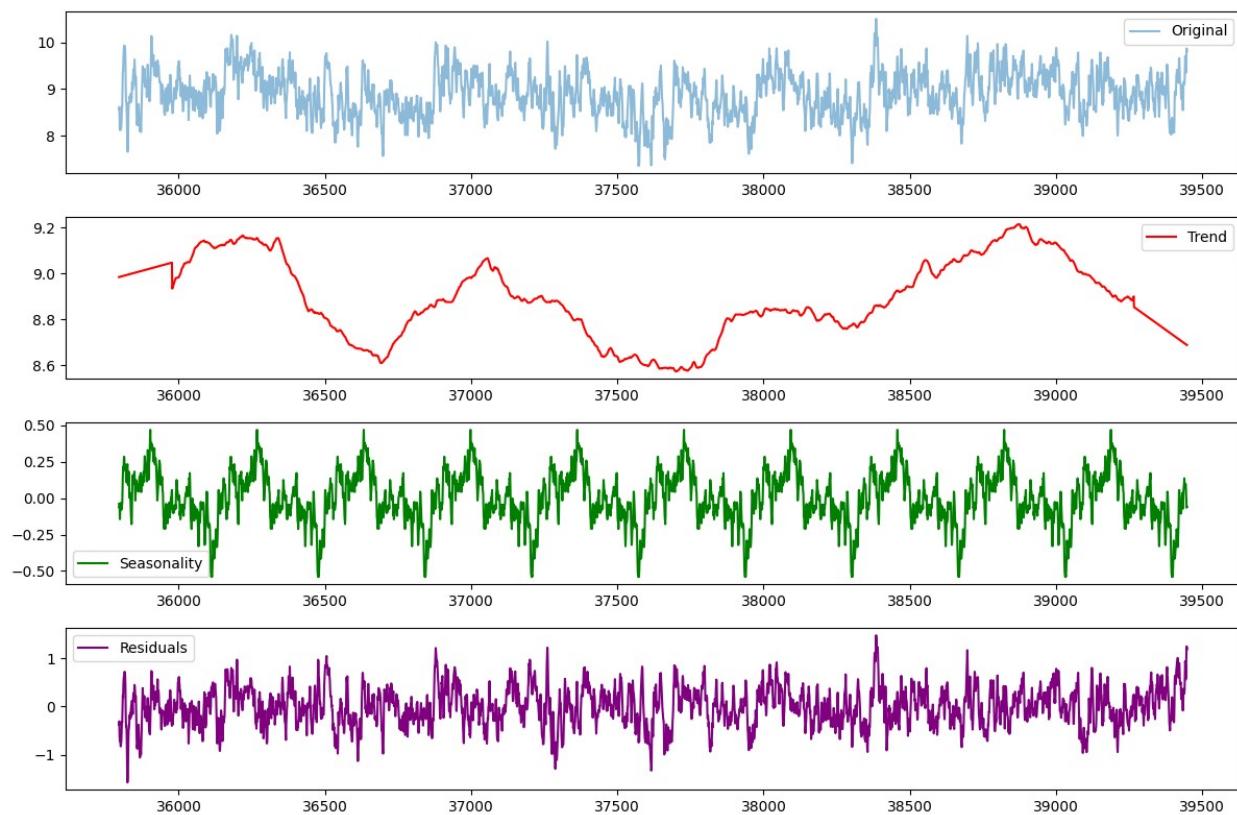
Analyzing period: 1980-1990
Number of Outliers: 17

Boxplot of Temperature (1980-1990)

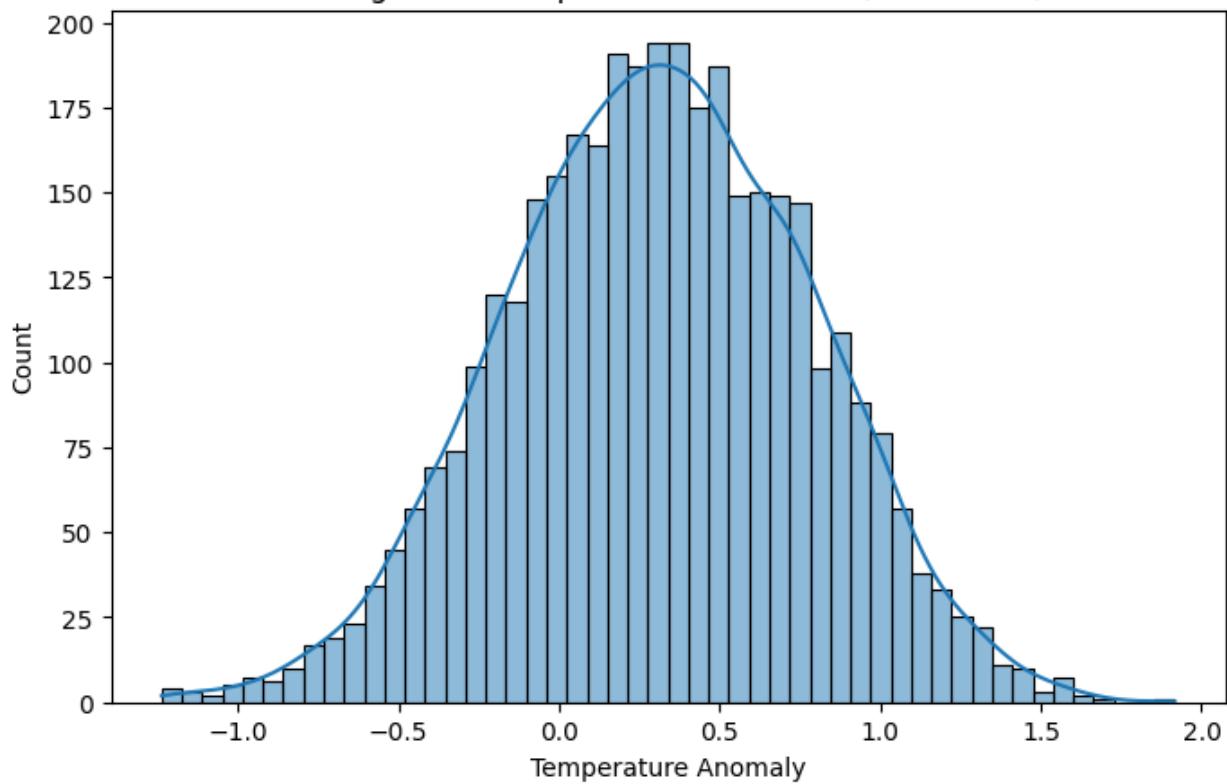


Trend Analysis (1980-1990)



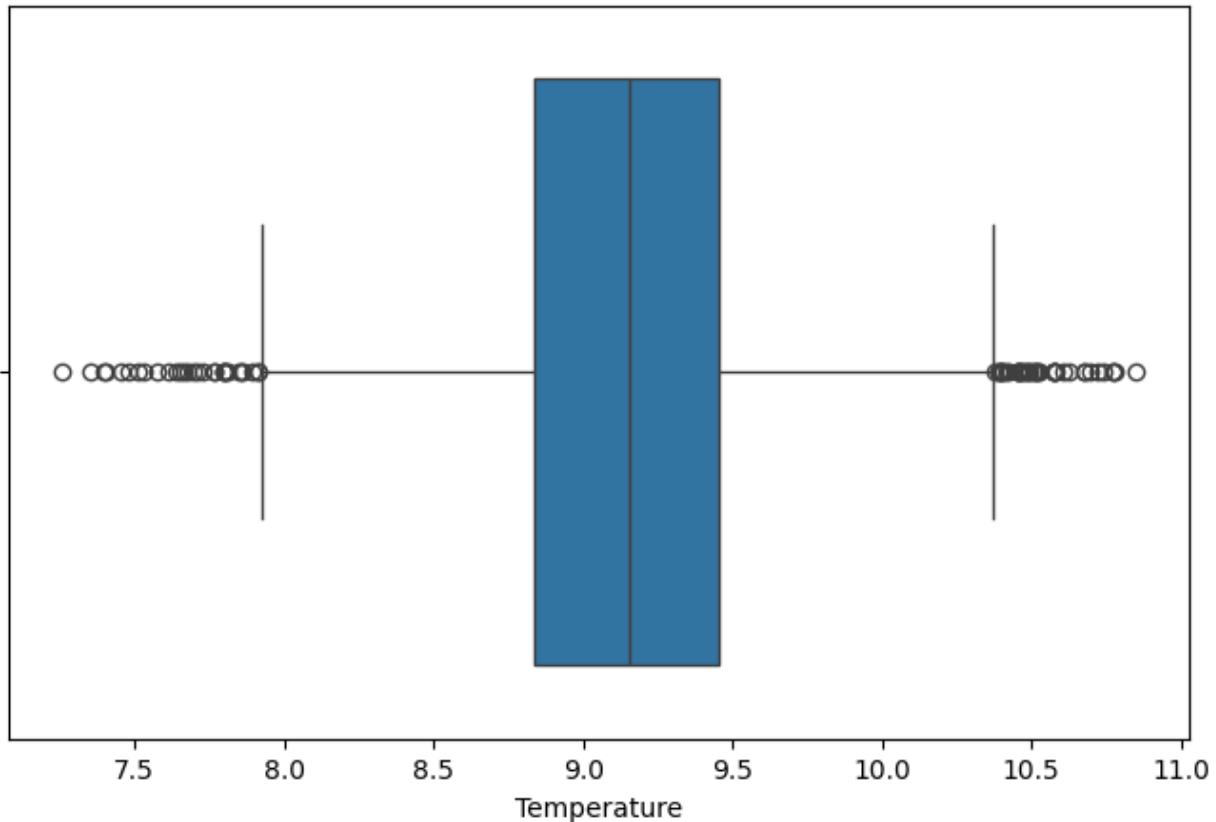


Histogram of Temperature Anomalies (1980-1990)

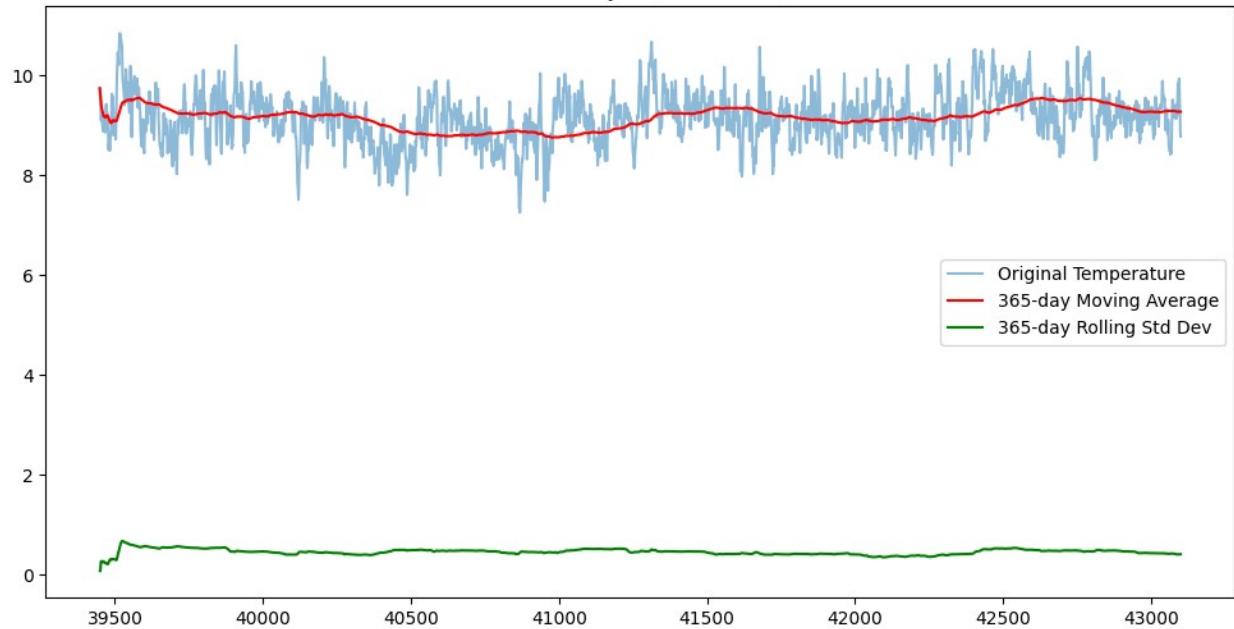


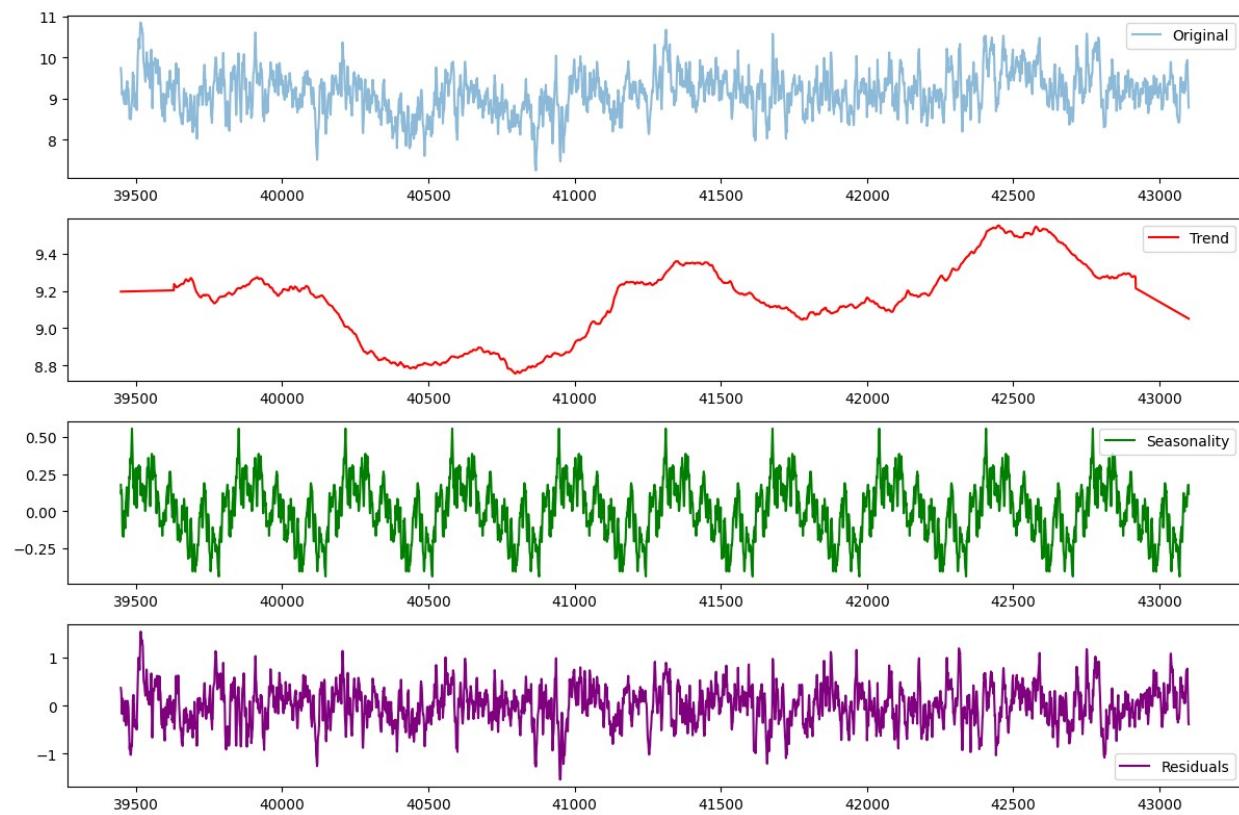
Analyzing period: 1990-2000
Number of Outliers: 72

Boxplot of Temperature (1990-2000)

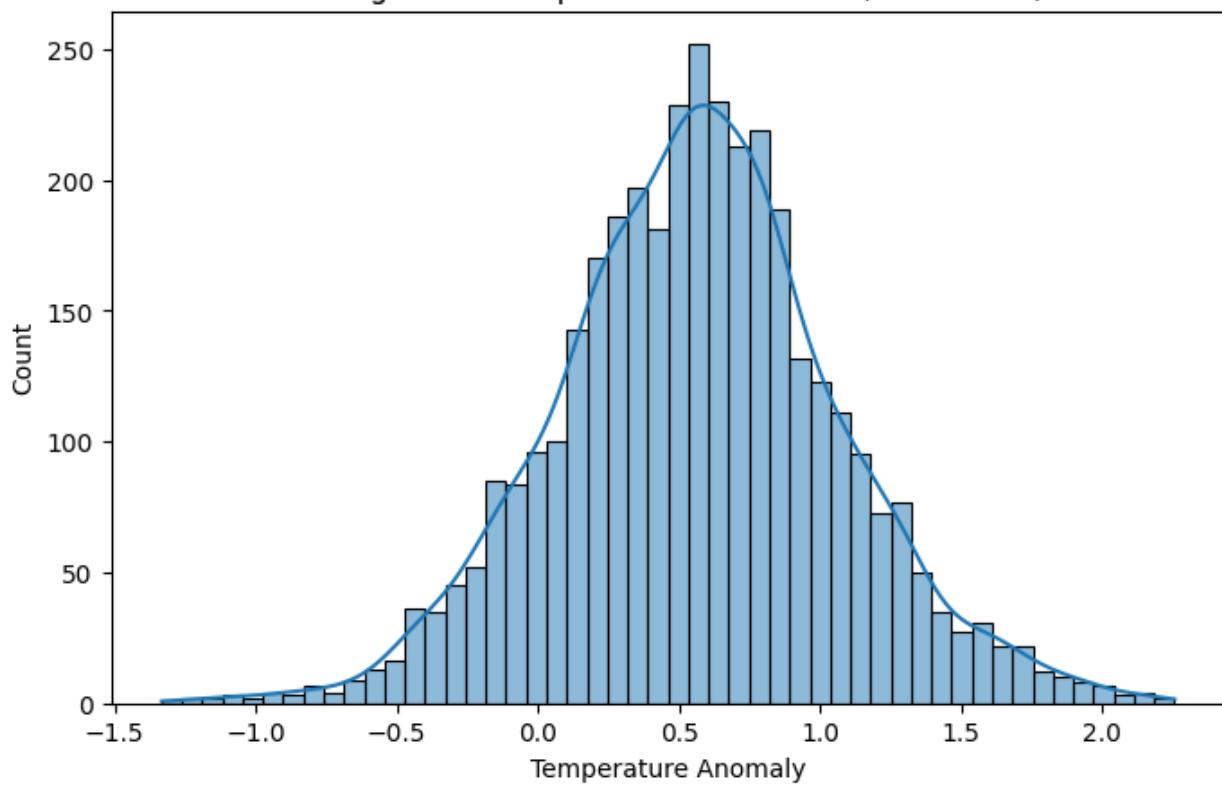


Trend Analysis (1990-2000)



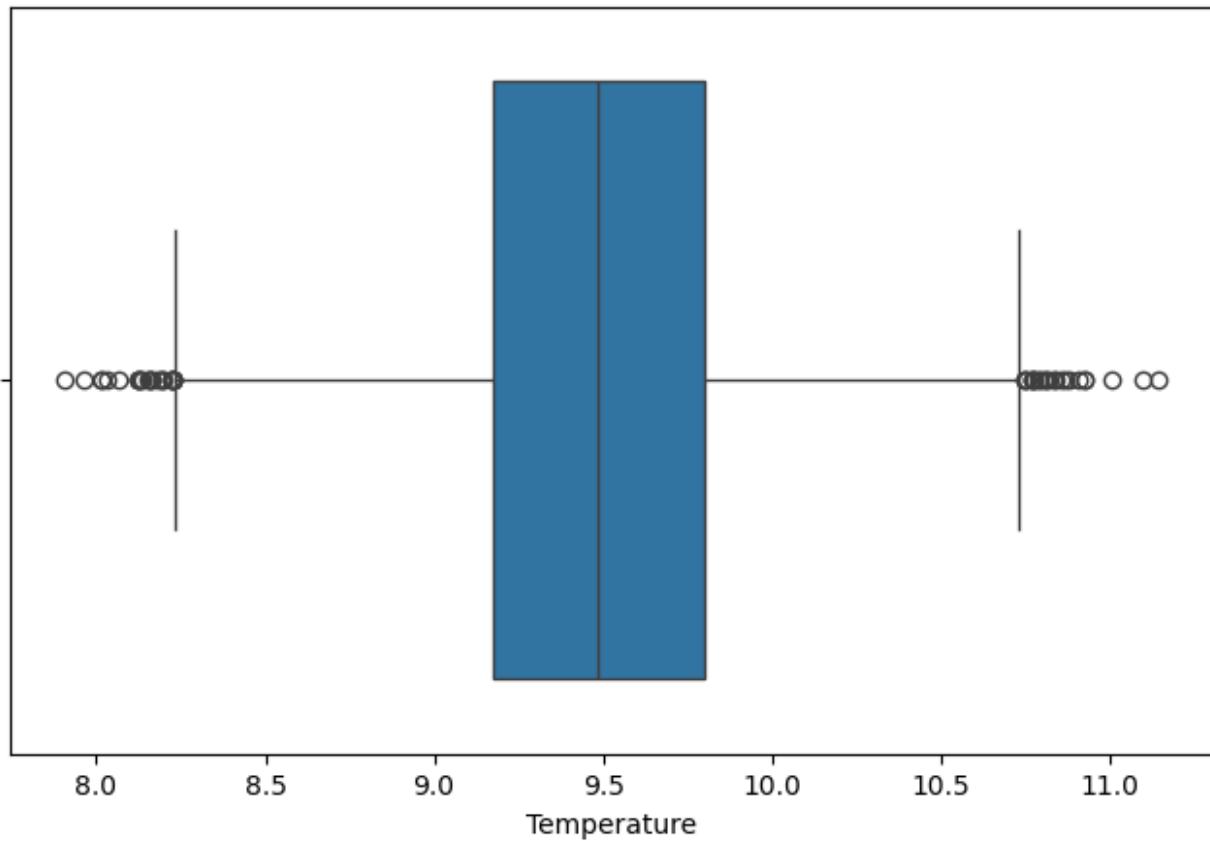


Histogram of Temperature Anomalies (1990-2000)

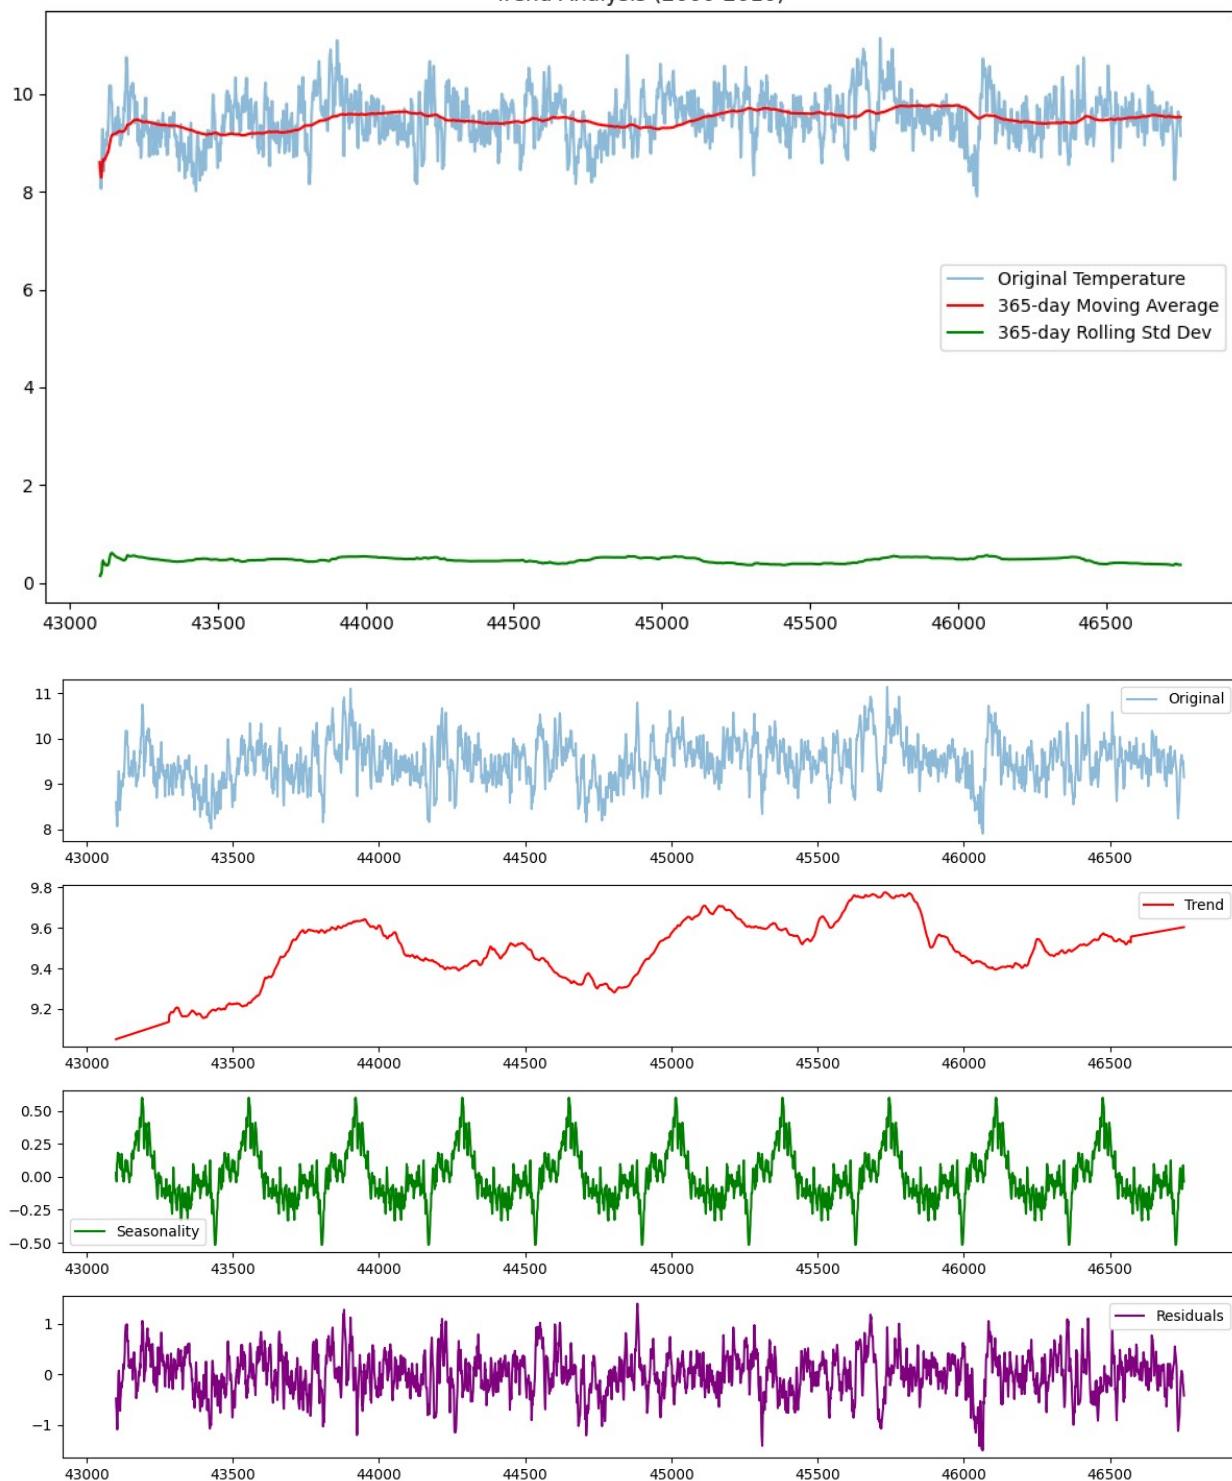


Analyzing period: 2000-2010
Number of Outliers: 48

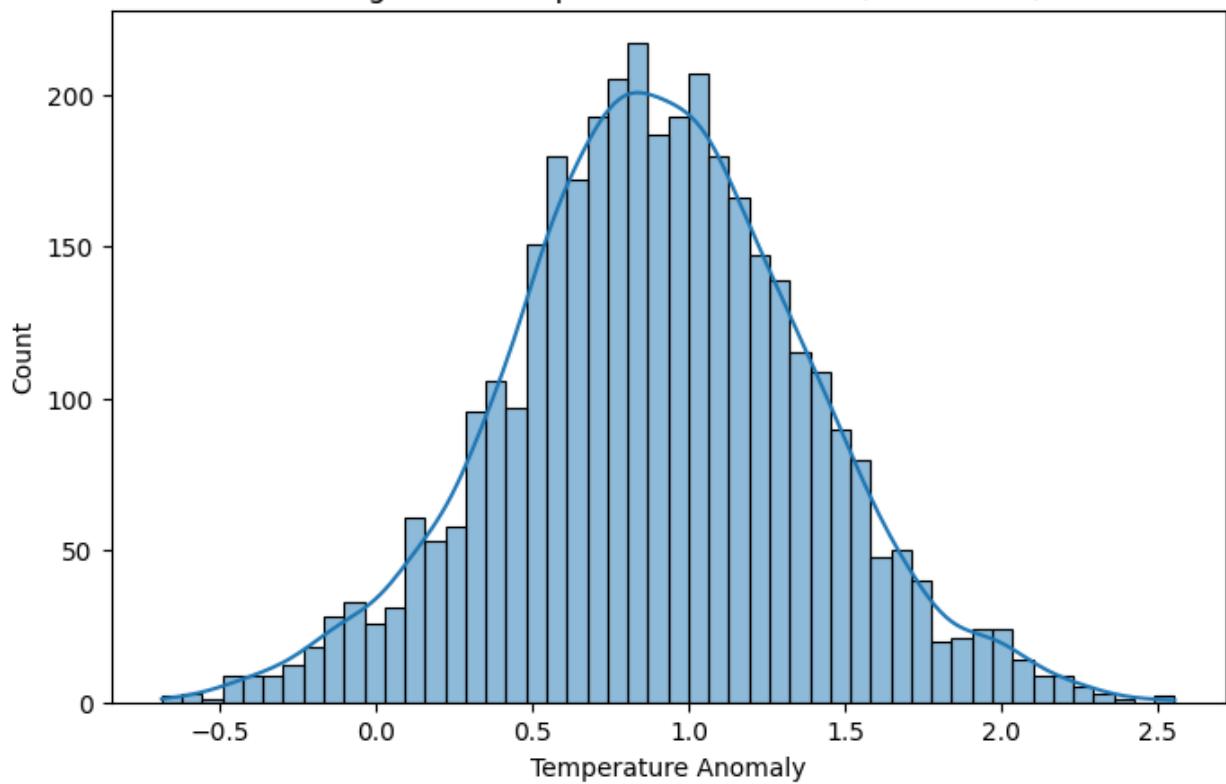
Boxplot of Temperature (2000-2010)



Trend Analysis (2000-2010)

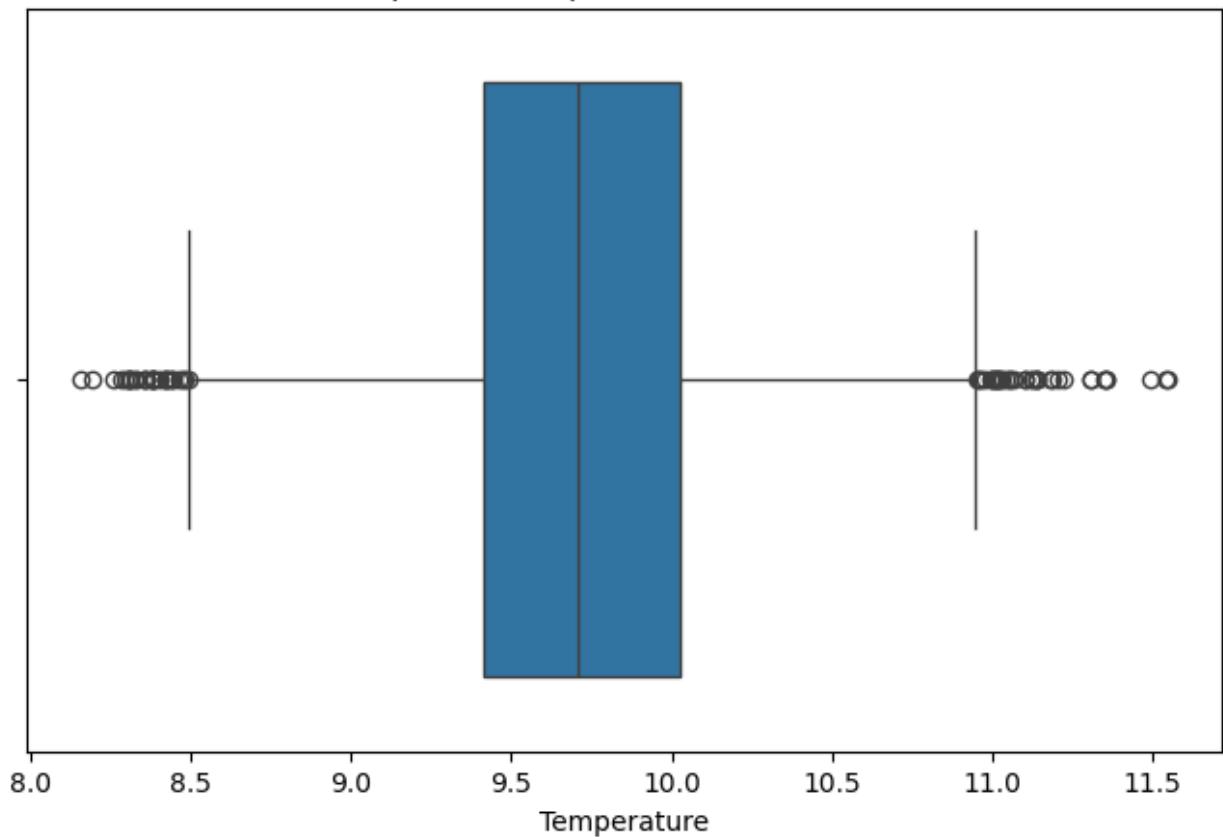


Histogram of Temperature Anomalies (2000-2010)

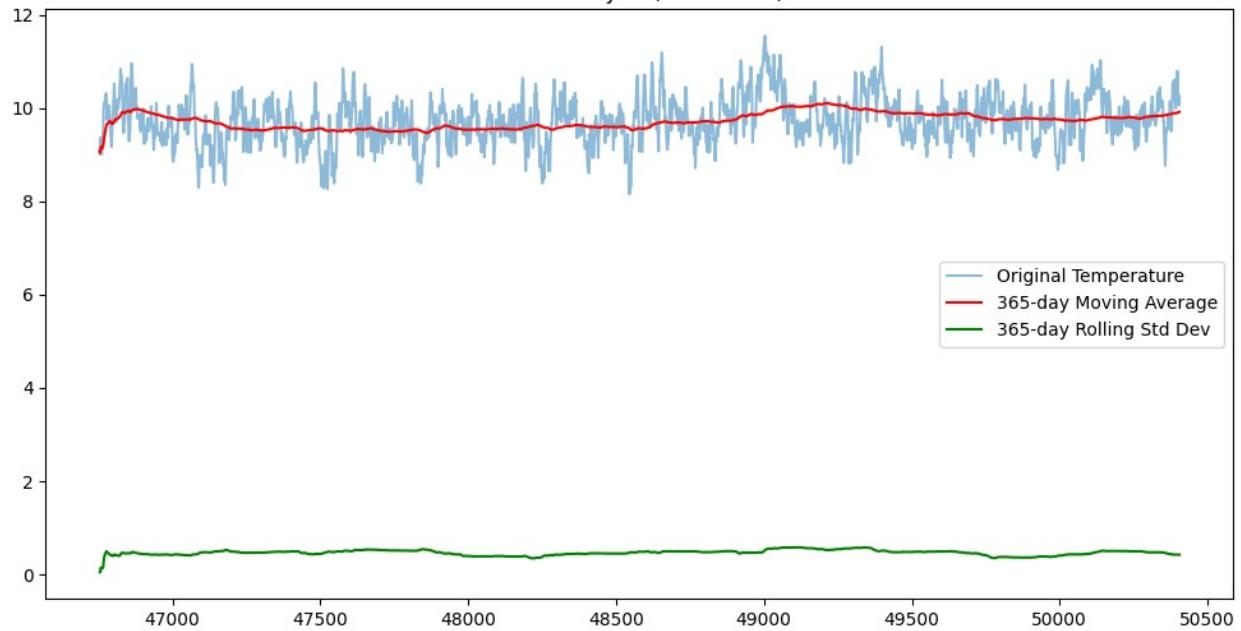


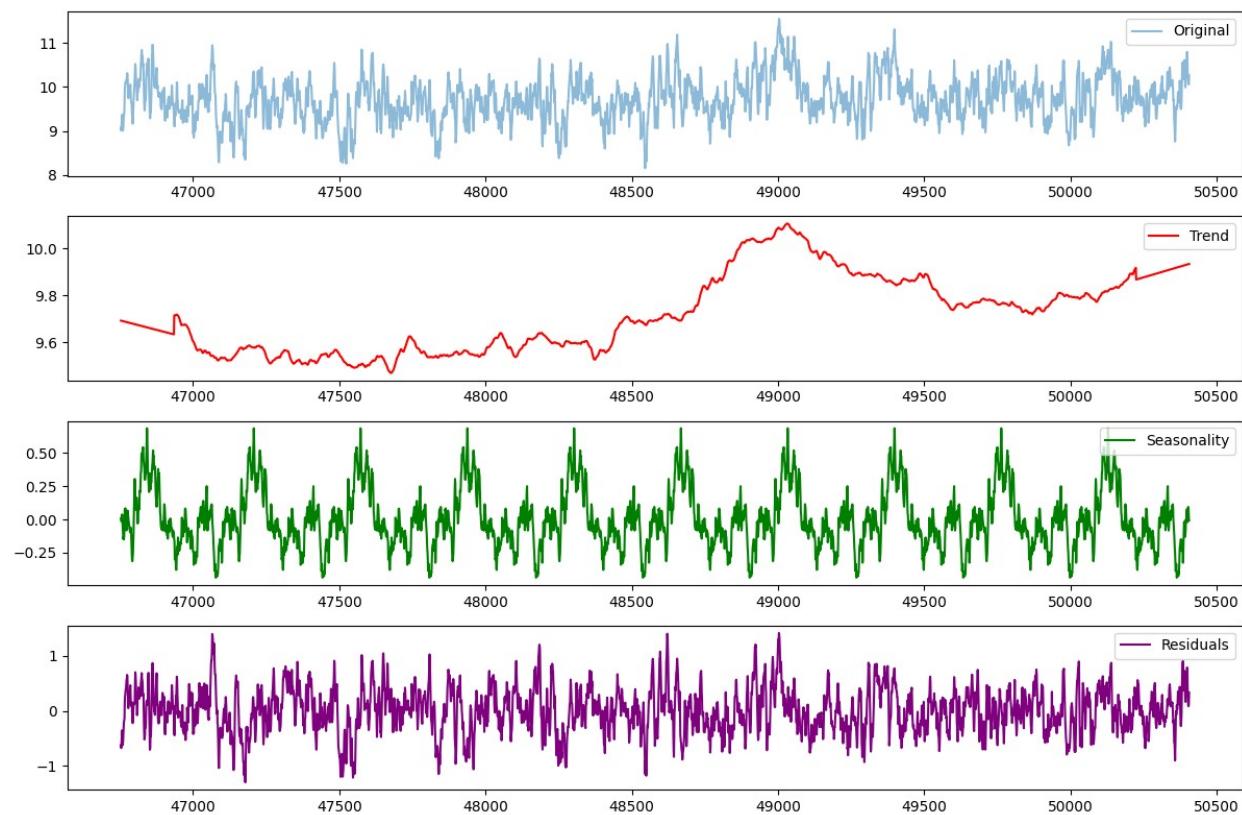
Analyzing period: 2010-2020
Number of Outliers: 67

Boxplot of Temperature (2010-2020)

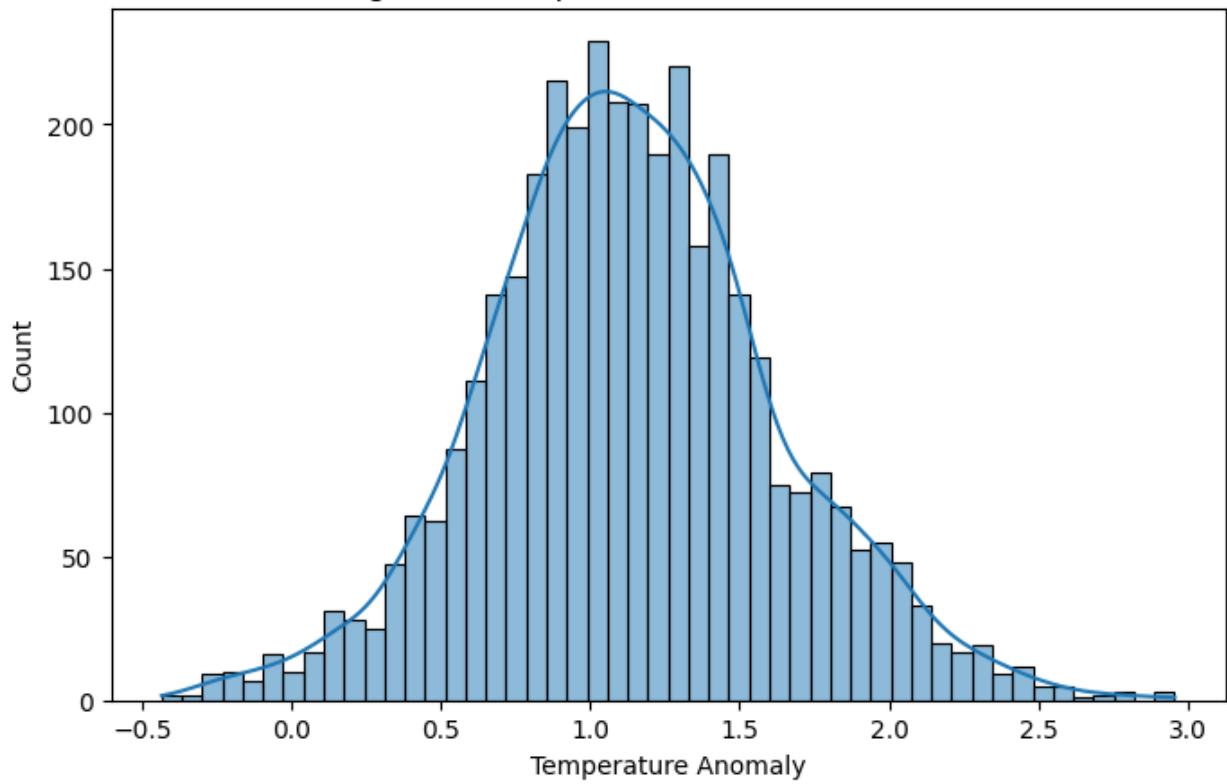


Trend Analysis (2010-2020)





Histogram of Temperature Anomalies (2010-2020)



Mean significance tests

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from scipy.stats import ttest_ind, f_oneway, shapiro
import scipy.stats as stats

# Convert 'Date' column to datetime (modify column name if needed)
if 'Date' in df.columns:
    df['Date'] = pd.to_datetime(df['Date'])
    df.set_index('Date', inplace=True)

# Drop missing values
df.dropna(subset=['Temperature'], inplace=True)
print(df.head())
### ---- 1. Hypothesis Testing ---- ###

# Splitting the data into two periods for t-test
early_period = df[df.index.year < 1950]['Temperature']
recent_period = df[df.index.year >= 2000]['Temperature']

# Independent t-test (assumes normality)
t_stat, p_value = ttest_ind(early_period, recent_period,
                            equal_var=False)
print(f"T-test results: t-statistic={t_stat:.4f}, p-
value={p_value:.4f}")

if p_value < 0.05:
    print("Reject H0: The mean temperature has significantly changed
over time.")
else:
    print("Fail to reject H0: No significant change in mean
temperature.")

# ANOVA test for multiple time periods (dividing data into 50-year
bins)
df['Year'] = df.index.year
df['Period'] = pd.cut(df['Year'], bins=[1880, 1930, 1980, 2020],
                      labels=['1880-1930', '1930-1980', '1980-2020'])

# Perform ANOVA
anova_groups = [df[df['Period'] == period]['Temperature'] for period
in df['Period'].unique()]
anova_stat, anova_p = f_oneway(*anova_groups)
print(f"ANOVA results: F-statistic={anova_stat:.4f}, p-
value={anova_p:.4f}")

if anova_p < 0.05:
```

```

        print("Reject H0: Mean temperatures differ significantly across
time periods.")
else:
    print("Fail to reject H0: No significant difference in mean
temperatures across periods.")

### ---- 2. Fitting a Multivariate Normal Distribution ---- ###

# Check normality assumption (Shapiro-Wilk Test on a sample)
sample = df['Temperature'].sample(500, random_state=42) # Take a
sample for normality check
shapiro_stat, shapiro_p = shapiro(sample)
print(f"Shapiro-Wilk Test: Statistic={shapiro_stat:.4f}, p-
value={shapiro_p:.4f}")

if shapiro_p < 0.05:
    print("Data is not normally distributed (p < 0.05). Consider
transformations.")
else:
    print("Data follows normal distribution (p >= 0.05).")

# QQ Plot for normality check
plt.figure(figsize=(6, 6))
stats.probplot(sample, dist="norm", plot=plt)
plt.title("QQ Plot of Temperature")
plt.show()

# Compute Covariance Matrix
cov_matrix = np.cov(df[['Temperature', 'Anomaly']].dropna().T)
print("Covariance Matrix:\n", cov_matrix)

```

	Year	Month	Day	Day of Year	Anomaly	Temperature	
Rolling_Mean	\						
0	1880	1	1	1	-0.692	7.898	7.898000
1	1880	1	2	2	-0.592	7.998	7.948000
2	1880	1	3	3	-0.673	7.917	7.937667
3	1880	1	4	4	-0.615	7.975	7.947000
4	1880	1	5	5	-0.681	7.909	7.939400

	Rolling_Std
0	NaN
1	NaN
2	NaN
3	NaN
4	NaN

```

-----
AttributeError                               Traceback (most recent call
last)
Cell In[8], line 19
  15 print(df.head())
  16 ### ---- 1. Hypothesis Testing ---- ###
  17
  18 # Splitting the data into two periods for t-test
--> 19 early_period = df[df.index.year < 1950]['Temperature']
  20 recent_period = df[df.index.year >= 2000]['Temperature']
  22 # Independent t-test (assumes normality)

AttributeError: 'RangeIndex' object has no attribute 'year'

```

10 years group study and 35 years group study

```

# Extract Year and compute yearly average temperature
df["Year"] = df.index.year
yearly_avg_temp = df.groupby("Year")[
    "Temperature"].mean().reset_index()

# Define group configurations
group_configs = {
    "10-Year Groups": (14, 10),
    "35-Year Groups": (4, 35),
}

# Dictionary to store all results
all_distribution_stats = {}

# Perform analysis for each group size
for group_name, (num_groups, group_size) in group_configs.items():

    # Create figure for subplots
    fig, axes = plt.subplots(num_groups, 2, figsize=(12, 5 * num_groups))

    # Dictionary to store statistics for the current group
    distribution_stats = {}

    # Perform distribution analysis for each group
    for i in range(num_groups):
        start_index = i * group_size
        end_index = start_index + group_size

        if end_index > len(yearly_avg_temp):
            break # Stop if out of range

```

```

# Extract temperature data for the group
group_data = yearly_avg_temp.iloc[start_index:end_index]
["Temperature"]
start_year = yearly_avg_temp.iloc[start_index]["Year"]
end_year = yearly_avg_temp.iloc[end_index - 1]["Year"]
group_label = f"{start_year}-{end_year}"

# Compute statistical metrics
stats_dict = {
    "Mean": np.mean(group_data),
    "Median": np.median(group_data),
    "Variance": np.var(group_data, ddof=1),
    "Skewness": stats.skew(group_data),
    "Kurtosis": stats.kurtosis(group_data),
}
distribution_stats[group_label] = stats_dict

# Plot Histogram & KDE
sns.histplot(group_data, bins=10, kde=True, ax=axes[i, 0],
color="skyblue")
axes[i, 0].set_title(f"Histogram & KDE: {group_label}")
axes[i, 0].set_xlabel("Temperature")

# Q-Q Plot (Normality Check)
stats.probplot(group_data, dist="norm", plot=axes[i, 1])
axes[i, 1].set_title(f"Q-Q Plot: {group_label}")

# Convert dictionary to DataFrame and store results
distribution_df = pd.DataFrame.from_dict(distribution_stats,
orient="index")
all_distribution_stats[group_name] = distribution_df

# Save statistics to CSV
filename = f"temperature_distribution_{group_name.replace(' ', '_').lower()}.csv"
distribution_df.to_csv(filename)

# Show plots
plt.tight_layout()
plt.show()

# Print DataFrame
print(f"\n Distribution Statistics for {group_name}:")
print(distribution_df)
print("\n" + "=" * 80 + "\n")

df["Year"] = df.index.year
yearly_avg_temp = df.groupby("Year")
["Temperature"].mean().reset_index()

```

```

# Define the number of groups and group size
num_groups = 14
group_size = 10

# Plot KDE curves for each group
plt.figure(figsize=(12, 6))

for i in range(num_groups):
    start_index = i * group_size
    end_index = start_index + group_size

    if end_index > len(yearly_avg_temp):
        break # Stop if out of range

    # Extract temperature data for the group
    group_data = yearly_avg_temp.iloc[start_index:end_index]
    ["Temperature"]
    start_year = yearly_avg_temp.iloc[start_index]["Year"]
    end_year = yearly_avg_temp.iloc[end_index - 1]["Year"]

    # Plot KDE for the group
    sns.kdeplot(group_data, label=f"{start_year}-{end_year}",
    linewidth=2)

# Add labels and title
plt.xlabel("Temperature", fontsize=12)
plt.ylabel("Density", fontsize=12)
plt.title("Kernel Density Estimation (KDE) of Temperature Distributions (14 Groups of 10 Years)", fontsize=14)
plt.legend(title="Year Ranges", bbox_to_anchor=(1.05, 1), loc='upper left')
plt.grid(True)

# Show plot
plt.show()

```

20 years group study

```

import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
import scipy.stats as stats

# Extract Year, Month, and Day
df["Year"] = df.index.year
df["Month"] = df.index.month

```

```

df["Day"] = df.index.day

# Define the number of groups and group size
num_groups = 7
group_size = 20

# Create an empty dictionary to store statistics
distribution_stats = {}

# Create plots
fig, axes = plt.subplots(num_groups, 3, figsize=(18, 5 * num_groups))
# Added an extra column for log-normality

# Perform distribution analysis for each group
for i in range(num_groups):
    start_index = i * group_size
    end_index = start_index + group_size

    if end_index > len(df):
        break # Stop if out of range

    # Extract daily temperature data for the group
    group_data = df.iloc[start_index:end_index]["Temperature"]
    group_label = f"Group {i+1} ({df.iloc[start_index]['Year']}-{df.iloc[start_index]['Month']}-{df.iloc[start_index]['Day']} to {df.iloc[end_index-1]['Year']}-{df.iloc[end_index-1]['Month']}-{df.iloc[end_index-1]['Day']})"

    # Compute descriptive statistics
    stats_dict = {
        "Mean": np.mean(group_data),
        "Median": np.median(group_data),
        "Variance": np.var(group_data, ddof=1),
        "Skewness": stats.skew(group_data),
        "Kurtosis": stats.kurtosis(group_data),
    }

    # **Normality Test** (Shapiro-Wilk)
    shapiro_test = stats.shapiro(group_data)
    stats_dict["Normality (Shapiro-Wilk) p-value"] =
shapiro_test.pvalue

    # **Log-Normality Test** (Kolmogorov-Smirnov Test)
    log_data = np.log(group_data[group_data > 0]) # Take log, avoid log(0) issues
    ks_test = stats.kstest(log_data, 'norm', args=(np.mean(log_data),
np.std(log_data)))
    stats_dict["Log-Normality (KS test) p-value"] = ks_test.pvalue

```

```

distribution_stats[group_label] = stats_dict

# **Histogram & KDE plot**
sns.histplot(group_data, bins=10, kde=True, ax=axes[i, 0],
color="skyblue")
axes[i, 0].set_title(f"Histogram & KDE: {group_label}")
axes[i, 0].set_xlabel("Temperature")

# **Q-Q Plot for Normality**
stats.probplot(group_data, dist="norm", plot=axes[i, 1])
axes[i, 1].set_title(f"Q-Q Plot: {group_label}")

# **Q-Q Plot for Log-Normality**
stats.probplot(log_data, dist="norm", plot=axes[i, 2])
axes[i, 2].set_title(f"Log-Normal Q-Q Plot: {group_label}")

# Save distribution statistics to a DataFrame
distribution_df = pd.DataFrame.from_dict(distribution_stats,
orient="index")

# Show plots
plt.tight_layout()
plt.show()

# Save statistics to CSV
distribution_df.to_csv("daily_temperature_distribution_study_with_norm
ality.csv")

# Print results
print(distribution_df)

# Define the number of groups and group size
num_groups = 7
group_size = 20

# Plot KDE curves for each group
plt.figure(figsize=(12, 6))

for i in range(num_groups):
    start_index = i * group_size
    end_index = start_index + group_size

    if end_index > len(yearly_avg_temp):
        break # Stop if out of range

    # Extract temperature data for the group
    group_data = yearly_avg_temp.iloc[start_index:end_index]
    [ "Temperature" ]
    start_year = yearly_avg_temp.iloc[start_index][ "Year" ]
    end_year = yearly_avg_temp.iloc[end_index - 1][ "Year" ]

```

```

# Plot KDE for the group
sns.kdeplot(group_data, label=f"{start_year}-{end_year}",
linewidth=2)

# Add labels and title
plt.xlabel("Temperature", fontsize=12)
plt.ylabel("Density", fontsize=12)
plt.title("Kernel Density Estimation (KDE) of Temperature Distributions (7 Groups of 20 Years)", fontsize=14)
plt.legend(title="Year Ranges")
plt.grid(True)

# Show plot
plt.show()

```

Growth Rate calculations

Annual growth rate and n- yearly growth rates (n=10,20,35,70)

```

import pandas as pd
import numpy as np

# Load data (Assuming df is already loaded and indexed by Date)

# Extract Year and calculate yearly average temperature
df["Year"] = df.index.year
yearly_avg_temp = df.groupby("Year")[
    "Temperature"].mean().reset_index()

# Function to calculate CAGR using two formulas
def calculate_cagr1(initial, final, years):
    return ((final / initial) ** (1 / years)) - 1 if years > 0 else
np.nan

def calculate_cagr2(initial, final, years):
    return ((final / initial)) - 1 if years > 0 else np.nan

# Define the group sizes and names
group_sizes = [10, 20, 35, 70] # Groups of 10, 20, 35, and 70 years
group_labels = ["10-Year Growth", "20-Year Growth", "35-Year Growth",
"70-Year Growth"]

all_growth_data = {}

```

```

# Loop through each group size and compute CAGR using both formulas
for group_size, label in zip(group_sizes, group_labels):
    num_groups = len(yearly_avg_temp) // group_size # Number of full groups
    growth_list1, growth_list2 = [], []

    for i in range(num_groups):
        start_index = i * group_size
        end_index = start_index + group_size - 1

        if end_index < len(yearly_avg_temp):
            temp_initial = yearly_avg_temp.iloc[start_index]
            temp_final = yearly_avg_temp.iloc[end_index]
        else:
            break # Stop if the range goes out of bounds

    # Store results
    all_growth_data[f"{label} - CAGR1"] = growth_list1
    all_growth_data[f"{label} - CAGR2"] = growth_list2

# Find the maximum length among the lists
max_length = max(map(len, all_growth_data.values()))

# Pad shorter lists with NaN to match the maximum length
for key in all_growth_data:
    all_growth_data[key] += [np.nan] * (max_length - len(all_growth_data[key]))

# Convert to DataFrame
growth_df = pd.DataFrame.from_dict(all_growth_data)

growth_df.index = [f"Group {i+1}" for i in range(len(growth_df))]

# Save to CSV
growth_df.to_csv("temperature1_growth_analysis.csv")

# Display results
print(growth_df)

```

Time Series Modelling

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from statsmodels.tsa.stattools import adfuller
from statsmodels.graphics.tsaplots import plot_acf, plot_pacf
from statsmodels.tsa.arima.model import ARIMA
from sklearn.metrics import mean_absolute_error, mean_squared_error
from pmдарима import auto_arima

# Exploratory Analysis
plt.figure(figsize=(14,6))
plt.plot(df['Temperature'], label='Daily Temperature')
plt.title('Daily Global Temperature Over Time')
plt.xlabel('Index')
plt.ylabel('Temperature')
plt.legend()
plt.savefig("daily_global_temperature.png") # Save with a meaningful filename
plt.close()

# ADF Test for Stationarity
def adf_test(series):
    result = adfuller(series.dropna())
    print(f"ADF Statistic: {result[0]}")
    print(f"p-value: {result[1]}")
    if result[1] > 0.05:
        print("Non-stationary series (p > 0.05)")
    else:
        print("Stationary series (p <= 0.05)")

print("ADF Test (Original Series):")
adf_test(df['Temperature'])

# Differencing
df['Temperature_Diff'] = df['Temperature'].diff()

print("\nADF Test (After Differencing):")
adf_test(df['Temperature_Diff'])

# ACF and PACF
fig, axes = plt.subplots(1, 2, figsize=(14,5))
plot_acf(df['Temperature_Diff'].dropna(), ax=axes[0])
plot_pacf(df['Temperature_Diff'].dropna(), ax=axes[1])
plt.savefig("acf_pacf_analysis.png") # Save with an appropriate filename
plt.close()
```

```

# Automatically select best ARIMA order
auto_model = auto_arima(
    df['Temperature'],
    seasonal=False,
    stepwise=True,
    suppress_warnings=True,
    error_action='ignore',
    trace=True,
    max_p=6,
    max_q=6
)

print(f"\nBest ARIMA Order Found: {auto_model.order}")

# Fit ARIMA Model
model = ARIMA(df['Temperature'], order=auto_model.order)
arima_result = model.fit()
print(arima_result.summary())

# Forecast next 365 days
forecast = arima_result.forecast(steps=365)

# Plot
plt.figure(figsize=(14,6))
plt.plot(df['Temperature'][-365:], label='Actual')
plt.plot(np.arange(len(df), len(df) + 365), forecast,
label='Forecast', linestyle='dashed')
plt.title('ARIMA Forecast vs Actual (Last 365 Days)')
plt.xlabel('Index')
plt.ylabel('Temperature')
plt.legend()
plt.savefig("arima_forecast_vs_actual.png") # Save with an
appropriate filename
plt.close()

# Evaluation
y_actual = df['Temperature'][-365:]
y_pred = forecast[:365]
mae = mean_absolute_error(y_actual, y_pred)
rmse = np.sqrt(mean_squared_error(y_actual, y_pred))
print(f"\nARIMA Performance on Full Dataset:")
print(f"MAE: {mae:.4f}")
print(f"RMSE: {rmse:.4f}")

ADF Test (Original Series):
ADF Statistic: -9.437660805688283
p-value: 4.991882507815218e-16
Stationary series (p <= 0.05)

```

```

ADF Test (After Differencing):
ADF Statistic: -43.762296113685935
p-value: 0.0
Stationary series (p <= 0.05)
Performing stepwise search to minimize aic
ARIMA(2,1,2)(0,0,0)[0] intercept : AIC=-57163.089, Time=29.44 sec
ARIMA(0,1,0)(0,0,0)[0] intercept : AIC=-48399.827, Time=2.20 sec
ARIMA(1,1,0)(0,0,0)[0] intercept : AIC=-53767.450, Time=1.42 sec
ARIMA(0,1,1)(0,0,0)[0] intercept : AIC=-54586.133, Time=3.08 sec
ARIMA(0,1,0)(0,0,0)[0] intercept : AIC=-48401.822, Time=1.34 sec
ARIMA(1,1,2)(0,0,0)[0] intercept : AIC=-54582.434, Time=8.30 sec
ARIMA(2,1,1)(0,0,0)[0] intercept : AIC=-56872.395, Time=29.94 sec
ARIMA(3,1,2)(0,0,0)[0] intercept : AIC=inf, Time=33.32 sec
ARIMA(2,1,3)(0,0,0)[0] intercept : AIC=inf, Time=38.43 sec
ARIMA(1,1,1)(0,0,0)[0] intercept : AIC=-54594.783, Time=3.82 sec
ARIMA(1,1,3)(0,0,0)[0] intercept : AIC=-57204.109, Time=32.06 sec
ARIMA(0,1,3)(0,0,0)[0] intercept : AIC=-54821.407, Time=8.26 sec
ARIMA(1,1,4)(0,0,0)[0] intercept : AIC=-57172.578, Time=39.37 sec
ARIMA(0,1,2)(0,0,0)[0] intercept : AIC=-54599.336, Time=6.09 sec
ARIMA(0,1,4)(0,0,0)[0] intercept : AIC=-55238.337, Time=18.25 sec
ARIMA(2,1,4)(0,0,0)[0] intercept : AIC=-57161.968, Time=45.42 sec
ARIMA(1,1,3)(0,0,0)[0] intercept : AIC=-57205.935, Time=13.19 sec
ARIMA(0,1,3)(0,0,0)[0] intercept : AIC=-54823.403, Time=4.57 sec
ARIMA(1,1,2)(0,0,0)[0] intercept : AIC=-54632.880, Time=10.79 sec
ARIMA(2,1,3)(0,0,0)[0] intercept : AIC=-57115.991, Time=21.29 sec
ARIMA(1,1,4)(0,0,0)[0] intercept : AIC=-57206.792, Time=15.39 sec
ARIMA(0,1,4)(0,0,0)[0] intercept : AIC=-55240.332, Time=4.13 sec
ARIMA(2,1,4)(0,0,0)[0] intercept : AIC=-57204.926, Time=15.86 sec
ARIMA(1,1,5)(0,0,0)[0] intercept : AIC=inf, Time=18.30 sec
ARIMA(0,1,5)(0,0,0)[0] intercept : AIC=-55650.498, Time=5.85 sec
ARIMA(2,1,5)(0,0,0)[0] intercept : AIC=-57202.826, Time=25.41 sec

```

Best model: ARIMA(1,1,4)(0,0,0)[0]

Total fit time: 435.555 seconds

Best ARIMA Order Found: (1, 1, 4)
SARIMAX Results

```

=====
=====
Dep. Variable: Temperature No. Observations: 50985
Model: ARIMA(1, 1, 4) Log Likelihood: 28609.396
Date: Tue, 08 Apr 2025 AIC: -
57206.792
Time: 15:58:53 BIC: -
57153.756
Sample: 0 HQIC: -
57190.194

```

```

- 50985

Covariance Type: opg
=====

=====
```

	coef	std err	z	P> z	[0.025
0.975]					
-----	-----	-----	-----	-----	-----
ar.L1	0.8650	0.003	256.401	0.000	0.858
0.872					
ma.L1	-0.5537	0.005	-105.126	0.000	-0.564
-0.543					
ma.L2	-0.3358	0.005	-69.416	0.000	-0.345
-0.326					
ma.L3	-0.0832	0.005	-16.923	0.000	-0.093
-0.074					
ma.L4	-0.0092	0.005	-1.920	0.055	-0.019
0.000					
sigma2	0.0191	0.000	175.469	0.000	0.019
0.019					

```

=====
```

Ljung-Box (L1) (Q):	0.00	Jarque-Bera (JB):
374.89		
Prob(Q):	0.98	Prob(JB):
0.00		
Heteroskedasticity (H):	1.07	Skew:
-0.03		
Prob(H) (two-sided):	0.00	Kurtosis:
3.42		

```

=====
```

Warnings:
[1] Covariance matrix calculated using the outer product of gradients (complex-step).

ARIMA Performance on Full Dataset:

MAE: 0.3041

RMSE: 0.3773

```

import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from statsmodels.tsa.stattools import adfuller
from statsmodels.graphics.tsaplots import plot_acf, plot_pacf
from sklearn.metrics import mean_absolute_error, mean_squared_error
```

```

from pmdarima import auto_arima

group_size = 20
start_year = 1880
end_year = 2020

def adf_test(series):
    result = adfuller(series.dropna())
    print(f"ADF Statistic: {result[0]}")
    print(f"p-value: {result[1]}")
    if result[1] > 0.05:
        print("Non-stationary series (p > 0.05)")
    else:
        print("Stationary series (p <= 0.05)")

results_list = []

for start in range(start_year, end_year, group_size):
    df_group = df[(df['Year'] >= start) & (df['Year'] < start + group_size)].copy()
    df_group.reset_index(drop=True, inplace=True)

    print(f"\nAnalyzing period: {start}-{start + group_size}")

    # Plot original temperature
    plt.figure(figsize=(14,6))
    plt.plot(df_group['Temperature'], label='Daily Temperature')
    plt.title(f'Daily Global Temperature ({start}-{start + group_size})')
    plt.xlabel('Day Index')
    plt.ylabel('Temperature')
    plt.legend()
    plt.show()

    # ADF Test
    adf_test(df_group['Temperature'])

    # Differencing
    df_group['Temperature_Diff'] = df_group['Temperature'].diff()
    adf_test(df_group['Temperature_Diff'].dropna())

    # ACF and PACF
    fig, axes = plt.subplots(1, 2, figsize=(14,5))
    plot_acf(df_group['Temperature_Diff'].dropna(), ax=axes[0])
    plot_pacf(df_group['Temperature_Diff'].dropna(), ax=axes[1])
    plt.show()

    # Auto ARIMA Order Selection
    auto_model = auto_arima(
        df_group['Temperature'],

```

```

        seasonal=False,
        stepwise=True,
        suppress_warnings=True,
        error_action='ignore',
        trace=True,
        max_p=6,
        max_q=6
    )

print(f"Selected Order (p,d,q): {auto_model.order}")

# Fit ARIMA using selected order
model = ARIMA(df_group['Temperature'], order=auto_model.order)
model_fit = model.fit()
print(model_fit.summary())

# Forecast next 4 values
forecast = model_fit.forecast(steps=4)

# Plot actual and forecast
plt.figure(figsize=(14,6))
plt.plot(df_group['Temperature'][-10:], label='Actual')
plt.plot(
    np.arange(len(df_group), len(df_group)+4),
    forecast,
    label='Forecast',
    linestyle='dashed'
)
plt.title(f'ARIMA Forecast vs Actual ({start}-{start + group_size})')
plt.legend()
plt.show()

# Performance metrics using actual last 4 values
y_actual = df_group['Temperature'][-4:].values
y_pred = forecast[:4]
mae = mean_absolute_error(y_actual, y_pred)
rmse = np.sqrt(mean_squared_error(y_actual, y_pred))
print(f'ARIMA MAE ({start}-{start + group_size}): {mae:.4f}, RMSE: {rmse:.4f}')

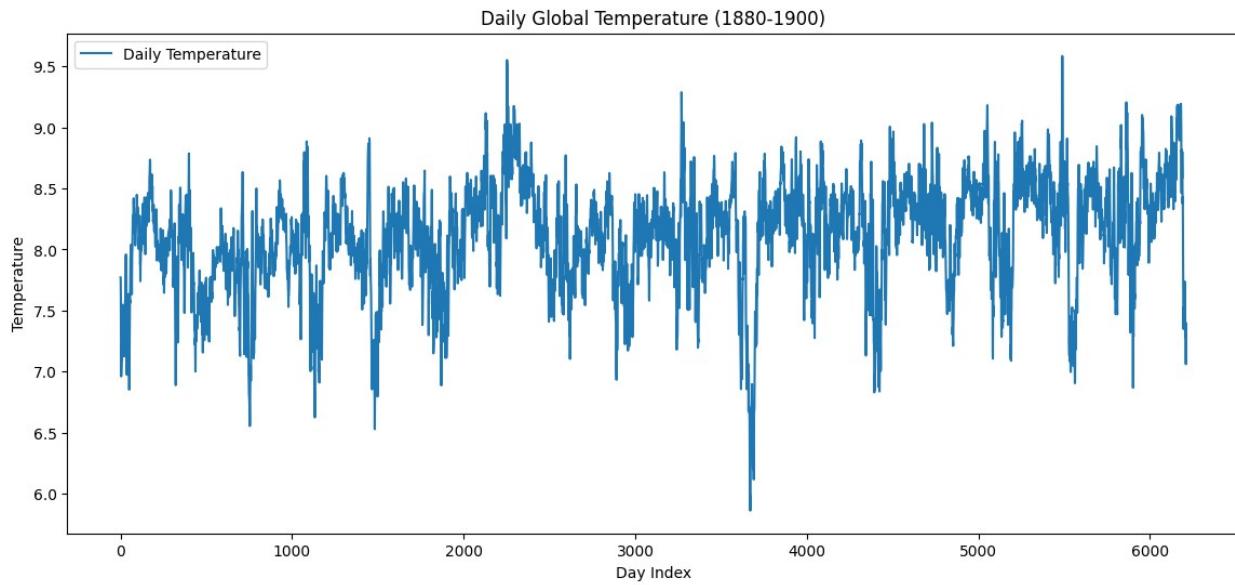
results_list.append({
    "Time Period": f'{start}-{start + group_size}',
    "Order": auto_model.order,
    "MAE": mae,
    "RMSE": rmse
})

# Final DataFrame of results
results_df = pd.DataFrame(results_list)

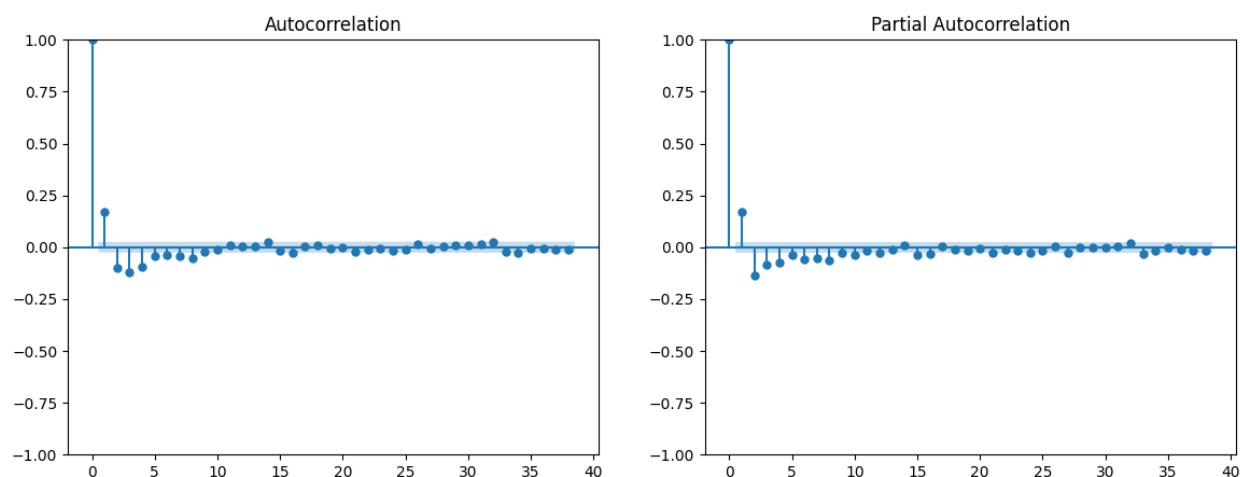
```

```
print("\nARIMA Performance Summary:")
print(results_df)
```

Analyzing period: 1880-1900



```
ADF Statistic: -9.505999565937447
p-value: 3.3461405645231796e-16
Stationary series (p <= 0.05)
ADF Statistic: -24.730374567814753
p-value: 0.0
Stationary series (p <= 0.05)
```



```
Performing stepwise search to minimize aic
ARIMA(2,1,2)(0,0,0)[0] intercept : AIC=-7356.225, Time=4.08 sec
```

```

ARIMA(0,1,0)(0,0,0)[0] intercept      : AIC=-6890.814, Time=0.34 sec
ARIMA(1,1,0)(0,0,0)[0] intercept      : AIC=-7074.713, Time=0.51 sec
ARIMA(0,1,1)(0,0,0)[0] intercept      : AIC=-7120.486, Time=0.57 sec
ARIMA(0,1,0)(0,0,0)[0] intercept      : AIC=-6892.813, Time=0.16 sec
ARIMA(1,1,2)(0,0,0)[0] intercept      : AIC=-7360.900, Time=3.05 sec
ARIMA(0,1,2)(0,0,0)[0] intercept      : AIC=-7150.586, Time=1.63 sec
ARIMA(1,1,1)(0,0,0)[0] intercept      : AIC=-7134.081, Time=1.28 sec
ARIMA(1,1,3)(0,0,0)[0] intercept      : AIC=-7362.650, Time=4.74 sec
ARIMA(0,1,3)(0,0,0)[0] intercept      : AIC=-7222.532, Time=1.04 sec
ARIMA(2,1,3)(0,0,0)[0] intercept      : AIC=-7357.826, Time=2.19 sec
ARIMA(1,1,4)(0,0,0)[0] intercept      : AIC=-7361.274, Time=4.74 sec
ARIMA(0,1,4)(0,0,0)[0] intercept      : AIC=-7295.500, Time=1.42 sec
ARIMA(2,1,4)(0,0,0)[0] intercept      : AIC=-7359.212, Time=5.32 sec
ARIMA(1,1,3)(0,0,0)[0]                : AIC=-7364.650, Time=1.21 sec
ARIMA(0,1,3)(0,0,0)[0]                : AIC=-7224.531, Time=0.35 sec
ARIMA(1,1,2)(0,0,0)[0]                : AIC=-7363.973, Time=0.88 sec
ARIMA(2,1,3)(0,0,0)[0]                : AIC=-7362.857, Time=2.27 sec
ARIMA(1,1,4)(0,0,0)[0]                : AIC=inf, Time=2.42 sec
ARIMA(0,1,2)(0,0,0)[0]                : AIC=-7152.585, Time=0.31 sec
ARIMA(0,1,4)(0,0,0)[0]                : AIC=-7297.499, Time=0.48 sec
ARIMA(2,1,2)(0,0,0)[0]                : AIC=-7364.120, Time=1.69 sec
ARIMA(2,1,4)(0,0,0)[0]                : AIC=-7362.469, Time=4.57 sec

```

Best model: ARIMA(1,1,3)(0,0,0)[0]

Total fit time: 45.300 seconds

Selected Order (p,d,q): (1, 1, 3)

SARIMAX Results

```
=====
=====
Dep. Variable:           Temperature    No. Observations:      6213
Model:                 ARIMA(1, 1, 3)    Log Likelihood:   3687.325
Date:                  Tue, 08 Apr 2025   AIC:                 -7364.650
Time:                  16:14:27         BIC:                 -7330.979
Sample:                   0   HQIC:                -7352.978
                           - 6213
```

Covariance Type: opg

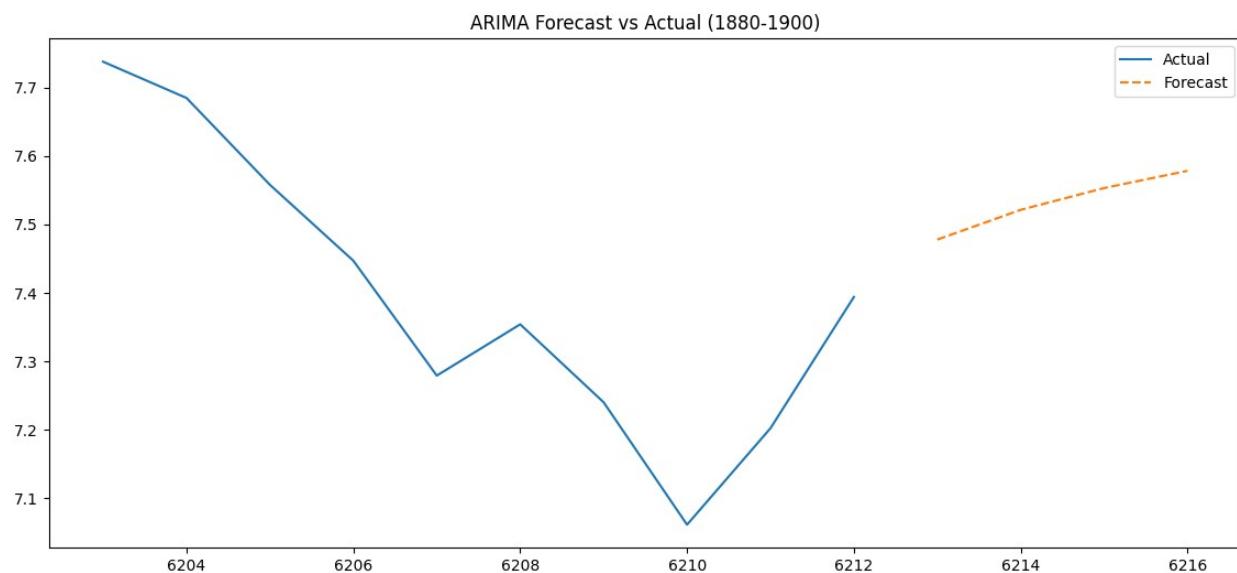
```
=====
=====
            coef    std err          z     P>|z|      [0.025
0.975]
```

ar.L1	0.7916	0.018	43.431	0.000	0.756
0.827					
ma.L1	-0.6332	0.021	-30.125	0.000	-0.674
-0.592					
ma.L2	-0.2433	0.013	-18.590	0.000	-0.269
-0.218					
ma.L3	-0.0317	0.014	-2.318	0.020	-0.059
-0.005					
sigma2	0.0179	0.000	65.123	0.000	0.017
0.018					

Ljung-Box (L1) (Q):	144.16	0.00	Jarque-Bera (JB):
Prob(Q):	0.00	0.99	Prob(JB):
Heteroskedasticity (H):	-0.06	1.38	Skew:
Prob(H) (two-sided):	3.74	0.00	Kurtosis:

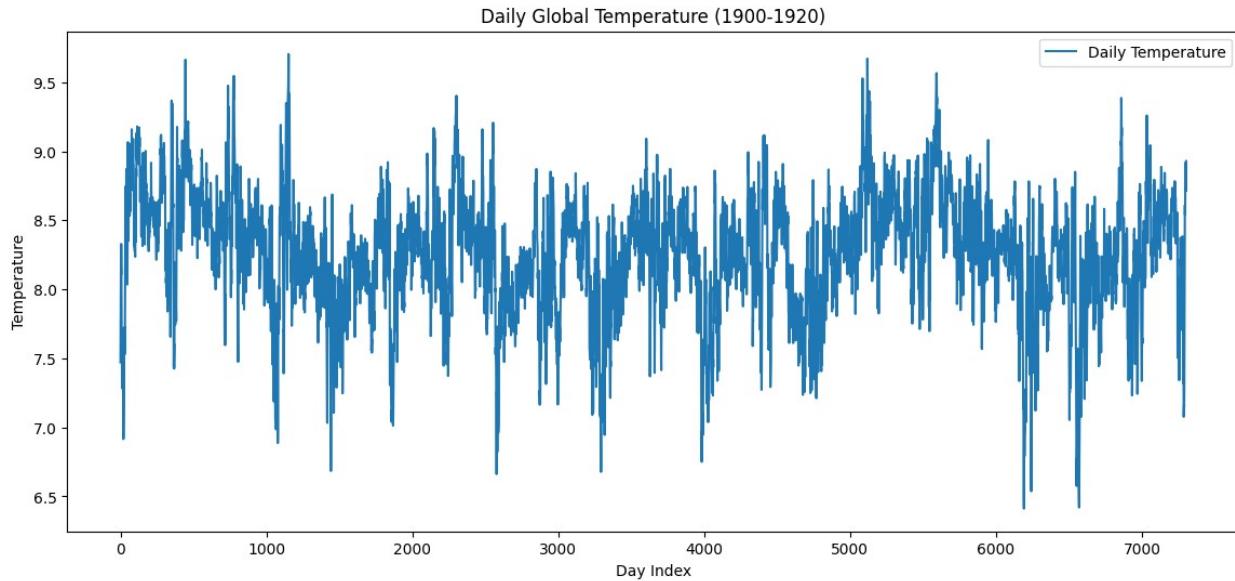
Warnings:

[1] Covariance matrix calculated using the outer product of gradients (complex-step).



ARIMA MAE (1880-1900): 0.3085, RMSE: 0.3264

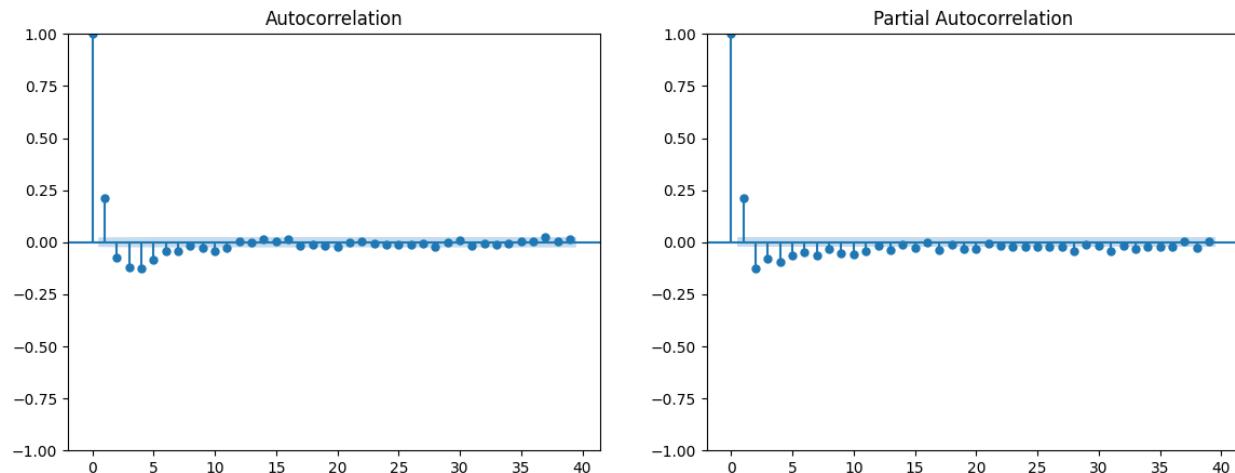
Analyzing period: 1900-1920



```

ADF Statistic: -11.319604391380265
p-value: 1.1830963854949672e-20
Stationary series (p <= 0.05)
ADF Statistic: -21.167372868090293
p-value: 0.0
Stationary series (p <= 0.05)

```



```

Performing stepwise search to minimize aic
ARIMA(2,1,2)(0,0,0)[0] intercept : AIC=-7798.864, Time=4.07 sec
ARIMA(0,1,0)(0,0,0)[0] intercept : AIC=-7008.392, Time=0.44 sec
ARIMA(1,1,0)(0,0,0)[0] intercept : AIC=-7349.110, Time=0.31 sec
ARIMA(0,1,1)(0,0,0)[0] intercept : AIC=-7415.822, Time=0.58 sec
ARIMA(0,1,0)(0,0,0)[0] intercept : AIC=-7010.382, Time=0.15 sec
ARIMA(1,1,2)(0,0,0)[0] intercept : AIC=inf, Time=5.81 sec
ARIMA(2,1,1)(0,0,0)[0] intercept : AIC=-7789.314, Time=5.10 sec
ARIMA(3,1,2)(0,0,0)[0] intercept : AIC=-7791.952, Time=4.65 sec

```

ARIMA(2,1,3)(0,0,0)[0]	intercept	:	AIC=-7789.912, Time=4.96 sec
ARIMA(1,1,1)(0,0,0)[0]	intercept	:	AIC=-7420.823, Time=1.64 sec
ARIMA(1,1,3)(0,0,0)[0]	intercept	:	AIC=-7794.653, Time=4.56 sec
ARIMA(3,1,1)(0,0,0)[0]	intercept	:	AIC=-7804.966, Time=5.13 sec
ARIMA(3,1,0)(0,0,0)[0]	intercept	:	AIC=-7511.321, Time=0.74 sec
ARIMA(4,1,1)(0,0,0)[0]	intercept	:	AIC=-7804.548, Time=5.47 sec
ARIMA(2,1,0)(0,0,0)[0]	intercept	:	AIC=-7466.929, Time=0.35 sec
ARIMA(4,1,0)(0,0,0)[0]	intercept	:	AIC=-7576.185, Time=0.99 sec
ARIMA(4,1,2)(0,0,0)[0]	intercept	:	AIC=-7799.927, Time=6.10 sec
ARIMA(3,1,1)(0,0,0)[0]		:	AIC=-7806.970, Time=1.35 sec
ARIMA(2,1,1)(0,0,0)[0]		:	AIC=-7791.304, Time=1.14 sec
ARIMA(3,1,0)(0,0,0)[0]		:	AIC=-7513.312, Time=0.40 sec
ARIMA(4,1,1)(0,0,0)[0]		:	AIC=-7806.861, Time=1.75 sec
ARIMA(3,1,2)(0,0,0)[0]		:	AIC=-7811.460, Time=2.74 sec
ARIMA(2,1,2)(0,0,0)[0]		:	AIC=-7805.344, Time=2.16 sec
ARIMA(4,1,2)(0,0,0)[0]		:	AIC=-7803.154, Time=2.85 sec
ARIMA(3,1,3)(0,0,0)[0]		:	AIC=-7801.584, Time=3.49 sec
ARIMA(2,1,3)(0,0,0)[0]		:	AIC=-7803.638, Time=3.47 sec
ARIMA(4,1,3)(0,0,0)[0]		:	AIC=-7797.230, Time=3.16 sec

Best model: ARIMA(3,1,2)(0,0,0)[0]

Total fit time: 73.559 seconds

Selected Order (p,d,q): (3, 1, 2)

SARIMAX Results

```
=====
=====
Dep. Variable: Temperature No. Observations: 7304
Model: ARIMA(3, 1, 2) Log Likelihood: 3911.730
Date: Tue, 08 Apr 2025 AIC: -7811.460
Time: 16:15:44 BIC: -7770.084
Sample: 0 HQIC: -7797.233
- 7304
Covariance Type: opg
```

	coef	std err	z	P> z	[0.025
0.975]					
-----	-----	-----	-----	-----	-----
ar.L1	1.6683	0.133	12.513	0.000	1.407
1.930					
ar.L2	-0.9191	0.149	-6.167	0.000	-1.211

```

-0.627
ar.L3      0.1943    0.035     5.615     0.000     0.126
0.262
ma.L1      -1.4803   0.135    -10.984    0.000    -1.744
-1.216
ma.L2      0.4939    0.130     3.797     0.000     0.239
0.749
sigma2     0.0201   0.000     69.269    0.000     0.020
0.021
=====
=====
```

```
Ljung-Box (L1) (Q):          0.01  Jarque-Bera (JB):
```

```
134.73
```

```
Prob(Q):
```

```
0.00
```

```
Heteroskedasticity (H):
```

```
-0.05
```

```
Prob(H) (two-sided):
```

```
3.66
```

=====

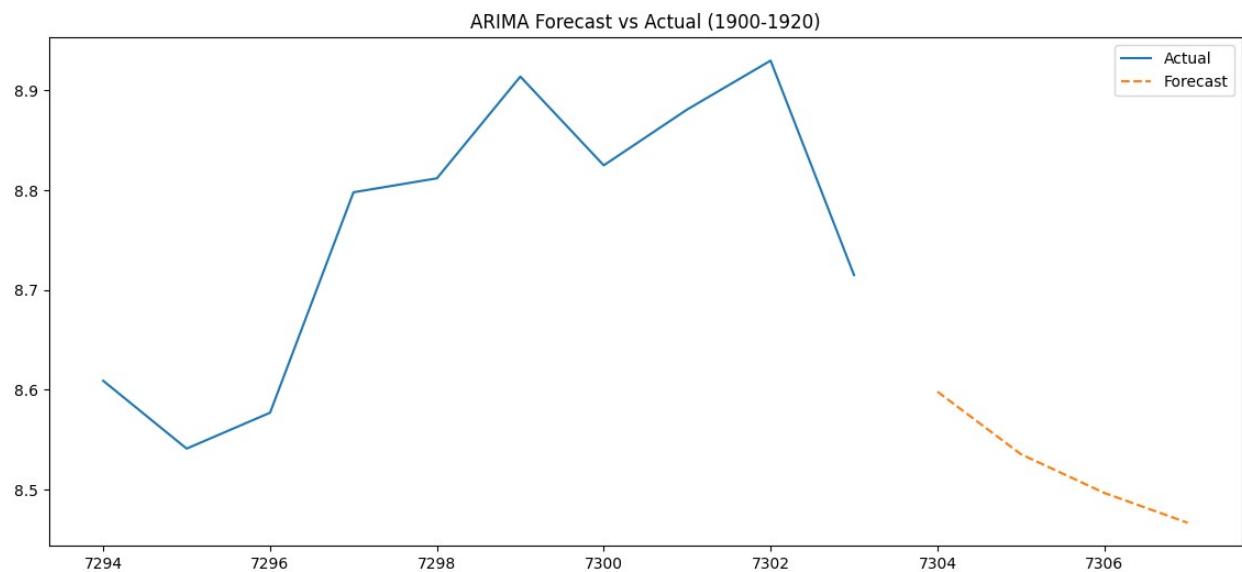
```
0.91  Prob(JB):
```

```
0.98  Skew:
```

```
0.67  Kurtosis:
```

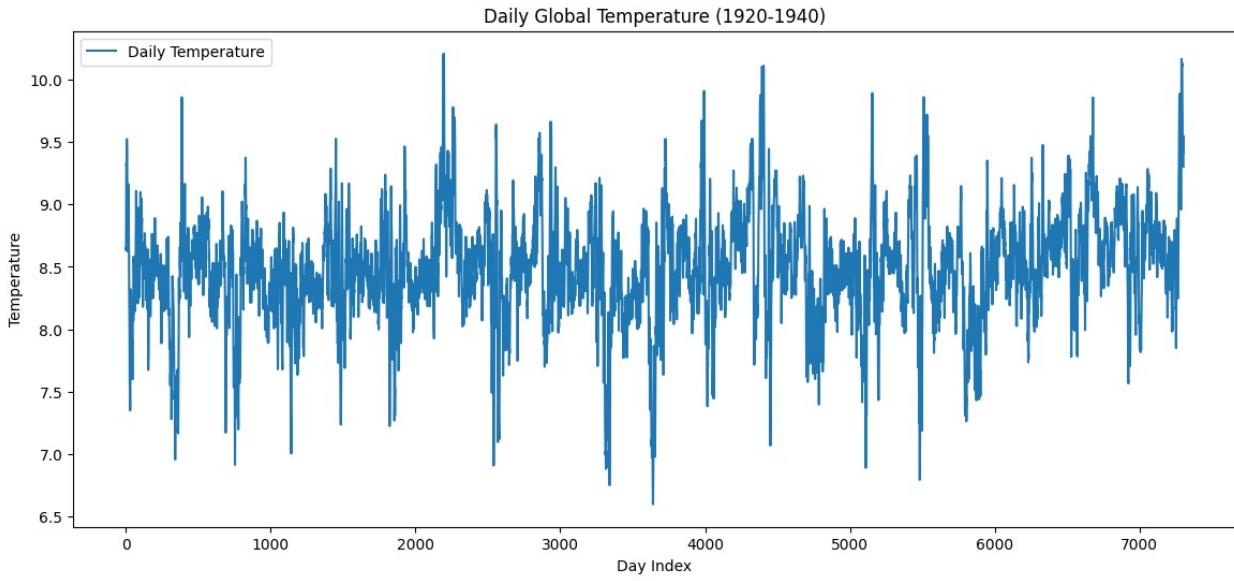
Warnings:

[1] Covariance matrix calculated using the outer product of gradients (complex-step).

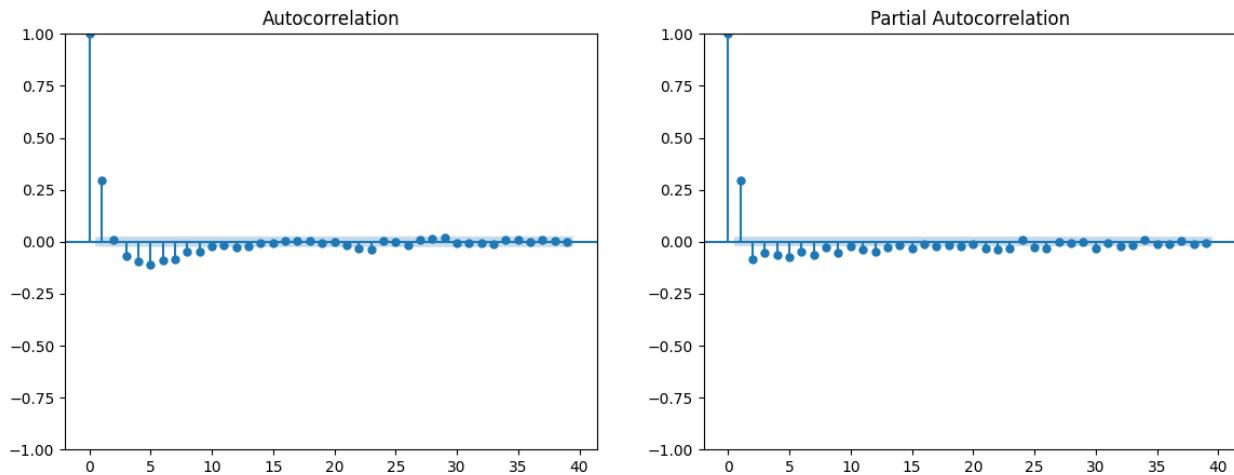


ARIMA MAE (1900-1920): 0.3136, RMSE: 0.3242

Analyzing period: 1920-1940



```
ADF Statistic: -13.478767214028714
p-value: 3.2724619072701088e-25
Stationary series (p <= 0.05)
ADF Statistic: -22.837800799250253
p-value: 0.0
Stationary series (p <= 0.05)
```



```
Performing stepwise search to minimize aic
ARIMA(2,1,2)(0,0,0)[0] intercept : AIC=-9226.554, Time=4.12 sec
ARIMA(0,1,0)(0,0,0)[0] intercept : AIC=-8286.312, Time=0.30 sec
ARIMA(1,1,0)(0,0,0)[0] intercept : AIC=-8955.796, Time=1.40 sec
ARIMA(0,1,1)(0,0,0)[0] intercept : AIC=-8989.146, Time=0.54 sec
ARIMA(0,1,0)(0,0,0)[0] : AIC=-8288.309, Time=0.14 sec
ARIMA(1,1,2)(0,0,0)[0] intercept : AIC=-8999.913, Time=3.03 sec
ARIMA(2,1,1)(0,0,0)[0] intercept : AIC=-9268.339, Time=3.10 sec
ARIMA(1,1,1)(0,0,0)[0] intercept : AIC=-8995.692, Time=0.78 sec
```

```

ARIMA(2,1,0)(0,0,0)[0] intercept : AIC=-9009.027, Time=0.29 sec
ARIMA(3,1,1)(0,0,0)[0] intercept : AIC=inf, Time=4.15 sec
ARIMA(3,1,0)(0,0,0)[0] intercept : AIC=-9025.861, Time=0.77 sec
ARIMA(3,1,2)(0,0,0)[0] intercept : AIC=inf, Time=6.89 sec
ARIMA(2,1,1)(0,0,0)[0] intercept : AIC=inf, Time=1.32 sec

```

Best model: ARIMA(2,1,1)(0,0,0)[0] intercept

Total fit time: 26.883 seconds

Selected Order (p,d,q): (2, 1, 1)

SARIMAX Results

```
=====
=====
Dep. Variable: Temperature No. Observations: 7305
Model: ARIMA(2, 1, 1) Log Likelihood: 4659.833
Date: Tue, 08 Apr 2025 AIC: -9311.666
Time: 16:16:14 BIC: -9284.081
Sample: 0 HQIC: -9302.181
- 7305
Covariance Type: opg
```

	coef	std err	z	P> z	[0.025
0.975]					
-----	-----	-----	-----	-----	-----
ar.L1	1.2597	0.010	128.984	0.000	1.241
1.279					
ar.L2	-0.3330	0.010	-33.819	0.000	-0.352
-0.314					
ma.L1	-0.9957	0.001	-862.638	0.000	-0.998
-0.993					
sigma2	0.0163	0.000	68.804	0.000	0.016
0.017					

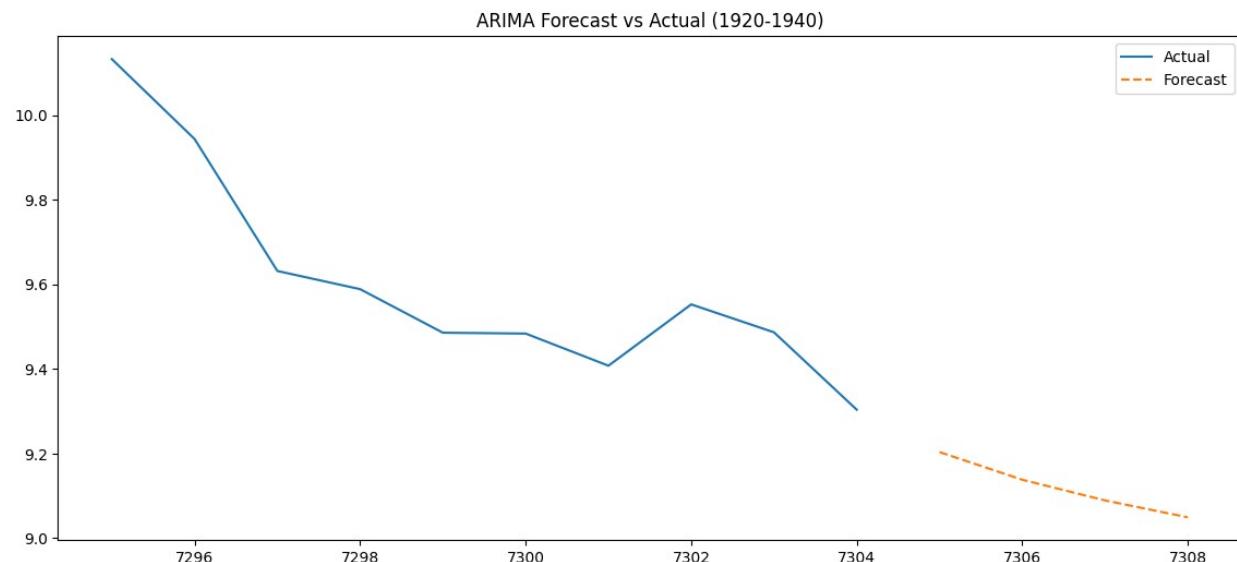
Ljung-Box (L1) (Q):	1.09	Jarque-Bera (JB):
124.05		
Prob(Q):	0.30	Prob(JB):
0.00		
Heteroskedasticity (H):	0.81	Skew:
-0.10		
Prob(H) (two-sided):	0.00	Kurtosis:

3.61

```
=====
```

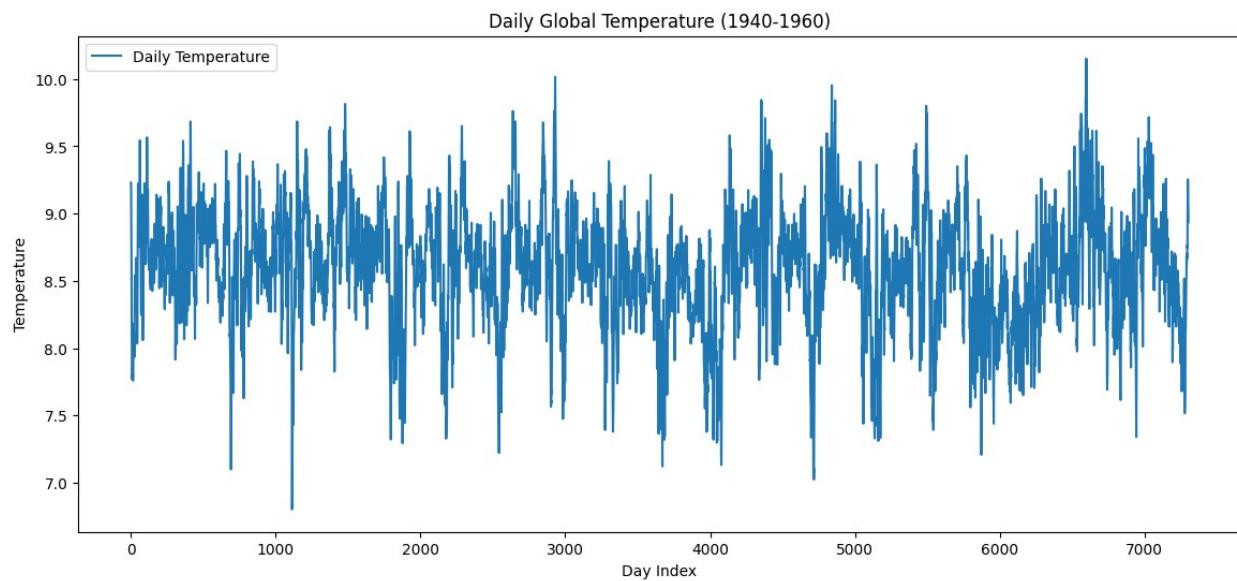
Warnings:

```
[1] Covariance matrix calculated using the outer product of gradients  
(complex-step).
```



ARIMA MAE (1920-1940): 0.3177, RMSE: 0.3303

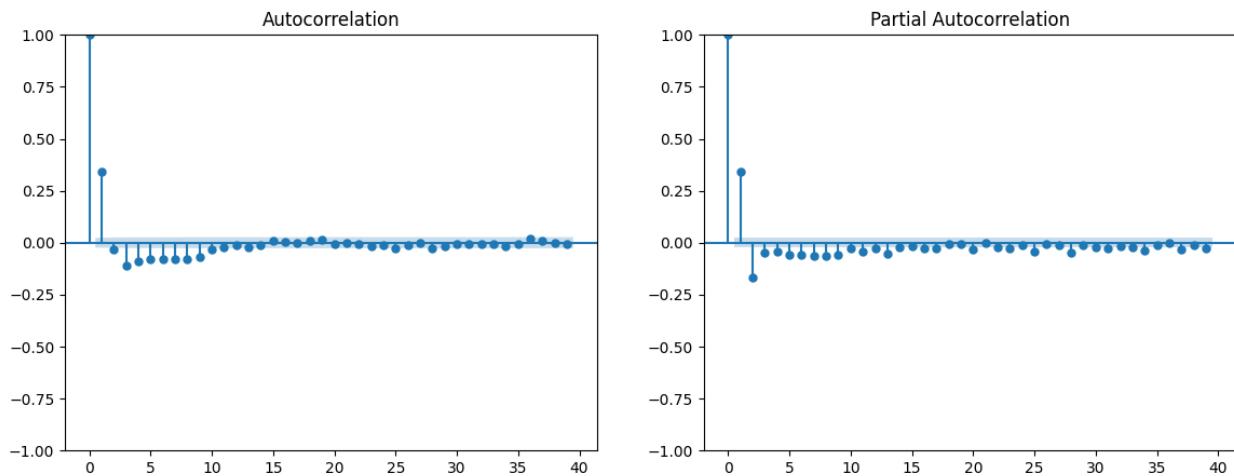
Analyzing period: 1940-1960



```

ADF Statistic: -12.27529371649404
p-value: 8.459582293576229e-23
Stationary series (p <= 0.05)
ADF Statistic: -20.324966838133054
p-value: 0.0
Stationary series (p <= 0.05)

```



```

Performing stepwise search to minimize aic
ARIMA(2,1,2)(0,0,0)[0] intercept : AIC=-9197.336, Time=3.74 sec
ARIMA(0,1,0)(0,0,0)[0] intercept : AIC=-7804.979, Time=0.31 sec
ARIMA(1,1,0)(0,0,0)[0] intercept : AIC=-8696.580, Time=0.67 sec
ARIMA(0,1,1)(0,0,0)[0] intercept : AIC=-8841.272, Time=0.58 sec
ARIMA(0,1,0)(0,0,0)[0] : AIC=-7806.978, Time=0.22 sec
ARIMA(1,1,2)(0,0,0)[0] intercept : AIC=-9110.173, Time=2.61 sec
ARIMA(2,1,1)(0,0,0)[0] intercept : AIC=-9131.656, Time=2.43 sec
ARIMA(3,1,2)(0,0,0)[0] intercept : AIC=-9165.646, Time=4.32 sec
ARIMA(2,1,3)(0,0,0)[0] intercept : AIC=inf, Time=5.45 sec
ARIMA(1,1,1)(0,0,0)[0] intercept : AIC=-8843.897, Time=0.91 sec
ARIMA(1,1,3)(0,0,0)[0] intercept : AIC=inf, Time=5.18 sec
ARIMA(3,1,1)(0,0,0)[0] intercept : AIC=-9201.021, Time=4.94 sec
ARIMA(3,1,0)(0,0,0)[0] intercept : AIC=-8910.869, Time=0.38 sec
ARIMA(4,1,1)(0,0,0)[0] intercept : AIC=-9186.042, Time=4.95 sec
ARIMA(2,1,0)(0,0,0)[0] intercept : AIC=-8896.629, Time=0.30 sec
ARIMA(4,1,0)(0,0,0)[0] intercept : AIC=-8922.922, Time=0.54 sec
ARIMA(4,1,2)(0,0,0)[0] intercept : AIC=-9106.094, Time=5.54 sec
ARIMA(3,1,1)(0,0,0)[0] : AIC=inf, Time=1.94 sec

```

```

Best model: ARIMA(3,1,1)(0,0,0)[0] intercept
Total fit time: 45.012 seconds
Selected Order (p,d,q): (3, 1, 1)

```

SARIMAX Results

```

=====
=====
```

```

Dep. Variable: Temperature No. Observations: 7305
7305
Model: ARIMA(3, 1, 1) Log Likelihood 4610.040
4610.040
Date: Tue, 08 Apr 2025 AIC 9210.079
9210.079
Time: 16:17:02 BIC 9175.599
9175.599
Sample: 0 HQIC 9198.223
9198.223
- 7305

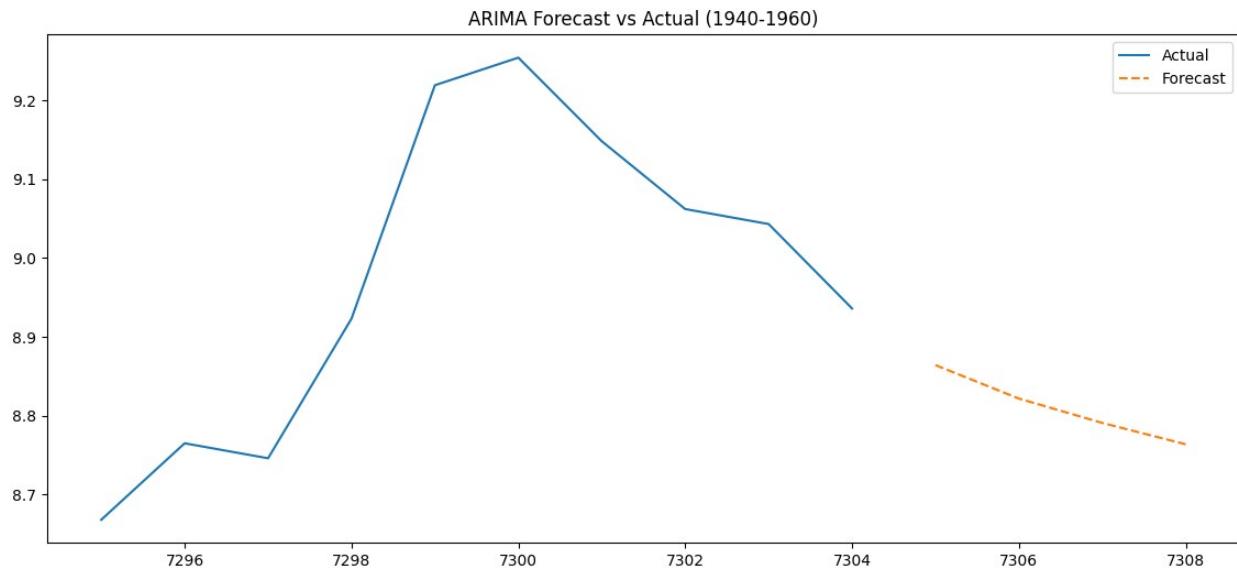
Covariance Type: opg
=====

=====      coef    std err      z      P>|z|      [0.025
0.975]
-----
ar.L1      1.3389    0.011    121.546      0.000      1.317
1.361
ar.L2     -0.5224    0.018    -29.676      0.000     -0.557
-0.488
ar.L3      0.1090    0.011      9.966      0.000      0.088
0.130
ma.L1     -0.9937    0.001    -685.111      0.000     -0.997
-0.991
sigma2     0.0166    0.000      65.655      0.000      0.016
0.017
=====

=====      Ljung-Box (L1) (Q): 0.00      Jarque-Bera (JB):
45.52
Prob(Q): 0.99      Prob(JB):
0.00
Heteroskedasticity (H): 1.14      Skew:
-0.06
Prob(H) (two-sided): 0.00      Kurtosis:
3.36
=====

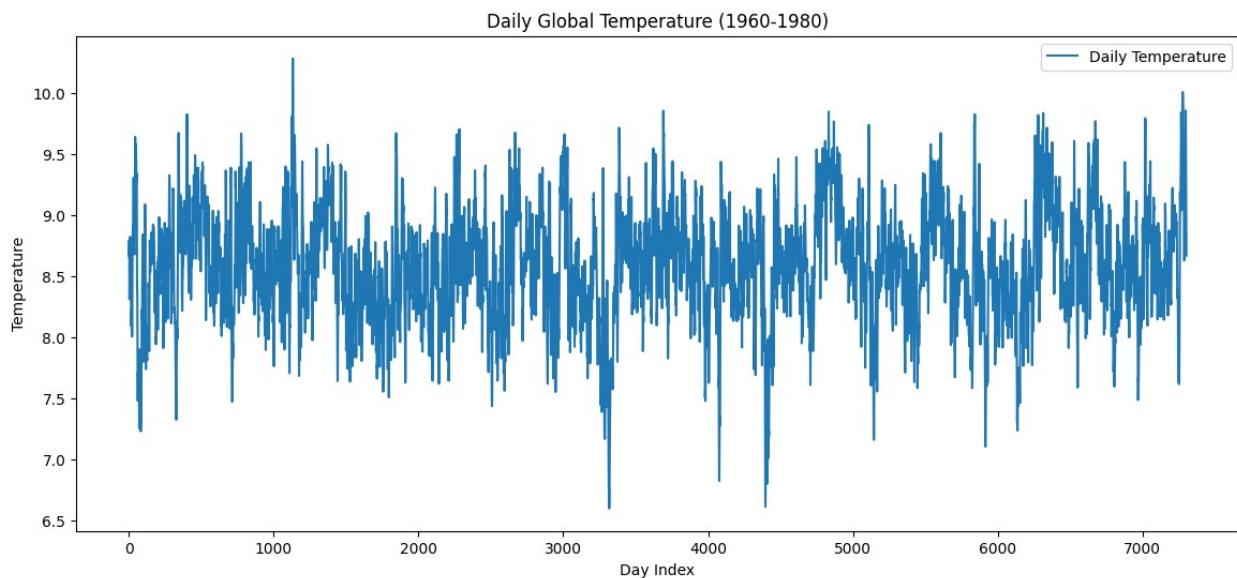
=====      Warnings:
[1] Covariance matrix calculated using the outer product of gradients
(complex-step).

```

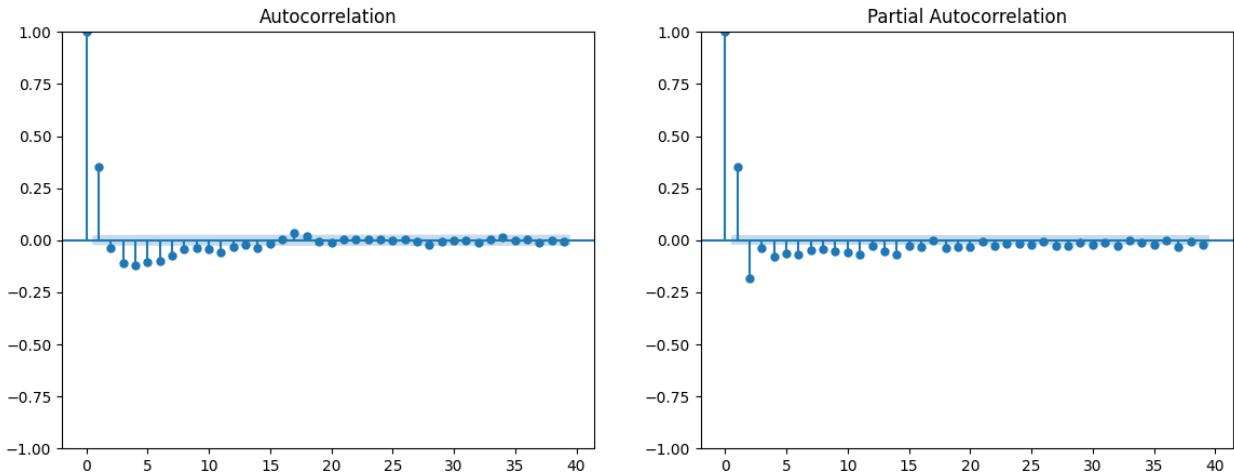


ARIMA MAE (1940-1960): 0.2372, RMSE: 0.2406

Analyzing period: 1960-1980



ADF Statistic: -11.993721933526885
 p-value: 3.4729483329151113e-22
 Stationary series (p <= 0.05)
 ADF Statistic: -20.037106267414817
 p-value: 0.0
 Stationary series (p <= 0.05)



```

Performing stepwise search to minimize aic
ARIMA(2,0,2)(0,0,0)[0] : AIC=-7113.495, Time=2.74 sec
ARIMA(0,0,0)(0,0,0)[0] : AIC=52160.627, Time=0.04 sec
ARIMA(1,0,0)(0,0,0)[0] : AIC=inf, Time=0.27 sec
ARIMA(0,0,1)(0,0,0)[0] : AIC=inf, Time=0.32 sec
ARIMA(1,0,2)(0,0,0)[0] : AIC=-7118.198, Time=0.73 sec
ARIMA(0,0,2)(0,0,0)[0] : AIC=32969.431, Time=0.49 sec
ARIMA(1,0,1)(0,0,0)[0] : AIC=-7117.974, Time=0.23 sec
ARIMA(1,0,3)(0,0,0)[0] : AIC=-7144.448, Time=0.96 sec
ARIMA(0,0,3)(0,0,0)[0] : AIC=25465.631, Time=0.91 sec
ARIMA(2,0,3)(0,0,0)[0] : AIC=-7114.261, Time=2.36 sec
ARIMA(1,0,4)(0,0,0)[0] : AIC=-7207.252, Time=0.78 sec
ARIMA(0,0,4)(0,0,0)[0] : AIC=19531.028, Time=1.38 sec
ARIMA(2,0,4)(0,0,0)[0] : AIC=-7140.346, Time=3.28 sec
ARIMA(1,0,5)(0,0,0)[0] : AIC=-7270.931, Time=1.07 sec
ARIMA(0,0,5)(0,0,0)[0] : AIC=15142.056, Time=2.12 sec
ARIMA(2,0,5)(0,0,0)[0] : AIC=-7203.939, Time=4.29 sec
ARIMA(1,0,6)(0,0,0)[0] : AIC=-7353.574, Time=1.52 sec
ARIMA(0,0,6)(0,0,0)[0] : AIC=11626.131, Time=2.00 sec
ARIMA(2,0,6)(0,0,0)[0] : AIC=-7536.162, Time=2.82 sec
ARIMA(3,0,6)(0,0,0)[0] : AIC=-7212.935, Time=5.59 sec
ARIMA(3,0,5)(0,0,0)[0] : AIC=-7155.627, Time=4.42 sec
ARIMA(2,0,6)(0,0,0)[0] intercept : AIC=-7550.224, Time=6.41 sec
ARIMA(1,0,6)(0,0,0)[0] intercept : AIC=-7542.015, Time=9.12 sec
ARIMA(2,0,5)(0,0,0)[0] intercept : AIC=-7542.308, Time=6.02 sec
ARIMA(3,0,6)(0,0,0)[0] intercept : AIC=-7555.912, Time=9.39 sec
ARIMA(3,0,5)(0,0,0)[0] intercept : AIC=-7554.213, Time=9.39 sec
ARIMA(4,0,6)(0,0,0)[0] intercept : AIC=-7556.616, Time=11.34 sec
ARIMA(4,0,5)(0,0,0)[0] intercept : AIC=-7556.238, Time=8.57 sec
ARIMA(5,0,6)(0,0,0)[0] intercept : AIC=-7556.134, Time=11.71 sec
ARIMA(5,0,5)(0,0,0)[0] intercept : AIC=-7543.517, Time=11.42 sec
ARIMA(4,0,6)(0,0,0)[0] : AIC=-7140.561, Time=5.66 sec

```

Best model: ARIMA(4,0,6)(0,0,0)[0] intercept

Total fit time: 127.404 seconds
Selected Order (p,d,q): (4, 0, 6)

SARIMAX Results

=====

=====

Dep. Variable:	Temperature	No. Observations:
7305		
Model:	ARIMA(4, 0, 6)	Log Likelihood
3785.042		
Date:	Tue, 08 Apr 2025	AIC
7546.084		
Time:	16:19:23	BIC
7463.328		
Sample:	0	HQIC
7517.629		
	- 7305	

Covariance Type: opg

=====

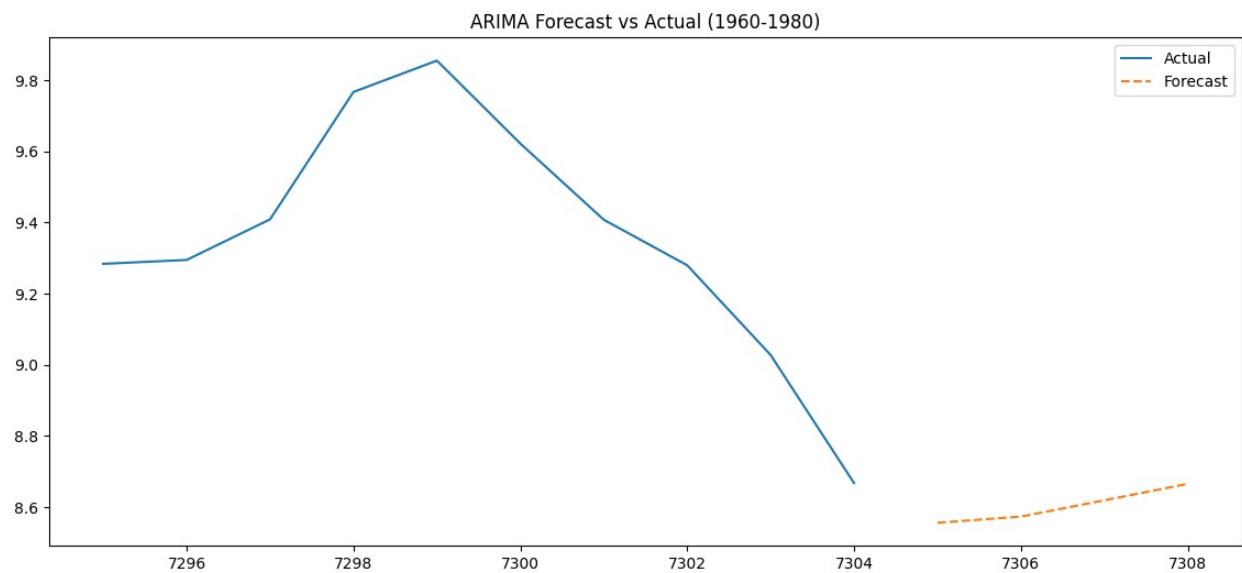
=====

	coef	std err	z	P> z	[0.025
0.975]					
-----	-----	-----	-----	-----	-----
const	8.5912	0.027	319.794	0.000	8.539
8.644					
ar.L1	0.4649	0.160	2.909	0.004	0.152
0.778					
ar.L2	0.2813	0.204	1.378	0.168	-0.119
0.681					
ar.L3	0.6666	0.206	3.233	0.001	0.262
1.071					
ar.L4	-0.4816	0.123	-3.909	0.000	-0.723
-0.240					
ma.L1	0.8999	0.160	5.629	0.000	0.587
1.213					
ma.L2	0.3656	0.166	2.203	0.028	0.040
0.691					
ma.L3	-0.5479	0.141	-3.883	0.000	-0.825
-0.271					
ma.L4	-0.3966	0.069	-5.774	0.000	-0.531
-0.262					
ma.L5	-0.1670	0.027	-6.194	0.000	-0.220
-0.114					
ma.L6	-0.0677	0.015	-4.508	0.000	-0.097
-0.038					
sigma2	0.0209	0.000	62.788	0.000	0.020
0.022					

```
=====
=====
Ljung-Box (L1) (Q):          0.64    Jarque-Bera (JB):
16.64                         0.42    Prob(JB):
Prob(Q):                      0.00
Heteroskedasticity (H):      1.05    Skew:
-0.04                         0.20    Kurtosis:
Prob(H) (two-sided):        3.22
=====
=====
```

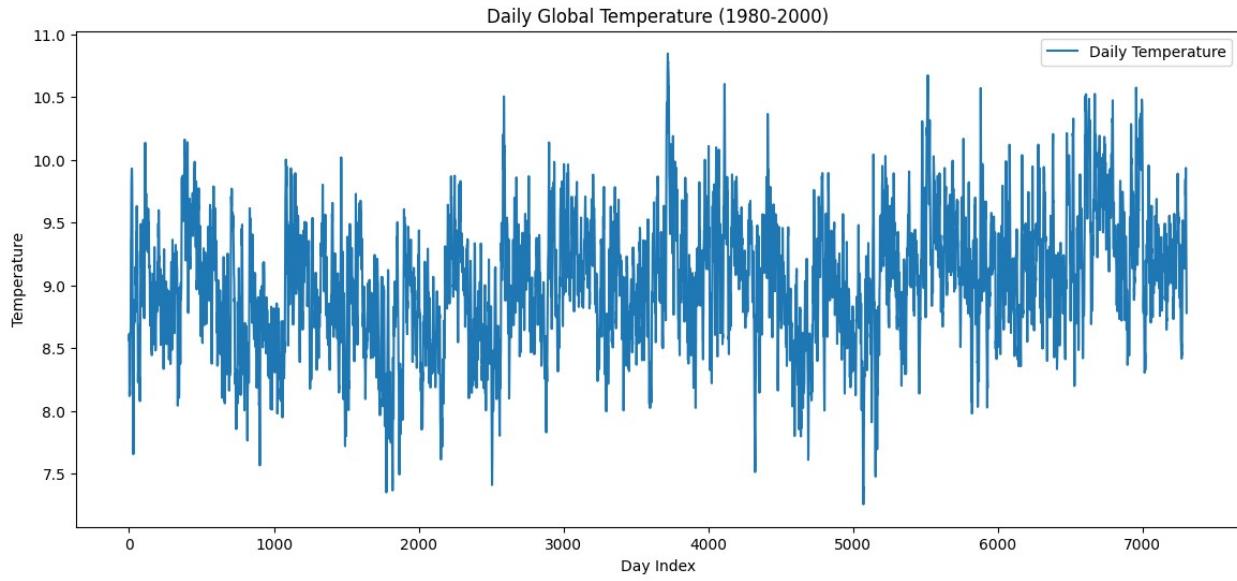
Warnings:

```
[1] Covariance matrix calculated using the outer product of gradients
(complex-step).
```

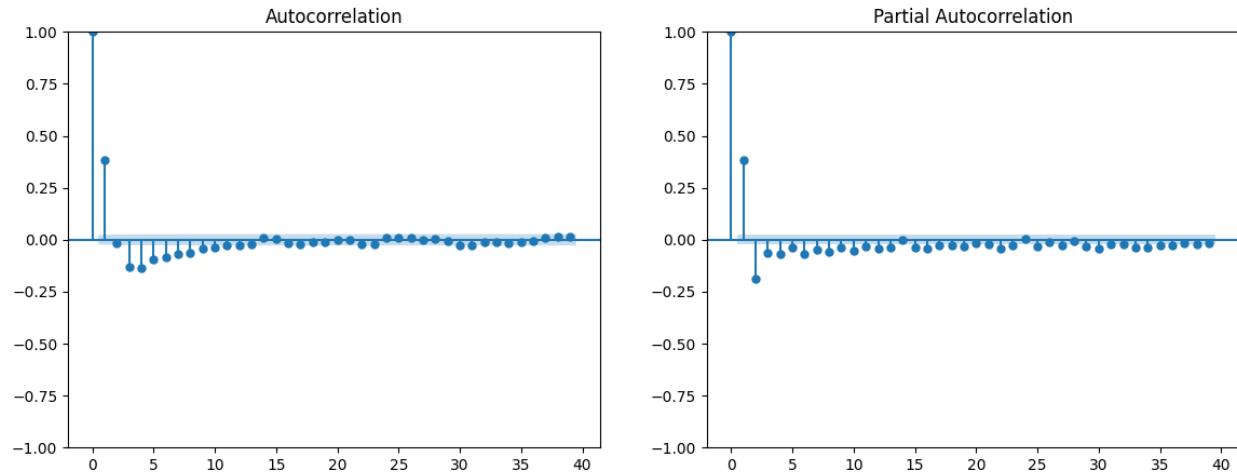


```
ARIMA MAE (1960-1980): 0.4921, RMSE: 0.5897
```

```
Analyzing period: 1980-2000
```



```
ADF Statistic: -7.820413456227168
p-value: 6.693545479622451e-12
Stationary series (p <= 0.05)
ADF Statistic: -21.99891719650499
p-value: 0.0
Stationary series (p <= 0.05)
```



```
Performing stepwise search to minimize aic
ARIMA(2,1,2)(0,0,0)[0] intercept : AIC=-7568.907, Time=3.75 sec
ARIMA(0,1,0)(0,0,0)[0] intercept : AIC=-5799.393, Time=0.41 sec
ARIMA(1,1,0)(0,0,0)[0] intercept : AIC=-6941.991, Time=0.76 sec
ARIMA(0,1,1)(0,0,0)[0] intercept : AIC=-7123.404, Time=0.53 sec
ARIMA(0,1,0)(0,0,0)[0] intercept : AIC=-5801.392, Time=0.15 sec
ARIMA(1,1,2)(0,0,0)[0] intercept : AIC=inf, Time=3.84 sec
ARIMA(2,1,1)(0,0,0)[0] intercept : AIC=-7483.240, Time=3.38 sec
ARIMA(3,1,2)(0,0,0)[0] intercept : AIC=-7505.095, Time=4.59 sec
```

```

ARIMA(2,1,3)(0,0,0)[0] intercept      : AIC=-7550.989, Time=5.04 sec
ARIMA(1,1,1)(0,0,0)[0] intercept      : AIC=-7136.362, Time=1.01 sec
ARIMA(1,1,3)(0,0,0)[0] intercept      : AIC=inf, Time=3.21 sec
ARIMA(3,1,1)(0,0,0)[0] intercept      : AIC=-7572.700, Time=4.65 sec
ARIMA(3,1,0)(0,0,0)[0] intercept      : AIC=-7229.854, Time=0.45 sec
ARIMA(4,1,1)(0,0,0)[0] intercept      : AIC=-7551.042, Time=4.83 sec
ARIMA(2,1,0)(0,0,0)[0] intercept      : AIC=-7204.426, Time=0.33 sec
ARIMA(4,1,0)(0,0,0)[0] intercept      : AIC=-7264.488, Time=0.55 sec
ARIMA(4,1,2)(0,0,0)[0] intercept      : AIC=-7562.450, Time=6.00 sec
ARIMA(3,1,1)(0,0,0)[0]                : AIC=-7575.347, Time=1.84 sec
ARIMA(2,1,1)(0,0,0)[0]                : AIC=-7485.235, Time=1.17 sec
ARIMA(3,1,0)(0,0,0)[0]                : AIC=-7231.853, Time=0.71 sec
ARIMA(4,1,1)(0,0,0)[0]                : AIC=-7573.350, Time=2.56 sec
ARIMA(3,1,2)(0,0,0)[0]                : AIC=-7573.293, Time=2.92 sec
ARIMA(2,1,0)(0,0,0)[0]                : AIC=-7206.426, Time=0.81 sec
ARIMA(2,1,2)(0,0,0)[0]                : AIC=-7571.295, Time=1.82 sec
ARIMA(4,1,0)(0,0,0)[0]                : AIC=-7266.487, Time=0.27 sec
ARIMA(4,1,2)(0,0,0)[0]                : AIC=-7573.069, Time=3.14 sec

```

Best model: ARIMA(3,1,1)(0,0,0)[0]

Total fit time: 58.757 seconds

Selected Order (p,d,q): (3, 1, 1)

SARIMAX Results

```

=====
=====
Dep. Variable: Temperature No. Observations: 7305
Model: ARIMA(3, 1, 1) Log Likelihood: 3792.673
Date: Tue, 08 Apr 2025 AIC: 7575.347
Time: 16:20:25 BIC: 7540.866
Sample: 0 HQIC: 7563.491
- 7305

```

Covariance Type: opg

```

=====
=====
            coef    std err          z      P>|z|      [0.025
0.975]
-----
ar.L1      1.3794    0.012    119.194      0.000      1.357
1.402
ar.L2     -0.5858    0.018    -31.970      0.000     -0.622
-0.550

```

ar.L3	0.1164	0.012	9.993	0.000	0.094
0.139					
ma.L1	-0.9875	0.002	-443.015	0.000	-0.992
-0.983					
sigma2	0.0207	0.000	62.051	0.000	0.020
0.021					

Ljung-Box (L1) (Q): 0.00 Jarque-Bera (JB):

5.47

Prob(Q):

0.07

Heteroskedasticity (H):

0.03

Prob(H) (two-sided):

3.11

Prob(JB):

1.00

Skew:

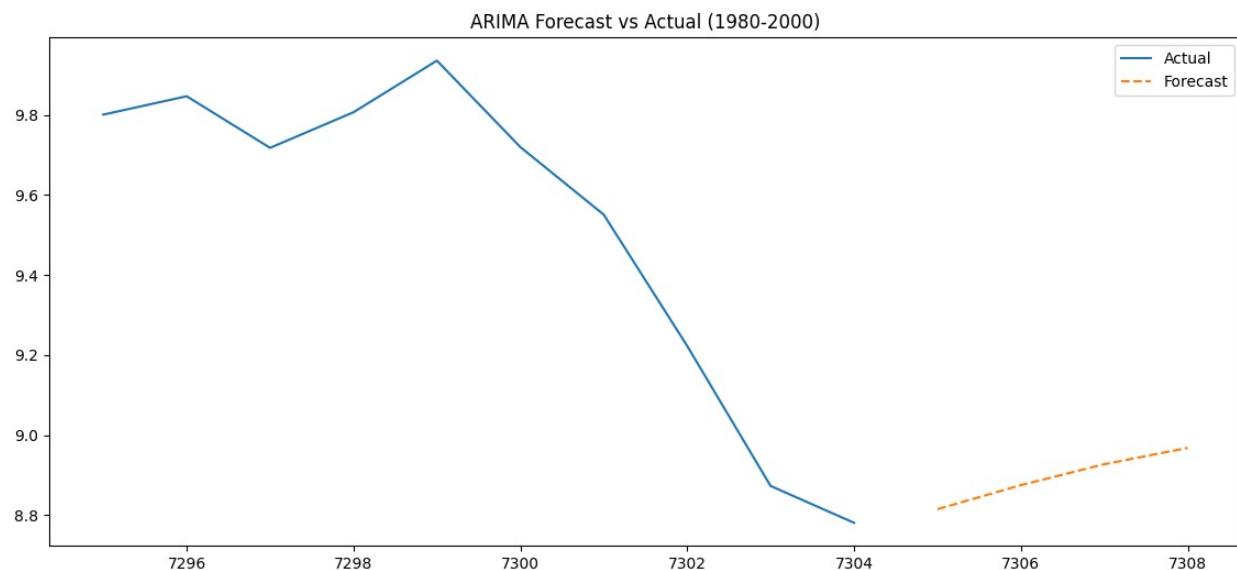
1.09

Kurtosis:

0.03

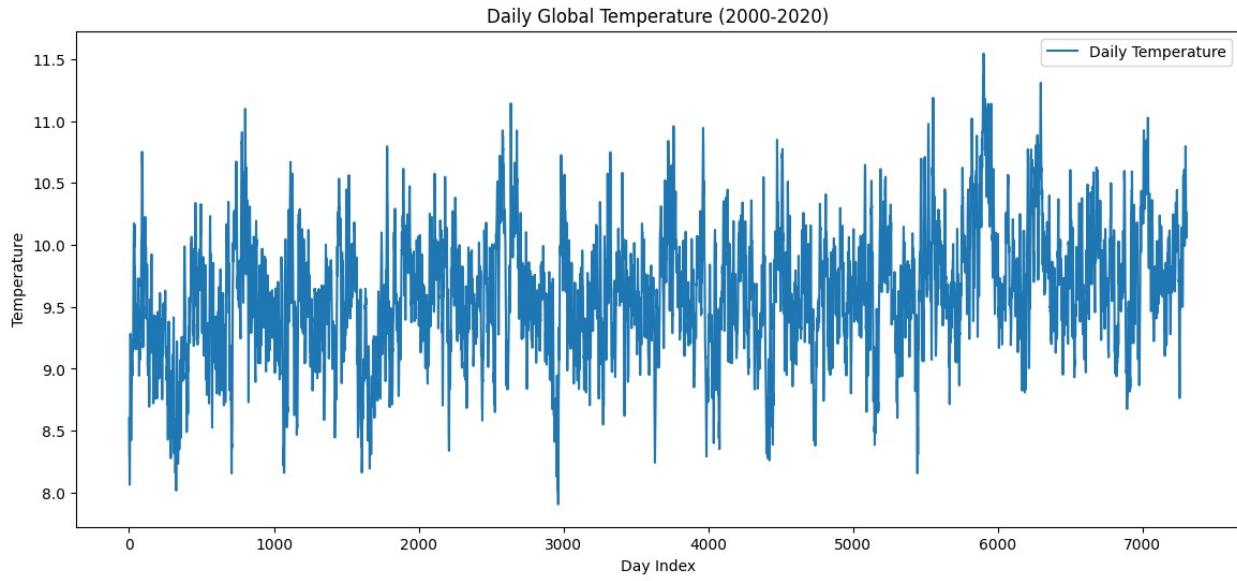
Warnings:

[1] Covariance matrix calculated using the outer product of gradients
(complex-step).



ARIMA MAE (1980-2000): 0.3313, RMSE: 0.4185

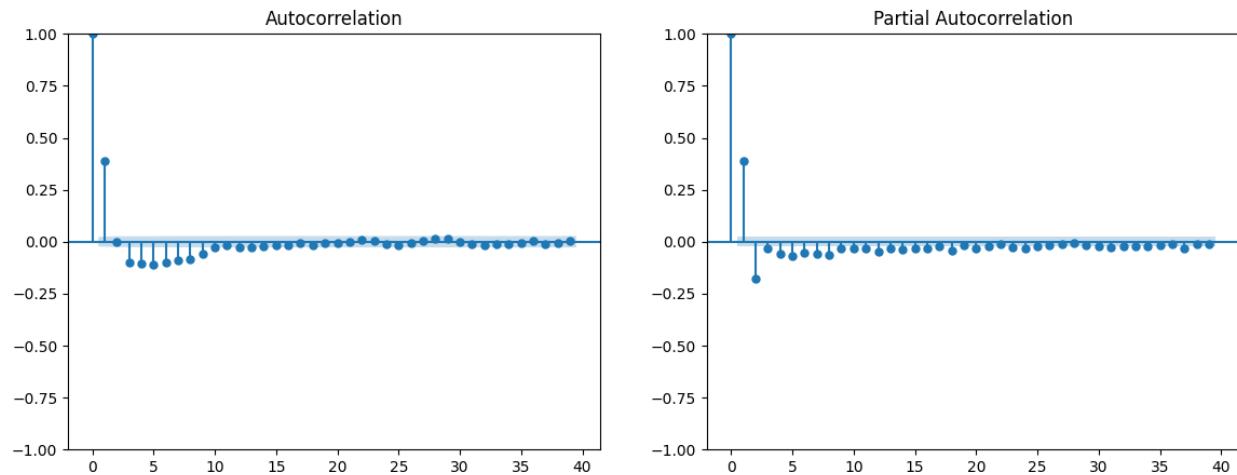
Analyzing period: 2000-2020



```

ADF Statistic: -14.245515559630434
p-value: 1.5119106986016173e-26
Stationary series (p <= 0.05)
ADF Statistic: -20.557143402726886
p-value: 0.0
Stationary series (p <= 0.05)

```



```

Performing stepwise search to minimize aic
ARIMA(2,1,2)(0,0,0)[0] intercept      : AIC=inf, Time=6.59 sec
ARIMA(0,1,0)(0,0,0)[0] intercept      : AIC=-6194.864, Time=0.32 sec
ARIMA(1,1,0)(0,0,0)[0] intercept      : AIC=-7387.833, Time=0.18 sec
ARIMA(0,1,1)(0,0,0)[0] intercept      : AIC=-7558.217, Time=0.52 sec
ARIMA(0,1,0)(0,0,0)[0]                : AIC=-6196.849, Time=0.19 sec
ARIMA(1,1,1)(0,0,0)[0] intercept      : AIC=-7577.079, Time=0.91 sec
ARIMA(2,1,1)(0,0,0)[0] intercept      : AIC=-7864.979, Time=3.34 sec
ARIMA(2,1,0)(0,0,0)[0] intercept      : AIC=-7624.647, Time=0.75 sec

```

```

ARIMA(3,1,1)(0,0,0)[0] intercept : AIC=inf, Time=4.58 sec
ARIMA(1,1,2)(0,0,0)[0] intercept : AIC=inf, Time=3.46 sec
ARIMA(3,1,0)(0,0,0)[0] intercept : AIC=-7630.980, Time=0.45 sec
ARIMA(3,1,2)(0,0,0)[0] intercept : AIC=-7880.396, Time=4.77 sec
ARIMA(4,1,2)(0,0,0)[0] intercept : AIC=-7927.260, Time=6.13 sec
ARIMA(4,1,1)(0,0,0)[0] intercept : AIC=-7879.431, Time=6.29 sec
ARIMA(5,1,2)(0,0,0)[0] intercept : AIC=-7924.450, Time=7.32 sec
ARIMA(4,1,3)(0,0,0)[0] intercept : AIC=-7948.695, Time=6.21 sec
ARIMA(3,1,3)(0,0,0)[0] intercept : AIC=-7948.723, Time=5.66 sec
ARIMA(2,1,3)(0,0,0)[0] intercept : AIC=inf, Time=5.12 sec
ARIMA(3,1,4)(0,0,0)[0] intercept : AIC=inf, Time=5.02 sec
ARIMA(2,1,4)(0,0,0)[0] intercept : AIC=-7951.121, Time=7.82 sec
ARIMA(1,1,4)(0,0,0)[0] intercept : AIC=-7952.044, Time=5.30 sec
ARIMA(0,1,4)(0,0,0)[0] intercept : AIC=-7628.879, Time=1.97 sec
ARIMA(1,1,3)(0,0,0)[0] intercept : AIC=inf, Time=6.44 sec
ARIMA(1,1,5)(0,0,0)[0] intercept : AIC=-7929.124, Time=6.24 sec
ARIMA(0,1,3)(0,0,0)[0] intercept : AIC=-7602.985, Time=1.69 sec
ARIMA(0,1,5)(0,0,0)[0] intercept : AIC=-7680.740, Time=2.74 sec
ARIMA(2,1,5)(0,0,0)[0] intercept : AIC=-7934.863, Time=7.41 sec
ARIMA(1,1,4)(0,0,0)[0] : AIC=inf, Time=2.97 sec

```

Best model: ARIMA(1,1,4)(0,0,0)[0] intercept

Total fit time: 110.410 seconds

Selected Order (p,d,q): (1, 1, 4)

SARIMAX Results

```
=====
=====
Dep. Variable: Temperature No. Observations: 7305
Model: ARIMA(1, 1, 4) Log Likelihood: 3987.527
Date: Tue, 08 Apr 2025 AIC: 7963.053
Time: 16:22:19 BIC: 7921.676
Sample: 0 HQIC: 7948.826
- 7305
```

Covariance Type: opg

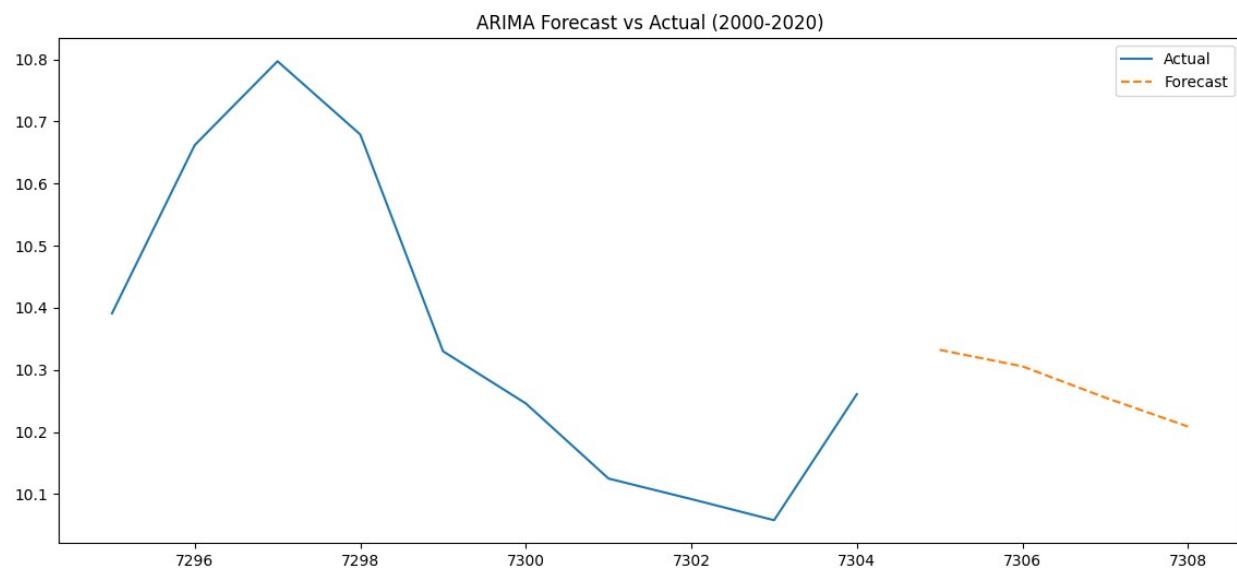
```
=====
=====
            coef    std err          z      P>|z|      [0.025
0.975]
-----
ar.L1      0.8811    0.007    118.404      0.000      0.867
0.896
```

ma.L1	-0.4752	0.013	-35.700	0.000	-0.501
-0.449					
ma.L2	-0.3923	0.013	-30.102	0.000	-0.418
-0.367					
ma.L3	-0.1173	0.013	-8.985	0.000	-0.143
-0.092					
ma.L4	-0.0098	0.013	-0.754	0.451	-0.035
0.016					
sigma2	0.0196	0.000	63.814	0.000	0.019
0.020					

Ljung-Box (L1) (Q):	0.00	Jarque-Bera (JB):
24.88		
Prob(Q):	0.99	Prob(JB):
0.00		
Heteroskedasticity (H):	0.96	Skew:
0.07		
Prob(H) (two-sided):	0.29	Kurtosis:
3.25		

Warnings:

[1] Covariance matrix calculated using the outer product of gradients (complex-step).



ARIMA MAE (2000-2020): 0.1676, RMSE: 0.1805

ARIMA Performance Summary:

Time Period	Order	MAE	RMSE
-------------	-------	-----	------

0	1880-1900	(1, 1, 3)	0.308529	0.326355
1	1900-1920	(3, 1, 2)	0.313560	0.324202
2	1920-1940	(2, 1, 1)	0.317736	0.330327
3	1940-1960	(3, 1, 1)	0.237159	0.240629
4	1960-1980	(4, 0, 6)	0.492115	0.589666
5	1980-2000	(3, 1, 1)	0.331345	0.418542
6	2000-2020	(1, 1, 4)	0.167644	0.180514

```

# todays work
==== Import Libraries ====
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
from statsmodels.tsa.stattools import adfuller
from statsmodels.graphics.tsaplots import plot_acf, plot_pacf
from pmdarima import auto_arima
from statsmodels.tsa.arima.model import ARIMA
from sklearn.metrics import mean_absolute_error, mean_squared_error

# === Step 1: Prepare Data ===
# Assuming df already has 'Year', 'Month', 'Day', 'Temperature', and
# 'Anomaly' columns

# Convert to datetime and set as index
df['Date'] = pd.to_datetime(df[['Year', 'Month', 'Day']])
df.set_index('Date', inplace=True)

# Retain only needed columns
df = df[['Temperature', 'Anomaly']]

# Add 'Year' column for aggregation
df['Year'] = df.index.year

# === Step 2: Aggregate to Annual Min, Max, Mean ===
annual_stats = df.groupby('Year')['Temperature'].agg(['min', 'max',
    'mean']).reset_index()
annual_stats.columns = ['Year', 'MinTemp', 'MaxTemp', 'MeanTemp']
print("Annual Temperature Summary:")
print(annual_stats.head())

# === Step 3: Define ADF Test Function ===
def adf_test(series, title=""):
    result = adfuller(series.dropna())
    print(f"\nADF Test - {title}")
    print(f"ADF Statistic: {result[0]:.4f}")
    print(f"p-value: {result[1]:.4f}")
    if result[1] > 0.05:
        print("=> Non-stationary (p > 0.05)")
    else:
        print("=> Stationary (p <= 0.05)")



```

```

# === Step 4: Define Forecasting Function ===
def model_forecast(series, title, forecast_years=5):
    print(f"\n== {title} ==")

    # Plot time series
    plt.figure(figsize=(10, 4))
    plt.plot(series.values, label='Actual')
    plt.title(f'{title} Over Years')
    plt.xlabel('Year Index')
    plt.ylabel('Temperature')
    plt.grid(True)
    plt.legend()
    plt.tight_layout()
    plt.show()

    # ADF Test
    adf_test(series, title)

    # Differencing
    series_diff = series.diff().dropna()

    # ACF and PACF
    fig, axes = plt.subplots(1, 2, figsize=(14, 4))
    plot_acf(series_diff, ax=axes[0])
    plot_pacf(series_diff, ax=axes[1])
    fig.suptitle(f'{title} - ACF & PACF')
    plt.tight_layout()
    plt.show()

    # Fit Auto ARIMA
    auto_model = auto_arima(series, seasonal=False, trace=True,
                           error_action='ignore',
                           suppress_warnings=True)
    print(f"Selected ARIMA Order: {auto_model.order}")

    # Fit ARIMA model
    model = ARIMA(series, order=auto_model.order)
    model_fit = model.fit()
    print(model_fit.summary())

    # Forecast
    forecast = model_fit.forecast(steps=forecast_years)
    forecast_index = np.arange(len(series), len(series) +
forecast_years)

    # Plot forecast
    plt.figure(figsize=(10, 4))
    plt.plot(series.values, label='Actual')

```

```

        plt.plot(forecast_index, forecast, label='Forecast',
linestyle='--', color='red')
        plt.title(f'{title} Forecast for {forecast_years} Years')
        plt.xlabel('Year Index')
        plt.ylabel('Temperature')
        plt.legend()
        plt.grid(True)
        plt.tight_layout()
        plt.show()

    return forecast

# === Step 5: Apply to All Three Series ===
min_forecast = model_forecast(annual_stats['MinTemp'], "Annual Minimum
Temperature")
max_forecast = model_forecast(annual_stats['MaxTemp'], "Annual Maximum
Temperature")
mean_forecast = model_forecast(annual_stats['MeanTemp'], "Annual Mean
Temperature")

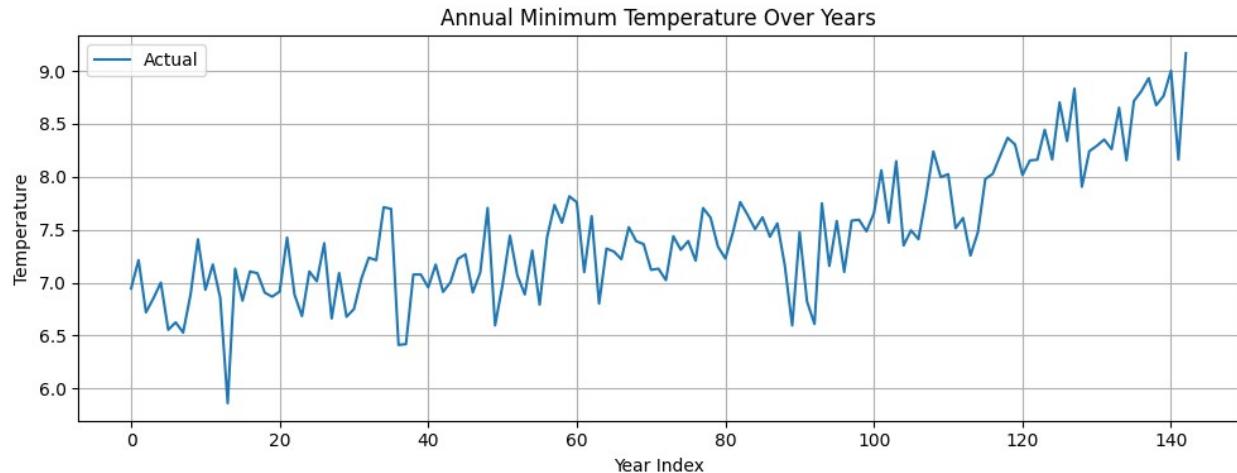
# === Step 6: Combine Forecasts into One Table ===
forecast_years = np.arange(annual_stats['Year'].iloc[-1] + 1,
annual_stats['Year'].iloc[-1] + 6)
forecast_df = pd.DataFrame({
    'Year': forecast_years,
    'Min Forecast': min_forecast,
    'Max Forecast': max_forecast,
    'Mean Forecast': mean_forecast
})

print("\nForecasted Annual Temperatures:")
print(forecast_df)

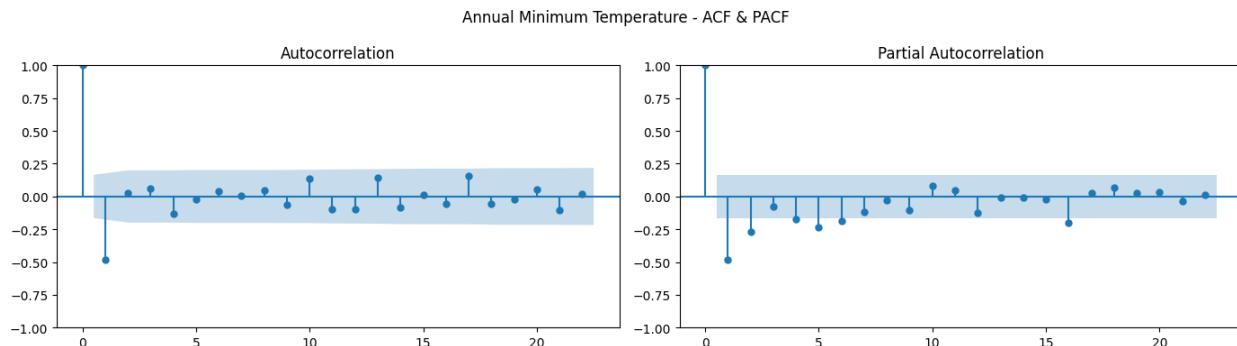
Annual Temperature Summary:
   Year  MinTemp  MaxTemp  MeanTemp
0  1880      6.945     8.666  7.995557
1  1881      7.213     9.006  8.255378
2  1882      6.720     9.325  8.116529
3  1883      6.852     8.735  7.995145
4  1884      7.001     8.786  7.801880

==== Annual Minimum Temperature ====

```



```
ADF Test - Annual Minimum Temperature
ADF Statistic: 0.1330
p-value: 0.9682
=> Non-stationary (p > 0.05)
```



```
Performing stepwise search to minimize aic
ARIMA(2,1,2)(0,0,0)[0] intercept : AIC=104.664, Time=0.34 sec
ARIMA(0,1,0)(0,0,0)[0] intercept : AIC=163.981, Time=0.03 sec
ARIMA(1,1,0)(0,0,0)[0] intercept : AIC=126.994, Time=0.01 sec
ARIMA(0,1,1)(0,0,0)[0] intercept : AIC=100.377, Time=0.06 sec
ARIMA(0,1,0)(0,0,0)[0]          : AIC=162.173, Time=0.01 sec
ARIMA(1,1,1)(0,0,0)[0] intercept : AIC=101.663, Time=0.08 sec
ARIMA(0,1,2)(0,0,0)[0] intercept : AIC=101.781, Time=0.09 sec
ARIMA(1,1,2)(0,0,0)[0] intercept : AIC=104.128, Time=0.09 sec
ARIMA(0,1,1)(0,0,0)[0]          : AIC=103.725, Time=0.03 sec
```

Best model: ARIMA(0,1,1)(0,0,0)[0] intercept

Total fit time: 0.754 seconds

Selected ARIMA Order: (0, 1, 1)

SARIMAX Results

```

=====
Dep. Variable: MinTemp   No. Observations: 143
Model: ARIMA(0, 1, 1)   Log Likelihood: -49.863
Date: Tue, 06 May 2025   AIC: 103.725
Time: 14:34:53   BIC: 109.637
Sample: 0   HQIC: 106.127
                           - 143

Covariance Type: opg
=====
```

	coef	std err	z	P> z	[0.025
0.975]					
-----	-----	-----	-----	-----	-----
ma.L1	-0.7660	0.048	-16.102	0.000	-0.859
-0.673					
sigma2	0.1174	0.014	8.355	0.000	0.090
0.145					

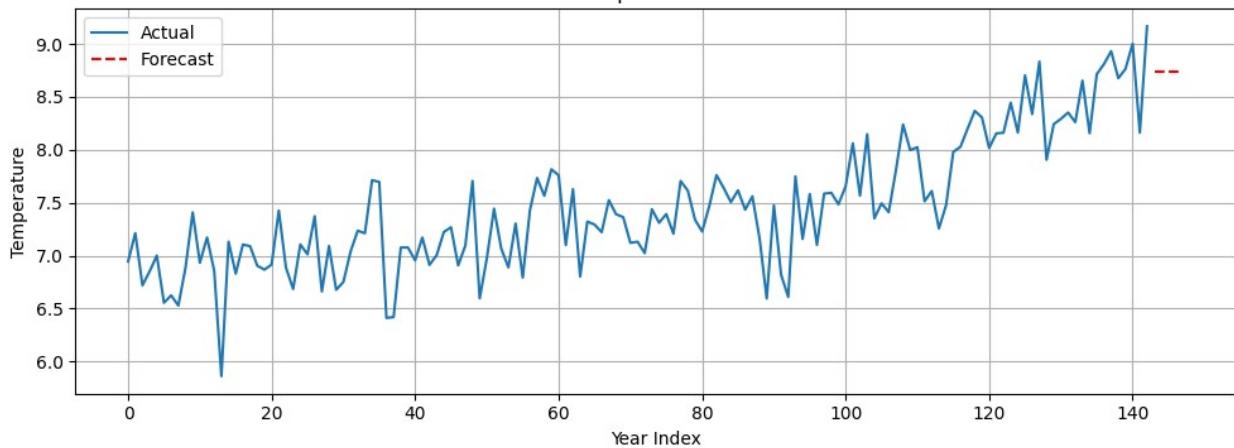
```

=====
Ljung-Box (L1) (Q): 0.00   Jarque-Bera (JB): 8.93
Prob(Q): 0.98   Prob(JB): 0.01
Heteroskedasticity (H): 0.94   Skew: -0.57
Prob(H) (two-sided): 0.82   Kurtosis: 3.44
=====
```

```

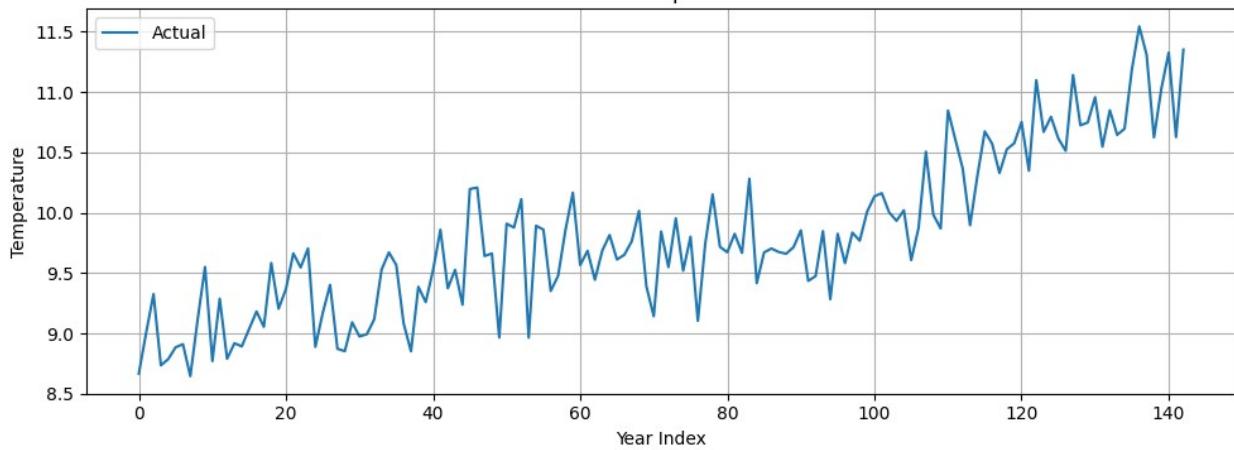
=====
Warnings:
[1] Covariance matrix calculated using the outer product of gradients
(complex-step).
=====
```

Annual Minimum Temperature Forecast for 5 Years



==== Annual Maximum Temperature ===

Annual Maximum Temperature Over Years

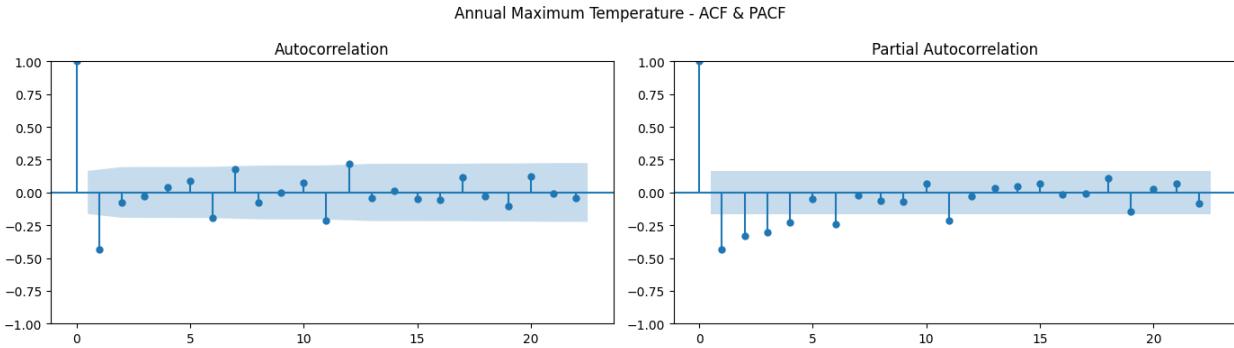


ADF Test - Annual Maximum Temperature

ADF Statistic: 0.8198

p-value: 0.9919

=> Non-stationary ($p > 0.05$)



```

Performing stepwise search to minimize aic
ARIMA(2,1,2)(0,0,0)[0] intercept : AIC=82.085, Time=0.34 sec
ARIMA(0,1,0)(0,0,0)[0] intercept : AIC=149.246, Time=0.02 sec
ARIMA(1,1,0)(0,0,0)[0] intercept : AIC=120.757, Time=0.02 sec
ARIMA(0,1,1)(0,0,0)[0] intercept : AIC=77.889, Time=0.06 sec
ARIMA(0,1,0)(0,0,0)[0]          : AIC=147.557, Time=0.01 sec
ARIMA(1,1,1)(0,0,0)[0] intercept : AIC=79.858, Time=0.11 sec
ARIMA(0,1,2)(0,0,0)[0] intercept : AIC=79.849, Time=0.21 sec
ARIMA(1,1,2)(0,0,0)[0] intercept : AIC=inf, Time=0.25 sec
ARIMA(0,1,1)(0,0,0)[0]          : AIC=84.407, Time=0.03 sec

```

Best model: ARIMA(0,1,1)(0,0,0)[0] intercept

Total fit time: 1.068 seconds

Selected ARIMA Order: (0, 1, 1)

SARIMAX Results

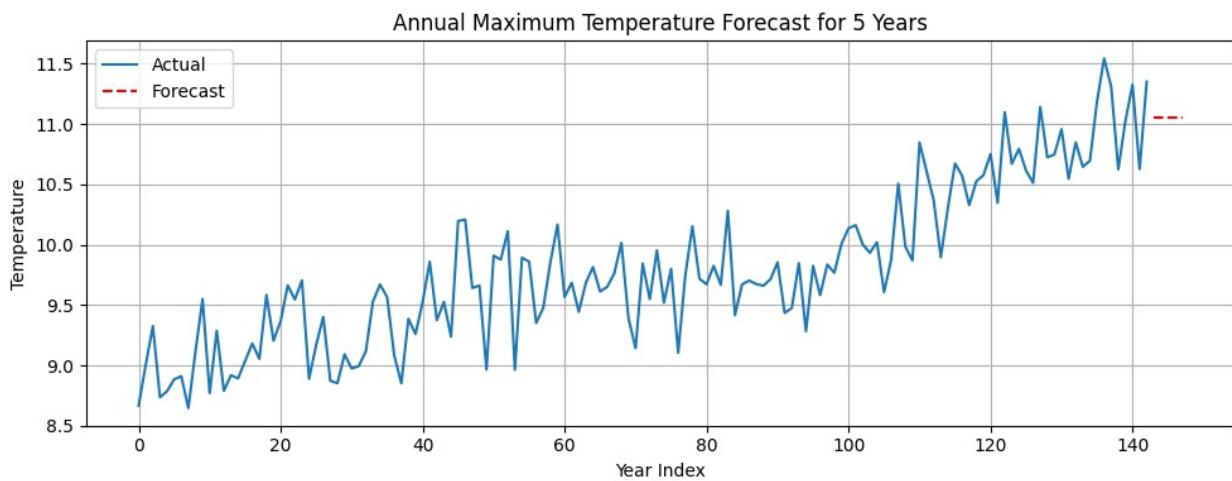
Dep. Variable:	MaxTemp	No. Observations:			
143					
Model:	ARIMA(0, 1, 1)	Log Likelihood			
-40.204					
Date:	Tue, 06 May 2025	AIC			
84.407					
Time:	14:34:55	BIC			
90.319					
Sample:	0	HQIC			
86.810					
	- 143				
Covariance Type:	opg				
	coef	std err	z	P> z	[0.025
0.975]					
ma.L1	-0.7663	0.055	-14.022	0.000	-0.873

```
-0.659  
sigma2          0.1025      0.012      8.216      0.000      0.078  
0.127
```

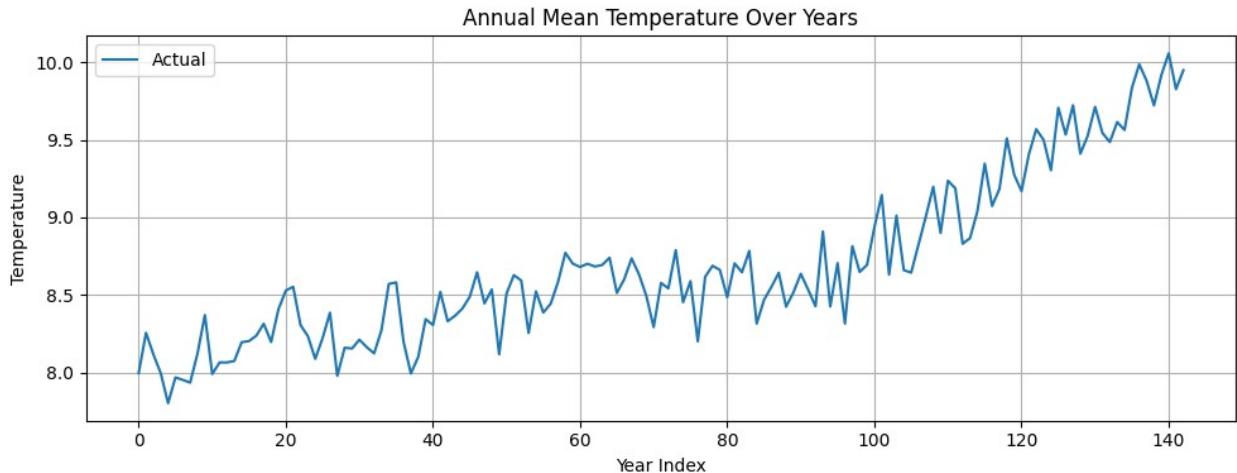
```
Ljung-Box (L1) (Q):           0.33   Jarque-Bera (JB):  
0.11  
Prob(Q):                   0.57   Prob(JB):  
0.94  
Heteroskedasticity (H):     0.97   Skew:  
-0.07  
Prob(H) (two-sided):        0.92   Kurtosis:  
2.95
```

Warnings:

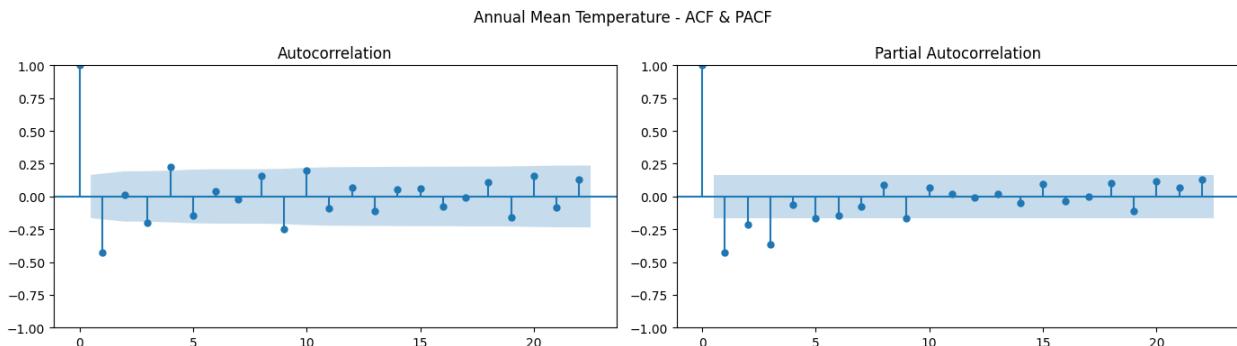
```
[1] Covariance matrix calculated using the outer product of gradients  
(complex-step).
```



```
==== Annual Mean Temperature ===
```



```
ADF Test - Annual Mean Temperature
ADF Statistic: 1.7800
p-value: 0.9983
=> Non-stationary (p > 0.05)
```



```
Performing stepwise search to minimize aic
ARIMA(2,1,2)(0,0,0)[0] intercept : AIC=-83.595, Time=0.22 sec
ARIMA(0,1,0)(0,0,0)[0] intercept : AIC=-32.737, Time=0.04 sec
ARIMA(1,1,0)(0,0,0)[0] intercept : AIC=-59.463, Time=0.04 sec
ARIMA(0,1,1)(0,0,0)[0] intercept : AIC=-84.205, Time=0.06 sec
ARIMA(0,1,0)(0,0,0)[0]          : AIC=-34.143, Time=0.01 sec
ARIMA(1,1,1)(0,0,0)[0] intercept : AIC=-82.957, Time=0.09 sec
ARIMA(0,1,2)(0,0,0)[0] intercept : AIC=-82.956, Time=0.09 sec
ARIMA(1,1,2)(0,0,0)[0] intercept : AIC=-82.726, Time=0.17 sec
ARIMA(0,1,1)(0,0,0)[0]          : AIC=-78.140, Time=0.02 sec
```

Best model: ARIMA(0,1,1)(0,0,0)[0] intercept

Total fit time: 0.737 seconds

Selected ARIMA Order: (0, 1, 1)

SARIMAX Results

```

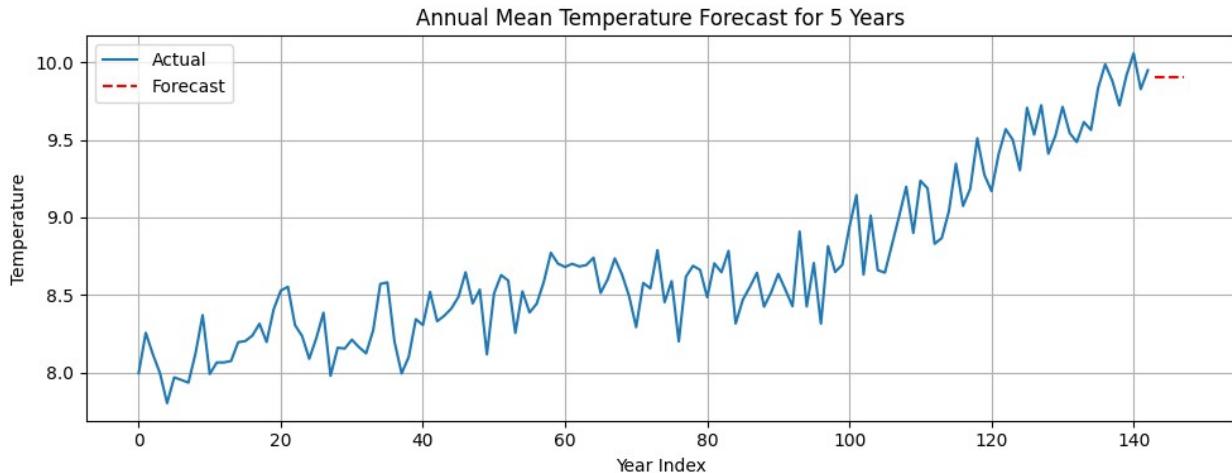
=====
Dep. Variable:           MeanTemp    No. Observations:      143
Model:                 ARIMA(0, 1, 1)    Log Likelihood   -41.070
Date:                  Tue, 06 May 2025   AIC             -78.140
Time:                  14:34:56        BIC             -72.228
Sample:                   0   HQIC            -75.738
                           - 143

Covariance Type: opg

=====
=====
```

	coef	std err	z	P> z	[0.025
0.975]					
-----	-----	-----	-----	-----	-----
ma.L1	-0.6489	0.067	-9.684	0.000	-0.780
-0.518					
sigma2	0.0327	0.004	7.340	0.000	0.024
0.041					

=====



Forecasted Annual Temperatures:

	Year	Min Forecast	Max Forecast	Mean Forecast
143	2023	8.734658	11.052061	9.905106
144	2024	8.734658	11.052061	9.905106
145	2025	8.734658	11.052061	9.905106
146	2026	8.734658	11.052061	9.905106
147	2027	8.734658	11.052061	9.905106

```

import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from statsmodels.tsa.stattools import adfuller
from statsmodels.graphics.tsaplots import plot_acf, plot_pacf
from statsmodels.tsa.arima.model import ARIMA
from sklearn.metrics import mean_absolute_error, mean_squared_error

# Convert to datetime if not already
if 'Date' in df.columns:
    df['Date'] = pd.to_datetime(df['Date'])
    df.set_index('Date', inplace=True)

# Set year column for grouping
df['Year'] = df.index.year

# Initialize parameters
group_size = 10
start_year = 1880
end_year = 2020

# Function to perform ADF test
def adf_test(series):
    result = adfuller(series.dropna())

```

```

    return result[0], result[1]

# Store results
mae_rmse_arima_results = []

# Loop over 20-year groups
for start in range(start_year, end_year, group_size):
    df_group = df[(df['Year'] >= start) & (df['Year'] < start + group_size)].copy()
    if len(df_group) < 40: # skip very small samples
        continue

    # ADF Test
    _, pval = adf_test(df_group['Temperature'])
    if pval > 0.05:
        df_group['Temperature_Diff'] = df_group['Temperature'].diff()
        diff_data = df_group['Temperature_Diff'].dropna()
    else:
        diff_data = df_group['Temperature']

    # Fit ARIMA (5,1,2) arbitrarily; in full work, ACF/PACF should guide p,d,q
    try:
        model = ARIMA(df_group['Temperature'], order=(5,1,2))
        arima_result = model.fit()

        # Forecast last 4 days
        forecast = arima_result.forecast(steps=4)
        y_actual = df_group['Temperature'][-4:]
        y_pred = forecast[:4]

        # Compute MAE and RMSE
        mae = mean_absolute_error(y_actual, y_pred)
        rmse = np.sqrt(mean_squared_error(y_actual, y_pred))
        mae_rmse_arima_results.append([f"{start}-{start + group_size}", mae, rmse])
    except:
        continue

# Create results DataFrame
results_arima_df = pd.DataFrame(mae_rmse_arima_results, columns=["Time Period", "ARIMA_MAE", "ARIMA_RMSE"])
results_arima_df

-----
AttributeError                                Traceback (most recent call
last)
Cell In[22], line 17

```

```

14     df.set_index('Date', inplace=True)
16 # Set year column for grouping
---> 17 df['Year'] = df.index.year
19 # Initialize parameters
20 group_size = 10

AttributeError: 'RangeIndex' object has no attribute 'year'

```

lstm models

```

import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns

from statsmodels.tsa.stattools import adfuller
from statsmodels.graphics.tsaplots import plot_acf, plot_pacf
from statsmodels.tsa.arima.model import ARIMA
from statsmodels.tsa.seasonal import seasonal_decompose

from sklearn.metrics import mean_absolute_error, mean_squared_error
from sklearn.preprocessing import MinMaxScaler
from sklearn.model_selection import train_test_split

import tensorflow as tf
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import LSTM, Dense

# Define group range
group_size = 10
start_year = 1880
end_year = 2020

# ADF Test
def adf_test(series):
    result = adfuller(series.dropna())
    print(f"ADF Statistic: {result[0]}")
    print(f"p-value: {result[1]}")
    if result[1] > 0.05:
        print("Non-stationary series (p > 0.05)")
    else:
        print("Stationary series (p <= 0.05)")

# Function to create sequences for LSTM
def create_sequences(data, seq_length=10):
    X, y = [], []
    for i in range(len(data) - seq_length):
        X.append(data[i:i + seq_length])

```

```

        y.append(data[i + seq_length])
    return np.array(X), np.array(y)

# Assuming df is already loaded and contains 'Year' and 'Temperature'
mae_rmse_results = []

for start in range(start_year, end_year, group_size):
    df_group = df[(df['Year'] >= start) & (df['Year'] < start + group_size)].copy()
    print(f"\nAnalyzing period: {start}-{start + group_size}")

    # Decompose time series
    decomposition = seasonal_decompose(df_group['Temperature'],
model='additive', period=365)
    fig, (ax1, ax2, ax3) = plt.subplots(3, 1, figsize=(14, 8))
    decomposition.trend.plot(ax=ax1, title='Trend')
    decomposition.seasonal.plot(ax=ax2, title='Seasonality')
    decomposition.resid.plot(ax=ax3, title='Residuals')
    plt.tight_layout()
    plt.show()

    # Normalize Data
    scaler = MinMaxScaler(feature_range=(0, 1))
    df_group['Temperature_Scaled'] =
scaler.fit_transform(df_group[['Temperature']])

    # Prepare Sequences for LSTM
    seq_length = 30 # Using 30 days of past data
    X, y = create_sequences(df_group['Temperature_Scaled'].values,
seq_length)

    # Reshape for LSTM
    X = X.reshape((X.shape[0], X.shape[1], 1))

    # Train-Test Split
    X_train, X_test, y_train, y_test = train_test_split(X, y,
test_size=0.2, shuffle=False)

    # Build LSTM Model
    model = Sequential([
        LSTM(50, activation='relu', return_sequences=True,
input_shape=(seq_length, 1)),
        LSTM(50, activation='relu'),
        Dense(1)
    ])
    model.compile(optimizer='adam', loss='mse')

    # Train Model
    model.fit(X_train, y_train, epochs=20, batch_size=16, verbose=1)

```

```

# Predictions
y_pred_lstm = model.predict(X_test)
y_pred_lstm = scaler.inverse_transform(y_pred_lstm)

# Evaluate Performance
y_test_actual = scaler.inverse_transform(y_test.reshape(-1, 1))
mae_lstm = mean_absolute_error(y_test_actual, y_pred_lstm)
rmse_lstm = np.sqrt(mean_squared_error(y_test_actual,
y_pred_lstm))

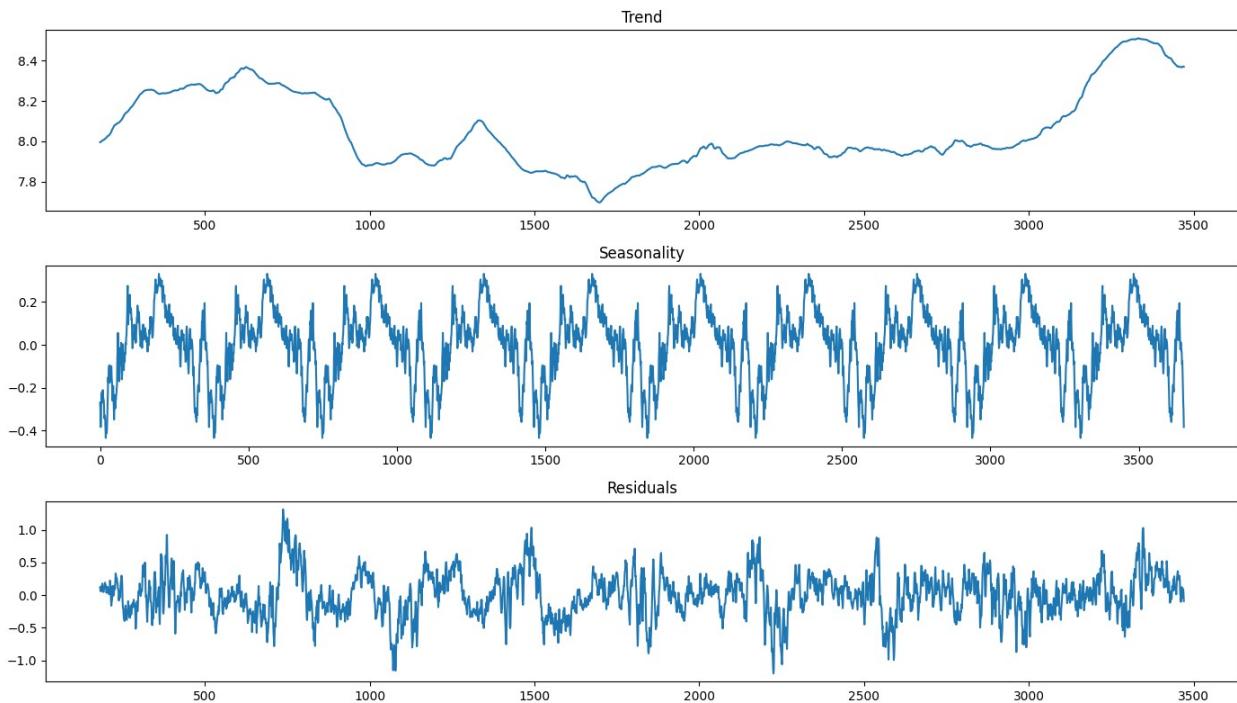
# Store Results
mae_rmse_results.append([f'{start}-{start + group_size}', 
mae_lstm, rmse_lstm])
print(f'LSTM MAE ({start}-{start + group_size}): {mae_lstm:.4f}, 
RMSE: {rmse_lstm:.4f}')

# Create DataFrame for Results
results_df = pd.DataFrame(mae_rmse_results, columns=["Time Period",
"MAE", "RMSE"])
print("\nSummary of LSTM Performance:")
print(results_df)

# Plot Results
plt.figure(figsize=(12, 6))
plt.plot(results_df["Time Period"], results_df["MAE"], marker='o',
label="MAE")
plt.plot(results_df["Time Period"], results_df["RMSE"], marker='s',
label="RMSE")
plt.xlabel("Time Period")
plt.ylabel("Error Value")
plt.title("MAE and RMSE for Each 10-Year Group using LSTM")
plt.xticks(rotation=45)
plt.legend()
plt.tight_layout()
plt.show()

```

Analyzing period: 1880-1890



```

Epoch 1/20
182/182 ━━━━━━━━ 5s 11ms/step - loss: 0.0464
Epoch 2/20
182/182 ━━━━━━━━ 2s 11ms/step - loss: 0.0066
Epoch 3/20
182/182 ━━━━━━━━ 2s 11ms/step - loss: 0.0063
Epoch 4/20
182/182 ━━━━━━━━ 2s 12ms/step - loss: 0.0044
Epoch 5/20
182/182 ━━━━━━━━ 2s 12ms/step - loss: 0.0032
Epoch 6/20
182/182 ━━━━━━━━ 2s 11ms/step - loss: 0.0023
Epoch 7/20
182/182 ━━━━━━━━ 2s 11ms/step - loss: 0.0020
Epoch 8/20
182/182 ━━━━━━━━ 2s 10ms/step - loss: 0.0020
Epoch 9/20
182/182 ━━━━━━━━ 2s 11ms/step - loss: 0.0021
Epoch 10/20
182/182 ━━━━━━━━ 2s 11ms/step - loss: 0.0016
Epoch 11/20
182/182 ━━━━━━━━ 2s 11ms/step - loss: 0.0019
Epoch 12/20
182/182 ━━━━━━━━ 2s 12ms/step - loss: 0.0016
Epoch 13/20
182/182 ━━━━━━━━ 2s 12ms/step - loss: 0.0016
Epoch 14/20
182/182 ━━━━━━━━ 2s 11ms/step - loss: 0.0017

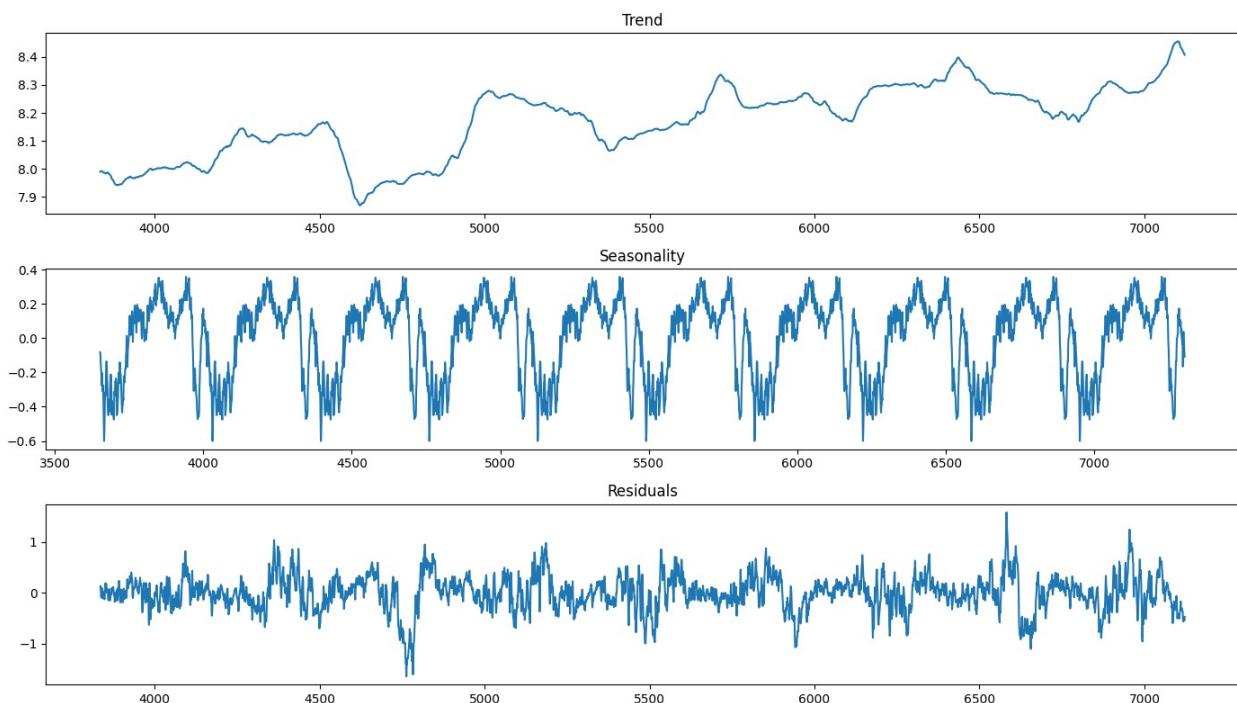
```

```

Epoch 15/20
182/182 ----- 2s 11ms/step - loss: 0.0017
Epoch 16/20
182/182 ----- 2s 11ms/step - loss: 0.0015
Epoch 17/20
182/182 ----- 2s 11ms/step - loss: 0.0016
Epoch 18/20
182/182 ----- 2s 9ms/step - loss: 0.0016
Epoch 19/20
182/182 ----- 2s 10ms/step - loss: 0.0015
Epoch 20/20
182/182 ----- 2s 11ms/step - loss: 0.0015
23/23 ----- 1s 14ms/step
LSTM MAE (1880-1890): 0.1122, RMSE: 0.1435

```

Analyzing period: 1890-1900



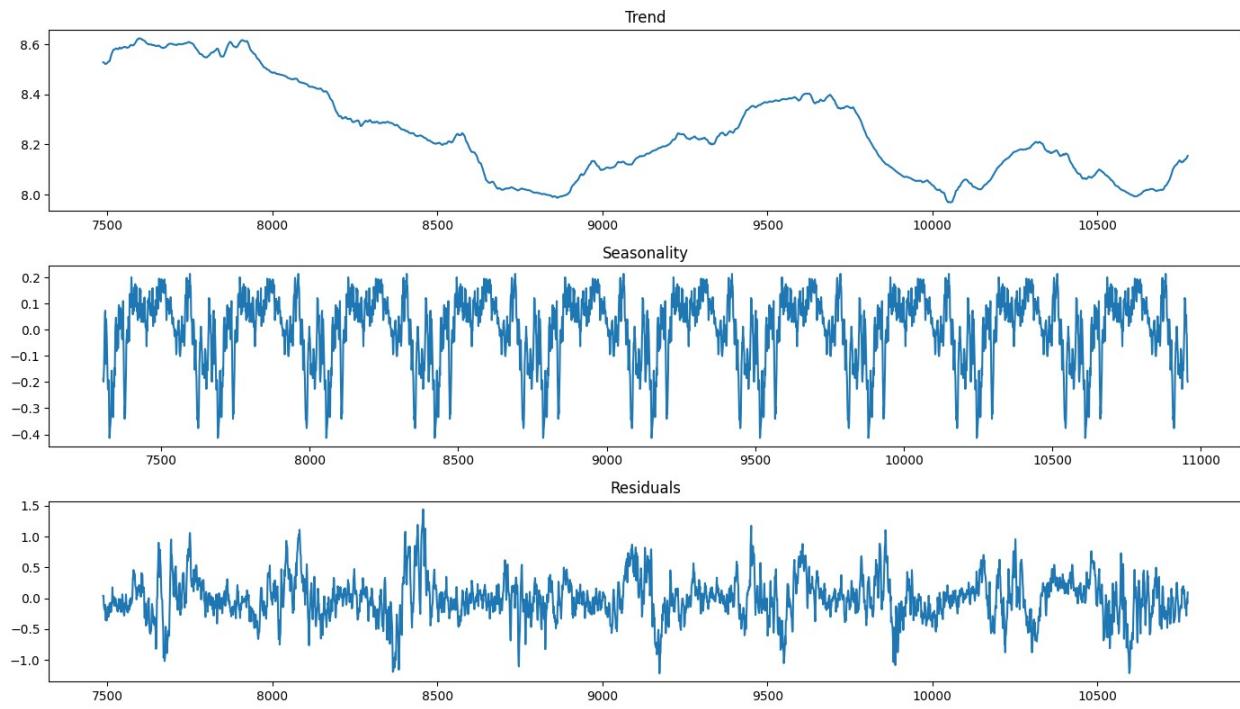
```

Epoch 1/20
182/182 ----- 4s 10ms/step - loss: 0.0894
Epoch 2/20
182/182 ----- 2s 10ms/step - loss: 0.0053
Epoch 3/20
182/182 ----- 2s 11ms/step - loss: 0.0045
Epoch 4/20
182/182 ----- 2s 11ms/step - loss: 0.0039
Epoch 5/20
182/182 ----- 2s 11ms/step - loss: 0.0030

```

```
Epoch 6/20
182/182 ━━━━━━━━ 2s 11ms/step - loss: 0.0026
Epoch 7/20
182/182 ━━━━━━━━ 2s 11ms/step - loss: 0.0022
Epoch 8/20
182/182 ━━━━━━━━ 2s 11ms/step - loss: 0.0021
Epoch 9/20
182/182 ━━━━━━━━ 2s 11ms/step - loss: 0.0015
Epoch 10/20
182/182 ━━━━━━━━ 2s 11ms/step - loss: 0.0017
Epoch 11/20
182/182 ━━━━━━━━ 2s 11ms/step - loss: 0.0016
Epoch 12/20
182/182 ━━━━━━━━ 2s 11ms/step - loss: 0.0017
Epoch 13/20
182/182 ━━━━━━━━ 2s 11ms/step - loss: 0.0015
Epoch 14/20
182/182 ━━━━━━━━ 2s 11ms/step - loss: 0.0016
Epoch 15/20
182/182 ━━━━━━━━ 2s 11ms/step - loss: 0.0016
Epoch 16/20
182/182 ━━━━━━━━ 2s 11ms/step - loss: 0.0015
Epoch 17/20
182/182 ━━━━━━━━ 2s 11ms/step - loss: 0.0015
Epoch 18/20
182/182 ━━━━━━━━ 2s 11ms/step - loss: 0.0015
Epoch 19/20
182/182 ━━━━━━━━ 2s 11ms/step - loss: 0.0020
Epoch 20/20
182/182 ━━━━━━━━ 2s 11ms/step - loss: 0.0015
23/23 ━━━━━━ 1s 14ms/step
LSTM MAE (1890-1900): 0.1150, RMSE: 0.1482
```

Analyzing period: 1900-1910



```

Epoch 1/20
182/182 ━━━━━━━━ 4s 10ms/step - loss: 0.0802
Epoch 2/20
182/182 ━━━━━━━━ 2s 10ms/step - loss: 0.0097
Epoch 3/20
182/182 ━━━━━━━━ 2s 11ms/step - loss: 0.0072
Epoch 4/20
182/182 ━━━━━━━━ 2s 11ms/step - loss: 0.0062
Epoch 5/20
182/182 ━━━━━━━━ 2s 11ms/step - loss: 0.0046
Epoch 6/20
182/182 ━━━━━━━━ 2s 11ms/step - loss: 0.0038
Epoch 7/20
182/182 ━━━━━━━━ 2s 11ms/step - loss: 0.0038
Epoch 8/20
182/182 ━━━━━━━━ 2s 11ms/step - loss: 0.0029
Epoch 9/20
182/182 ━━━━━━━━ 2s 11ms/step - loss: 0.0025
Epoch 10/20
182/182 ━━━━━━━━ 2s 11ms/step - loss: 0.0025
Epoch 11/20
182/182 ━━━━━━━━ 2s 11ms/step - loss: 0.0026
Epoch 12/20
182/182 ━━━━━━━━ 2s 11ms/step - loss: 0.0022
Epoch 13/20
182/182 ━━━━━━━━ 2s 11ms/step - loss: 0.0022
Epoch 14/20
182/182 ━━━━━━━━ 2s 11ms/step - loss: 0.0024

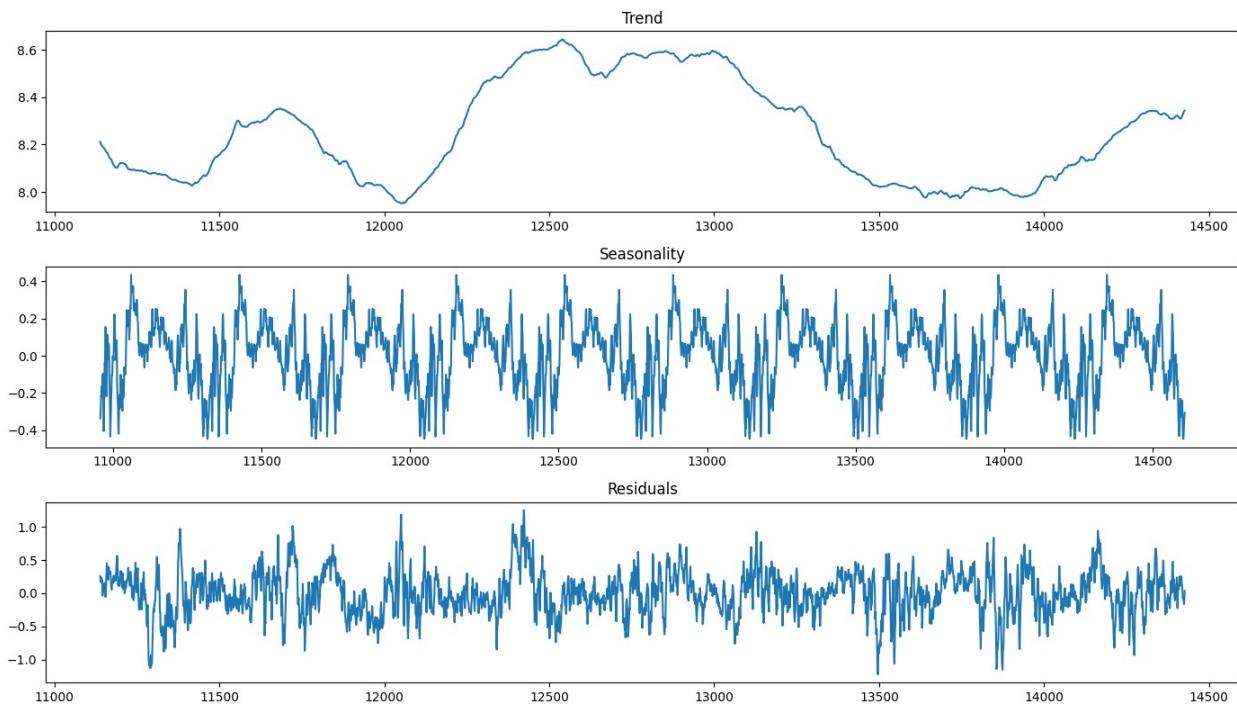
```

```

Epoch 15/20
182/182 ━━━━━━━━ 2s 11ms/step - loss: 0.0024
Epoch 16/20
182/182 ━━━━━━━━ 2s 11ms/step - loss: 0.0025
Epoch 17/20
182/182 ━━━━━━━━ 2s 11ms/step - loss: 0.0027
Epoch 18/20
182/182 ━━━━━━━━ 2s 11ms/step - loss: 0.0024
Epoch 19/20
182/182 ━━━━━━━━ 2s 11ms/step - loss: 0.0024
Epoch 20/20
182/182 ━━━━━━━━ 2s 11ms/step - loss: 0.0024
23/23 ━━━━━━ 1s 35ms/step
LSTM MAE (1900-1910): 0.1270, RMSE: 0.1642

```

Analyzing period: 1910-1920



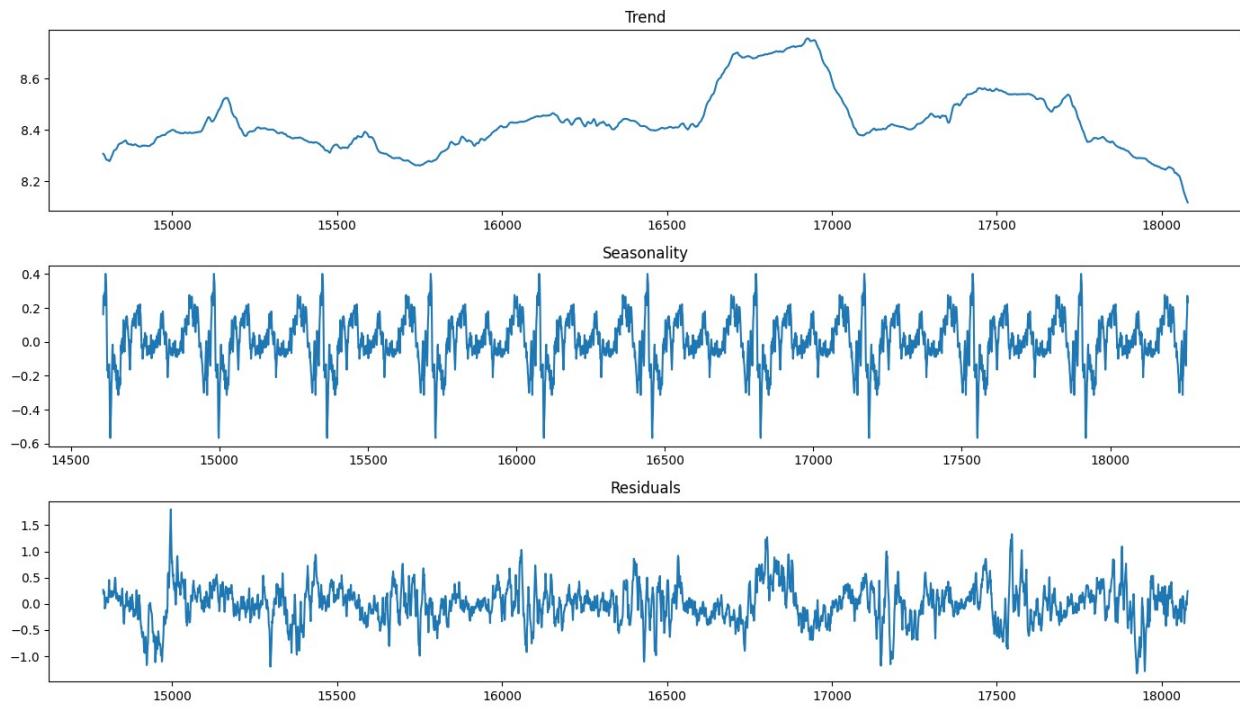
```

Epoch 1/20
182/182 ━━━━━━━━ 4s 10ms/step - loss: 0.0872
Epoch 2/20
182/182 ━━━━━━━━ 2s 10ms/step - loss: 0.0095
Epoch 3/20
182/182 ━━━━━━━━ 2s 10ms/step - loss: 0.0086
Epoch 4/20
182/182 ━━━━━━━━ 2s 11ms/step - loss: 0.0069
Epoch 5/20
182/182 ━━━━━━━━ 2s 11ms/step - loss: 0.0063

```

```
Epoch 6/20
182/182 ━━━━━━━━ 2s 11ms/step - loss: 0.0048
Epoch 7/20
182/182 ━━━━━━━━ 2s 11ms/step - loss: 0.0031
Epoch 8/20
182/182 ━━━━━━━━ 2s 11ms/step - loss: 0.0024
Epoch 9/20
182/182 ━━━━━━━━ 2s 11ms/step - loss: 0.0022
Epoch 10/20
182/182 ━━━━━━━━ 2s 11ms/step - loss: 0.0023
Epoch 11/20
182/182 ━━━━━━━━ 2s 11ms/step - loss: 0.0020
Epoch 12/20
182/182 ━━━━━━━━ 2s 11ms/step - loss: 0.0019
Epoch 13/20
182/182 ━━━━━━━━ 2s 11ms/step - loss: 0.0019
Epoch 14/20
182/182 ━━━━━━━━ 2s 12ms/step - loss: 0.0024
Epoch 15/20
182/182 ━━━━━━━━ 2s 12ms/step - loss: 0.0020
Epoch 16/20
182/182 ━━━━━━━━ 2s 12ms/step - loss: 0.0019
Epoch 17/20
182/182 ━━━━━━━━ 2s 12ms/step - loss: 0.0023
Epoch 18/20
182/182 ━━━━━━━━ 2s 12ms/step - loss: 0.0021
Epoch 19/20
182/182 ━━━━━━━━ 2s 10ms/step - loss: 0.0021
Epoch 20/20
182/182 ━━━━━━━━ 2s 12ms/step - loss: 0.0020
23/23 ━━━━━━━━ 0s 12ms/step
LSTM MAE (1910-1920): 0.1140, RMSE: 0.1437
```

Analyzing period: 1920-1930



```

Epoch 1/20
182/182 ━━━━━━━━ 4s 13ms/step - loss: 0.0709
Epoch 2/20
182/182 ━━━━━━━━ 2s 11ms/step - loss: 0.0073
Epoch 3/20
182/182 ━━━━━━━━ 2s 12ms/step - loss: 0.0059
Epoch 4/20
182/182 ━━━━━━━━ 2s 11ms/step - loss: 0.0042
Epoch 5/20
182/182 ━━━━━━━━ 2s 11ms/step - loss: 0.0032
Epoch 6/20
182/182 ━━━━━━━━ 2s 11ms/step - loss: 0.0024
Epoch 7/20
182/182 ━━━━━━━━ 2s 10ms/step - loss: 0.0019
Epoch 8/20
182/182 ━━━━━━━━ 2s 11ms/step - loss: 0.0022
Epoch 9/20
182/182 ━━━━━━━━ 2s 12ms/step - loss: 0.0017
Epoch 10/20
182/182 ━━━━━━━━ 2s 12ms/step - loss: 0.0017
Epoch 11/20
182/182 ━━━━━━━━ 2s 12ms/step - loss: 0.0021
Epoch 12/20
182/182 ━━━━━━━━ 2s 13ms/step - loss: 0.0015
Epoch 13/20
182/182 ━━━━━━━━ 2s 13ms/step - loss: 0.0017
Epoch 14/20
182/182 ━━━━━━━━ 2s 12ms/step - loss: 0.0015

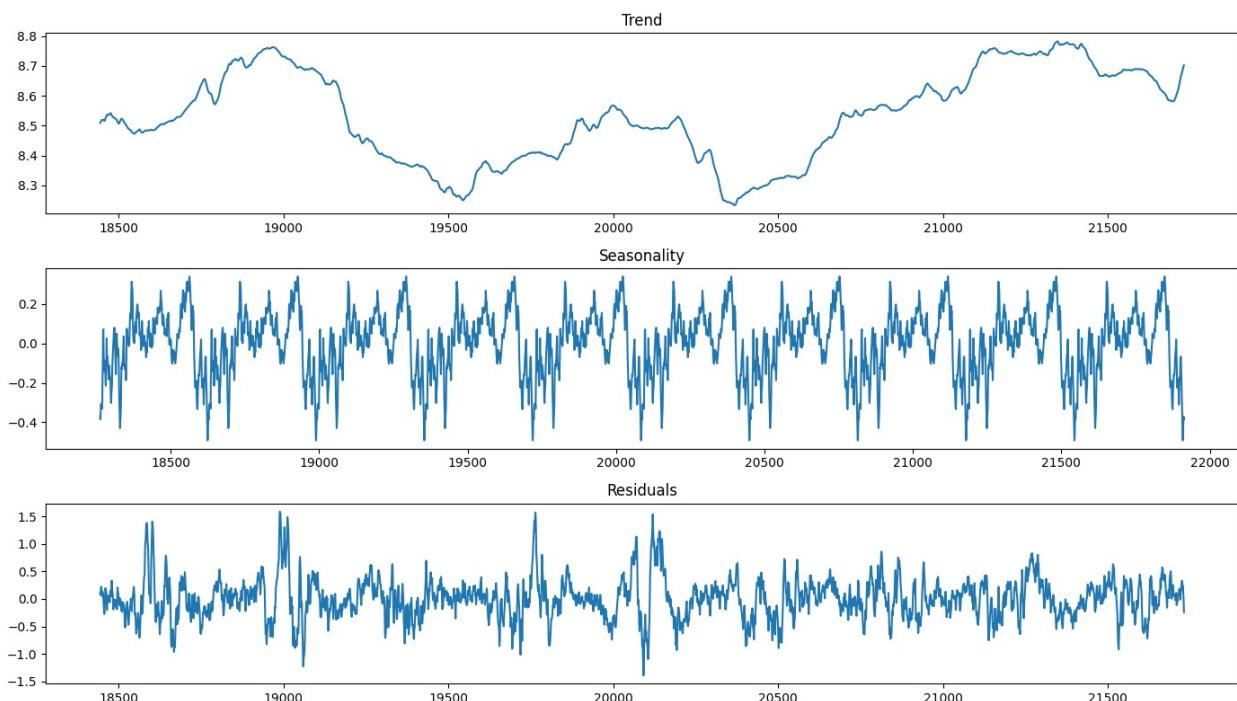
```

```

Epoch 15/20
182/182 ━━━━━━━━ 2s 12ms/step - loss: 0.0016
Epoch 16/20
182/182 ━━━━━━━━ 2s 11ms/step - loss: 0.0015
Epoch 17/20
182/182 ━━━━━━━━ 2s 10ms/step - loss: 0.0016
Epoch 18/20
182/182 ━━━━━━━━ 2s 10ms/step - loss: 0.0016
Epoch 19/20
182/182 ━━━━━━━━ 2s 10ms/step - loss: 0.0014
Epoch 20/20
182/182 ━━━━━━━━ 2s 11ms/step - loss: 0.0016
23/23 ━━━━━━ 1s 16ms/step
LSTM MAE (1920-1930): 0.1097, RMSE: 0.1429

```

Analyzing period: 1930-1940



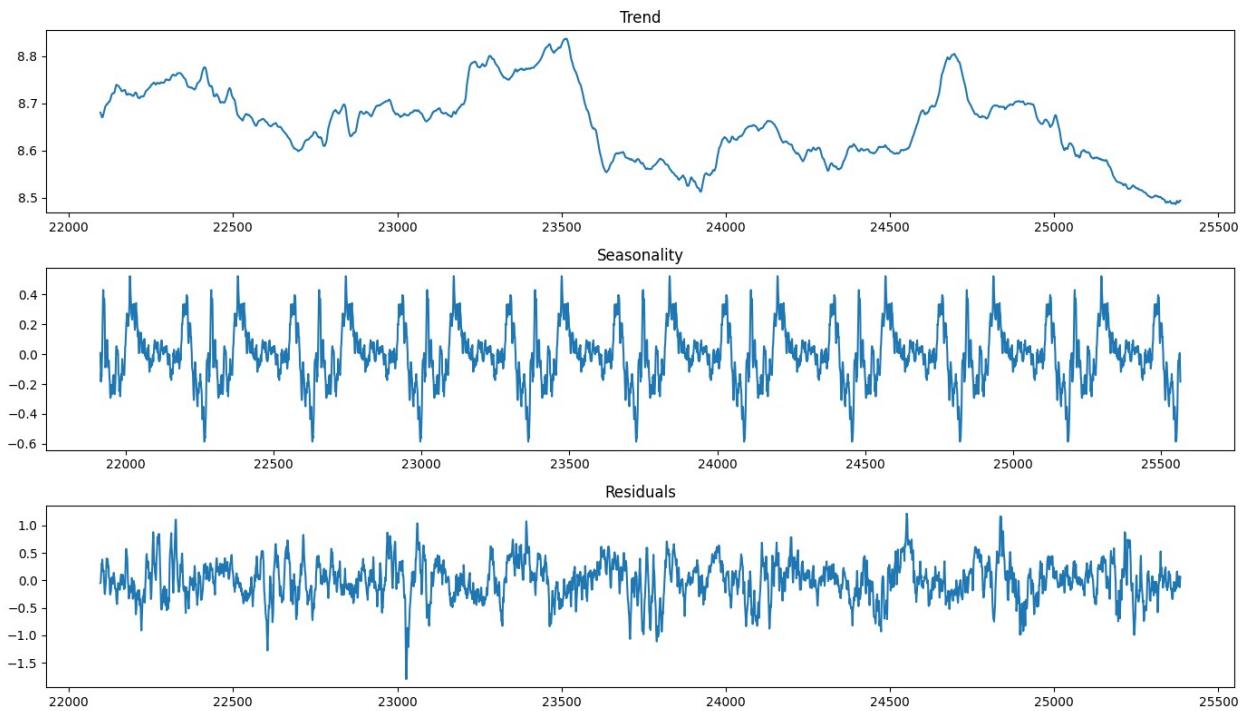
```

Epoch 1/20
182/182 ━━━━━━━━ 5s 12ms/step - loss: 0.0624
Epoch 2/20
182/182 ━━━━━━━━ 2s 13ms/step - loss: 0.0102
Epoch 3/20
182/182 ━━━━━━━━ 2s 11ms/step - loss: 0.0092
Epoch 4/20
182/182 ━━━━━━━━ 2s 10ms/step - loss: 0.0062
Epoch 5/20
182/182 ━━━━━━━━ 2s 10ms/step - loss: 0.0052

```

```
Epoch 6/20
182/182 ━━━━━━━━ 2s 12ms/step - loss: 0.0041
Epoch 7/20
182/182 ━━━━━━ 2s 12ms/step - loss: 0.0038
Epoch 8/20
182/182 ━━━━━━ 2s 11ms/step - loss: 0.0024
Epoch 9/20
182/182 ━━━━ 2s 10ms/step - loss: 0.0024
Epoch 10/20
182/182 ━━━━ 2s 10ms/step - loss: 0.0020
Epoch 11/20
182/182 ━━━━ 2s 11ms/step - loss: 0.0022
Epoch 12/20
182/182 ━━━━ 2s 12ms/step - loss: 0.0023
Epoch 13/20
182/182 ━━━━ 2s 11ms/step - loss: 0.0017
Epoch 14/20
182/182 ━━━━ 2s 11ms/step - loss: 0.0017
Epoch 15/20
182/182 ━━━━ 2s 12ms/step - loss: 0.0017
Epoch 16/20
182/182 ━━━━ 2s 12ms/step - loss: 0.0017
Epoch 17/20
182/182 ━━━━ 2s 12ms/step - loss: 0.0016
Epoch 18/20
182/182 ━━━━ 2s 12ms/step - loss: 0.0015
Epoch 19/20
182/182 ━━━━ 2s 12ms/step - loss: 0.0015
Epoch 20/20
182/182 ━━━━ 2s 10ms/step - loss: 0.0016
23/23 ━━━━ 0s 12ms/step
LSTM MAE (1930-1940): 0.0949, RMSE: 0.1206
```

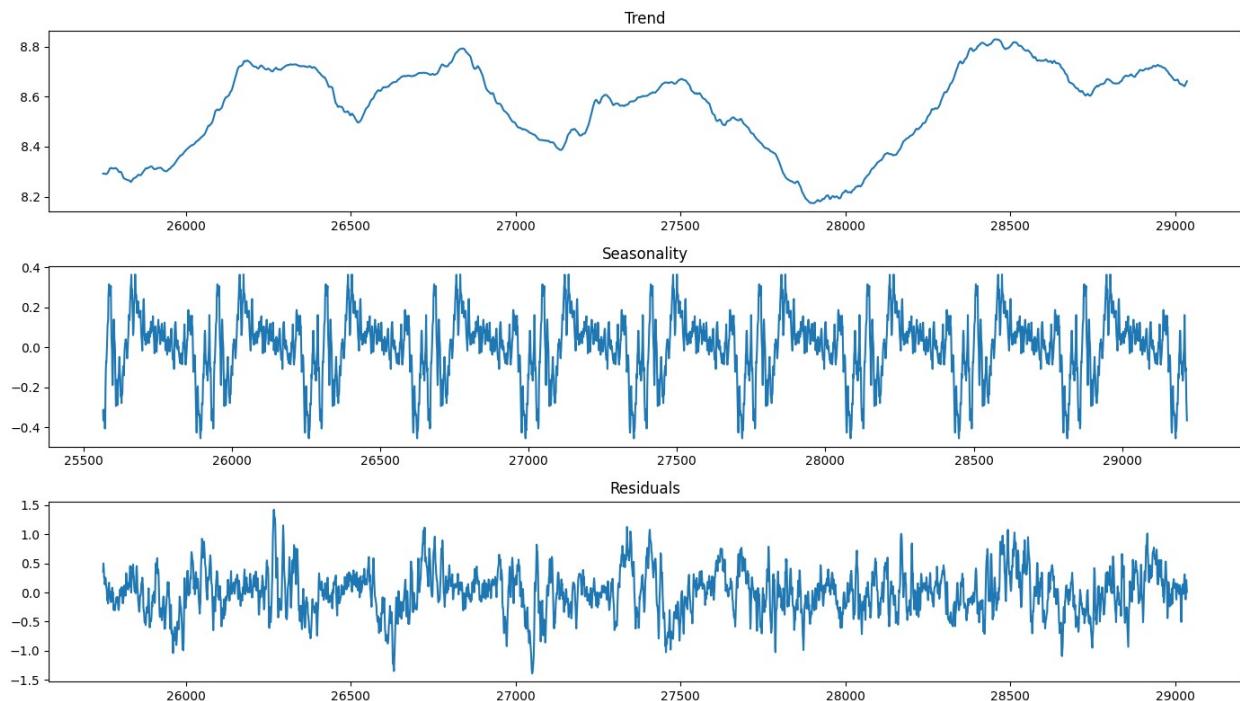
Analyzing period: 1940-1950



Epoch 1/20	182/182	4s 10ms/step - loss: 0.0917
Epoch 2/20	182/182	2s 11ms/step - loss: 0.0099
Epoch 3/20	182/182	2s 11ms/step - loss: 0.0078
Epoch 4/20	182/182	2s 10ms/step - loss: 0.0054
Epoch 5/20	182/182	2s 11ms/step - loss: 0.0041
Epoch 6/20	182/182	2s 12ms/step - loss: 0.0030
Epoch 7/20	182/182	3s 14ms/step - loss: 0.0026
Epoch 8/20	182/182	2s 12ms/step - loss: 0.0026
Epoch 9/20	182/182	2s 11ms/step - loss: 0.0022
Epoch 10/20	182/182	2s 10ms/step - loss: 0.0020
Epoch 11/20	182/182	2s 10ms/step - loss: 0.0018
Epoch 12/20	182/182	2s 11ms/step - loss: 0.0019
Epoch 13/20	182/182	2s 12ms/step - loss: 0.0023
Epoch 14/20	182/182	2s 14ms/step - loss: 0.0018

```
Epoch 15/20
182/182 ━━━━━━━━ 2s 10ms/step - loss: 0.0018
Epoch 16/20
182/182 ━━━━━━━━ 2s 10ms/step - loss: 0.0018
Epoch 17/20
182/182 ━━━━━━━━ 2s 10ms/step - loss: 0.0016
Epoch 18/20
182/182 ━━━━━━━━ 2s 11ms/step - loss: 0.0016
Epoch 19/20
182/182 ━━━━━━━━ 2s 12ms/step - loss: 0.0019
Epoch 20/20
182/182 ━━━━━━━━ 2s 12ms/step - loss: 0.0019
23/23 ━━━━━━ 1s 14ms/step
LSTM MAE (1940-1950): 0.1159, RMSE: 0.1465
```

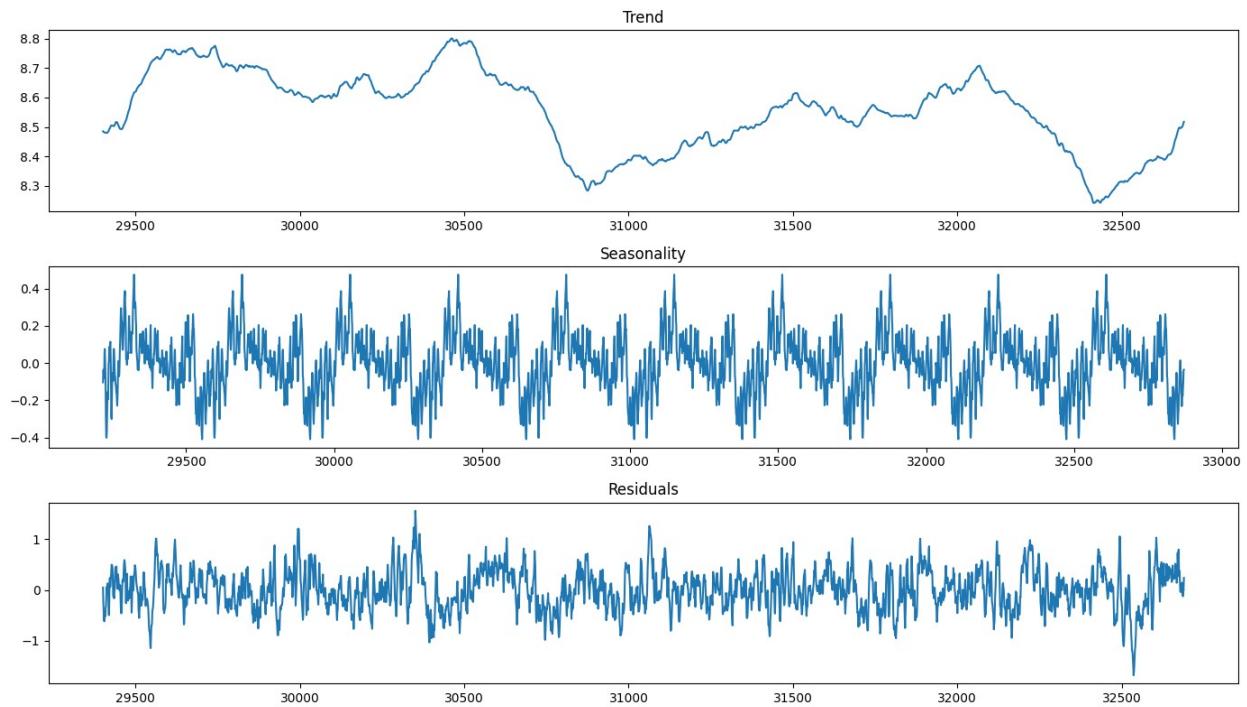
Analyzing period: 1950-1960



```
Epoch 1/20
182/182 ━━━━━━━━ 4s 9ms/step - loss: 0.0692
Epoch 2/20
182/182 ━━━━━━━━ 2s 10ms/step - loss: 0.0111
Epoch 3/20
182/182 ━━━━━━━━ 2s 11ms/step - loss: 0.0076
Epoch 4/20
182/182 ━━━━━━━━ 2s 11ms/step - loss: 0.0055
Epoch 5/20
182/182 ━━━━━━━━ 2s 11ms/step - loss: 0.0045
```

```
Epoch 6/20
182/182 ━━━━━━━━ 2s 11ms/step - loss: 0.0035
Epoch 7/20
182/182 ━━━━━━━━ 2s 11ms/step - loss: 0.0027
Epoch 8/20
182/182 ━━━━━━━━ 2s 11ms/step - loss: 0.0025
Epoch 9/20
182/182 ━━━━━━━━ 2s 11ms/step - loss: 0.0022
Epoch 10/20
182/182 ━━━━━━━━ 2s 11ms/step - loss: 0.0021
Epoch 11/20
182/182 ━━━━━━━━ 2s 11ms/step - loss: 0.0022
Epoch 12/20
182/182 ━━━━━━━━ 2s 11ms/step - loss: 0.0021
Epoch 13/20
182/182 ━━━━━━━━ 2s 11ms/step - loss: 0.0021
Epoch 14/20
182/182 ━━━━━━━━ 2s 11ms/step - loss: 0.0018
Epoch 15/20
182/182 ━━━━━━━━ 2s 11ms/step - loss: 0.0019
Epoch 16/20
182/182 ━━━━━━━━ 2s 11ms/step - loss: 0.0019
Epoch 17/20
182/182 ━━━━━━━━ 2s 11ms/step - loss: 0.0019
Epoch 18/20
182/182 ━━━━━━━━ 2s 11ms/step - loss: 0.0018
Epoch 19/20
182/182 ━━━━━━━━ 2s 11ms/step - loss: 0.0021
Epoch 20/20
182/182 ━━━━━━━━ 2s 11ms/step - loss: 0.0019
23/23 ━━━━━━━━ 1s 14ms/step
LSTM MAE (1950-1960): 0.1342, RMSE: 0.1706
```

Analyzing period: 1960-1970



```

Epoch 1/20
182/182 ━━━━━━━━ 4s 11ms/step - loss: 0.0446
Epoch 2/20
182/182 ━━━━━━━━ 2s 11ms/step - loss: 0.0102
Epoch 3/20
182/182 ━━━━━━━━ 2s 11ms/step - loss: 0.0090
Epoch 4/20
182/182 ━━━━━━━━ 2s 11ms/step - loss: 0.0086
Epoch 5/20
182/182 ━━━━━━━━ 2s 11ms/step - loss: 0.0076
Epoch 6/20
182/182 ━━━━━━━━ 2s 11ms/step - loss: 0.0068
Epoch 7/20
182/182 ━━━━━━━━ 2s 11ms/step - loss: 0.0048
Epoch 8/20
182/182 ━━━━━━━━ 2s 11ms/step - loss: 0.0038
Epoch 9/20
182/182 ━━━━━━━━ 2s 11ms/step - loss: 0.0031
Epoch 10/20
182/182 ━━━━━━━━ 2s 12ms/step - loss: 0.0025
Epoch 11/20
182/182 ━━━━━━━━ 2s 13ms/step - loss: 0.0023
Epoch 12/20
182/182 ━━━━━━━━ 2s 11ms/step - loss: 0.0019
Epoch 13/20
182/182 ━━━━━━━━ 2s 11ms/step - loss: 0.0018
Epoch 14/20
182/182 ━━━━━━━━ 2s 11ms/step - loss: 0.0018

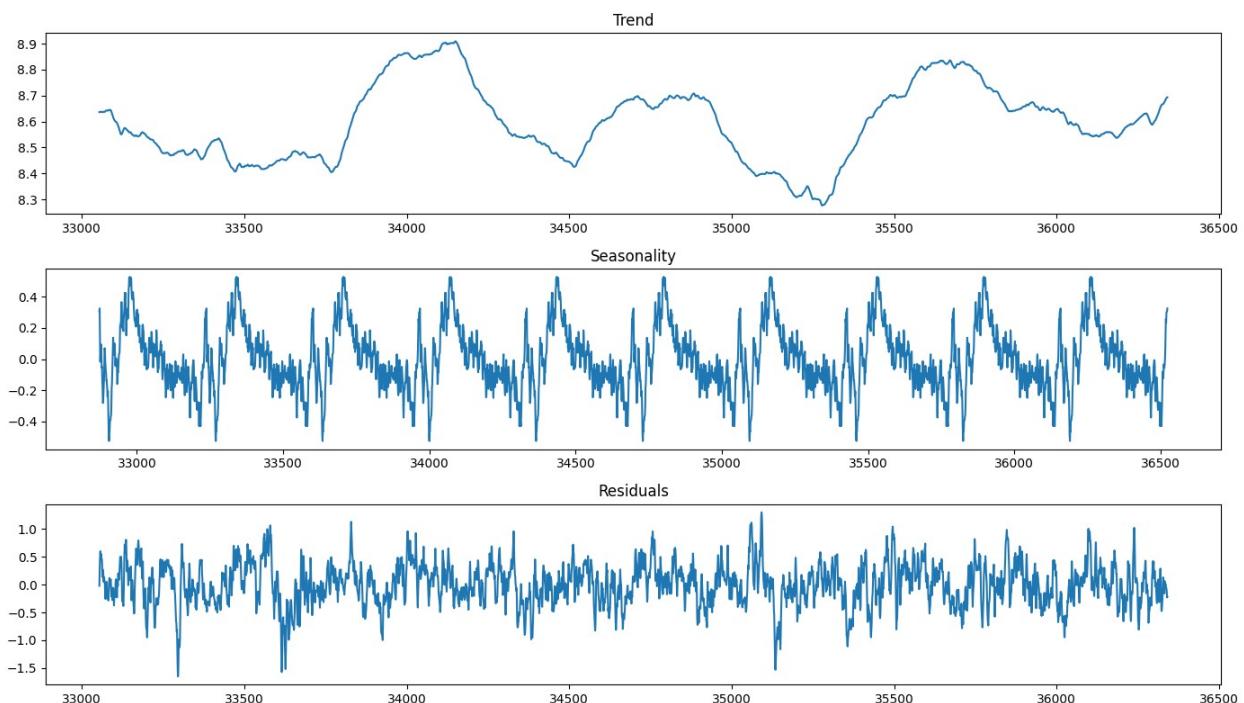
```

```

Epoch 15/20
182/182 ━━━━━━━━ 2s 10ms/step - loss: 0.0017
Epoch 16/20
182/182 ━━━━━━━━ 2s 9ms/step - loss: 0.0016
Epoch 17/20
182/182 ━━━━━━━━ 2s 10ms/step - loss: 0.0017
Epoch 18/20
182/182 ━━━━━━━━ 2s 11ms/step - loss: 0.0017
Epoch 19/20
182/182 ━━━━━━━━ 2s 11ms/step - loss: 0.0016
Epoch 20/20
182/182 ━━━━━━━━ 2s 11ms/step - loss: 0.0017
23/23 ━━━━━━ 1s 14ms/step
LSTM MAE (1960-1970): 0.1244, RMSE: 0.1596

```

Analyzing period: 1970-1980



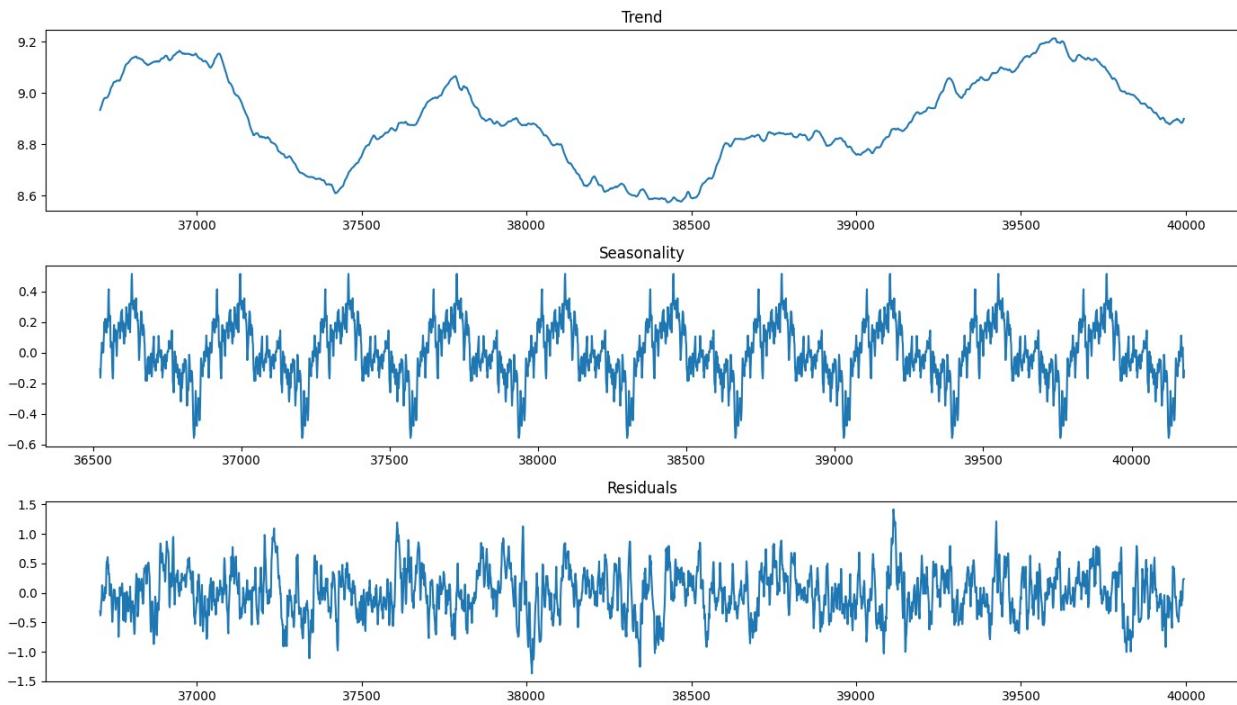
```

Epoch 1/20
182/182 ━━━━━━━━ 4s 11ms/step - loss: 0.0824
Epoch 2/20
182/182 ━━━━━━━━ 2s 11ms/step - loss: 0.0103
Epoch 3/20
182/182 ━━━━━━━━ 2s 11ms/step - loss: 0.0076
Epoch 4/20
182/182 ━━━━━━━━ 2s 11ms/step - loss: 0.0059
Epoch 5/20
182/182 ━━━━━━━━ 2s 11ms/step - loss: 0.0042

```

```
Epoch 6/20
182/182 ━━━━━━━━ 2s 11ms/step - loss: 0.0038
Epoch 7/20
182/182 ━━━━━━━━ 2s 11ms/step - loss: 0.0031
Epoch 8/20
182/182 ━━━━━━━━ 2s 11ms/step - loss: 0.0026
Epoch 9/20
182/182 ━━━━━━━━ 2s 11ms/step - loss: 0.0020
Epoch 10/20
182/182 ━━━━━━━━ 2s 11ms/step - loss: 0.0031
Epoch 11/20
182/182 ━━━━━━━━ 2s 11ms/step - loss: 0.0024
Epoch 12/20
182/182 ━━━━━━━━ 2s 11ms/step - loss: 0.0020
Epoch 13/20
182/182 ━━━━━━━━ 2s 11ms/step - loss: 0.0020
Epoch 14/20
182/182 ━━━━━━━━ 2s 11ms/step - loss: 0.0020
Epoch 15/20
182/182 ━━━━━━━━ 2s 12ms/step - loss: 0.0022
Epoch 16/20
182/182 ━━━━━━━━ 2s 11ms/step - loss: 0.0020
Epoch 17/20
182/182 ━━━━━━━━ 2s 11ms/step - loss: 0.0022
Epoch 18/20
182/182 ━━━━━━━━ 2s 11ms/step - loss: 0.0019
Epoch 19/20
182/182 ━━━━━━━━ 2s 12ms/step - loss: 0.0021
Epoch 20/20
182/182 ━━━━━━━━ 2s 12ms/step - loss: 0.0020
23/23 ━━━━━━━━ 1s 36ms/step
LSTM MAE (1970-1980): 0.1208, RMSE: 0.1495
```

Analyzing period: 1980-1990



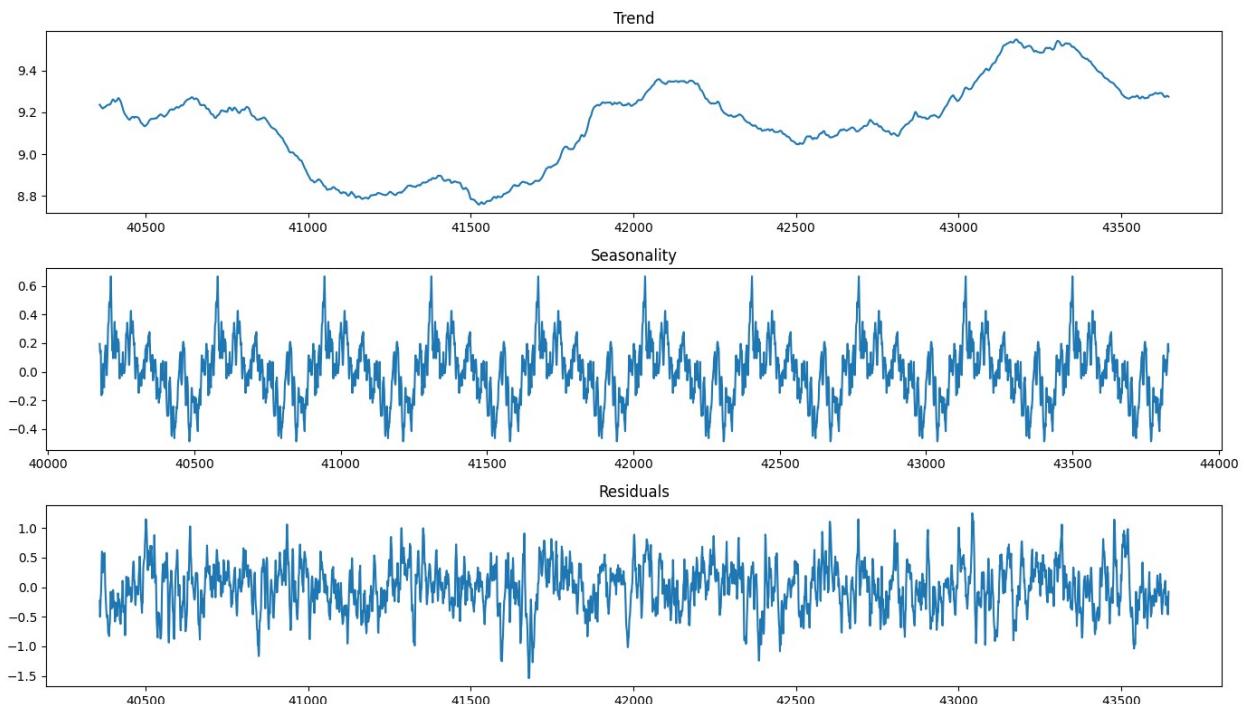
```

Epoch 1/20
182/182 ━━━━━━━━ 4s 9ms/step - loss: 0.0805
Epoch 2/20
182/182 ━━━━━━━━ 2s 10ms/step - loss: 0.0123
Epoch 3/20
182/182 ━━━━━━━━ 2s 11ms/step - loss: 0.0077
Epoch 4/20
182/182 ━━━━━━━━ 2s 12ms/step - loss: 0.0058
Epoch 5/20
182/182 ━━━━━━━━ 2s 10ms/step - loss: 0.0036
Epoch 6/20
182/182 ━━━━━━━━ 2s 11ms/step - loss: 0.0033
Epoch 7/20
182/182 ━━━━━━━━ 2s 10ms/step - loss: 0.0033
Epoch 8/20
182/182 ━━━━━━━━ 2s 10ms/step - loss: 0.0027
Epoch 9/20
182/182 ━━━━━━━━ 2s 10ms/step - loss: 0.0029
Epoch 10/20
182/182 ━━━━━━━━ 2s 10ms/step - loss: 0.0023
Epoch 11/20
182/182 ━━━━━━━━ 2s 10ms/step - loss: 0.0021
Epoch 12/20
182/182 ━━━━━━━━ 2s 11ms/step - loss: 0.0022
Epoch 13/20
182/182 ━━━━━━━━ 2s 11ms/step - loss: 0.0023
Epoch 14/20
182/182 ━━━━━━━━ 2s 11ms/step - loss: 0.0025

```

```
Epoch 15/20
182/182 ----- 2s 11ms/step - loss: 0.0023
Epoch 16/20
182/182 ----- 2s 11ms/step - loss: 0.0022
Epoch 17/20
182/182 ----- 2s 11ms/step - loss: 0.0022
Epoch 18/20
182/182 ----- 2s 11ms/step - loss: 0.0022
Epoch 19/20
182/182 ----- 2s 11ms/step - loss: 0.0023
Epoch 20/20
182/182 ----- 2s 11ms/step - loss: 0.0021
23/23 ----- 1s 14ms/step
LSTM MAE (1980-1990): 0.1397, RMSE: 0.1738
```

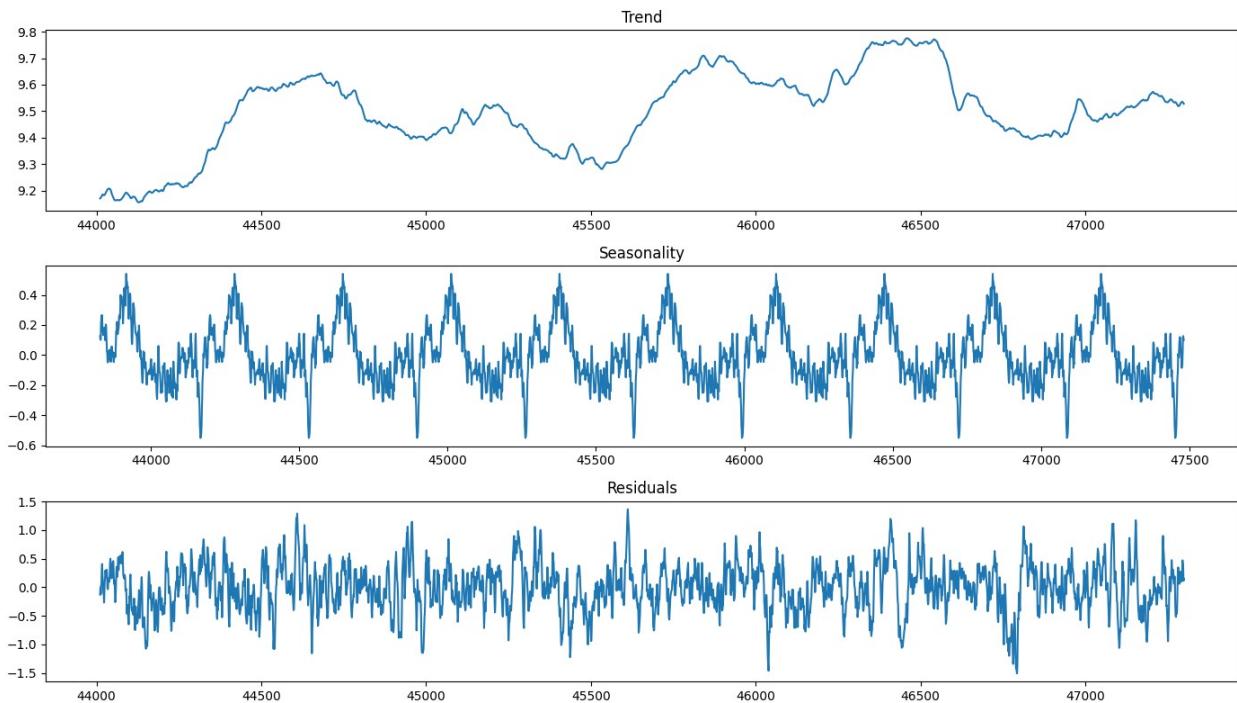
Analyzing period: 1990-2000



```
Epoch 1/20
182/182 ----- 4s 11ms/step - loss: 0.0741
Epoch 2/20
182/182 ----- 2s 11ms/step - loss: 0.0112
Epoch 3/20
182/182 ----- 2s 11ms/step - loss: 0.0089
Epoch 4/20
182/182 ----- 2s 11ms/step - loss: 0.0073
Epoch 5/20
182/182 ----- 2s 11ms/step - loss: 0.0048
```

```
Epoch 6/20
182/182 ----- 2s 11ms/step - loss: 0.0038
Epoch 7/20
182/182 ----- 2s 11ms/step - loss: 0.0032
Epoch 8/20
182/182 ----- 2s 12ms/step - loss: 0.0029
Epoch 9/20
182/182 ----- 2s 12ms/step - loss: 0.0026
Epoch 10/20
182/182 ----- 2s 11ms/step - loss: 0.0021
Epoch 11/20
182/182 ----- 2s 10ms/step - loss: 0.0024
Epoch 12/20
182/182 ----- 2s 10ms/step - loss: 0.0021
Epoch 13/20
182/182 ----- 2s 11ms/step - loss: 0.0020
Epoch 14/20
182/182 ----- 2s 11ms/step - loss: 0.0019
Epoch 15/20
182/182 ----- 2s 11ms/step - loss: 0.0019
Epoch 16/20
182/182 ----- 2s 11ms/step - loss: 0.0021
Epoch 17/20
182/182 ----- 2s 11ms/step - loss: 0.0028
Epoch 18/20
182/182 ----- 2s 11ms/step - loss: 0.0020
Epoch 19/20
182/182 ----- 2s 11ms/step - loss: 0.0019
Epoch 20/20
182/182 ----- 2s 11ms/step - loss: 0.0019
23/23 ----- 1s 14ms/step
LSTM MAE (1990-2000): 0.1223, RMSE: 0.1556
```

Analyzing period: 2000-2010



```

Epoch 1/20
182/182 ━━━━━━━━ 4s 11ms/step - loss: 0.0526
Epoch 2/20
182/182 ━━━━━━━━ 2s 11ms/step - loss: 0.0107
Epoch 3/20
182/182 ━━━━━━━━ 2s 11ms/step - loss: 0.0090
Epoch 4/20
182/182 ━━━━━━━━ 2s 11ms/step - loss: 0.0059
Epoch 5/20
182/182 ━━━━━━━━ 2s 11ms/step - loss: 0.0042
Epoch 6/20
182/182 ━━━━━━━━ 2s 11ms/step - loss: 0.0032
Epoch 7/20
182/182 ━━━━━━━━ 2s 11ms/step - loss: 0.0028
Epoch 8/20
182/182 ━━━━━━━━ 2s 11ms/step - loss: 0.0024
Epoch 9/20
182/182 ━━━━━━━━ 2s 12ms/step - loss: 0.0021
Epoch 10/20
182/182 ━━━━━━━━ 2s 12ms/step - loss: 0.0022
Epoch 11/20
182/182 ━━━━━━━━ 2s 11ms/step - loss: 0.0022
Epoch 12/20
182/182 ━━━━━━━━ 2s 11ms/step - loss: 0.0021
Epoch 13/20
182/182 ━━━━━━━━ 2s 11ms/step - loss: 0.0021
Epoch 14/20
182/182 ━━━━━━━━ 2s 10ms/step - loss: 0.0021

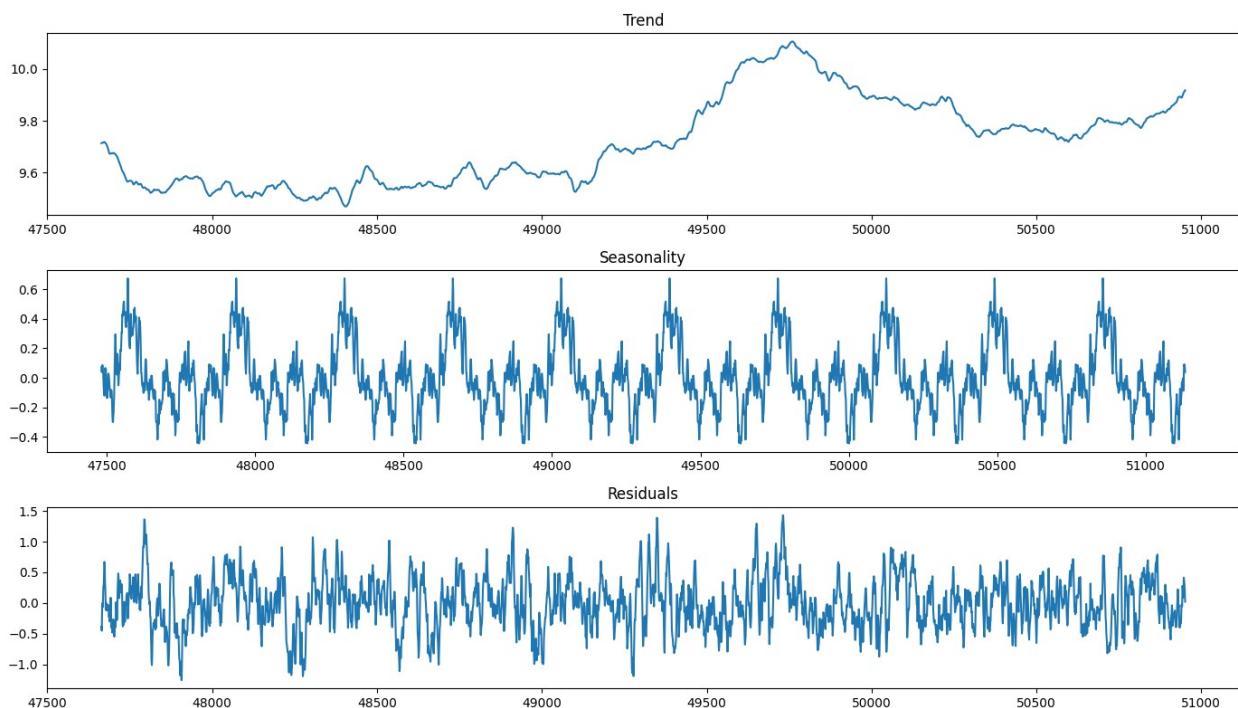
```

```

Epoch 15/20
182/182 ----- 2s 9ms/step - loss: 0.0021
Epoch 16/20
182/182 ----- 2s 9ms/step - loss: 0.0022
Epoch 17/20
182/182 ----- 2s 11ms/step - loss: 0.0021
Epoch 18/20
182/182 ----- 2s 11ms/step - loss: 0.0021
Epoch 19/20
182/182 ----- 2s 11ms/step - loss: 0.0026
Epoch 20/20
182/182 ----- 2s 11ms/step - loss: 0.0021
23/23 ----- 1s 15ms/step
LSTM MAE (2000-2010): 0.1109, RMSE: 0.1401

```

Analyzing period: 2010-2020



```

Epoch 1/20
182/182 ----- 4s 11ms/step - loss: 0.0759
Epoch 2/20
182/182 ----- 2s 10ms/step - loss: 0.0119
Epoch 3/20
182/182 ----- 2s 11ms/step - loss: 0.0085
Epoch 4/20
182/182 ----- 2s 11ms/step - loss: 0.0051
Epoch 5/20
182/182 ----- 2s 11ms/step - loss: 0.0033

```

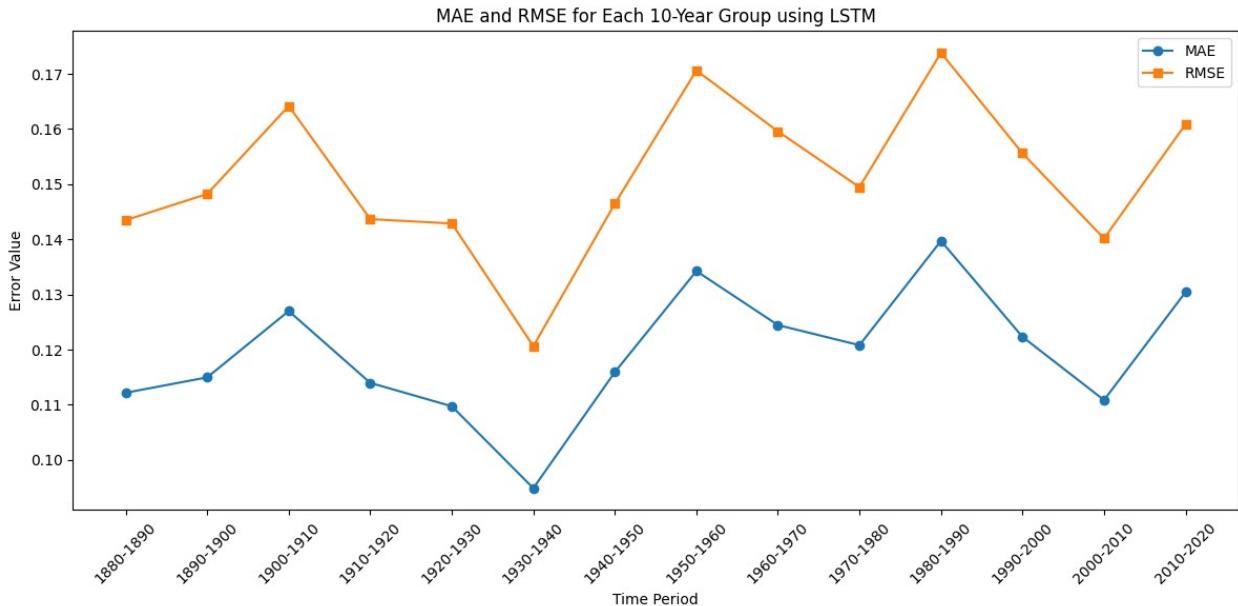
```

Epoch 6/20
182/182 ━━━━━━━━ 2s 11ms/step - loss: 0.0028
Epoch 7/20
182/182 ━━━━━━━━ 2s 11ms/step - loss: 0.0040
Epoch 8/20
182/182 ━━━━━━━━ 2s 11ms/step - loss: 0.0020
Epoch 9/20
182/182 ━━━━━━━━ 2s 11ms/step - loss: 0.0025
Epoch 10/20
182/182 ━━━━━━━━ 2s 11ms/step - loss: 0.0019
Epoch 11/20
182/182 ━━━━━━━━ 2s 11ms/step - loss: 0.0022
Epoch 12/20
182/182 ━━━━━━━━ 2s 11ms/step - loss: 0.0020
Epoch 13/20
182/182 ━━━━━━━━ 2s 11ms/step - loss: 0.0020
Epoch 14/20
182/182 ━━━━━━━━ 2s 11ms/step - loss: 0.0021
Epoch 15/20
182/182 ━━━━━━━━ 2s 11ms/step - loss: 0.0021
Epoch 16/20
182/182 ━━━━━━━━ 2s 11ms/step - loss: 0.0018
Epoch 17/20
182/182 ━━━━━━━━ 2s 11ms/step - loss: 0.0019
Epoch 18/20
182/182 ━━━━━━━━ 2s 11ms/step - loss: 0.0020
Epoch 19/20
182/182 ━━━━━━━━ 2s 11ms/step - loss: 0.0019
Epoch 20/20
182/182 ━━━━━━━━ 2s 11ms/step - loss: 0.0022
23/23 ━━━━━━━━ 1s 14ms/step
LSTM MAE (2010-2020): 0.1304, RMSE: 0.1609

```

Summary of LSTM Performance:

	Time Period	MAE	RMSE
0	1880-1890	0.112176	0.143471
1	1890-1900	0.114968	0.148225
2	1900-1910	0.126996	0.164180
3	1910-1920	0.113969	0.143658
4	1920-1930	0.109726	0.142871
5	1930-1940	0.094874	0.120625
6	1940-1950	0.115922	0.146455
7	1950-1960	0.134238	0.170619
8	1960-1970	0.124446	0.159572
9	1970-1980	0.120801	0.149459
10	1980-1990	0.139709	0.173818
11	1990-2000	0.122313	0.155600
12	2000-2010	0.110852	0.140144
13	2010-2020	0.130430	0.160865



```
#yearly Avg

# Calculate yearly average
df_yearly = df.groupby('Year')[['Temperature']].mean().reset_index()

# Calculate 5-year average
df['5YearGroup'] = (df['Year'] // 5) * 5
df_5year = df.groupby('5YearGroup')[['Temperature']].mean().reset_index()

from statsmodels.tsa.seasonal import seasonal_decompose

def train_lstm_on_series(series, label=""):
    # ----- Decomposition -----
    try:
        decomposition = seasonal_decompose(series, model='additive',
period=1)
        fig, axs = plt.subplots(3, 1, figsize=(12, 8))
        decomposition.trend.plot(ax=axs[0], title='Trend')
        decomposition.seasonal.plot(ax=axs[1], title='Seasonality')
        decomposition.resid.plot(ax=axs[2], title='Residuals')
        plt.suptitle(f"Decomposition Analysis - {label}", fontsize=14)
        plt.tight_layout()
        plt.show()
    except Exception as e:
        print(f"Decomposition failed: {e}")

    # ----- Normalize -----
    scaler = MinMaxScaler()
    series_scaled = scaler.fit_transform(series.values.reshape(-1, 1))
```

```

# ----- Create Sequences -----
seq_length = 5
X, y = create_sequences(series_scaled, seq_length)
X = X.reshape((X.shape[0], X.shape[1], 1))

# ----- Train-test Split -----
X_train, X_test, y_train, y_test = train_test_split(X, y,
test_size=0.2, shuffle=False)

# ----- Build and Train LSTM -----
model = Sequential([
    LSTM(50, activation='relu', return_sequences=True,
input_shape=(seq_length, 1)),
    LSTM(50, activation='relu'),
    Dense(1)
])
model.compile(optimizer='adam', loss='mse')
model.fit(X_train, y_train, epochs=50, batch_size=8, verbose=1)

# ----- Prediction -----
y_pred = model.predict(X_test)
y_pred_inv = scaler.inverse_transform(y_pred)
y_test_inv = scaler.inverse_transform(y_test.reshape(-1, 1))

# ----- Evaluation -----
mae = mean_absolute_error(y_test_inv, y_pred_inv)
rmse = np.sqrt(mean_squared_error(y_test_inv, y_pred_inv))

print(f"\nLSTM Performance ({label}):")
print(f"MAE: {mae:.4f}")
print(f"RMSE: {rmse:.4f}")

# ----- Plot Prediction -----
plt.figure(figsize=(10, 4))
plt.plot(y_test_inv, label='Actual')
plt.plot(y_pred_inv, label='Predicted')
plt.title(f'LSTM Prediction - {label}')
plt.legend()
plt.tight_layout()
plt.show()

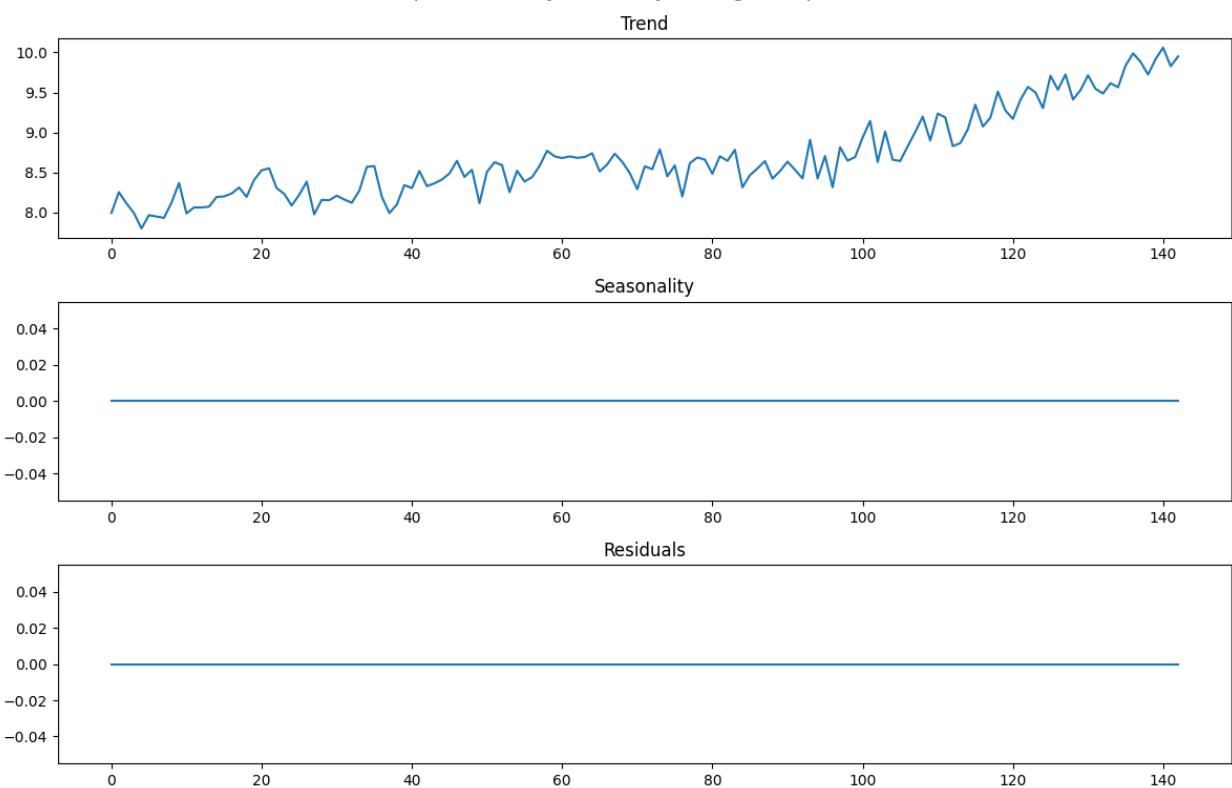
return model

# Yearly model
train_lstm_on_series(df_yearly['Temperature'], label="Yearly Average
Temperature")

# 5-Year model
train_lstm_on_series(df_5year['Temperature'], label="5-Year Average
Temperature")

```

Decomposition Analysis - Yearly Average Temperature



Epoch 1/50

14/14 ━━━━━━━━ 5s 5ms/step - loss: 0.0979

Epoch 2/50

14/14 ━━━━━━━━ 0s 5ms/step - loss: 0.0583

Epoch 3/50

14/14 ━━━━━━━━ 0s 6ms/step - loss: 0.0172

Epoch 4/50

14/14 ━━━━━━━━ 0s 5ms/step - loss: 0.0092

Epoch 5/50

14/14 ━━━━━━━━ 0s 6ms/step - loss: 0.0100

Epoch 6/50

14/14 ━━━━━━━━ 0s 6ms/step - loss: 0.0094

Epoch 7/50

14/14 ━━━━━━━━ 0s 6ms/step - loss: 0.0087

Epoch 8/50

14/14 ━━━━━━━━ 0s 5ms/step - loss: 0.0064

Epoch 9/50

14/14 ━━━━━━━━ 0s 5ms/step - loss: 0.0077

Epoch 10/50

14/14 ━━━━━━━━ 0s 5ms/step - loss: 0.0078

Epoch 11/50

14/14 ━━━━━━━━ 0s 5ms/step - loss: 0.0065

Epoch 12/50

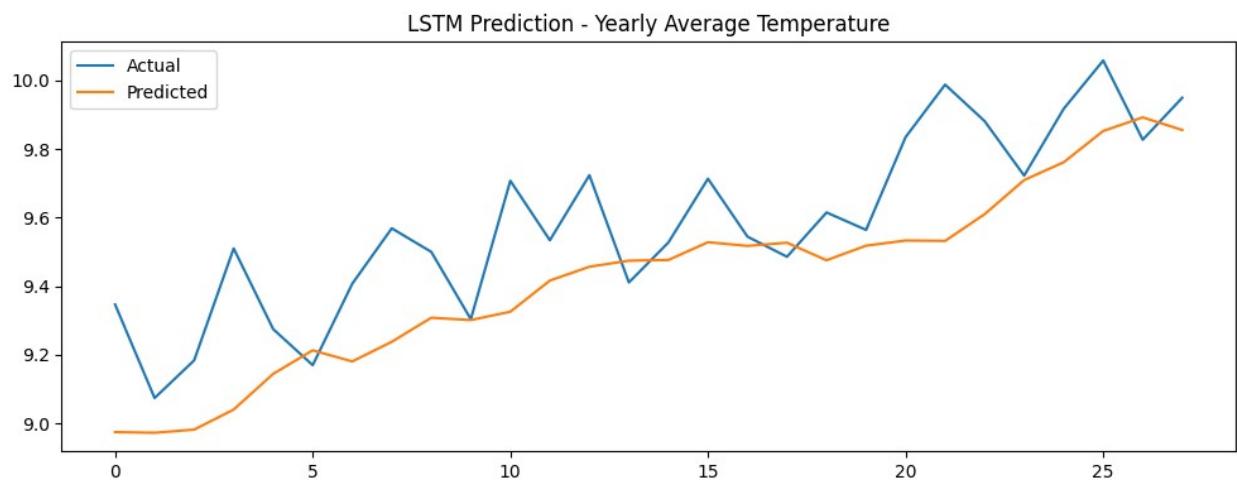
14/14 ━━━━━━━━ 0s 9ms/step - loss: 0.0078

```
Epoch 13/50
14/14 ━━━━━━━━ 0s 6ms/step - loss: 0.0073
Epoch 14/50
14/14 ━━━━━━━━ 0s 6ms/step - loss: 0.0085
Epoch 15/50
14/14 ━━━━━━━━ 0s 6ms/step - loss: 0.0087
Epoch 16/50
14/14 ━━━━━━━━ 0s 6ms/step - loss: 0.0070
Epoch 17/50
14/14 ━━━━━━━━ 0s 5ms/step - loss: 0.0070
Epoch 18/50
14/14 ━━━━━━━━ 0s 6ms/step - loss: 0.0056
Epoch 19/50
14/14 ━━━━━━━━ 0s 6ms/step - loss: 0.0066
Epoch 20/50
14/14 ━━━━━━━━ 0s 6ms/step - loss: 0.0067
Epoch 21/50
14/14 ━━━━━━━━ 0s 6ms/step - loss: 0.0069
Epoch 22/50
14/14 ━━━━━━━━ 0s 6ms/step - loss: 0.0071
Epoch 23/50
14/14 ━━━━━━━━ 0s 6ms/step - loss: 0.0064
Epoch 24/50
14/14 ━━━━━━━━ 0s 6ms/step - loss: 0.0078
Epoch 25/50
14/14 ━━━━━━━━ 0s 6ms/step - loss: 0.0061
Epoch 26/50
14/14 ━━━━━━━━ 0s 6ms/step - loss: 0.0070
Epoch 27/50
14/14 ━━━━━━━━ 0s 7ms/step - loss: 0.0071
Epoch 28/50
14/14 ━━━━━━━━ 0s 6ms/step - loss: 0.0088
Epoch 29/50
14/14 ━━━━━━━━ 0s 6ms/step - loss: 0.0070
Epoch 30/50
14/14 ━━━━━━━━ 0s 5ms/step - loss: 0.0072
Epoch 31/50
14/14 ━━━━━━━━ 0s 5ms/step - loss: 0.0052
Epoch 32/50
14/14 ━━━━━━━━ 0s 5ms/step - loss: 0.0074
Epoch 33/50
14/14 ━━━━━━━━ 0s 5ms/step - loss: 0.0076
Epoch 34/50
14/14 ━━━━━━━━ 0s 6ms/step - loss: 0.0075
Epoch 35/50
14/14 ━━━━━━━━ 0s 6ms/step - loss: 0.0064
Epoch 36/50
14/14 ━━━━━━━━ 0s 5ms/step - loss: 0.0082
Epoch 37/50
```

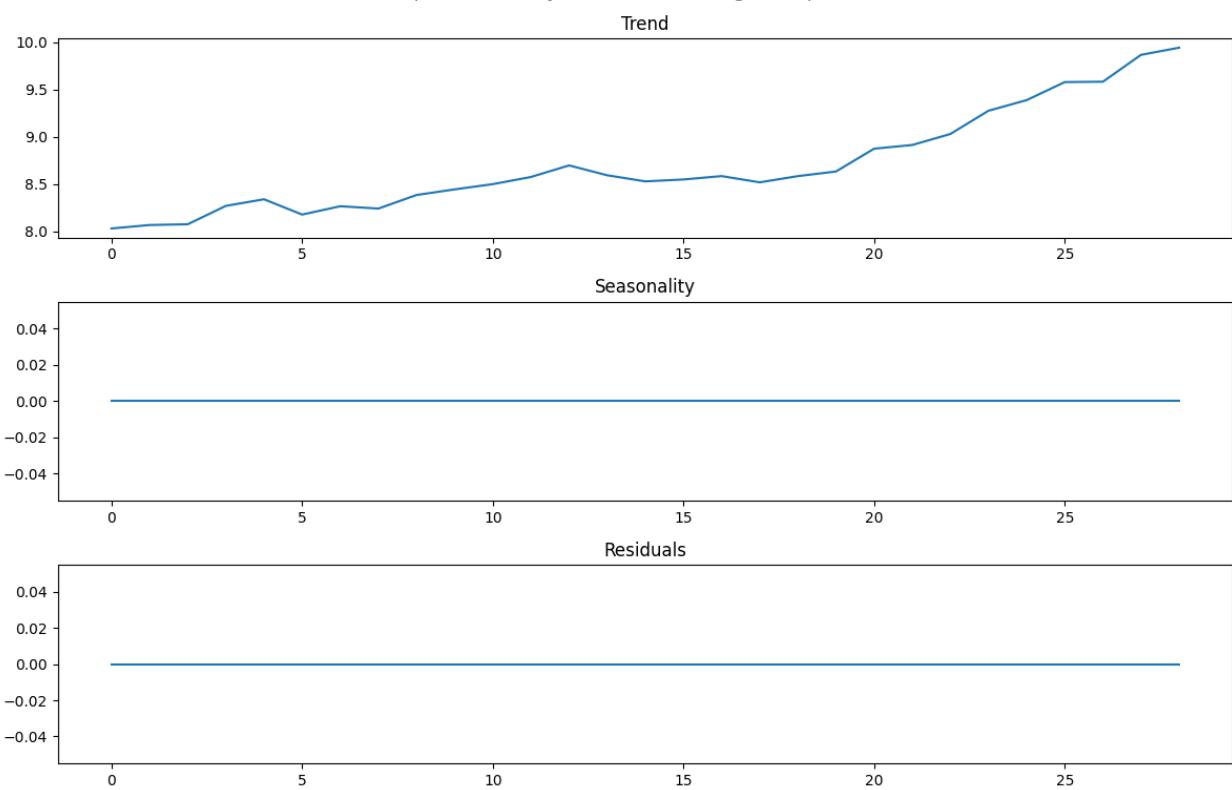
```
14/14 ━━━━━━━━━━ 0s 6ms/step - loss: 0.0069
Epoch 38/50
14/14 ━━━━━━━━━━ 0s 6ms/step - loss: 0.0075
Epoch 39/50
14/14 ━━━━━━━━━━ 0s 6ms/step - loss: 0.0061
Epoch 40/50
14/14 ━━━━━━━━━━ 0s 6ms/step - loss: 0.0065
Epoch 41/50
14/14 ━━━━━━━━━━ 0s 6ms/step - loss: 0.0078
Epoch 42/50
14/14 ━━━━━━━━━━ 0s 8ms/step - loss: 0.0087
Epoch 43/50
14/14 ━━━━━━━━━━ 0s 6ms/step - loss: 0.0068
Epoch 44/50
14/14 ━━━━━━━━━━ 0s 6ms/step - loss: 0.0081
Epoch 45/50
14/14 ━━━━━━━━━━ 0s 6ms/step - loss: 0.0064
Epoch 46/50
14/14 ━━━━━━━━━━ 0s 5ms/step - loss: 0.0057
Epoch 47/50
14/14 ━━━━━━━━━━ 0s 5ms/step - loss: 0.0067
Epoch 48/50
14/14 ━━━━━━━━━━ 0s 5ms/step - loss: 0.0062
Epoch 49/50
14/14 ━━━━━━━━━━ 0s 6ms/step - loss: 0.0068
Epoch 50/50
14/14 ━━━━━━━━━━ 0s 6ms/step - loss: 0.0081
1/1 ━━━━━━━━━━ 0s 481ms/step
```

LSTM Performance (Yearly Average Temperature):

MAE: 0.1769
RMSE: 0.2216



Decomposition Analysis - 5-Year Average Temperature



Epoch 1/50

3/3 ━━━━━━ 5s 9ms/step - loss: 0.0973

Epoch 2/50

3/3 ━━━━━━ 0s 7ms/step - loss: 0.0981

Epoch 3/50

3/3 ━━━━━━ 0s 6ms/step - loss: 0.0842

Epoch 4/50

3/3 ━━━━━━ 0s 5ms/step - loss: 0.0719

Epoch 5/50

3/3 ━━━━━━ 0s 7ms/step - loss: 0.0699

Epoch 6/50

3/3 ━━━━━━ 0s 5ms/step - loss: 0.0544

Epoch 7/50

3/3 ━━━━━━ 0s 5ms/step - loss: 0.0480

Epoch 8/50

3/3 ━━━━━━ 0s 6ms/step - loss: 0.0334

Epoch 9/50

3/3 ━━━━━━ 0s 8ms/step - loss: 0.0260

Epoch 10/50

3/3 ━━━━━━ 0s 7ms/step - loss: 0.0134

Epoch 11/50

3/3 ━━━━━━ 0s 7ms/step - loss: 0.0130

Epoch 12/50

3/3 ━━━━━━ 0s 7ms/step - loss: 0.0132

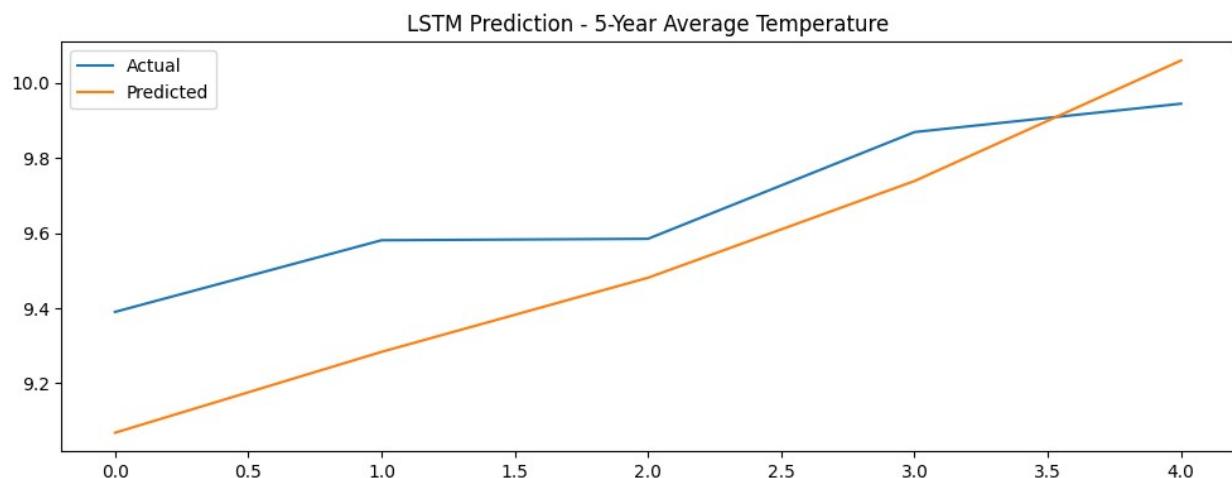
```
Epoch 13/50
3/3 ━━━━━━━━ 0s 8ms/step - loss: 0.0125
Epoch 14/50
3/3 ━━━━━━ 0s 10ms/step - loss: 0.0157
Epoch 15/50
3/3 ━━━━ 0s 7ms/step - loss: 0.0131
Epoch 16/50
3/3 ━━ 0s 7ms/step - loss: 0.0130
Epoch 17/50
3/3 ━ 0s 7ms/step - loss: 0.0096
Epoch 18/50
3/3 0s 6ms/step - loss: 0.0111
Epoch 19/50
3/3 0s 8ms/step - loss: 0.0144
Epoch 20/50
3/3 0s 7ms/step - loss: 0.0096
Epoch 21/50
3/3 0s 9ms/step - loss: 0.0101
Epoch 22/50
3/3 0s 6ms/step - loss: 0.0098
Epoch 23/50
3/3 0s 7ms/step - loss: 0.0121
Epoch 24/50
3/3 0s 8ms/step - loss: 0.0105
Epoch 25/50
3/3 0s 5ms/step - loss: 0.0103
Epoch 26/50
3/3 0s 4ms/step - loss: 0.0101
Epoch 27/50
3/3 0s 7ms/step - loss: 0.0096
Epoch 28/50
3/3 0s 6ms/step - loss: 0.0086
Epoch 29/50
3/3 0s 7ms/step - loss: 0.0111
Epoch 30/50
3/3 0s 8ms/step - loss: 0.0085
Epoch 31/50
3/3 0s 5ms/step - loss: 0.0086
Epoch 32/50
3/3 0s 8ms/step - loss: 0.0081
Epoch 33/50
3/3 0s 6ms/step - loss: 0.0084
Epoch 34/50
3/3 0s 7ms/step - loss: 0.0088
Epoch 35/50
3/3 0s 8ms/step - loss: 0.0093
Epoch 36/50
3/3 0s 6ms/step - loss: 0.0074
Epoch 37/50
```

```
3/3 ━━━━━━━━━━ 0s 10ms/step - loss: 0.0084
Epoch 38/50
3/3 ━━━━━━━━ 0s 7ms/step - loss: 0.0078
Epoch 39/50
3/3 ━━━━━━ 0s 8ms/step - loss: 0.0090
Epoch 40/50
3/3 ━━━━ 0s 6ms/step - loss: 0.0079
Epoch 41/50
3/3 ━━ 0s 7ms/step - loss: 0.0081
Epoch 42/50
3/3 ━ 0s 8ms/step - loss: 0.0068
Epoch 43/50
3/3 0s 7ms/step - loss: 0.0068
Epoch 44/50
3/3 0s 8ms/step - loss: 0.0079
Epoch 45/50
3/3 0s 7ms/step - loss: 0.0057
Epoch 46/50
3/3 0s 6ms/step - loss: 0.0081
Epoch 47/50
3/3 0s 7ms/step - loss: 0.0066
Epoch 48/50
3/3 0s 6ms/step - loss: 0.0071
Epoch 49/50
3/3 0s 11ms/step - loss: 0.0065
Epoch 50/50
3/3 0s 7ms/step - loss: 0.0070
1/1 ━ 0s 468ms/step
```

LSTM Performance (5-Year Average Temperature):

MAE: 0.1937

RMSE: 0.2159



```
<Sequential name=sequential_48, built=True>
```

```

def forecast_future(model, series, years, scaler, seq_length=5):
    """Forecast future temperature values"""
    predictions = []
    input_seq = scaler.transform(series[-seq_length:]).values.reshape(-1, 1).tolist()

    for _ in range(years):
        # Prepare input for prediction
        X_input = np.array(input_seq[-seq_length:]).reshape((1, seq_length, 1))
        next_pred = model.predict(X_input, verbose=0)[0][0]
        predictions.append(next_pred)
        input_seq.append([next_pred]) # keep the sequence growing

    # Inverse scale predictions
    predictions_inv =
    scaler.inverse_transform(np.array(predictions).reshape(-1, 1)).flatten()
    return predictions_inv

# Train on full data for forecasting
scaler = MinMaxScaler()
series_scaled =
scaler.fit_transform(df_yearly['Temperature'].values.reshape(-1, 1))

# Re-train model on entire dataset
X_full, y_full = create_sequences(series_scaled, seq_length=5)
X_full = X_full.reshape((X_full.shape[0], X_full.shape[1], 1))
model_full = Sequential([
    LSTM(50, activation='relu', return_sequences=True, input_shape=(5, 1)),
    LSTM(50, activation='relu'),
    Dense(1)
])
model_full.compile(optimizer='adam', loss='mse')
model_full.fit(X_full, y_full, epochs=50, batch_size=8, verbose=1)

# Forecast next 10 years
future_years = list(range(df_yearly['Year'].iloc[-1] + 1,
df_yearly['Year'].iloc[-1] + 11))
future_predictions = forecast_future(model_full,
df_yearly['Temperature'], 10, scaler)

# Plot
plt.figure(figsize=(10, 4))
plt.plot(df_yearly['Year'], df_yearly['Temperature'],
label='Historical')
plt.plot(future_years, future_predictions, label='Forecast (2021–2030)', linestyle='--', marker='o')
plt.xlabel("Year")

```

```
plt.ylabel("Avg Temperature")
plt.title("Yearly Average Temperature Forecast (Next 10 Years)")
plt.legend()
plt.tight_layout()
plt.show()

# Print values
forecast_df = pd.DataFrame({'Year': future_years, 'Predicted
Temperature': future_predictions})
print("\nForecasted Temperatures (2021–2030):")
print(forecast_df)

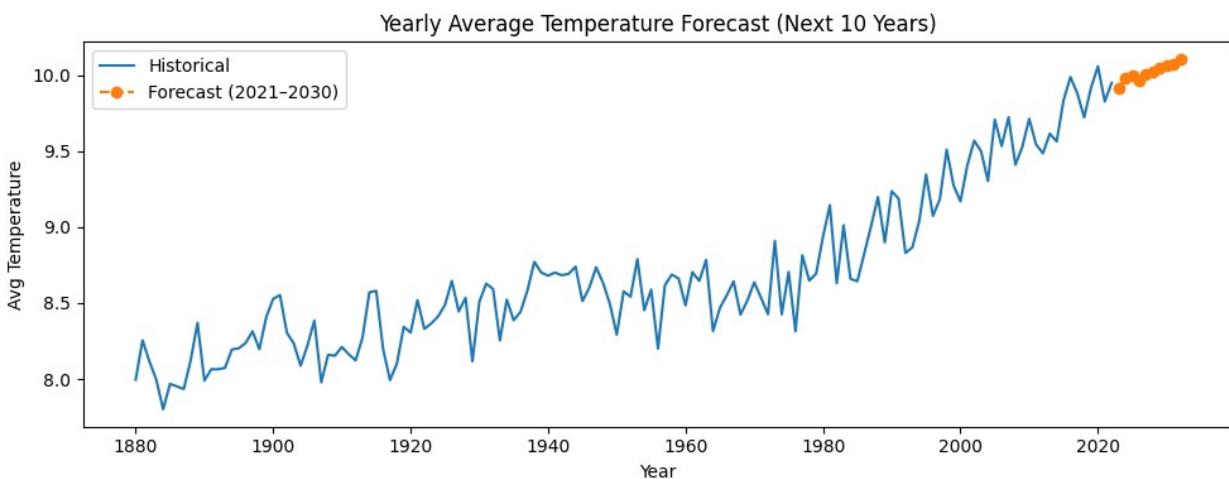
Epoch 1/50
18/18 ━━━━━━━━ 2s 2ms/step - loss: 0.2090
Epoch 2/50
18/18 ━━━━━━━━ 0s 3ms/step - loss: 0.1328
Epoch 3/50
18/18 ━━━━━━━━ 0s 3ms/step - loss: 0.0302
Epoch 4/50
18/18 ━━━━━━━━ 0s 3ms/step - loss: 0.0153
Epoch 5/50
18/18 ━━━━━━━━ 0s 3ms/step - loss: 0.0102
Epoch 6/50
18/18 ━━━━━━━━ 0s 3ms/step - loss: 0.0076
Epoch 7/50
18/18 ━━━━━━━━ 0s 3ms/step - loss: 0.0076
Epoch 8/50
18/18 ━━━━━━━━ 0s 3ms/step - loss: 0.0072
Epoch 9/50
18/18 ━━━━━━━━ 0s 3ms/step - loss: 0.0077
Epoch 10/50
18/18 ━━━━━━━━ 0s 3ms/step - loss: 0.0070
Epoch 11/50
18/18 ━━━━━━━━ 0s 3ms/step - loss: 0.0077
Epoch 12/50
18/18 ━━━━━━━━ 0s 3ms/step - loss: 0.0076
Epoch 13/50
18/18 ━━━━━━━━ 0s 3ms/step - loss: 0.0068
Epoch 14/50
18/18 ━━━━━━━━ 0s 3ms/step - loss: 0.0072
Epoch 15/50
18/18 ━━━━━━━━ 0s 3ms/step - loss: 0.0064
Epoch 16/50
18/18 ━━━━━━━━ 0s 3ms/step - loss: 0.0070
Epoch 17/50
18/18 ━━━━━━━━ 0s 3ms/step - loss: 0.0075
Epoch 18/50
18/18 ━━━━━━━━ 0s 3ms/step - loss: 0.0069
Epoch 19/50
18/18 ━━━━━━━━ 0s 3ms/step - loss: 0.0084
```

```
Epoch 20/50
18/18 ━━━━━━━━ 0s 4ms/step - loss: 0.0057
Epoch 21/50
18/18 ━━━━━━━━ 0s 3ms/step - loss: 0.0064
Epoch 22/50
18/18 ━━━━━━━━ 0s 3ms/step - loss: 0.0066
Epoch 23/50
18/18 ━━━━━━━━ 0s 3ms/step - loss: 0.0082
Epoch 24/50
18/18 ━━━━━━━━ 0s 3ms/step - loss: 0.0082
Epoch 25/50
18/18 ━━━━━━━━ 0s 3ms/step - loss: 0.0070
Epoch 26/50
18/18 ━━━━━━━━ 0s 3ms/step - loss: 0.0074
Epoch 27/50
18/18 ━━━━━━━━ 0s 3ms/step - loss: 0.0072
Epoch 28/50
18/18 ━━━━━━━━ 0s 3ms/step - loss: 0.0077
Epoch 29/50
18/18 ━━━━━━━━ 0s 3ms/step - loss: 0.0072
Epoch 30/50
18/18 ━━━━━━━━ 0s 3ms/step - loss: 0.0075
Epoch 31/50
18/18 ━━━━━━━━ 0s 3ms/step - loss: 0.0070
Epoch 32/50
18/18 ━━━━━━━━ 0s 5ms/step - loss: 0.0065
Epoch 33/50
18/18 ━━━━━━━━ 0s 3ms/step - loss: 0.0061
Epoch 34/50
18/18 ━━━━━━━━ 0s 3ms/step - loss: 0.0060
Epoch 35/50
18/18 ━━━━━━━━ 0s 3ms/step - loss: 0.0077
Epoch 36/50
18/18 ━━━━━━━━ 0s 3ms/step - loss: 0.0065
Epoch 37/50
18/18 ━━━━━━━━ 0s 3ms/step - loss: 0.0065
Epoch 38/50
18/18 ━━━━━━━━ 0s 3ms/step - loss: 0.0068
Epoch 39/50
18/18 ━━━━━━━━ 0s 3ms/step - loss: 0.0080
Epoch 40/50
18/18 ━━━━━━━━ 0s 4ms/step - loss: 0.0082
Epoch 41/50
18/18 ━━━━━━━━ 0s 3ms/step - loss: 0.0046
Epoch 42/50
18/18 ━━━━━━━━ 0s 3ms/step - loss: 0.0067
Epoch 43/50
18/18 ━━━━━━━━ 0s 3ms/step - loss: 0.0067
Epoch 44/50
```

```

18/18 ━━━━━━━━━━ 0s 3ms/step - loss: 0.0075
Epoch 45/50
18/18 ━━━━━━━━ 0s 4ms/step - loss: 0.0063
Epoch 46/50
18/18 ━━━━━━ 0s 3ms/step - loss: 0.0081
Epoch 47/50
18/18 ━━━━ 0s 4ms/step - loss: 0.0073
Epoch 48/50
18/18 ━━ 0s 3ms/step - loss: 0.0067
Epoch 49/50
18/18 ━ 0s 3ms/step - loss: 0.0070
Epoch 50/50
18/18 0s 3ms/step - loss: 0.0069
WARNING:tensorflow:5 out of the last 26 calls to <function
TensorFlowTrainer.make_predict_function.<locals>.one_step_on_data_distributed at 0x000001192C49E980> triggered tf.function retracing.
Tracing is expensive and the excessive number of tracings could be due to (1) creating @tf.function repeatedly in a loop, (2) passing tensors with different shapes, (3) passing Python objects instead of tensors. For (1), please define your @tf.function outside of the loop. For (2), @tf.function has reduce_retracing=True option that can avoid unnecessary retracing. For (3), please refer to
https://www.tensorflow.org/guide/function#controlling\_retracing and
https://www.tensorflow.org/api\_docs/python/tf/function for more details.

```



Forecasted Temperatures (2021–2030):

	Year	Predicted Temperature
0	2023	9.916318
1	2024	9.985366
2	2025	9.995470
3	2026	9.963815
4	2027	10.008984

```

5 2028           10.019996
6 2029           10.049272
7 2030           10.060635
8 2031           10.074349
9 2032           10.104771

#maximum

df_yearly_max = df.groupby('Year')[['Temperature']].max().reset_index()

# Calculate 5-year average
df['5YearGroup'] = (df['Year'] // 5) * 5
df_5year_max = df.groupby('5YearGroup')[['Temperature']].max().reset_index()

def train_lstm_on_series(series, label=""):
    # Normalize
    scaler = MinMaxScaler()
    series_scaled = scaler.fit_transform(series.values.reshape(-1, 1))

    # Create sequences
    seq_length = 5
    X, y = create_sequences(series_scaled, seq_length)
    X = X.reshape((X.shape[0], X.shape[1], 1))

    # Train-test split
    X_train, X_test, y_train, y_test = train_test_split(X, y,
test_size=0.2, shuffle=False)

    # Model
    model = Sequential([
        LSTM(50, activation='relu', return_sequences=True,
input_shape=(seq_length, 1)),
        LSTM(50, activation='relu'),
        Dense(1)
    ])
    model.compile(optimizer='adam', loss='mse')
    model.fit(X_train, y_train, epochs=50, batch_size=8, verbose=1)

    # Predict and inverse transform
    y_pred = model.predict(X_test)
    y_pred_inv = scaler.inverse_transform(y_pred)
    y_test_inv = scaler.inverse_transform(y_test.reshape(-1, 1))

    # Metrics
    mae = mean_absolute_error(y_test_inv, y_pred_inv)
    rmse = np.sqrt(mean_squared_error(y_test_inv, y_pred_inv))

    print(f"\nLSTM Performance ({label}):")
    print(f"MAE: {mae:.4f}")

```

```

print(f"RMSE: {rmse:.4f}")

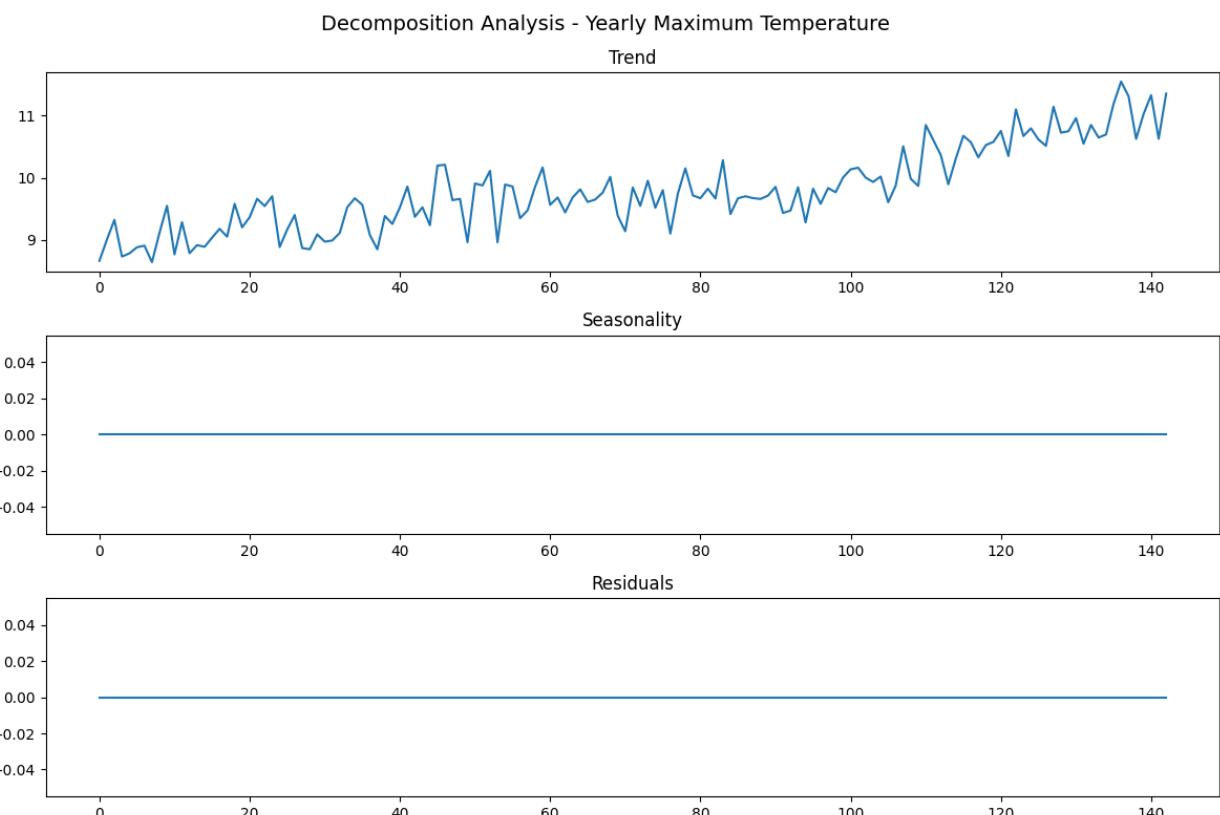
# Plot
plt.figure(figsize=(10, 4))
plt.plot(y_test_inv, label='Actual')
plt.plot(y_pred_inv, label='Predicted')
plt.title(f"LSTM Prediction - {label}")
plt.legend()
plt.tight_layout()
plt.show()

return model

# Yearly model
train_lstm_on_series(df_yearly_max['Temperature'], label="Yearly Maximum Temperature")

# 5-Year model
train_lstm_on_series(df_5year_max['Temperature'], label="5-Year Maximum Temperature")

```



```

Epoch 1/50
14/14 ━━━━━━━━━━━━ 4s 5ms/step - loss: 0.1002
Epoch 2/50

```

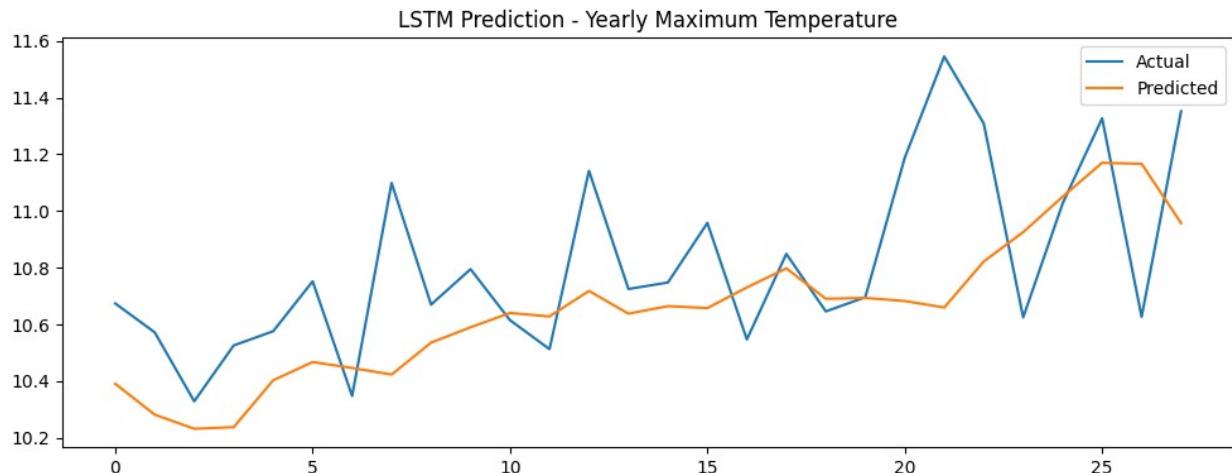
14/14 ━━━━━━━━ 0s 5ms/step - loss: 0.0487
Epoch 3/50
14/14 ━━━━━━ 0s 5ms/step - loss: 0.0214
Epoch 4/50
14/14 ━━━━━━ 0s 5ms/step - loss: 0.0155
Epoch 5/50
14/14 ━━━━━━ 0s 5ms/step - loss: 0.0177
Epoch 6/50
14/14 ━━━━━━ 0s 5ms/step - loss: 0.0170
Epoch 7/50
14/14 ━━━━━━ 0s 11ms/step - loss: 0.0188
Epoch 8/50
14/14 ━━━━━━ 0s 6ms/step - loss: 0.0139
Epoch 9/50
14/14 ━━━━━━ 0s 6ms/step - loss: 0.0143
Epoch 10/50
14/14 ━━━━━━ 0s 5ms/step - loss: 0.0130
Epoch 11/50
14/14 ━━━━━━ 0s 6ms/step - loss: 0.0120
Epoch 12/50
14/14 ━━━━━━ 0s 6ms/step - loss: 0.0131
Epoch 13/50
14/14 ━━━━━━ 0s 6ms/step - loss: 0.0117
Epoch 14/50
14/14 ━━━━━━ 0s 6ms/step - loss: 0.0138
Epoch 15/50
14/14 ━━━━━━ 0s 5ms/step - loss: 0.0115
Epoch 16/50
14/14 ━━━━━━ 0s 5ms/step - loss: 0.0113
Epoch 17/50
14/14 ━━━━━━ 0s 6ms/step - loss: 0.0149
Epoch 18/50
14/14 ━━━━━━ 0s 6ms/step - loss: 0.0132
Epoch 19/50
14/14 ━━━━━━ 0s 5ms/step - loss: 0.0108
Epoch 20/50
14/14 ━━━━━━ 0s 5ms/step - loss: 0.0144
Epoch 21/50
14/14 ━━━━━━ 0s 7ms/step - loss: 0.0115
Epoch 22/50
14/14 ━━━━━━ 0s 7ms/step - loss: 0.0138
Epoch 23/50
14/14 ━━━━━━ 0s 6ms/step - loss: 0.0132
Epoch 24/50
14/14 ━━━━━━ 0s 6ms/step - loss: 0.0134
Epoch 25/50
14/14 ━━━━━━ 0s 6ms/step - loss: 0.0151
Epoch 26/50
14/14 ━━━━━━ 0s 6ms/step - loss: 0.0115

```
Epoch 27/50
14/14 ━━━━━━━━ 0s 6ms/step - loss: 0.0147
Epoch 28/50
14/14 ━━━━━━━━ 0s 5ms/step - loss: 0.0130
Epoch 29/50
14/14 ━━━━━━━━ 0s 5ms/step - loss: 0.0110
Epoch 30/50
14/14 ━━━━━━━━ 0s 6ms/step - loss: 0.0121
Epoch 31/50
14/14 ━━━━━━━━ 0s 6ms/step - loss: 0.0121
Epoch 32/50
14/14 ━━━━━━━━ 0s 6ms/step - loss: 0.0134
Epoch 33/50
14/14 ━━━━━━━━ 0s 6ms/step - loss: 0.0130
Epoch 34/50
14/14 ━━━━━━━━ 0s 5ms/step - loss: 0.0146
Epoch 35/50
14/14 ━━━━━━━━ 0s 6ms/step - loss: 0.0117
Epoch 36/50
14/14 ━━━━━━━━ 0s 6ms/step - loss: 0.0128
Epoch 37/50
14/14 ━━━━━━━━ 0s 6ms/step - loss: 0.0115
Epoch 38/50
14/14 ━━━━━━━━ 0s 11ms/step - loss: 0.0116
Epoch 39/50
14/14 ━━━━━━━━ 0s 6ms/step - loss: 0.0122
Epoch 40/50
14/14 ━━━━━━━━ 0s 6ms/step - loss: 0.0101
Epoch 41/50
14/14 ━━━━━━━━ 0s 6ms/step - loss: 0.0119
Epoch 42/50
14/14 ━━━━━━━━ 0s 6ms/step - loss: 0.0119
Epoch 43/50
14/14 ━━━━━━━━ 0s 6ms/step - loss: 0.0121
Epoch 44/50
14/14 ━━━━━━━━ 0s 6ms/step - loss: 0.0118
Epoch 45/50
14/14 ━━━━━━━━ 0s 6ms/step - loss: 0.0149
Epoch 46/50
14/14 ━━━━━━━━ 0s 6ms/step - loss: 0.0139
Epoch 47/50
14/14 ━━━━━━━━ 0s 6ms/step - loss: 0.0128
Epoch 48/50
14/14 ━━━━━━━━ 0s 5ms/step - loss: 0.0119
Epoch 49/50
14/14 ━━━━━━━━ 0s 6ms/step - loss: 0.0135
Epoch 50/50
14/14 ━━━━━━━━ 0s 6ms/step - loss: 0.0143
1/1 ━━━━━━━━ 0s 500ms/step
```

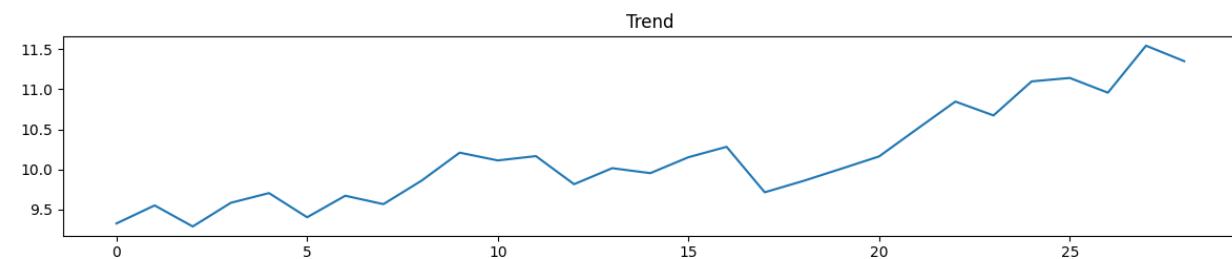
LSTM Performance (Yearly Maximum Temperature):

MAE: 0.2550

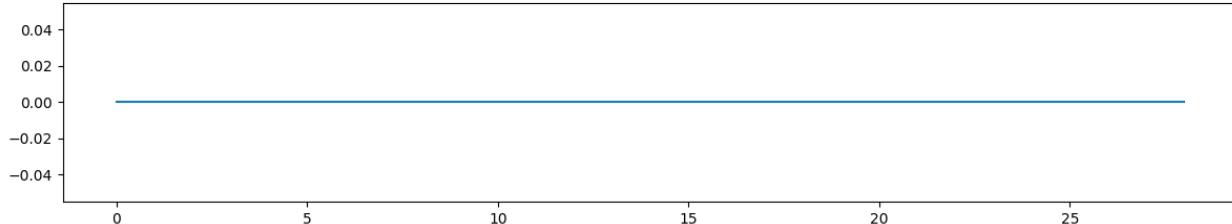
RMSE: 0.3311



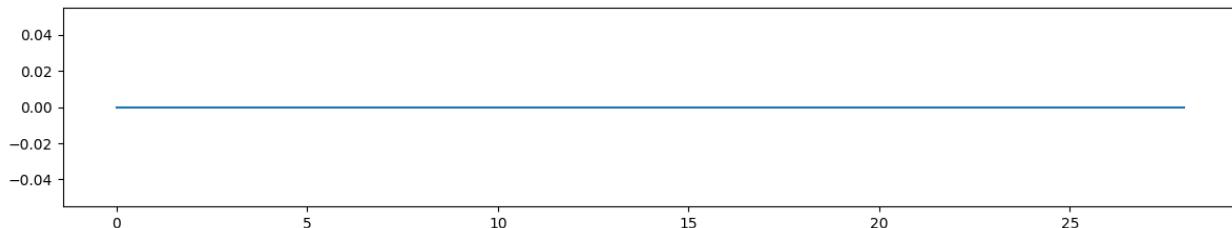
Decomposition Analysis - 5-Year Maximum Temperature



Seasonality



Residuals



Epoch 1/50

3/3 ━━━━━━━━ 5s 6ms/step - loss: 0.1466

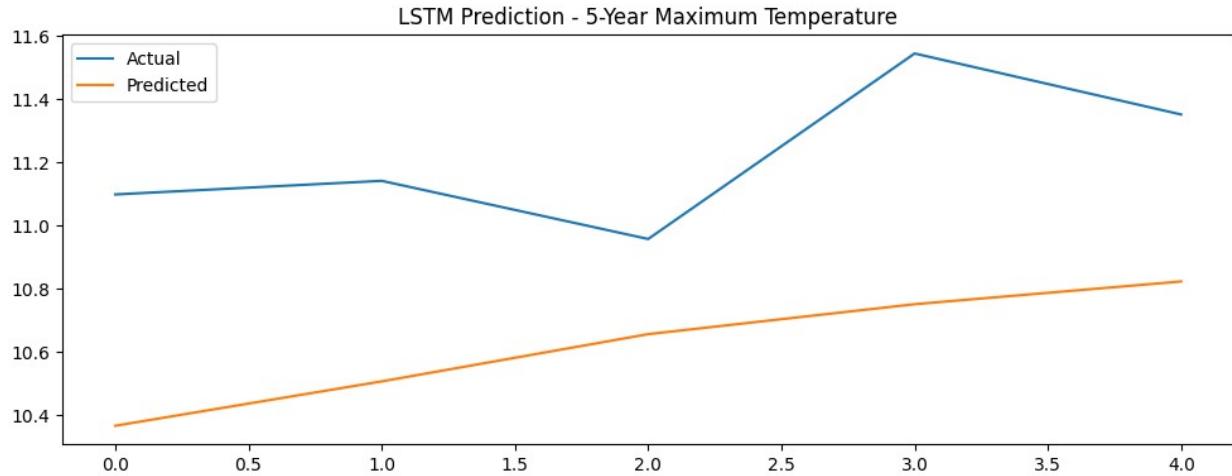
Epoch 2/50

3/3 ━━━━━━━━ 0s 6ms/step - loss: 0.1295

```
Epoch 3/50
3/3 ━━━━━━━━ 0s 7ms/step - loss: 0.1180
Epoch 4/50
3/3 ━━━━━━━━ 0s 7ms/step - loss: 0.1011
Epoch 5/50
3/3 ━━━━━━━━ 0s 6ms/step - loss: 0.1012
Epoch 6/50
3/3 ━━━━━━━━ 0s 6ms/step - loss: 0.0860
Epoch 7/50
3/3 ━━━━━━━━ 0s 5ms/step - loss: 0.0713
Epoch 8/50
3/3 ━━━━━━━━ 0s 6ms/step - loss: 0.0573
Epoch 9/50
3/3 ━━━━━━━━ 0s 7ms/step - loss: 0.0433
Epoch 10/50
3/3 ━━━━━━━━ 0s 6ms/step - loss: 0.0354
Epoch 11/50
3/3 ━━━━━━━━ 0s 7ms/step - loss: 0.0292
Epoch 12/50
3/3 ━━━━━━━━ 0s 6ms/step - loss: 0.0193
Epoch 13/50
3/3 ━━━━━━━━ 0s 6ms/step - loss: 0.0203
Epoch 14/50
3/3 ━━━━━━━━ 0s 6ms/step - loss: 0.0217
Epoch 15/50
3/3 ━━━━━━━━ 0s 8ms/step - loss: 0.0234
Epoch 16/50
3/3 ━━━━━━━━ 0s 7ms/step - loss: 0.0185
Epoch 17/50
3/3 ━━━━━━━━ 0s 7ms/step - loss: 0.0196
Epoch 18/50
3/3 ━━━━━━━━ 0s 5ms/step - loss: 0.0196
Epoch 19/50
3/3 ━━━━━━━━ 0s 6ms/step - loss: 0.0199
Epoch 20/50
3/3 ━━━━━━━━ 0s 6ms/step - loss: 0.0192
Epoch 21/50
3/3 ━━━━━━━━ 0s 6ms/step - loss: 0.0221
Epoch 22/50
3/3 ━━━━━━━━ 0s 8ms/step - loss: 0.0209
Epoch 23/50
3/3 ━━━━━━━━ 0s 7ms/step - loss: 0.0182
Epoch 24/50
3/3 ━━━━━━━━ 0s 5ms/step - loss: 0.0204
Epoch 25/50
3/3 ━━━━━━━━ 0s 6ms/step - loss: 0.0229
Epoch 26/50
3/3 ━━━━━━━━ 0s 13ms/step - loss: 0.0195
Epoch 27/50
```

```
3/3 ━━━━━━━━━━ 0s 6ms/step - loss: 0.0182
Epoch 28/50
3/3 ━━━━━━━━ 0s 6ms/step - loss: 0.0176
Epoch 29/50
3/3 ━━━━━━ 0s 7ms/step - loss: 0.0235
Epoch 30/50
3/3 ━━━━ 0s 6ms/step - loss: 0.0233
Epoch 31/50
3/3 ━━ 0s 6ms/step - loss: 0.0200
Epoch 32/50
3/3 ━ 0s 4ms/step - loss: 0.0170
Epoch 33/50
3/3 0s 9ms/step - loss: 0.0202
Epoch 34/50
3/3 0s 7ms/step - loss: 0.0181
Epoch 35/50
3/3 0s 8ms/step - loss: 0.0200
Epoch 36/50
3/3 0s 5ms/step - loss: 0.0169
Epoch 37/50
3/3 0s 4ms/step - loss: 0.0157
Epoch 38/50
3/3 0s 8ms/step - loss: 0.0172
Epoch 39/50
3/3 0s 7ms/step - loss: 0.0209
Epoch 40/50
3/3 0s 6ms/step - loss: 0.0149
Epoch 41/50
3/3 0s 8ms/step - loss: 0.0190
Epoch 42/50
3/3 0s 9ms/step - loss: 0.0202
Epoch 43/50
3/3 0s 6ms/step - loss: 0.0169
Epoch 44/50
3/3 0s 7ms/step - loss: 0.0176
Epoch 45/50
3/3 0s 7ms/step - loss: 0.0192
Epoch 46/50
3/3 0s 7ms/step - loss: 0.0177
Epoch 47/50
3/3 0s 7ms/step - loss: 0.0195
Epoch 48/50
3/3 0s 5ms/step - loss: 0.0185
Epoch 49/50
3/3 0s 6ms/step - loss: 0.0182
Epoch 50/50
3/3 0s 7ms/step - loss: 0.0214
1/1 ━━━━━━━━ 0s 394ms/step
```

LSTM Performance (5-Year Maximum Temperature):
MAE: 0.5979
RMSE: 0.6225



```
<Sequential name=sequential_50, built=True>

# Train on full data for forecasting
scaler = MinMaxScaler()
series_scaled =
scaler.fit_transform(df_yearly_max['Temperature'].values.reshape(-1,
1))

# Re-train model on entire dataset
X_full, y_full = create_sequences(series_scaled, seq_length=5)
X_full = X_full.reshape((X_full.shape[0], X_full.shape[1], 1))
model_full = Sequential([
    LSTM(50, activation='relu', return_sequences=True, input_shape=(5,
1)),
    LSTM(50, activation='relu'),
    Dense(1)
])
model_full.compile(optimizer='adam', loss='mse')
model_full.fit(X_full, y_full, epochs=50, batch_size=8, verbose=1)

# Forecast next 10 years
future_years = list(range(df_yearly_max['Year'].iloc[-1] + 1,
df_yearly_max['Year'].iloc[-1] + 11))
future_predictions = forecast_future(model_full,
df_yearly_max['Temperature'], 10, scaler)

# Plot
plt.figure(figsize=(10, 4))
plt.plot(df_yearly_max['Year'], df_yearly_max['Temperature'],
label='Historical')
```

```
plt.plot(future_years, future_predictions, label='Forecast (2021–2030)', linestyle='--', marker='o')
plt.xlabel("Year")
plt.ylabel("Avg Temperature")
plt.title("Yearly Maximum Temperature Forecast (Next 10 Years)")
plt.legend()
plt.tight_layout()
plt.show()

# Print values
forecast_df_max = pd.DataFrame({'Year': future_years, 'Predicted Temperature': future_predictions})
print("\nForecasted Temperatures (2021–2030):")
print(forecast_df_max)

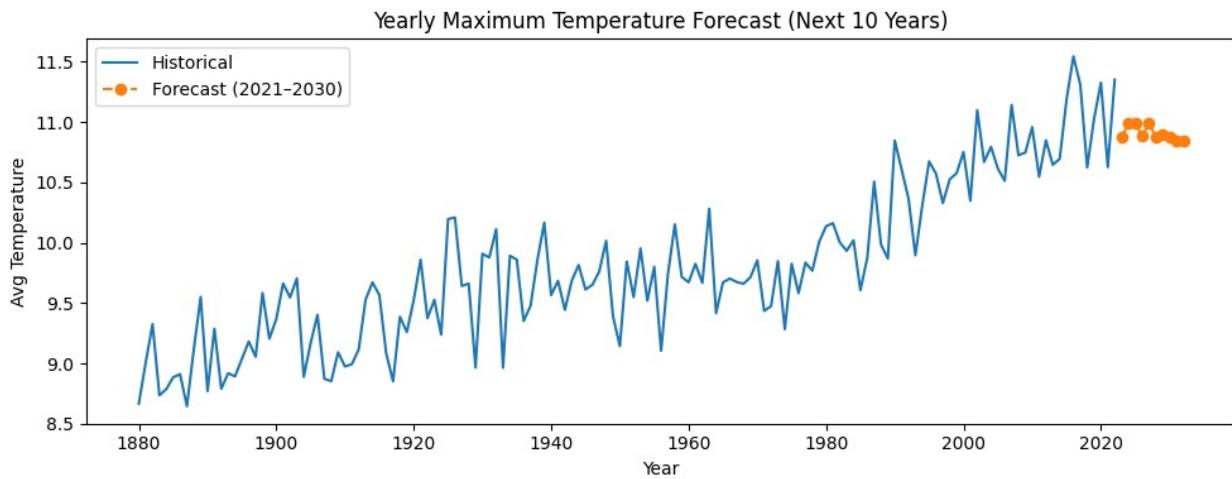
Epoch 1/50
18/18 ━━━━━━━━ 5s 5ms/step - loss: 0.1838
Epoch 2/50
18/18 ━━━━━━━━ 0s 6ms/step - loss: 0.0912
Epoch 3/50
18/18 ━━━━━━━━ 0s 6ms/step - loss: 0.0277
Epoch 4/50
18/18 ━━━━━━━━ 0s 6ms/step - loss: 0.0219
Epoch 5/50
18/18 ━━━━━━━━ 0s 6ms/step - loss: 0.0216
Epoch 6/50
18/18 ━━━━━━━━ 0s 7ms/step - loss: 0.0155
Epoch 7/50
18/18 ━━━━━━━━ 0s 6ms/step - loss: 0.0139
Epoch 8/50
18/18 ━━━━━━━━ 0s 6ms/step - loss: 0.0129
Epoch 9/50
18/18 ━━━━━━━━ 0s 6ms/step - loss: 0.0122
Epoch 10/50
18/18 ━━━━━━━━ 0s 6ms/step - loss: 0.0126
Epoch 11/50
18/18 ━━━━━━━━ 0s 6ms/step - loss: 0.0136
Epoch 12/50
18/18 ━━━━━━━━ 0s 6ms/step - loss: 0.0102
Epoch 13/50
18/18 ━━━━━━━━ 0s 6ms/step - loss: 0.0116
Epoch 14/50
18/18 ━━━━━━━━ 0s 6ms/step - loss: 0.0122
Epoch 15/50
18/18 ━━━━━━━━ 0s 6ms/step - loss: 0.0110
Epoch 16/50
18/18 ━━━━━━━━ 0s 6ms/step - loss: 0.0122
Epoch 17/50
18/18 ━━━━━━━━ 0s 6ms/step - loss: 0.0139
Epoch 18/50
```

```
18/18 ━━━━━━━━━━ 0s 5ms/step - loss: 0.0121
Epoch 19/50
18/18 ━━━━━━━━ 0s 6ms/step - loss: 0.0108
Epoch 20/50
18/18 ━━━━━━ 0s 6ms/step - loss: 0.0118
Epoch 21/50
18/18 ━━━━ 0s 6ms/step - loss: 0.0132
Epoch 22/50
18/18 ━━ 0s 6ms/step - loss: 0.0143
Epoch 23/50
18/18 ━ 0s 6ms/step - loss: 0.0126
Epoch 24/50
18/18 0s 8ms/step - loss: 0.0111
Epoch 25/50
18/18 0s 7ms/step - loss: 0.0102
Epoch 26/50
18/18 0s 7ms/step - loss: 0.0106
Epoch 27/50
18/18 0s 7ms/step - loss: 0.0110
Epoch 28/50
18/18 0s 6ms/step - loss: 0.0126
Epoch 29/50
18/18 0s 6ms/step - loss: 0.0129
Epoch 30/50
18/18 0s 5ms/step - loss: 0.0111
Epoch 31/50
18/18 0s 6ms/step - loss: 0.0122
Epoch 32/50
18/18 0s 5ms/step - loss: 0.0138
Epoch 33/50
18/18 0s 6ms/step - loss: 0.0114
Epoch 34/50
18/18 0s 6ms/step - loss: 0.0107
Epoch 35/50
18/18 0s 6ms/step - loss: 0.0136
Epoch 36/50
18/18 0s 6ms/step - loss: 0.0123
Epoch 37/50
18/18 0s 5ms/step - loss: 0.0122
Epoch 38/50
18/18 0s 6ms/step - loss: 0.0112
Epoch 39/50
18/18 0s 6ms/step - loss: 0.0147
Epoch 40/50
18/18 0s 6ms/step - loss: 0.0116
Epoch 41/50
18/18 0s 7ms/step - loss: 0.0129
Epoch 42/50
18/18 0s 6ms/step - loss: 0.0135
```

```

Epoch 43/50
18/18 ━━━━━━━━ 0s 5ms/step - loss: 0.0120
Epoch 44/50
18/18 ━━━━━━━━ 0s 6ms/step - loss: 0.0110
Epoch 45/50
18/18 ━━━━━━━━ 0s 5ms/step - loss: 0.0125
Epoch 46/50
18/18 ━━━━━━━━ 0s 5ms/step - loss: 0.0131
Epoch 47/50
18/18 ━━━━━━━━ 0s 9ms/step - loss: 0.0140
Epoch 48/50
18/18 ━━━━━━━━ 0s 6ms/step - loss: 0.0120
Epoch 49/50
18/18 ━━━━━━━━ 0s 6ms/step - loss: 0.0151
Epoch 50/50
18/18 ━━━━━━━━ 0s 6ms/step - loss: 0.0106

```



Forecasted Temperatures (2021–2030):

	Year	Predicted Temperature
0	2023	10.872462
1	2024	10.988500
2	2025	10.987010
3	2026	10.886519
4	2027	10.990862
5	2028	10.874409
6	2029	10.892254
7	2030	10.871459
8	2031	10.845276
9	2032	10.844179

#minimum TS

```

# Calculate yearly min
df_yearly_min = df.groupby('Year')[['Temperature']].min().reset_index()

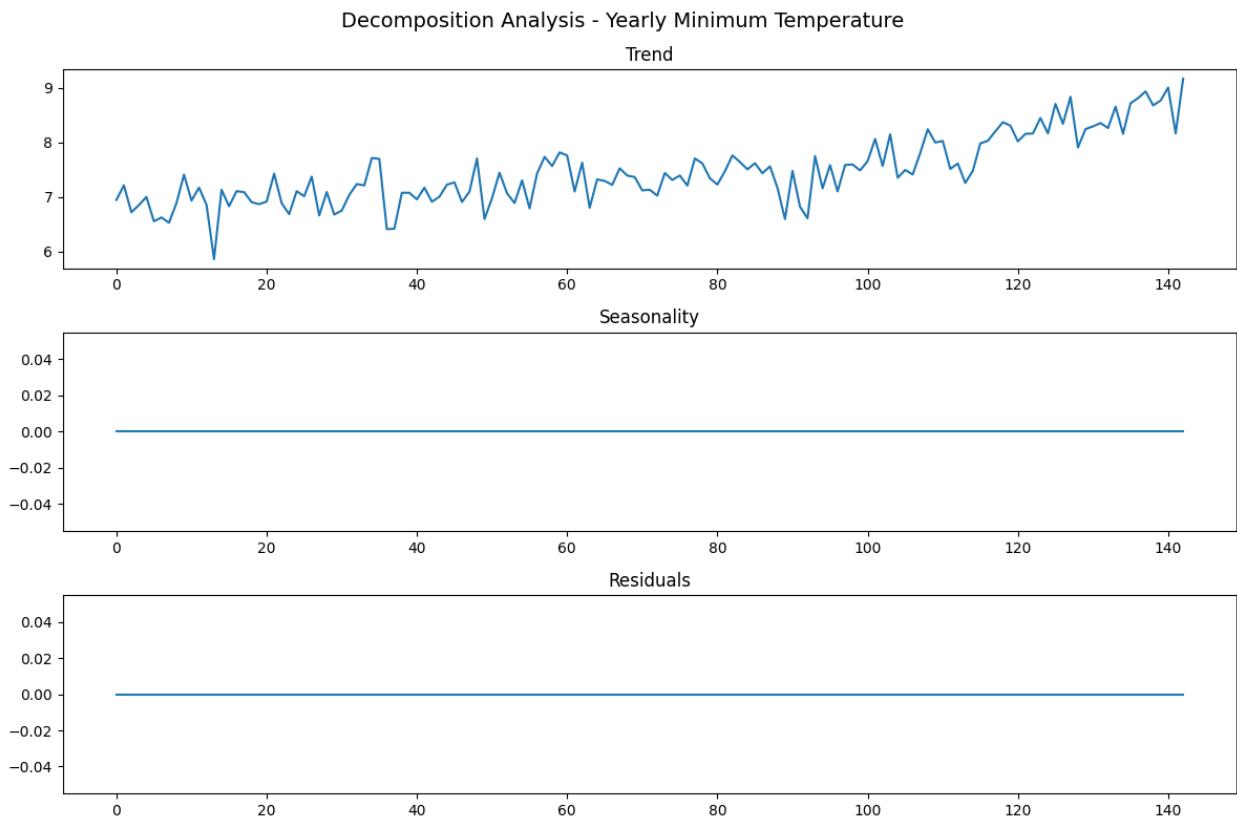
# Calculate 5-year min
df['5YearGroup'] = (df['Year'] // 5) * 5
df_5year_min = df.groupby('5YearGroup')[['Temperature']].min().reset_index()
df_yearly_min.tail()

      Year  Temperature
138  2018       8.676
139  2019       8.764
140  2020       9.002
141  2021       8.162
142  2022       9.166

# Yearly model
train_lstm_on_series(df_yearly_min['Temperature'], label="Yearly Minimum Temperature")

# 5-Year model
train_lstm_on_series(df_5year_min['Temperature'], label="5-Year Minimum Temperature")

```



```
Epoch 1/50
14/14 ━━━━━━━━ 5s 5ms/step - loss: 0.1750
Epoch 2/50
14/14 ━━━━━━━━ 0s 5ms/step - loss: 0.1084
Epoch 3/50
14/14 ━━━━━━━━ 0s 5ms/step - loss: 0.0191
Epoch 4/50
14/14 ━━━━━━━━ 0s 5ms/step - loss: 0.0126
Epoch 5/50
14/14 ━━━━━━━━ 0s 5ms/step - loss: 0.0154
Epoch 6/50
14/14 ━━━━━━━━ 0s 5ms/step - loss: 0.0134
Epoch 7/50
14/14 ━━━━━━━━ 0s 5ms/step - loss: 0.0109
Epoch 8/50
14/14 ━━━━━━━━ 0s 5ms/step - loss: 0.0131
Epoch 9/50
14/14 ━━━━━━━━ 0s 5ms/step - loss: 0.0140
Epoch 10/50
14/14 ━━━━━━━━ 0s 5ms/step - loss: 0.0131
Epoch 11/50
14/14 ━━━━━━━━ 0s 5ms/step - loss: 0.0119
Epoch 12/50
14/14 ━━━━━━━━ 0s 5ms/step - loss: 0.0154
Epoch 13/50
14/14 ━━━━━━━━ 0s 5ms/step - loss: 0.0118
Epoch 14/50
14/14 ━━━━━━━━ 0s 6ms/step - loss: 0.0115
Epoch 15/50
14/14 ━━━━━━━━ 0s 5ms/step - loss: 0.0119
Epoch 16/50
14/14 ━━━━━━━━ 0s 5ms/step - loss: 0.0102
Epoch 17/50
14/14 ━━━━━━━━ 0s 5ms/step - loss: 0.0104
Epoch 18/50
14/14 ━━━━━━━━ 0s 5ms/step - loss: 0.0124
Epoch 19/50
14/14 ━━━━━━━━ 0s 5ms/step - loss: 0.0121
Epoch 20/50
14/14 ━━━━━━━━ 0s 5ms/step - loss: 0.0112
Epoch 21/50
14/14 ━━━━━━━━ 0s 5ms/step - loss: 0.0113
Epoch 22/50
14/14 ━━━━━━━━ 0s 5ms/step - loss: 0.0113
Epoch 23/50
14/14 ━━━━━━━━ 0s 5ms/step - loss: 0.0144
Epoch 24/50
14/14 ━━━━━━━━ 0s 6ms/step - loss: 0.0108
Epoch 25/50
14/14 ━━━━━━━━ 0s 5ms/step - loss: 0.0148
```

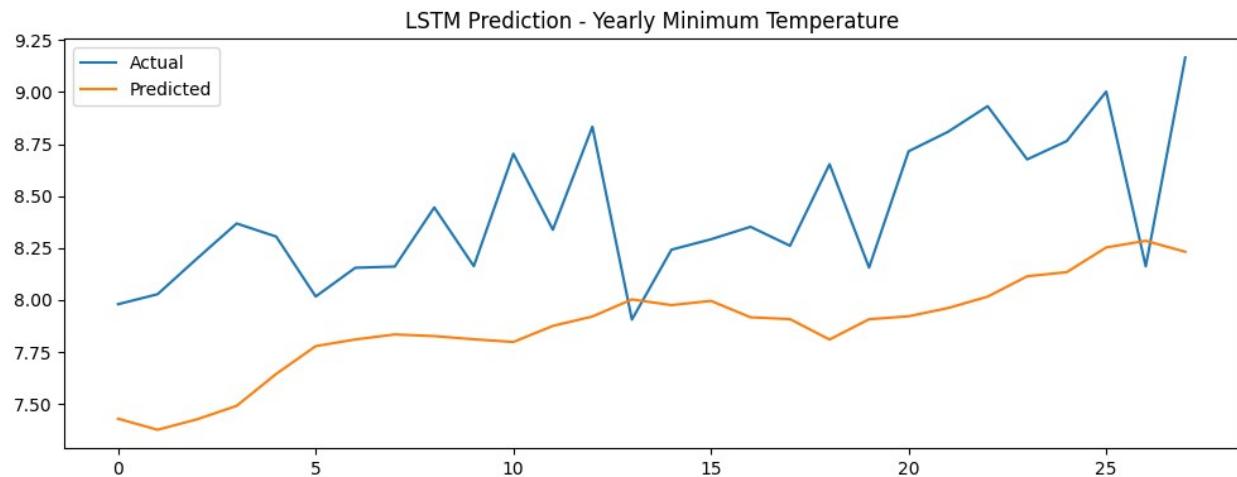
Epoch 26/50
14/14 ━━━━━━━━ 0s 5ms/step - loss: 0.0133
Epoch 27/50
14/14 ━━━━━━━━ 0s 5ms/step - loss: 0.0137
Epoch 28/50
14/14 ━━━━━━━━ 0s 5ms/step - loss: 0.0129
Epoch 29/50
14/14 ━━━━━━━━ 0s 5ms/step - loss: 0.0110
Epoch 30/50
14/14 ━━━━━━━━ 0s 6ms/step - loss: 0.0129
Epoch 31/50
14/14 ━━━━━━━━ 0s 5ms/step - loss: 0.0115
Epoch 32/50
14/14 ━━━━━━━━ 0s 5ms/step - loss: 0.0109
Epoch 33/50
14/14 ━━━━━━━━ 0s 5ms/step - loss: 0.0135
Epoch 34/50
14/14 ━━━━━━━━ 0s 5ms/step - loss: 0.0120
Epoch 35/50
14/14 ━━━━━━━━ 0s 5ms/step - loss: 0.0117
Epoch 36/50
14/14 ━━━━━━━━ 0s 5ms/step - loss: 0.0095
Epoch 37/50
14/14 ━━━━━━━━ 0s 5ms/step - loss: 0.0128
Epoch 38/50
14/14 ━━━━━━━━ 0s 5ms/step - loss: 0.0096
Epoch 39/50
14/14 ━━━━━━━━ 0s 5ms/step - loss: 0.0100
Epoch 40/50
14/14 ━━━━━━━━ 0s 5ms/step - loss: 0.0135
Epoch 41/50
14/14 ━━━━━━━━ 0s 6ms/step - loss: 0.0101
Epoch 42/50
14/14 ━━━━━━━━ 0s 6ms/step - loss: 0.0118
Epoch 43/50
14/14 ━━━━━━━━ 0s 6ms/step - loss: 0.0134
Epoch 44/50
14/14 ━━━━━━━━ 0s 6ms/step - loss: 0.0137
Epoch 45/50
14/14 ━━━━━━━━ 0s 5ms/step - loss: 0.0131
Epoch 46/50
14/14 ━━━━━━━━ 0s 5ms/step - loss: 0.0091
Epoch 47/50
14/14 ━━━━━━━━ 0s 6ms/step - loss: 0.0124
Epoch 48/50
14/14 ━━━━━━━━ 0s 5ms/step - loss: 0.0103
Epoch 49/50
14/14 ━━━━━━━━ 0s 5ms/step - loss: 0.0118
Epoch 50/50
14/14 ━━━━━━━━ 0s 5ms/step - loss: 0.0100

1/1 ————— 0s 419ms/step

LSTM Performance (Yearly Minimum Temperature):

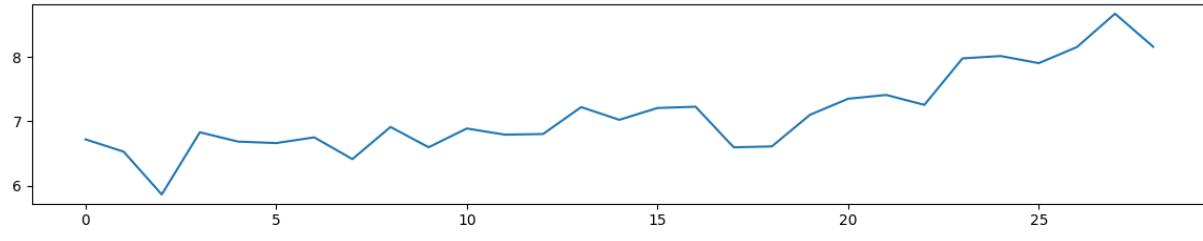
MAE: 0.5630

RMSE: 0.6206

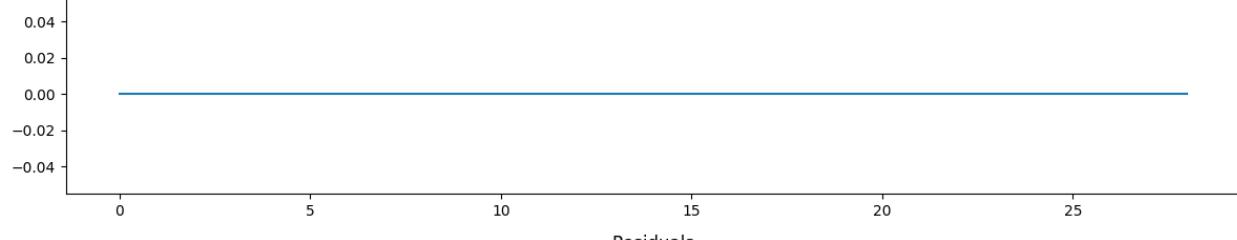


Decomposition Analysis - 5-Year Minimum Temperature

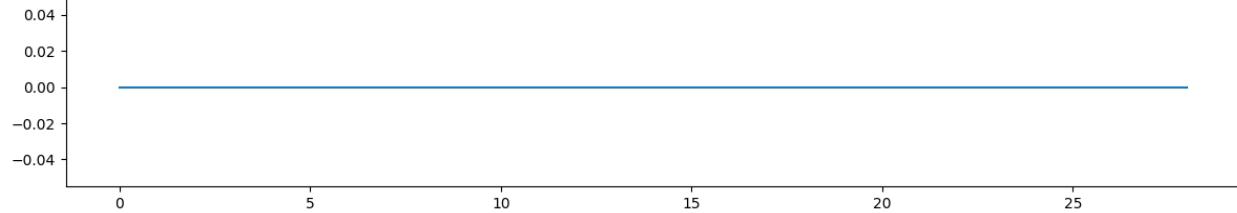
Trend



Seasonality



Residuals



Epoch 1/50

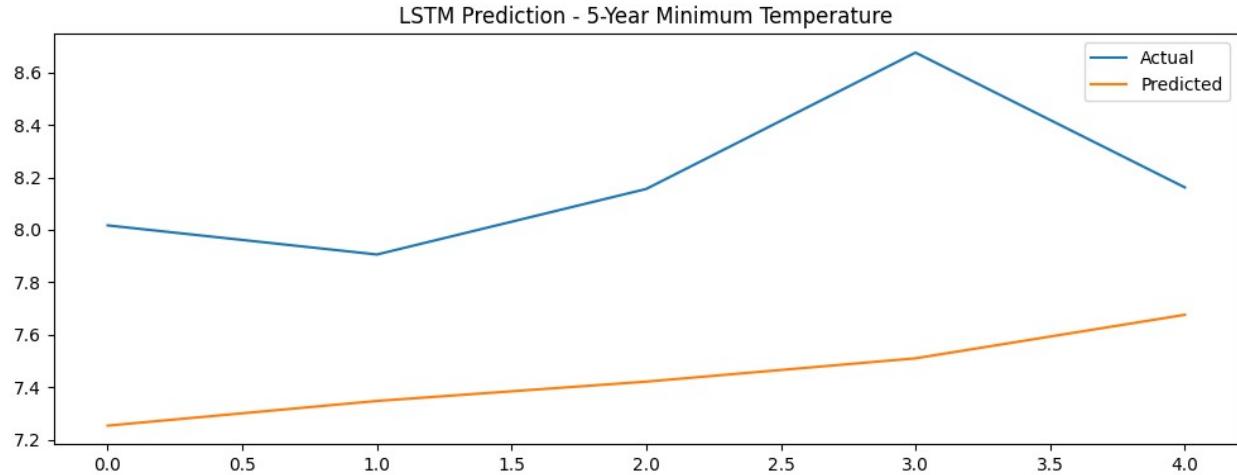
3/3 ————— 4s 9ms/step - loss: 0.1909

Epoch 2/50

```
3/3 ━━━━━━━━━━ 0s 8ms/step - loss: 0.1558
Epoch 3/50
3/3 ━━━━━━ 0s 5ms/step - loss: 0.1475
Epoch 4/50
3/3 ━━━━ 0s 6ms/step - loss: 0.1236
Epoch 5/50
3/3 ━━ 0s 7ms/step - loss: 0.1185
Epoch 6/50
3/3 ━ 0s 6ms/step - loss: 0.1015
Epoch 7/50
3/3 0s 6ms/step - loss: 0.1024
Epoch 8/50
3/3 0s 6ms/step - loss: 0.0780
Epoch 9/50
3/3 0s 5ms/step - loss: 0.0653
Epoch 10/50
3/3 0s 6ms/step - loss: 0.0356
Epoch 11/50
3/3 0s 5ms/step - loss: 0.0259
Epoch 12/50
3/3 0s 9ms/step - loss: 0.0134
Epoch 13/50
3/3 0s 6ms/step - loss: 0.0149
Epoch 14/50
3/3 0s 7ms/step - loss: 0.0200
Epoch 15/50
3/3 0s 7ms/step - loss: 0.0205
Epoch 16/50
3/3 0s 6ms/step - loss: 0.0161
Epoch 17/50
3/3 0s 9ms/step - loss: 0.0168
Epoch 18/50
3/3 0s 6ms/step - loss: 0.0127
Epoch 19/50
3/3 0s 6ms/step - loss: 0.0128
Epoch 20/50
3/3 0s 5ms/step - loss: 0.0176
Epoch 21/50
3/3 0s 10ms/step - loss: 0.0178
Epoch 22/50
3/3 0s 7ms/step - loss: 0.0157
Epoch 23/50
3/3 0s 7ms/step - loss: 0.0132
Epoch 24/50
3/3 0s 6ms/step - loss: 0.0153
Epoch 25/50
3/3 0s 7ms/step - loss: 0.0115
Epoch 26/50
3/3 0s 7ms/step - loss: 0.0128
```

```
Epoch 27/50
3/3 ━━━━━━━━ 0s 6ms/step - loss: 0.0117
Epoch 28/50
3/3 ━━━━━━━━ 0s 6ms/step - loss: 0.0156
Epoch 29/50
3/3 ━━━━━━━━ 0s 6ms/step - loss: 0.0136
Epoch 30/50
3/3 ━━━━━━━━ 0s 7ms/step - loss: 0.0124
Epoch 31/50
3/3 ━━━━━━━━ 0s 7ms/step - loss: 0.0163
Epoch 32/50
3/3 ━━━━━━━━ 0s 8ms/step - loss: 0.0126
Epoch 33/50
3/3 ━━━━━━━━ 0s 5ms/step - loss: 0.0144
Epoch 34/50
3/3 ━━━━━━━━ 0s 7ms/step - loss: 0.0144
Epoch 35/50
3/3 ━━━━━━━━ 0s 9ms/step - loss: 0.0150
Epoch 36/50
3/3 ━━━━━━━━ 0s 6ms/step - loss: 0.0156
Epoch 37/50
3/3 ━━━━━━━━ 0s 6ms/step - loss: 0.0164
Epoch 38/50
3/3 ━━━━━━━━ 0s 7ms/step - loss: 0.0136
Epoch 39/50
3/3 ━━━━━━━━ 0s 6ms/step - loss: 0.0163
Epoch 40/50
3/3 ━━━━━━━━ 0s 9ms/step - loss: 0.0135
Epoch 41/50
3/3 ━━━━━━━━ 0s 8ms/step - loss: 0.0146
Epoch 42/50
3/3 ━━━━━━━━ 0s 6ms/step - loss: 0.0144
Epoch 43/50
3/3 ━━━━━━━━ 0s 4ms/step - loss: 0.0125
Epoch 44/50
3/3 ━━━━━━━━ 0s 6ms/step - loss: 0.0137
Epoch 45/50
3/3 ━━━━━━━━ 0s 6ms/step - loss: 0.0137
Epoch 46/50
3/3 ━━━━━━━━ 0s 6ms/step - loss: 0.0129
Epoch 47/50
3/3 ━━━━━━━━ 0s 7ms/step - loss: 0.0129
Epoch 48/50
3/3 ━━━━━━━━ 0s 4ms/step - loss: 0.0130
Epoch 49/50
3/3 ━━━━━━━━ 0s 4ms/step - loss: 0.0167
Epoch 50/50
3/3 ━━━━━━━━ 0s 8ms/step - loss: 0.0124
1/1 ━━━━━━━━ 0s 404ms/step
```

```
LSTM Performance (5-Year Minimum Temperature):
MAE: 0.7418
RMSE: 0.7786
```



```
<Sequential name=sequential_52, built=True>

# Train on full data for forecasting
scaler = MinMaxScaler()
series_scaled =
scaler.fit_transform(df_yearly_min['Temperature'].values.reshape(-1,
1))

# Re-train model on entire dataset
X_full, y_full = create_sequences(series_scaled, seq_length=5)
X_full = X_full.reshape((X_full.shape[0], X_full.shape[1], 1))
model_full = Sequential([
    LSTM(50, activation='relu', return_sequences=True, input_shape=(5,
1)),
    LSTM(50, activation='relu'),
    Dense(1)
])
model_full.compile(optimizer='adam', loss='mse')
model_full.fit(X_full, y_full, epochs=50, batch_size=8, verbose=1)

# Forecast next 10 years
future_years = list(range(df_yearly_min['Year'].iloc[-1] + 1,
df_yearly_min['Year'].iloc[-1] + 11))
future_predictions = forecast_future(model_full,
df_yearly_min['Temperature'], 10, scaler)

# Plot
plt.figure(figsize=(10, 4))
plt.plot(df_yearly_min['Year'], df_yearly_min['Temperature'],
```

```
label='Historical')
plt.plot(future_years, future_predictions, label='Forecast (2021–
2030)', linestyle='--', marker='o')
plt.xlabel("Year")
plt.ylabel("Avg Temperature")
plt.title("Yearly Minimum Temperature Forecast (Next 10 Years)")
plt.legend()
plt.tight_layout()
plt.show()

# Print values
forecast_df = pd.DataFrame({'Year': future_years, 'Predicted
Temperature': future_predictions})
print("\nForecasted Temperatures (2021–2030):")
print(forecast_df)

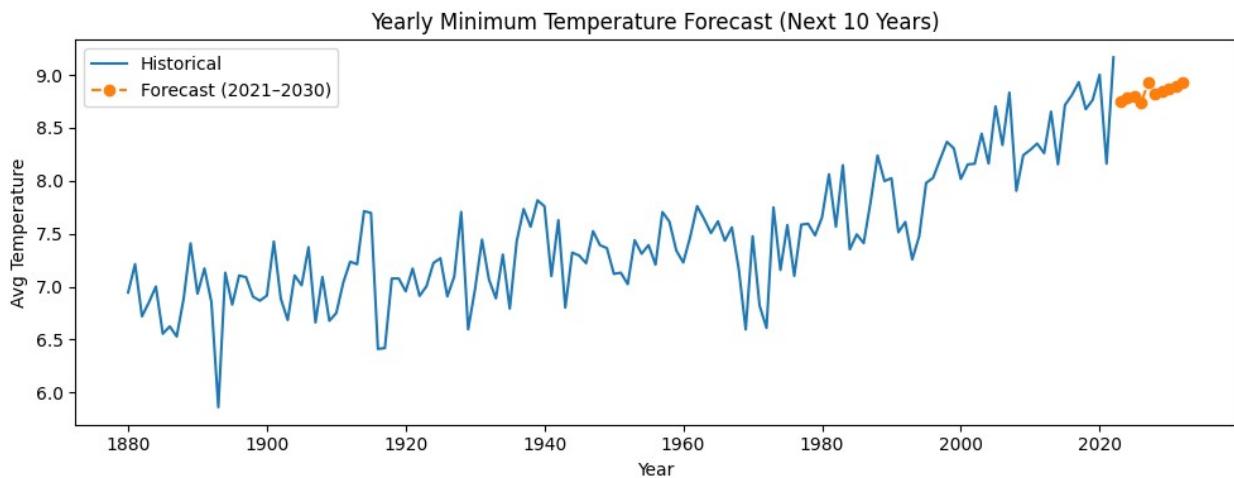
Epoch 1/50
18/18 ━━━━━━━━ 5s 5ms/step - loss: 0.2459
Epoch 2/50
18/18 ━━━━━━ 0s 5ms/step - loss: 0.1163
Epoch 3/50
18/18 ━━━━ 0s 5ms/step - loss: 0.0219
Epoch 4/50
18/18 ━━ 0s 5ms/step - loss: 0.0172
Epoch 5/50
18/18 ━ 0s 5ms/step - loss: 0.0159
Epoch 6/50
18/18 0s 5ms/step - loss: 0.0149
Epoch 7/50
18/18 0s 5ms/step - loss: 0.0169
Epoch 8/50
18/18 0s 5ms/step - loss: 0.0156
Epoch 9/50
18/18 0s 5ms/step - loss: 0.0148
Epoch 10/50
18/18 0s 5ms/step - loss: 0.0152
Epoch 11/50
18/18 0s 5ms/step - loss: 0.0121
Epoch 12/50
18/18 0s 5ms/step - loss: 0.0105
Epoch 13/50
18/18 0s 5ms/step - loss: 0.0137
Epoch 14/50
18/18 0s 5ms/step - loss: 0.0125
Epoch 15/50
18/18 0s 5ms/step - loss: 0.0120
Epoch 16/50
18/18 0s 5ms/step - loss: 0.0137
Epoch 17/50
18/18 0s 5ms/step - loss: 0.0157
```

```
Epoch 18/50
18/18 ━━━━━━━━ 0s 5ms/step - loss: 0.0139
Epoch 19/50
18/18 ━━━━━━━━ 0s 5ms/step - loss: 0.0113
Epoch 20/50
18/18 ━━━━━━━━ 0s 5ms/step - loss: 0.0134
Epoch 21/50
18/18 ━━━━━━━━ 0s 5ms/step - loss: 0.0125
Epoch 22/50
18/18 ━━━━━━━━ 0s 5ms/step - loss: 0.0136
Epoch 23/50
18/18 ━━━━━━━━ 0s 5ms/step - loss: 0.0168
Epoch 24/50
18/18 ━━━━━━━━ 0s 5ms/step - loss: 0.0117
Epoch 25/50
18/18 ━━━━━━━━ 0s 5ms/step - loss: 0.0117
Epoch 26/50
18/18 ━━━━━━━━ 0s 5ms/step - loss: 0.0117
Epoch 27/50
18/18 ━━━━━━━━ 0s 5ms/step - loss: 0.0148
Epoch 28/50
18/18 ━━━━━━━━ 0s 5ms/step - loss: 0.0130
Epoch 29/50
18/18 ━━━━━━━━ 0s 5ms/step - loss: 0.0121
Epoch 30/50
18/18 ━━━━━━━━ 0s 5ms/step - loss: 0.0133
Epoch 31/50
18/18 ━━━━━━━━ 0s 5ms/step - loss: 0.0137
Epoch 32/50
18/18 ━━━━━━━━ 0s 5ms/step - loss: 0.0116
Epoch 33/50
18/18 ━━━━━━━━ 0s 5ms/step - loss: 0.0142
Epoch 34/50
18/18 ━━━━━━━━ 0s 6ms/step - loss: 0.0127
Epoch 35/50
18/18 ━━━━━━━━ 0s 5ms/step - loss: 0.0126
Epoch 36/50
18/18 ━━━━━━━━ 0s 5ms/step - loss: 0.0104
Epoch 37/50
18/18 ━━━━━━━━ 0s 5ms/step - loss: 0.0113
Epoch 38/50
18/18 ━━━━━━━━ 0s 6ms/step - loss: 0.0157
Epoch 39/50
18/18 ━━━━━━━━ 0s 6ms/step - loss: 0.0141
Epoch 40/50
18/18 ━━━━━━━━ 0s 5ms/step - loss: 0.0158
Epoch 41/50
18/18 ━━━━━━━━ 0s 6ms/step - loss: 0.0125
Epoch 42/50
```

```

18/18 ━━━━━━━━━━ 0s 8ms/step - loss: 0.0111
Epoch 43/50
18/18 ━━━━━━━━ 0s 5ms/step - loss: 0.0123
Epoch 44/50
18/18 ━━━━━━ 0s 5ms/step - loss: 0.0130
Epoch 45/50
18/18 ━━━━ 0s 6ms/step - loss: 0.0141
Epoch 46/50
18/18 ━━ 0s 6ms/step - loss: 0.0133
Epoch 47/50
18/18 ━ 0s 6ms/step - loss: 0.0162
Epoch 48/50
18/18 0s 6ms/step - loss: 0.0166
Epoch 49/50
18/18 0s 7ms/step - loss: 0.0143
Epoch 50/50
18/18 0s 7ms/step - loss: 0.0118

```



Forecasted Temperatures (2021–2030):

	Year	Predicted Temperature
0	2023	8.751654
1	2024	8.786212
2	2025	8.795724
3	2026	8.736594
4	2027	8.928153
5	2028	8.822813
6	2029	8.849542
7	2030	8.869045
8	2031	8.889892
9	2032	8.934145

```
# Filter data for the last decade
```

```
last_decade_df = df[(df['Year'] >= 2010) & (df['Year'] < 2020)].copy()
```

```

# Normalize temperature
scaler = MinMaxScaler(feature_range=(0, 1))
last_decade_df['Temperature_Scaled'] =
scaler.fit_transform(last_decade_df[['Temperature']])

# Prepare sequences
seq_length = 30
data = last_decade_df['Temperature_Scaled'].values

# Use last 30 days for prediction
last_sequence = data[-seq_length:].reshape(1, seq_length, 1)

# Load your last trained model (or retrain here if needed)

# Forecast next 10 days recursively
future_predictions = []

current_sequence = last_sequence.copy()

for _ in range(10):
    next_pred = model.predict(current_sequence)[0, 0]
    future_predictions.append(next_pred)

    # Update the sequence with the new prediction
    current_sequence = np.append(current_sequence[:, 1:, :], [[[next_pred]]], axis=1)

# Inverse transform predictions to original scale
future_predictions_actual =
scaler.inverse_transform(np.array(future_predictions).reshape(-1, 1))

# Show predictions
for i, temp in enumerate(future_predictions_actual.flatten(),
start=1):
    print(f"Day {i}: Predicted Temp = {temp:.3f}°C")

# Plotting
plt.figure(figsize=(10, 4))
plt.plot(range(1, 11), future_predictions_actual, marker='o',
linestyle='--')
plt.title("LSTM Forecast: Next 10 Days of Temperature")
plt.xlabel("Day")
plt.ylabel("Predicted Temperature (°C)")
plt.grid(True)
plt.tight_layout()
plt.show()

import pandas as pd
import numpy as np
import matplotlib.pyplot as plt

```

```

import seaborn as sns
from statsmodels.tsa.stattools import adfuller
from statsmodels.graphics.tsaplots import plot_acf, plot_pacf
from statsmodels.tsa.arima.model import ARIMA
from sklearn.metrics import mean_absolute_error, mean_squared_error
import tensorflow as tf
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import LSTM, Dense
from sklearn.preprocessing import MinMaxScaler
from sklearn.model_selection import train_test_split

# Set group size to 35 years
group_size = 35
start_year = 1880
end_year = 2020

def adf_test(series):
    result = adfuller(series.dropna())
    print(f"ADF Statistic: {result[0]}")
    print(f"p-value: {result[1]}")
    if result[1] > 0.05:
        print("Non-stationary series (p > 0.05)")
    else:
        print("Stationary series (p <= 0.05)")

def create_sequences(data, seq_length=30):
    X, y = [], []
    for i in range(len(data) - seq_length):
        X.append(data[i:i+seq_length])
        y.append(data[i+seq_length])
    return np.array(X), np.array(y)

mae_rmse_results = []

for start in range(start_year, end_year, group_size):
    df_group = df[(df['Year'] >= start) & (df['Year'] < start + group_size)].copy()
    if df_group.empty:
        continue

    print(f"\nAnalyzing period: {start}-{start + group_size}")

    # Normalize Data
    scaler = MinMaxScaler(feature_range=(0,1))
    df_group['Temperature_Scaled'] =
    scaler.fit_transform(df_group[['Temperature']])

    # Prepare Sequences for LSTM
    seq_length = 30
    X, y = create_sequences(df_group['Temperature_Scaled'].values,

```

```

seq_length)

# Reshape input for LSTM
X = X.reshape((X.shape[0], X.shape[1], 1))

# Train-Test Split
X_train, X_test, y_train, y_test = train_test_split(X, y,
test_size=0.2, shuffle=False)

# Build LSTM Model
model = Sequential([
    LSTM(50, activation='relu', return_sequences=True,
input_shape=(seq_length, 1)),
    LSTM(50, activation='relu'),
    Dense(1)
])

model.compile(optimizer='adam', loss='mse')

# Train Model
model.fit(X_train, y_train, epochs=20, batch_size=16, verbose=1)

# Predictions
y_pred_lstm = model.predict(X_test)
y_pred_lstm = scaler.inverse_transform(y_pred_lstm)

# Evaluate Performance
mae_lstm =
mean_absolute_error(scaler.inverse_transform(y_test.reshape(-1,1)),
y_pred_lstm)
rmse_lstm =
np.sqrt(mean_squared_error(scaler.inverse_transform(y_test.reshape(
-1,1)), y_pred_lstm))

mae_rmse_results.append([f'{start}-{start + group_size}',
mae_lstm, rmse_lstm])
print(f'LSTM MAE {start}-{start + group_size}: {mae_lstm:.4f},'
RMSE: {rmse_lstm:.4f} ')

# Create DataFrame for results
results_df = pd.DataFrame(mae_rmse_results, columns=["Time Period",
"MAE", "RMSE"])
print(results_df)

# Plot Results
plt.figure(figsize=(12, 6))
plt.plot(results_df["Time Period"], results_df["MAE"], marker='o',
label="MAE")
plt.plot(results_df["Time Period"], results_df["RMSE"], marker='s',

```

```

label="RMSE")
plt.xlabel("Time Period")
plt.ylabel("Error Value")
plt.title("MAE and RMSE for Each 35-Year Group")
plt.xticks(rotation=45)
plt.legend()
plt.tight_layout()
plt.show()

```

Model MAE (Lower is better) RMSE (Lower is better)
 Pros Cons ARIMA 0.3038, 0.3772 Works well for small data, interpretable
 Fails for non-linear patterns LSTM 0.1130, 0.1429 Captures complex relationships, good for long-term forecasting
 Computationally expensive, needs large data

```

from statsmodels.tsa.seasonal import seasonal_decompose

# Decompose the time series
decomposition = seasonal_decompose(df['Temperature'],
model='additive', period=365) # Assuming yearly seasonality

# Plot components
fig, (ax1, ax2, ax3) = plt.subplots(3, 1, figsize=(14,8))
decomposition.trend.plot(ax=ax1, title='Trend')
decomposition.seasonal.plot(ax=ax2, title='Seasonality')
decomposition.resid.plot(ax=ax3, title='Residuals')
plt.tight_layout()
plt.show()

import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from statsmodels.tsa.stattools import adfuller
from statsmodels.graphics.tsaplots import plot_acf, plot_pacf
from statsmodels.tsa.arima.model import ARIMA
from sklearn.metrics import mean_absolute_error, mean_squared_error
import tensorflow as tf
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import LSTM, Dense
from sklearn.preprocessing import MinMaxScaler
from sklearn.model_selection import train_test_split
from statsmodels.tsa.seasonal import seasonal_decompose

group_size = 20
start_year = 1880
end_year = 2020

def adf_test(series):
    result = adfuller(series.dropna())

```

```

print(f"ADF Statistic: {result[0]}")
print(f"p-value: {result[1]}")
if result[1] > 0.05:
    print("Non-stationary series (p > 0.05)")
else:
    print("Stationary series (p <= 0.05)")

def create_sequences(data, seq_length=10):
    X, y = [], []
    for i in range(len(data) - seq_length):
        X.append(data[i:i+seq_length])
        y.append(data[i+seq_length])
    return np.array(X), np.array(y)

mae_rmse_results = []

for start in range(start_year, end_year, group_size):
    df_group = df[(df['Year'] >= start) & (df['Year'] < start + group_size)].copy()
    print(f"Analyzing period: {start}-{start + group_size}")

    # Decomposition
    decomposition = seasonal_decompose(df_group['Temperature'],
                                         model='additive', period=365)
    fig, (ax1, ax2, ax3) = plt.subplots(3, 1, figsize=(14,8))
    decomposition.trend.plot(ax=ax1, title='Trend')
    decomposition.seasonal.plot(ax=ax2, title='Seasonality')
    decomposition.resid.plot(ax=ax3, title='Residuals')
    plt.tight_layout()
    plt.show()

    # Normalize Data
    scaler = MinMaxScaler(feature_range=(0,1))
    df_group['Temperature_Scaled'] =
    scaler.fit_transform(df_group[['Temperature']])

    # Prepare Sequences for LSTM
    seq_length = 30 # Using 30 days of past data
    X, y = create_sequences(df_group['Temperature_Scaled'].values,
                           seq_length)

    # Train-Test Split
    X_train, X_test, y_train, y_test = train_test_split(X, y,
                                                       test_size=0.2, shuffle=False)

    # Build LSTM Model
    model = Sequential([
        LSTM(50, activation='relu', return_sequences=True,
             input_shape=(seq_length, 1)),
        LSTM(50, activation='relu'),

```

```

        Dense(1)
    ])

model.compile(optimizer='adam', loss='mse')

# Train Model
model.fit(X_train, y_train, epochs=20, batch_size=16, verbose=1)

# Predictions
y_pred_lstm = model.predict(X_test)
y_pred_lstm = scaler.inverse_transform(y_pred_lstm)

# Evaluate Performance
mae_lstm =
mean_absolute_error(scaler.inverse_transform(y_test.reshape(-1,1)),
y_pred_lstm)
rmse_lstm =
np.sqrt(mean_squared_error(scaler.inverse_transform(y_test.reshape(
-1,1)), y_pred_lstm))

mae_rmse_results.append([f'{start}-{start + group_size}',
mae_lstm, rmse_lstm])
print(f'LSTM MAE {start}-{start + group_size}: {mae_lstm:.4f}, RMSE: {rmse_lstm:.4f}')

# Create DataFrame for results
results_df = pd.DataFrame(mae_rmse_results, columns=["Time Period",
"MAE", "RMSE"])
print(results_df)

# Plot Results
plt.figure(figsize=(12, 6))
plt.plot(results_df["Time Period"], results_df["MAE"], marker='o',
label="MAE")
plt.plot(results_df["Time Period"], results_df["RMSE"], marker='s',
label="RMSE")
plt.xlabel("Time Period")
plt.ylabel("Error Value")
plt.title("MAE and RMSE for Each 10-Year Group")
plt.xticks(rotation=45)
plt.legend()
plt.show()

import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from statsmodels.tsa.stattools import adfuller
from statsmodels.graphics.tsaplots import plot_acf, plot_pacf
from statsmodels.tsa.arima.model import ARIMA

```

```

from sklearn.metrics import mean_absolute_error, mean_squared_error
import tensorflow as tf
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import LSTM, Dense
from sklearn.preprocessing import MinMaxScaler
from sklearn.model_selection import train_test_split
from statsmodels.tsa.seasonal import seasonal_decompose

group_size = 35
start_year = 1880
end_year = 2020

def adf_test(series):
    result = adfuller(series.dropna())
    print(f"ADF Statistic: {result[0]}")
    print(f"p-value: {result[1]}")
    if result[1] > 0.05:
        print("Non-stationary series (p > 0.05)")
    else:
        print("Stationary series (p <= 0.05)")

def create_sequences(data, seq_length=10):
    X, y = [], []
    for i in range(len(data) - seq_length):
        X.append(data[i:i+seq_length])
        y.append(data[i+seq_length])
    return np.array(X), np.array(y)

mae_rmse_results = []

for start in range(start_year, end_year, group_size):
    df_group = df[(df['Year'] >= start) & (df['Year'] < start + group_size)].copy()
    print(f"Analyzing period: {start}-{start + group_size}")

    # Decomposition
    decomposition = seasonal_decompose(df_group['Temperature'], model='additive', period=365)
    fig, (ax1, ax2, ax3) = plt.subplots(3, 1, figsize=(14,8))
    decomposition.trend.plot(ax=ax1, title='Trend')
    decomposition.seasonal.plot(ax=ax2, title='Seasonality')
    decomposition.resid.plot(ax=ax3, title='Residuals')
    plt.tight_layout()
    plt.show()

    # Normalize Data
    scaler = MinMaxScaler(feature_range=(0,1))
    df_group['Temperature_Scaled'] =
    scaler.fit_transform(df_group[['Temperature']])

```

```

# Prepare Sequences for LSTM
seq_length = 30 # Using 30 days of past data
X, y = create_sequences(df_group['Temperature_Scaled'].values,
seq_length)

# Train-Test Split
X_train, X_test, y_train, y_test = train_test_split(X, y,
test_size=0.2, shuffle=False)

# Build LSTM Model
model = Sequential([
    LSTM(50, activation='relu', return_sequences=True,
input_shape=(seq_length, 1)),
    LSTM(50, activation='relu'),
    Dense(1)
])

model.compile(optimizer='adam', loss='mse')

# Train Model
model.fit(X_train, y_train, epochs=20, batch_size=16, verbose=1)

# Predictions
y_pred_lstm = model.predict(X_test)
y_pred_lstm = scaler.inverse_transform(y_pred_lstm)

# Evaluate Performance
mae_lstm =
mean_absolute_error(scaler.inverse_transform(y_test.reshape(-1,1)),
y_pred_lstm)
rmse_lstm =
np.sqrt(mean_squared_error(scaler.inverse_transform(y_test.reshape(
-1,1)), y_pred_lstm))

mae_rmse_results.append([f'{start}-{start + group_size}',
mae_lstm, rmse_lstm])
print(f'LSTM MAE {start}-{start + group_size}: {mae_lstm:.4f},\nRMSE: {rmse_lstm:.4f}')

# Create DataFrame for results
results_df = pd.DataFrame(mae_rmse_results, columns=["Time Period",
"MAE", "RMSE"])
print(results_df)

# Plot Results
plt.figure(figsize=(12, 6))
plt.plot(results_df["Time Period"], results_df["MAE"], marker='o',
label="MAE")
plt.plot(results_df["Time Period"], results_df["RMSE"], marker='s',

```

```

label="RMSE")
plt.xlabel("Time Period")
plt.ylabel("Error Value")
plt.title("MAE and RMSE for Each 10-Year Group")
plt.xticks(rotation=45)
plt.legend()
plt.show()

from statsmodels.graphics.tsaplots import plot_acf

# ACF Plot
plt.figure(figsize=(10,5))
plot_acf(df['Temperature'], lags=365)
plt.title("Autocorrelation Function (ACF) Plot")
plt.show()

import seaborn as sns

df['Month'] = df.index.month # Extract month
plt.figure(figsize=(12,6))
sns.boxplot(x='Month', y='Temperature', data=df)
plt.title("Monthly Temperature Variation")
plt.xlabel("Month")
plt.ylabel("Temperature")
plt.show()

```

Forecasting

```

combined_df = results_df.copy()
combined_df['ARIMA_MAE'] = results_arima_df['MAE']
combined_df['ARIMA_RMSE'] = results_arima_df['RMSE']

plt.figure(figsize=(14, 6))
plt.plot(combined_df['Time Period'], combined_df['MAE'], label='LSTM MAE', marker='o')
plt.plot(combined_df['Time Period'], combined_df['ARIMA_MAE'], label='ARIMA MAE', marker='s')
plt.plot(combined_df['Time Period'], combined_df['RMSE'], label='LSTM RMSE', marker='o', linestyle='--')
plt.plot(combined_df['Time Period'], combined_df['ARIMA_RMSE'], label='ARIMA RMSE', marker='s', linestyle='--')
plt.xticks(rotation=45)
plt.title("Model Performance Comparison: LSTM vs. ARIMA")
plt.xlabel("Time Period")
plt.ylabel("Error")

```

```
plt.legend()
plt.grid(True)
plt.tight_layout()
plt.show()

# Fit ARIMA on entire dataset
model_arima = ARIMA(df['Temperature'], order=(5,1,2))
result_arima = model_arima.fit()

# Forecast next 30 days
forecast_arima = result_arima.forecast(steps=30)

# Plot
plt.figure(figsize=(10, 5))
plt.plot(df['Temperature'], label='Original')
plt.plot(pd.date_range(df.index[-1], periods=30, freq='D'),
forecast_arima, label='ARIMA Forecast', color='green')
plt.title("ARIMA Forecast for Next 30 Days")
plt.xlabel("Date")
plt.ylabel("Temperature")
plt.legend()
plt.show()
```